

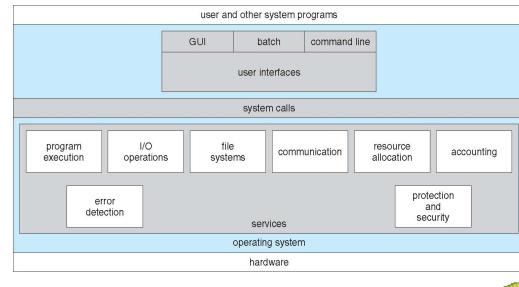
Chapter 2: System Structures



Operating System Concepts – 9th Edition

Silberschatz, Galvin and Gagne ©2013

A View of Operating System Services



Operating System Services (Cont)

- Another set of OS functions exists for ensuring the efficient operation of the system itself via resource sharing
 - **Resource allocation** - When multiple users or multiple jobs running concurrently, resources must be allocated to each of them
 - Many types of resources - Some (such as CPU cycles, main memory, and file storage) may have special allocation code, others (such as I/O devices) may have general request and release code
 - **Accounting** - To keep track of which users use how much and what kinds of computer resources
 - **Protection and security** - The owners of information stored in a multiuser or networked computer system may want to control use of that information, concurrent processes should not interfere with each other
 - **Protection** involves ensuring that all access to system resources is controlled
 - **Security** of the system from outsiders requires user authentication, extends to defending external I/O devices from invalid access attempts
 - If a system is to be protected and secure, precautions must be instituted throughout it. A chain is only as strong as its weakest link.

Operating System Concepts – 9th Edition

2.7

Silberschatz, Galvin and Gagne ©2013

Chapter 2: Operating System Structures

- Operating System Services
- System Calls
- System Programs
- Operating System Design and Implementation
- Operating System Structure
- Virtual Machines

Operating System Concepts – 9th Edition

2.2

Silberschatz, Galvin and Gagne ©2013

Operating System Services

- One set of operating-system services provides functions that are helpful to the user:
 - **User interface** - Almost all operating systems have a user interface (UI)
 - Varies between Command-Line (CLI), Graphics User Interface (GUI), Batch
 - **Program execution** - The system must be able to load a program into memory and to run that program, end execution, either normally or abnormally (indicating error)
 - **I/O operations** - A running program may require I/O, which may involve a file or an I/O device
 - **File-system manipulation** - The file system is of particular interest. Obviously, programs need to read and write files and directories, create and delete them, search them, list file information, permission management.

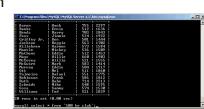
Operating System Concepts – 9th Edition

2.5

Silberschatz, Galvin and Gagne ©2013

User Operating System Interface - CLI

- Command Line Interface (CLI) or **command interpreter** allows direct command entry
 - Sometimes implemented in kernel, sometimes by systems program
 - Sometimes multiple flavors implemented – **shells**
 - Primarily fetches a command from user and executes it
 - Sometimes commands built-in, sometimes just names of programs
 - If the latter, adding new features doesn't require shell modification



Operating System Concepts – 9th Edition

2.8

Silberschatz, Galvin and Gagne ©2013

Objectives

- To describe the services an operating system provides to users, processes, and other systems
- To discuss the various ways of structuring an operating system

Operating System Concepts – 9th Edition

2.3

Silberschatz, Galvin and Gagne ©2013

Operating System Services (Cont)

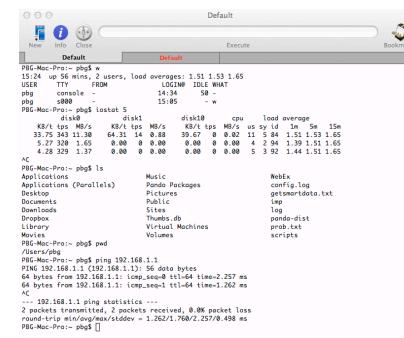
- One set of operating-system services provides functions that are helpful to the user (Cont):
 - **Communications** – Processes may exchange information, on the same computer or between computers over a network
 - Communications may be via shared memory or through message passing (packets moved by the OS)
 - **Error detection** – OS needs to be constantly aware of possible errors
 - May occur in the CPU and memory hardware, in I/O devices, in user program
 - For each type of error, OS should take the appropriate action to ensure correct and consistent computing
 - Debugging facilities can greatly enhance the user's and programmer's abilities to efficiently use the system

Operating System Concepts – 9th Edition

2.6

Silberschatz, Galvin and Gagne ©2013

Bourne Shell Command Interpreter



Operating System Concepts – 9th Edition

2.9

Silberschatz, Galvin and Gagne ©2013



User Operating System Interface - GUI

- User-friendly [desktop](#) metaphor interface
 - Usually mouse, keyboard, and monitor
 - [Icons](#) represent files, programs, actions, etc
 - Various mouse buttons over objects in the interface cause various actions (provide information, options, execute function, open directory (known as a [folder](#)))
 - Invented at Xerox PARC

- Many systems now include both CLI and GUI interfaces
 - Microsoft Windows is GUI with CLI "command" shell
 - Apple Mac OS X is "Aqua" GUI interface with UNIX kernel underneath and shells available
 - Unix and Linux have CLI with optional GUI interfaces (CDE, KDE, GNOME)

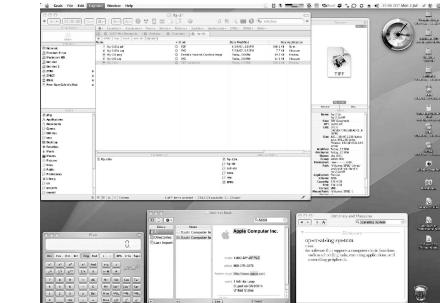
Operating System Concepts – 9th Edition

2.10

Silberschatz, Galvin and Gagne ©2013



The Mac OS X GUI

Operating System Concepts – 9th Edition

2.12

Silberschatz, Galvin and Gagne ©2013



System Calls

- Programming interface to the services provided by the OS
- Typically written in a high-level language (C or C++)
- Mostly accessed by programs via a high-level [Application Program Interface \(API\)](#) rather than direct system call use
- Three common APIs
 - Win32 API for Windows
 - POSIX API for POSIX-based systems (including virtually all versions of UNIX, Linux, and Mac OS X)
 - Java API for the Java virtual machine (JVM)
- Why use APIs rather than system calls?
 - Hide the complex details of the system call from users
 - Program portability

Operating System Concepts – 9th Edition

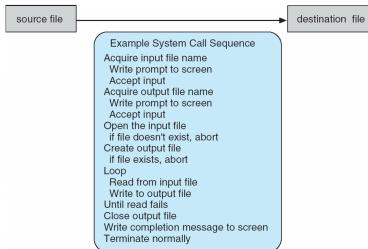
2.13

Silberschatz, Galvin and Gagne ©2013



Example of System Calls

- System call sequence to copy the contents of one file to another file

Operating System Concepts – 9th Edition

2.14

Silberschatz, Galvin and Gagne ©2013



Example of Standard API

EXAMPLE OF STANDARD API

As an example of a standard API, consider the `read()` function that is available in UNIX and Linux systems. The API for this function is obtained from the man page by invoking the command

`man read`

on the command line. A description of this API appears below:

```
#include <unistd.h>
ssize_t read(int fd, void *buf, size_t count)
    return value
        function name
            parameters
```

A program that uses the `read()` function must include the `unistd.h` header file, as this file defines the `ssize_t` and `size_t` data types (among other things). The parameters passed to `read()` are as follows:

- `int fd`—the file descriptor to be read
- `void *buf`—a buffer where the data will be read into
- `size_t count`—the maximum number of bytes to be read into the buffer

On a successful read, the number of bytes read is returned. A return value of 0 indicates end of file. If an error occurs, `read()` returns -1.

Operating System Concepts – 9th Edition

2.15

Silberschatz, Galvin and Gagne ©2013



System Call Implementation

- Typically, a number associated with each system call
 - System-call interface maintains a table indexed according to these numbers
- The system call interface invokes intended system call in OS kernel and returns status of the system call and any return values
- The caller need know nothing about how the system call is implemented
 - Just needs to obey API and understand what OS will do as a result call
 - Most details of OS interface hidden from programmer by API

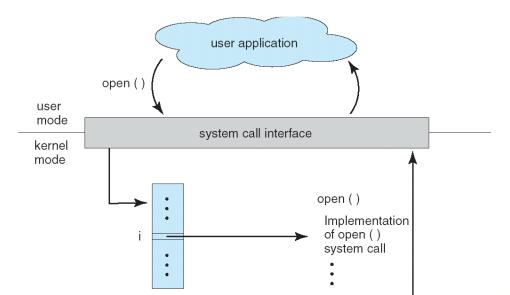
Operating System Concepts – 9th Edition

2.16

Silberschatz, Galvin and Gagne ©2013



API – System Call – OS Relationship

Operating System Concepts – 9th Edition

2.17

Silberschatz, Galvin and Gagne ©2013



System Call Parameter Passing

- Often, more information is required than simply the identity of desired system call
 - Exact type and amount of information vary according to OS and call
- Three general methods used to pass parameters to the OS
 - Simplest: pass the parameters in registers
 - In some cases, may be more parameters than registers
 - Parameters stored in a block, or table, in memory, and address of block passed as a parameter in a register
 - This approach taken by Linux and Solaris
 - Parameters placed, or [pushed](#), onto the [stack](#) by the program and [popped](#) off the stack by the operating system
 - Block and stack methods do not limit the number or length of parameters being passed

Operating System Concepts – 9th Edition

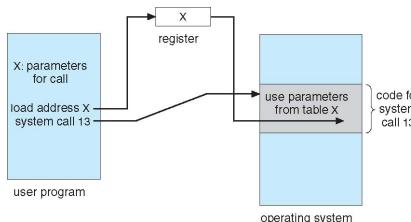
2.18

Silberschatz, Galvin and Gagne ©2013





Parameter Passing via Table

Operating System Concepts - 9th Edition

2.19

Silberschatz, Galvin and Gagne ©2013



Types of System Calls

Process control

- end, abort
- load, execute
- create process, terminate process
- get process attributes, set process attributes
- wait for time
- wait event, signal event
- allocate and free memory

File management

- create file, delete file
- open, close file
- read, write, reposition
- get and set file attributes

Operating System Concepts - 9th Edition

2.20

Silberschatz, Galvin and Gagne ©2013



Device management

- request device, release device
- read, write, reposition
- get device attributes, set device attributes
- logically attach or detach devices

Information maintenance

- get time or date, set time or date
- get system data, set system data
- get and set process, file, or device attributes

Operating System Concepts - 9th Edition

2.21

Silberschatz, Galvin and Gagne ©2013



Types of System Calls (Cont.)

Communications

- create, delete communication connection
- send, receive messages if **message passing model** to **host name** or **process name**
 - ▶ From client to server
- **Shared-memory model** create and gain access to memory regions
- transfer status information
- attach and detach remote devices

Protection

- Control access to resources
- Get and set permissions
- Allow and deny user access

Operating System Concepts - 9th Edition

2.22

Silberschatz, Galvin and Gagne ©2013



Examples of Windows and Unix System Calls

	Windows	Unix
Process Control	CreateProcess() ExitProcess() WaitForSingleObject()	fork() exit() wait()
File Manipulation	CreateFile() ReadFile() WriteFile() CloseHandle()	open() read() write() close()
Device Manipulation	SetConsoleMode() ReadConsole() WriteConsole()	ioctl() read() write()
Information Maintenance	GetCurrentProcessID() SetTimer() Sleep()	getpid() alarm() sleep()
Communication	CreatePipe() CreateFileMapping() MapViewOfFile()	pipe() shmget() map()
Protection	SetFileSecurity() InitializeSecurityDescriptor() SetSecurityDescriptorGroup()	chmod() unmap() chown()

Operating System Concepts - 9th Edition

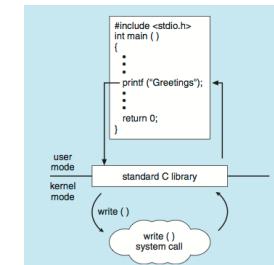
2.23

Silberschatz, Galvin and Gagne ©2013



Standard C Library Example

- C program invoking printf() library call, which calls write() system call

Operating System Concepts - 9th Edition

2.24

Silberschatz, Galvin and Gagne ©2013



System Programs

System programs provide a convenient environment for program development and execution. They can be divided into:

- File manipulation
- Status information sometimes stored in a File modification
- Programming language support
- Program loading and execution
- Communications
- Background services
- Application programs

Most users' view of an operation system is defined by system programs, not the actual system calls

Operating System Concepts - 9th Edition

2.25

Silberschatz, Galvin and Gagne ©2013



System Programs

Provide a convenient environment for program development and execution

- Some of them are simply user interfaces to system calls; others are considerably more complex

File management - Create, delete, copy, rename, print, dump, list, and generally manipulate files and directories

Status information

- Some ask the system for info - date, time, amount of available memory, disk space, number of users
- Others provide detailed performance, logging, and debugging information
- Typically, these programs format and print the output to the terminal or other output devices
- Some systems implement a **registry** - used to store and retrieve configuration information

Operating System Concepts - 9th Edition

2.26

Silberschatz, Galvin and Gagne ©2013



System Programs (Cont.)

File modification

- Text editors to create and modify files
- Special commands to search contents of files or perform transformations of the text

Programming-language support

- Compilers, assemblers, debuggers and interpreters sometimes provided

Program loading and execution

- Absolute loaders, relocatable loaders, linkage editors, and overlay-loaders, debugging systems for higher-level and machine language

Communications

- Provide the mechanism for creating virtual connections among processes, users, and computer systems

- Allow users to send messages to one another's screens, browse web pages, send electronic-mail messages, log in remotely, transfer files from one machine to another

Operating System Concepts - 9th Edition

2.27

Silberschatz, Galvin and Gagne ©2013





System Programs (Cont.)

- **Background Services**
 - Launch at boot time
 - Some for system startup, then terminate
 - Some from system boot to shutdown
 - Provide facilities like disk checking, process scheduling, error logging, printing
 - Run in user context not kernel context
 - Known as [services](#), [subsystems](#), [daemons](#)

- **Application programs**
 - Don't pertain to system
 - Run by users
 - Not typically considered part of OS
 - Launched by command line, mouse click, finger poke

Operating System Concepts – 9th Edition

2.28

Silberschatz, Galvin and Gagne ©2013



Operating System Design and Implementation

- Design and Implementation of OS not “solvable”, but some approaches have proven successful
- Internal structure of different Operating Systems can vary widely
- Start by defining goals and specifications
- Affected by choice of hardware, type of system
- **User** goals and **System** goals
 - User goals – operating system should be convenient to use, easy to learn, reliable, safe, and fast
 - System goals – operating system should be easy to design, implement, and maintain, as well as flexible, reliable, error-free, and efficient

Operating System Concepts – 9th Edition

2.29

Silberschatz, Galvin and Gagne ©2013



- Important principle to separate
 - Policy:** *What* will be done?
 - Mechanism:** *How* to do it?

- Mechanisms determine how to do something, policies decide what will be done
 - The separation of policy from mechanism is a very important principle, it allows maximum flexibility if policy decisions are to be changed later

- Specifying and designing OS is highly creative task of [software engineering](#)

Operating System Concepts – 9th Edition

2.30

Silberschatz, Galvin and Gagne ©2013



Implementation

- Much variation
 - Early OSes in assembly language
 - Then system programming languages like Algol, PL/I
 - Now C, C++
- Actually usually a mix of languages
 - Lowest levels in assembly
 - Main body in C
 - Systems programs in C, C++, scripting languages like PERL, Python, shell scripts
- More high-level language easier to [port](#) to other hardware
 - But slower

Operating System Concepts – 9th Edition

2.31

Silberschatz, Galvin and Gagne ©2013



Operating System Structure

- General-purpose OS is very large program
- Various ways to structure one as follows

Operating System Concepts – 9th Edition

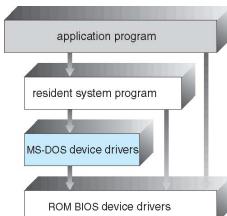
2.32

Silberschatz, Galvin and Gagne ©2013



Simple Structure

- I.e. MS-DOS – written to provide the most functionality in the least space
 - Not divided into modules
 - Although MS-DOS has some structure, its interfaces and levels of functionality are not well separated



Silberschatz, Galvin and Gagne ©2013



Layered Approach

- The operating system is divided into a number of layers (levels), each built on top of lower layers.
 - The bottom layer (layer 0), is the hardware;
 - The highest (layer N) is the user interface.
- With modularity, layers are selected such that each uses functions (operations) and services of only lower-level layers

Operating System Concepts – 9th Edition

2.34

Silberschatz, Galvin and Gagne ©2013



Layered Approach

- The operating system is divided into a number of layers (levels), each built on top of lower layers.
 - The bottom layer (layer 0), is the hardware; the highest (layer N) is the user interface.
- With modularity, layers are selected such that each uses functions (operations) and services of only lower-level layers

Operating System Concepts – 9th Edition

2.35

Silberschatz, Galvin and Gagne ©2013



UNIX

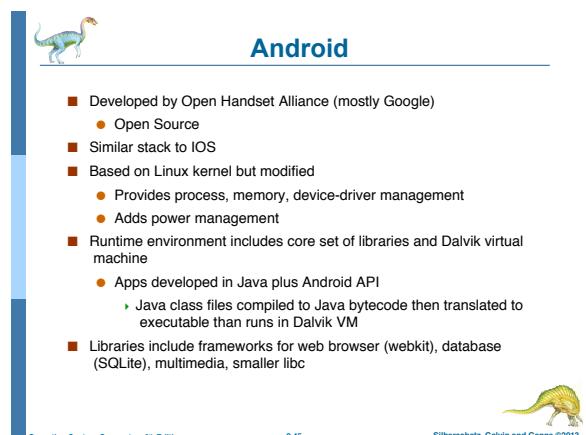
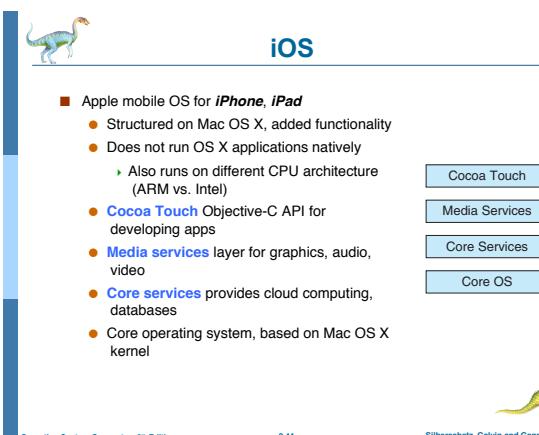
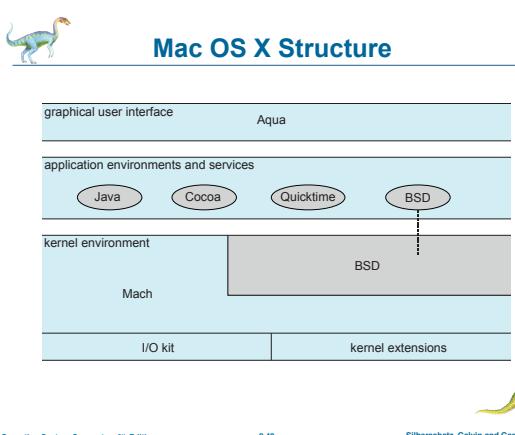
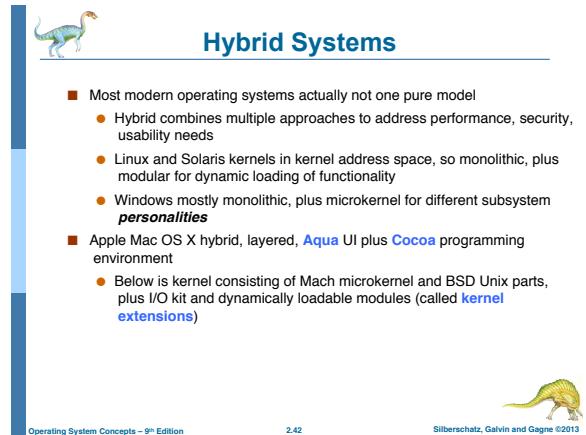
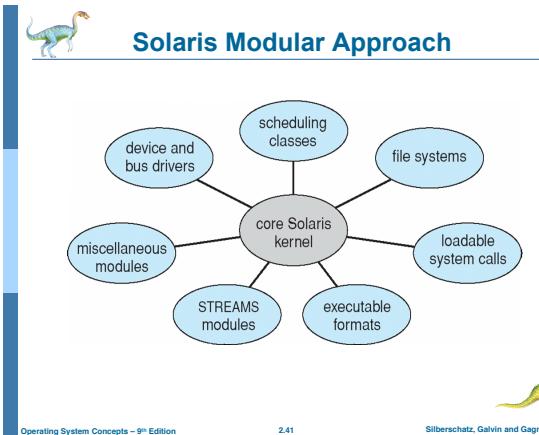
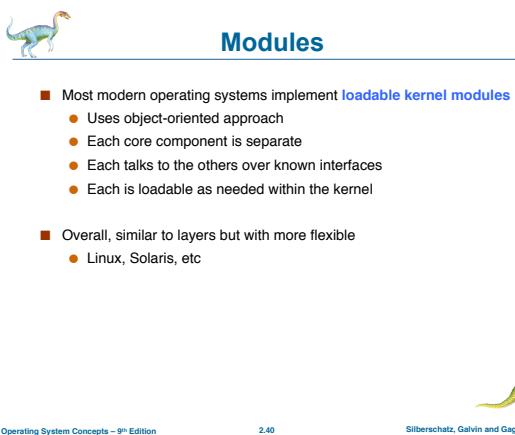
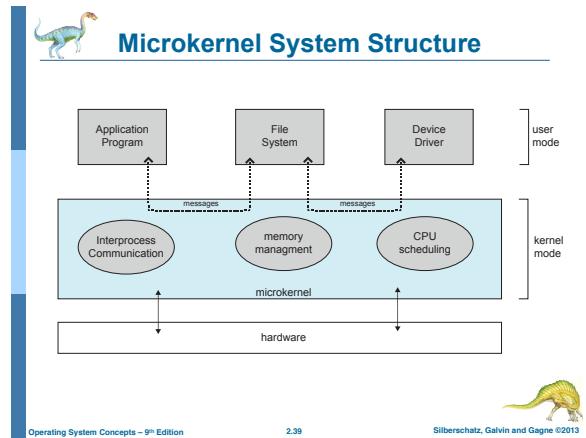
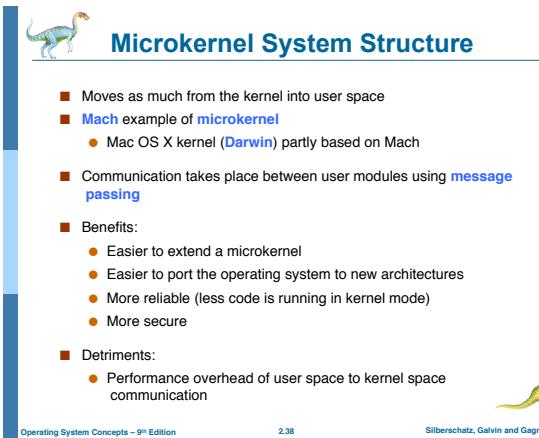
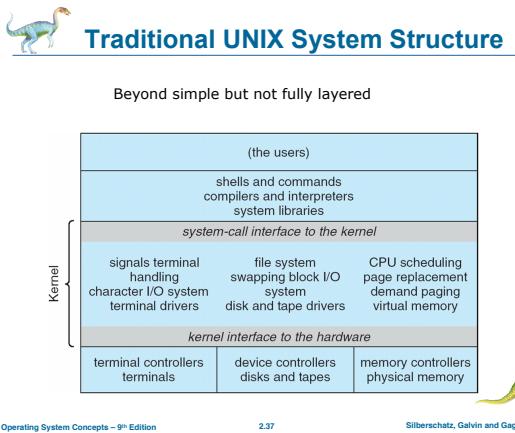
- UNIX – limited by hardware functionality, the original UNIX operating system had limited structuring. The UNIX OS consists of two separable parts
 - Systems programs
 - The kernel
 - Consists of everything below the system-call interface and above the physical hardware
 - Provides the file system, CPU scheduling, memory management, and other operating-system functions; a large number of functions for one level

Operating System Concepts – 9th Edition

2.36

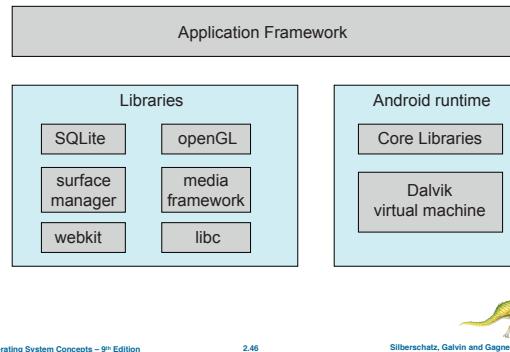
Silberschatz, Galvin and Gagne ©2013







Android Architecture

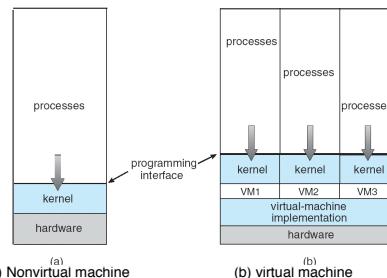
Operating System Concepts – 9th Edition

2.46

Silberschatz, Galvin and Gagne ©2013



Virtual Machines (Cont)

Operating System Concepts – 9th Edition

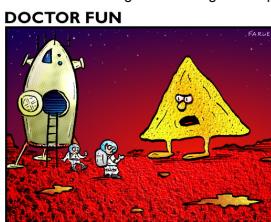
2.49

Silberschatz, Galvin and Gagne ©2013



Nachos: Virtual OS Environment

- You will be working with **Nachos**
 - Simulation environment
 - Hardware, interrupts, I/O
 - Execution of User Programs running on this platform



"This is the planet where nachos rule."

Operating System Concepts – 9th Edition

2.52

Silberschatz, Galvin and Gagne ©2013



Virtual Machines

- The **virtual machine** creates the illusion of multiple processes, each executing on its own processor with its own (virtual) memory
- Software emulation of an abstract machine
 - Make it look like hardware has features you want
 - Programs from one hardware & OS on another one
- Programming simplicity
 - Each process thinks it has all memory/CPU time
 - Each process thinks it owns all devices
 - Different Devices appear to have same interface
- Fault Isolation
 - Processes unable to directly impact other processes
 - Bugs cannot crash whole machine
- Protection and Portability
 - Java interface safe and stable across many platforms

Operating System Concepts – 9th Edition

2.47

Silberschatz, Galvin and Gagne ©2013



Virtual Machines (Cont)

- Allows operating systems to run applications within other OSes
 - Vast and growing industry
- Emulation** used when source CPU type different from target type (i.e. PowerPC to Intel x86)
 - Generally slowest method
 - When computer language not compiled to native code – **Interpretation**
- Virtualization** – OS natively compiled for CPU, running **guest** OSes also natively compiled
 - Consider VMware running WinXP guests, each running applications, all on native WinXP **host** OS
 - VMM** provides virtualization services

Operating System Concepts – 9th Edition

2.48

Silberschatz, Galvin and Gagne ©2013



Virtual Machines (Cont)

- Use cases involve laptops and desktops running multiple OSes for exploration or compatibility
 - Apple laptop running Mac OS X host, Windows as a guest
 - Developing apps for multiple OSes without having multiple systems
 - QA testing applications without having multiple systems
 - Executing and managing compute environments within data centers

Operating System Concepts – 9th Edition

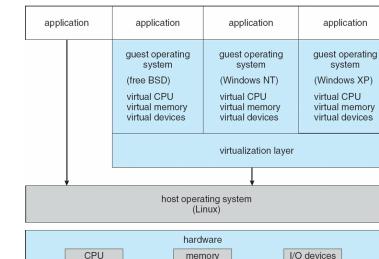
2.50

Silberschatz, Galvin and Gagne ©2013



Virtual Machines (Cont): Layers of OS

- Useful for OS development
 - When OS crashes, restricted to one VM
 - Can aid testing programs on other OS

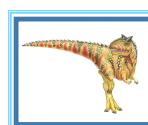
Operating System Concepts – 9th Edition

2.51

Silberschatz, Galvin and Gagne ©2013



End of Chapter 2



Silberschatz, Galvin and Gagne ©2013

