


## Clustering: $k$ -means

- dataset  $D$
- object  $\mathbf{x} = (x_1, x_2, \dots, x_d)$ 
  - every  $\mathbf{x}$  is a point in a  $d$ -dimensional space

**Clustering:** groups the data into clusters

- objects in the same cluster have high similarity
- objects in different clusters are dissimilar to each other



- unlike classification, there is **no class attribute** in clustering
  - this is quite common in large databases, because assigning class labels to a large number of objects can be very costly
- **unsupervised** learning (on the other hand, classification is **supervised** learning)

# Applications of Clustering

- business
  - discover distinct groups of customers based on their purchasing patterns
- biology
  - derive plant and animal taxonomies
- geographical data
  - group houses in a city according to house type and geographical location
- outlier detection
  - find credit card transactions that are not ordinary (fraud detection)




# Dissimilarity

- measures how **different** two objects are (lower when objects are more alike)
- minimum dissimilarity is often 0

$\mathbf{x}_1 = (x_{11}, x_{12}, \dots, x_{1d})$  and  $\mathbf{x}_2 = (x_{21}, x_{22}, \dots, x_{2d})$

- **Euclidean distance**:  $dist(\mathbf{x}_1, \mathbf{x}_2) = \sqrt{\sum_{i=1}^d (x_{1i} - x_{2i})^2}$
- **Manhattan distance**:  $dist(\mathbf{x}_1, \mathbf{x}_2) = \sum_{i=1}^d |x_{1i} - x_{2i}|$



- **standardization** is necessary if the scales of the attributes vary considerably

- measures how **alike** two objects are
  - higher when objects are more alike
  - often falls in the range  $[0, 1]$

## Example

cosine similarity:  $\cos(\mathbf{x}_1, \mathbf{x}_2) = \frac{\mathbf{x}_1 \cdot \mathbf{x}_2}{\|\mathbf{x}_1\| \cdot \|\mathbf{x}_2\|}$

- $\mathbf{x}_1 = (3, 2, 0, 5, 0, 0, 0, 2, 0, 0)$   
 $\mathbf{x}_2 = (1, 0, 0, 0, 0, 0, 0, 1, 0, 2)$
- cosine similarity is: 0.31

# Similarity Measures for Binary Attributes

- $\mathbf{x}_1$  and  $\mathbf{x}_2$ : two objects containing  $d$  **binary** attributes
  - $f_{00}$ : number of attributes with  $\mathbf{x}_1$  is 0 and  $\mathbf{x}_2$  is 0
  - $f_{01}$ : number of attributes with  $\mathbf{x}_1$  is 0 and  $\mathbf{x}_2$  is 1
  - $f_{10}$ : number of attributes with  $\mathbf{x}_1$  is 1 and  $\mathbf{x}_2$  is 0
  - $f_{11}$ : number of attributes with  $\mathbf{x}_1$  is 1 and  $\mathbf{x}_2$  is 1
- simple matching coefficient:  $SMC = \frac{f_{11} + f_{00}}{f_{01} + f_{10} + f_{00} + f_{11}} = \frac{f_{11} + f_{00}}{d}$


## Example

- $\mathbf{x}_1 = (1, 0, 0, 0, 0, 0, 0, 0, 0, 0)$ ,  $\mathbf{x}_2 = (0, 0, 0, 0, 0, 0, 1, 0, 0, 1)$
- $f_{00} = 7, f_{01} = 2, f_{10} = 1, f_{11} = 0$ ;  $SMC = \frac{0+7}{2+1+0+7} = 0.7$

# Partitioning Methods

## Basic idea

- organize the  $N$  objects into  $k$  partitions ( $k < N$ ), where each partition represents a **cluster**
  - objects within a cluster are similar, whereas objects of different clusters are dissimilar



- the clusters are formed to minimize an **objective criterion**, such as a **distance function**

Will focus on the following partitioning method

- **$k$ -means**

# $k$ -means

- $C_i$ : a cluster;  $N_i$ : number of points in  $C_i$
- **mean** of this cluster:  $\mathbf{c}_i = \frac{1}{N_i} \sum_{\mathbf{x} \in C_i} \mathbf{x}$ 
  - can be regarded as the **centroid** or **center of gravity** of the cluster

## $k$ -means algorithm







**input:** dataset  $D$ , number of clusters  $k$

**output:**  $k$  clusters

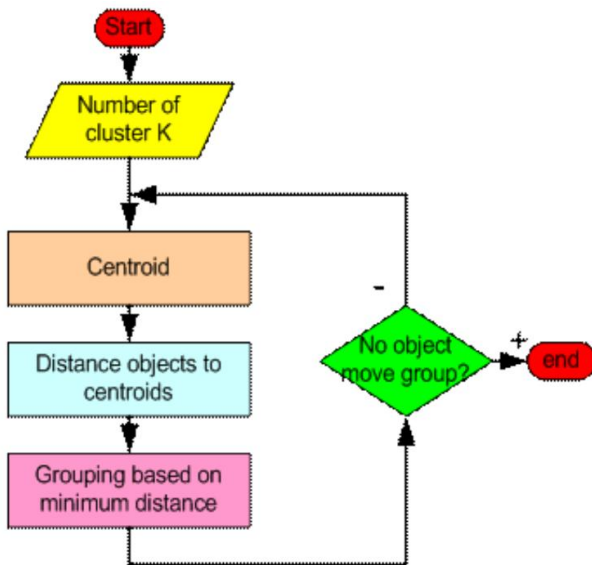
1. randomly select  $k$  points from  $D$  as the **initial** means
2. **repeat**
3.     form  $k$  clusters by assigning each point to its **closest** mean
4.     **recompute** the mean for every cluster
5. **until** no changes in the mean
6. **return** the  $k$  clusters



# Example



demo



# Issues: Local Minimum

- **objective criterion** that  $k$ -means attempts to minimize

(sum squared error) 
$$SSE = \sum_{i=1}^k \sum_{\mathbf{x} \in C_i} \|\mathbf{x} - \mathbf{c}_i\|^2$$

- converges to a **local minimum** → suboptimal clustering




- different initial starts of the means gives different final answers  
→ **initial selection** is very important

how to alleviate this problem?

- 1 multiple runs
- 2 pick the solution with minimum SSE

# Issues: Implicit Assumption on Cluster Shape

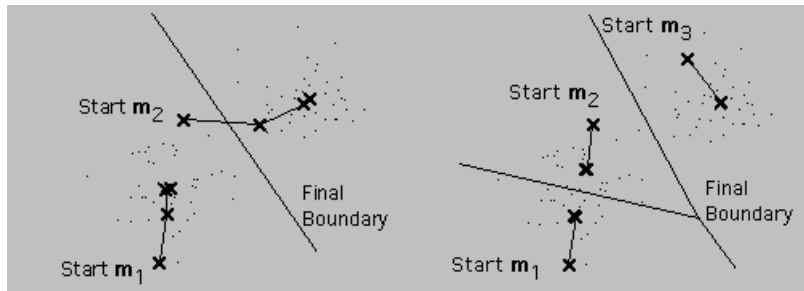


- can get wrong results when clusters have other shapes

# Issues: Number of Clusters

You have to pick the number of clusters


- in general, clustering result depends on  $k$



# Issues: Outliers

- the means may be **fictitious** (i.e., non-existent in the dataset)
- $k$ -means is sensitive to **outliers**

## Example



- by removing points 1 and 2, we obtain much tighter clusters