

Heterogeneous Parallel Programming

COMP4901D


CUDA Example: Convolution

Overview

- Convolution
 - Math operation
 - Baseline CUDA implementation
 - Tiling and reuse in convolution

Convolution: an Math Operation

- Given two functions, produce a third function
 - The output function is a modified version of one of the two input functions
 - The output function is on the area overlap of the two functions
 - The output function is given on the amount one of the input functions translates



<http://en.wikipedia.org/wiki/Convolution>

Applications of Convolution


- Signal, image, and video processing
 - Filtering out noises to show the big picture
 - Sharpening boundaries/edges
- Science and engineering computation
 - Fluid dynamics simulation
 - partial differential equations

Computation of Convolution

- An array operation where each output data element is a weighted sum of a collection of neighboring input elements
- The weights used in the weighted sum calculation are defined by an input mask array, commonly referred to as the **convolution kernel**
 - We will refer to these mask arrays as **convolution masks** to avoid confusion with CUDA kernel programs.
 - The same convolution mask is typically used for all elements of the array.

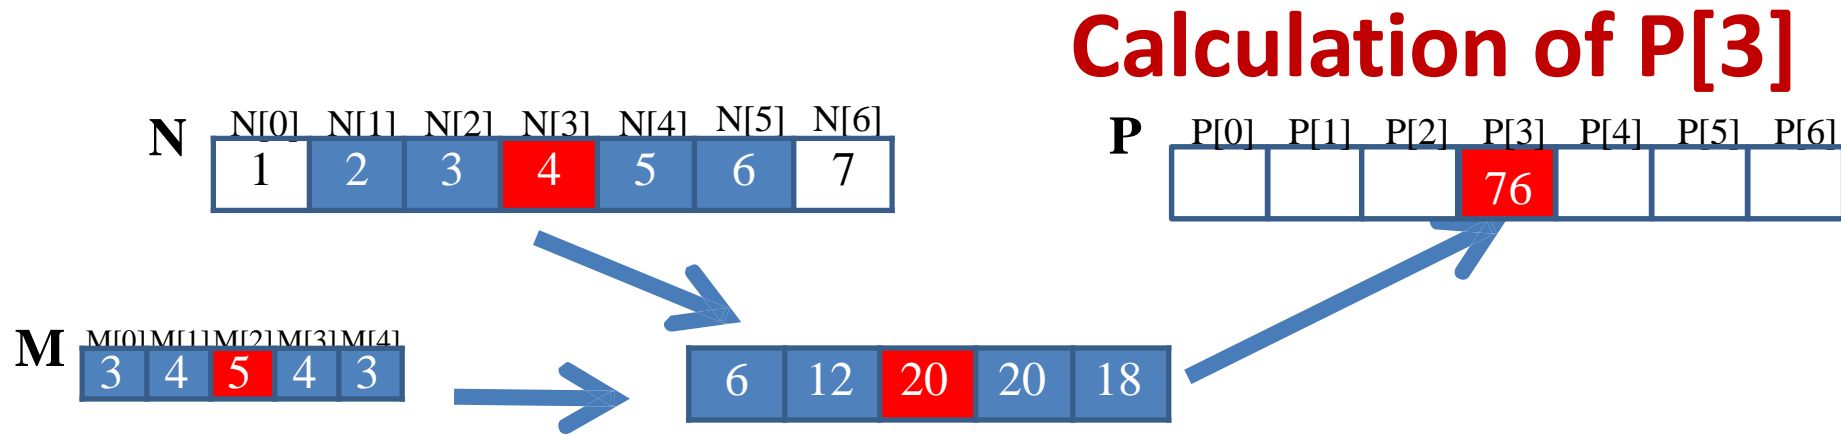
1D Convolution Example

Calculation of P[2]




- Commonly used for audio processing
 - Mask size is usually an odd number of elements for symmetry (5 in this example)

1D Convolution Example (Cont.)



1D Convolution Example (Cont.)

1D Convolution Boundary Condition




- Calculation of output elements near the boundaries (beginning and end) of the input array need to deal with “ghost” elements.
- Different policies (0, replicates of boundary values, etc.) may be used to fill in the boundaries.


1D Convolution Example in CUDA

```
__global__ void convolution_1D_basic_kernel
(float *N, float *M, float *P, int Mask_Width, int Width)
{
    int i = blockIdx.x*blockDim.x + threadIdx.x;
    float Pvalue = 0;
    int N_start_point = i - (Mask_Width/2);
    for (int j = 0; j < Mask_Width; j++) {
        if (N_start_point + j >= 0 && N_start_point + j < Width) {
            Pvalue += N[N_start_point + j] * M[j];
        }
    }
    P[i] = Pvalue;
}
```


2D Convolution Example



2D Convolution Example: Boundaries



Tiling Convolution: Output Elements



Each thread block calculates an output tile


Each output tile width is O_TILE_WIDTH

For each thread,

$\text{index_o} = \text{blockIdx.x} * \text{O_TILE_WIDTH} + \text{threadIdx.x}$

O_TILE_WIDTH is 4 in this example

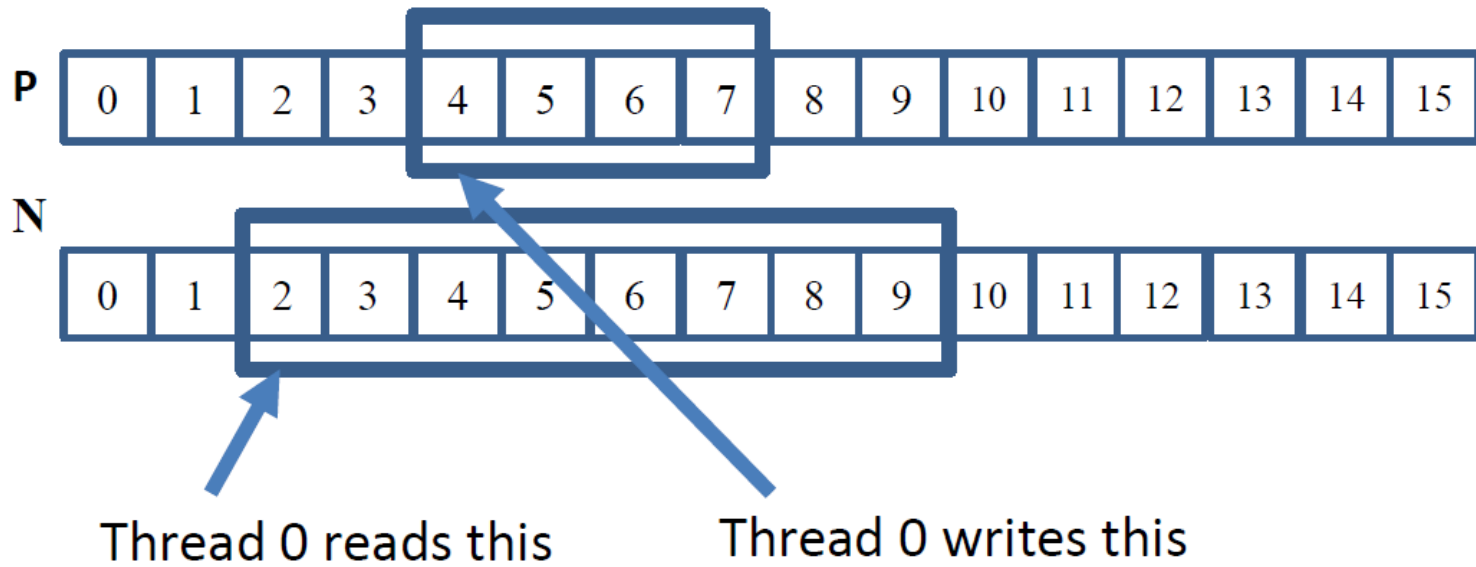
Tiling Convolution: Input Elements



Each input tile has all values needed to calculate the corresponding output tile.

Size each thread block to cover input tiles
blockDim.x is 8 in this example

Thread to Input Data Mapping



For each thread,
 $\text{index}_i = \text{index}_o - n$

where n is $\text{Mask_Width} / 2$. n is 2 in this example

All Threads Participate in Loading Input

```
float output = 0.0f;

if((index_i >= 0)  && (index_i < Width) ) {
    Ns[tx] = N[index_i];
}
else{
    Ns[tx] = 0.0f;
}
```

Only Some Threads Calculate Output


```
if (threadIdx.x < O_TILE_WIDTH) {  
    output = 0.0f;  
    for(j = 0; j < Mask_Width; j++) {  
        output += M[j] * Ns[j+threadIdx.x];  
    }  
    P[index_o] = output;  
}
```

Only Threads 0 through O_TILE_WIDTH-1 participate in calculation of output.

Setting Block Size

- `#define O_TILE_WIDTH 1020`
- `#define BLOCK_WIDTH (O_TILE_WIDTH + 4)`
- `dim3 dimBlock(BLOCK_WIDTH,1, 1);`
- `dim3 dimGrid((Width-1)/O_TILE_WIDTH+1, 1, 1)`
- The Mask_Width is 5 in this example
- In general, block width should be
output tile width + (mask width-1)

Shared Memory Data Reuse



Element 2 is used by thread 4 (1X)

Element 3 is used by threads 4, 5 (2X)

Element 4 is used by threads 4, 5, 6 (3X)

Element 5 is used by threads 4, 5, 6, 7 (4X)


Element 6 is used by threads 4, 5, 6, 7 (4X)

Element 7 is used by threads 5, 6, 7 (3X)


Element 8 is used by threads 6, 7 (2X)

Element 9 is used by thread 7 (1X)

Ghost Cells



Example 2D Image Matrix with Padding



- Example: a 3X3 matrix padded into a 3X4 matrix

Height is 3

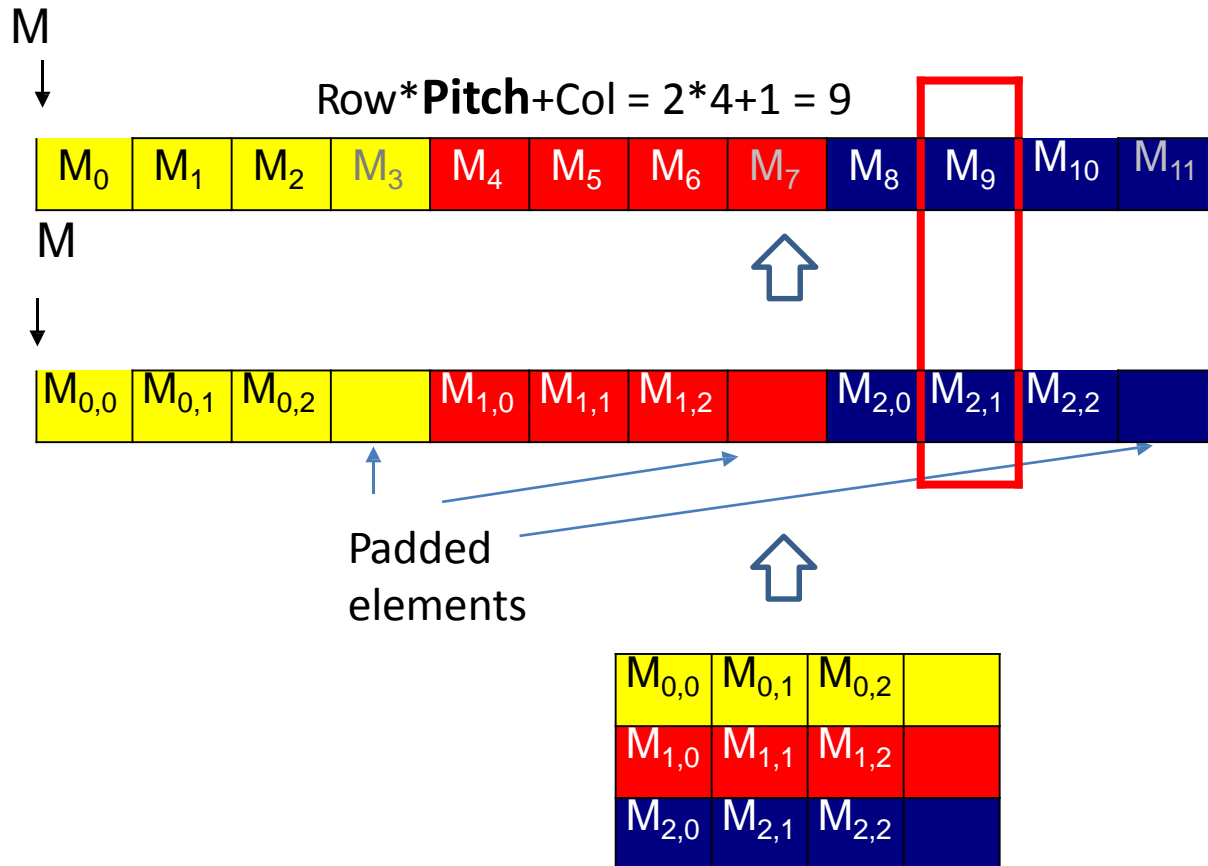
Width is 3

Channels is 1

Pitch is 4

- It is sometimes desirable to pad each row of a 2D matrix to multiples of DRAM bursts
 - So each row starts at the DRAM burst boundary
 - Effectively adding columns
 - This is usually done automatically by matrix allocation function
 - Pitch can be different for different hardware

Row-Major Layout with Pitch




Using Constant Memory for Mask

- Mask is used by all threads but not modified in the convolution kernel
 - All threads in a warp access the same locations at each point in time
- CUDA devices provide constant memory whose contents are aggressively cached
 - Cached values are broadcast to all threads in a warp
 - Effectively magnifies memory bandwidth without consuming shared memory
- Use of **const __restrict__** qualifiers for the mask parameter informs the compiler that it is eligible for constant caching, for example:

```
__global__ void convolution_2D_kernel(float *P, float *N,  
height, width, channels,  
const float * __restrict__ M) {
```

Shifting from Output Coordinates to Input Coordinate

```
int tx = threadIdx.x;  
int ty = threadIdx.y;  
int row_o = blockIdx.y * O_TILE_WIDTH + ty;  
int col_o = blockIdx.x * O_TILE_WIDTH + tx;  
int row_i = row_o - 2;  
int col_i = col_o - 2;
```



Taking Care of Boundaries

```
if((row_i >=0) && (row_i < height) &&
    (col_i >= 0) && (col_i < width) ) {
    Ns[ty][tx] = data[row_i*width + col_i];
} else{
    Ns[ty][tx] = 0.0f;
}
```



Use of width here is OK
since pitch is set to width
in this example.

Threads within the Output Tile Calculate

```
float output = 0.0f;
if(ty < O_TILE_WIDTH && tx < O_TILE_WIDTH){
    for(i = 0; i < MASK_WIDTH; i++) {
        for(j = 0; j < MASK_WIDTH; j++) {
            output += M[i][j] * Ns[i+ty][j+tx];
        }
    }
}
```

Threads within the Matrix Boundary

Output

```
if(row_o < height && col_o < width)  
    data[row_o*width + col_o] = output;
```

An Eight-Element Convolution Tile

N_ds



P

Mask_Width



For Mask_Width = 5, we load $8+5-1 = 12$ elements (12 memory l

Each Output Element Uses 5 Input Elements

N_ds



Mask_Width is 5

P



P[8] uses N[6], N[7], N[8], N[9], N[10]

P[9] uses N[7], N[8], N[9], N[10], N[11]

P[10] uses N[8], N[9], N[10], N[11], N[12]

...

P[14] uses N[12], N[13], N[14], N[15], N[16]

P[15] uses N[13], N[14], N[15], N[16], N[17]

A Simple Way of Calculating Tiling Benefit

- $(8+5-1)=12$ elements loaded
- $8*5$ global memory accesses replaced by shared memory accesses
- This gives a bandwidth reduction of $40/12=3.3$


Memory Access Reduction of 1D Tiled Convolution

- $O_TILE_WIDTH + MASK_WIDTH - 1$ elements loaded for each input tile
- $O_TILE_WIDTH * MASK_WIDTH$ global memory accesses replaced by shared memory accesses
- This gives a reduction factor of

$$(O_TILE_WIDTH * MASK_WIDTH) / (O_TILE_WIDTH + MASK_WIDTH - 1)$$

This ignores ghost elements in edge tiles.

Another Way of Looking at Reuse



N[6] is used by P[8] (1X)

N[7] is used by P[8], P[9] (2X)

N[8] is used by P[8], P[9], P[10] (3X)

N[9] is used by P[8], P[9], P[10], P[11] (4X)

N[10] is used by P[8], P[9], P[10], P[11], P[12] (5X)

... (5X)

N[14] is used by P[12], P[13], P[14], P[15] (4X)

N[15] is used by P[13], P[14], P[15] (3X)

Another Way of Calculating Memory Access Reduction

The total number of global memory accesses (to the $(8+5-1)=12$ N elements) replaced by shared memory accesses is:

$$\begin{aligned} &1 + 2 + 3 + 4 + 5 * (8 - 5 + 1) + 4 + 3 + 2 + 1 \\ &= 10 + 20 + 10 \\ &= 40 \end{aligned}$$

So the reduction is:

$$40/12 = 3.3$$

Another Way of Getting the Same Formula

The total number of global memory accesses to the input tile can be calculated as

$$\begin{aligned} & 1 + 2 + \dots + \text{MASK_WIDTH} - 1 + \text{MASK_WIDTH} * (\text{O_TILE_WIDTH} - \text{MASK_WIDTH} + 1) + \text{MASK_WIDTH} - 1 + \dots + 2 + 1 \\ &= \text{MASK_WIDTH} * (\text{MASK_WIDTH} - 1) + \text{MASK_WIDTH} * (\text{O_TILE_WIDTH} - \text{MASK_WIDTH} + 1) \\ &= \text{MASK_WIDTH} * \text{O_TILE_WIDTH} \end{aligned}$$

For a total of $\text{O_TILE_WIDTH} + \text{MASK_WIDTH} - 1$ input tile elements

Memory Access Reduction: 2D Tiled Convolution

- $(O_TILE_WIDTH + MASK_WIDTH - 1)^2$ input elements need to be loaded into shared memory
- The calculation of each output element needs to access $MASK_WIDTH^2$ input elements
- $O_TILE_WIDTH^2 * MASK_WIDTH^2$ global memory accesses are converted into shared memory accesses
- The reduction ratio is

$$O_TILE_WIDTH^2 * MASK_WIDTH^2 / (O_TILE_WIDTH + MASK_WIDTH - 1)^2$$

Summary

- Convolution is widely used in computation-intensive applications.
- Parallel convolution can be done in CUDA.
- Tiling with shared memory improves memory performance of parallel convolution.