

Data parallelism and graphics

Interactivity

- Interactive rendering applications
 - Video game, architectural modeling, 3D visualization (as opposed to offline rendering, e.g., movies)
- Requirement: Must be interactive
- How fast is enough?
- Monitors refresh at 60 Hertz (60 Hz)
 - Renders 60 images in each second
- Display 6 frames per second (6 FPS)
 - Each image will be repeated 10 times!
 - Testing shows this is the bare minimum


<http://www.realtimerendering.com/udacity/?load=demo/unit1-fps.js>

Real-time graphics

- Demands are often much higher
- 30-60 FPS are common thresholds
 - Tied to the 60Hz *refresh rate*
Each image being displayed once or twice
 - 60 FPS is upper limit
- If game takes 100ms to generate the frame
 - 10 FPS result
 - Display will still refresh at 60Hz but it will overwrite with same image

Screen resolution

- Pixel: A point in your screen
- Holds an RGB color value
- Higher screen resolution
leads to larger number of pixels
leads to better image
leads to lower higher processing cost
leads to lower FPS!



Pixel throughput

- A screen resolution of 1024 by 768 (modest!)
- A frame rate of 60 FPS
- How many pixels are we drawing per second?
 - A: 47,185,920
 - In each second the GPU has to calculate the color of a pixel over 47 million times!
 - A lot of computation
 - And that's a lower bound
 - Objects can overlap and be drawn on top of others

How can GPUs be so fast?!


A: Parallel processing

Data Parallelism for Graphics

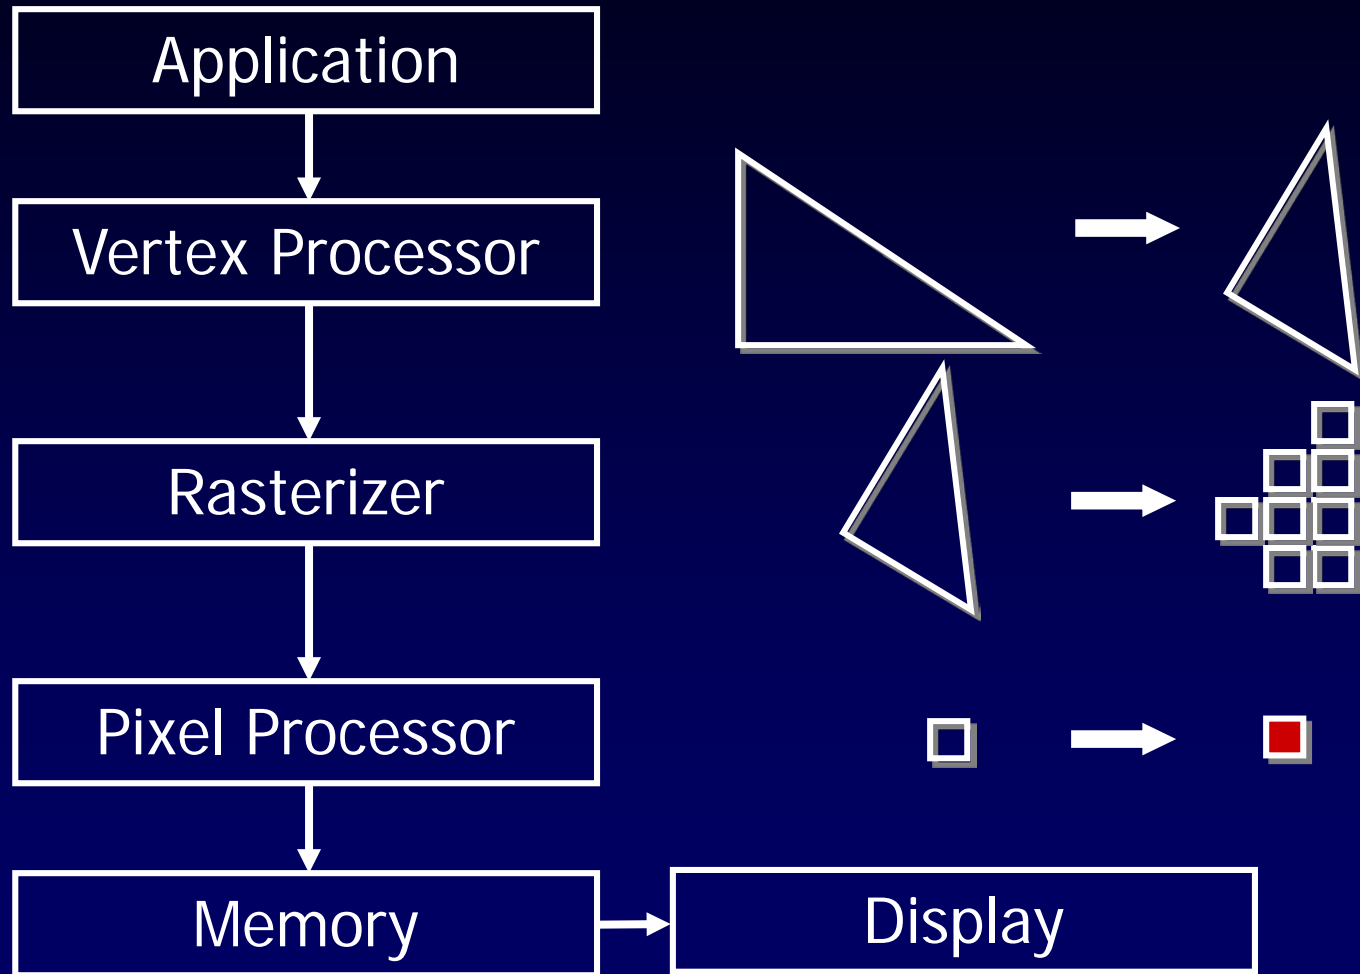
- GPUs are designed for graphics
 - Most popular consumer need for highly parallel computation
- GPUs process *independent* vertices & pixels
 - Temporary registers are zeroed
 - No shared or static data (not in graphics APIs)
 - No read-modify-write buffers
- Data-parallel processing
 - GPUs architecture is ALU-heavy
 - Multiple vertex & pixel pipelines

Today's GPUs

- Efficient architecture
- 1000+ parallel processors!
 - Working on vertices
 - Working on pixels
- SIMD processing
 - Must execute same code
 - Ok for graphics applications
- Improving at faster rate than CPUs



GPU processing at a glance





Vertex Processor

From object coordinates to screen coordinates


Geometry

- Composed of two parts
 - Vertex data in a vertex array/buffer
 - Primitives in an index/element array/buffer
- Vertices must include a position
 - Color and other optional parameters
- Primitives reference one or more vertices
 - Triangles, lines, points
- Geometry transformations act on vertices
 - And as a consequence change primitives too









Vertex transformation

Object or
Model
Coordinates

World
Coordinates

Eye or
Camera
Coordinates

Clip Space

Device
Coordinates

Window or
Screen
Coordinates






Modeling
transformation

Viewing
transformation


Projection
transformation

Perspective
division

Viewport
transformation



Rasterizer




From vertices and primitives to pixels

Pixel processor

Coloring the geometry

Use a popular illumination model

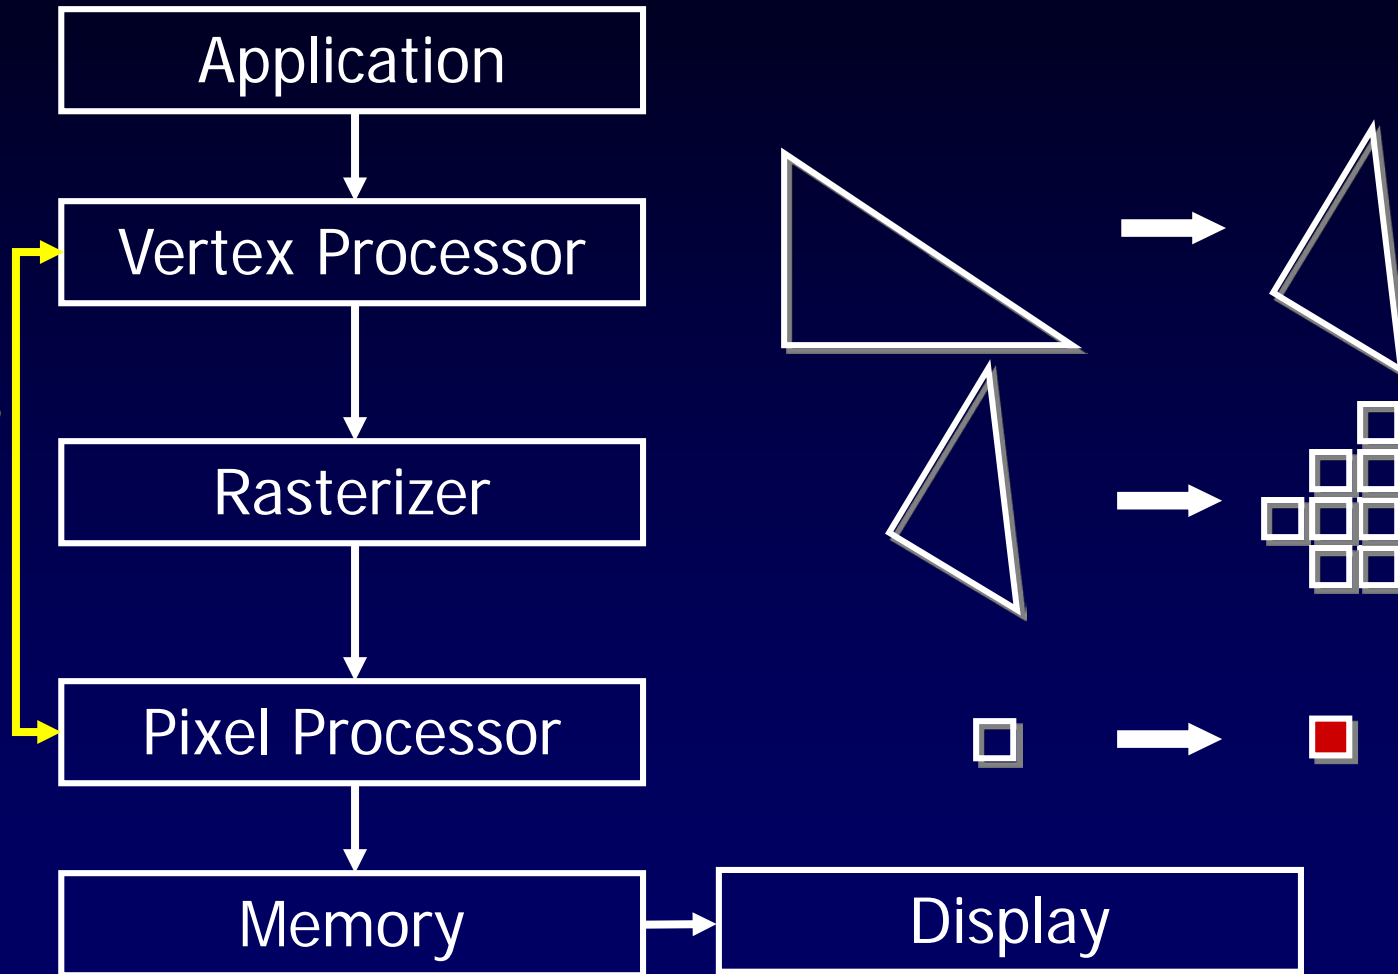


Pixel processor

Applying texture images



Programmability



Unified rendering pipeline

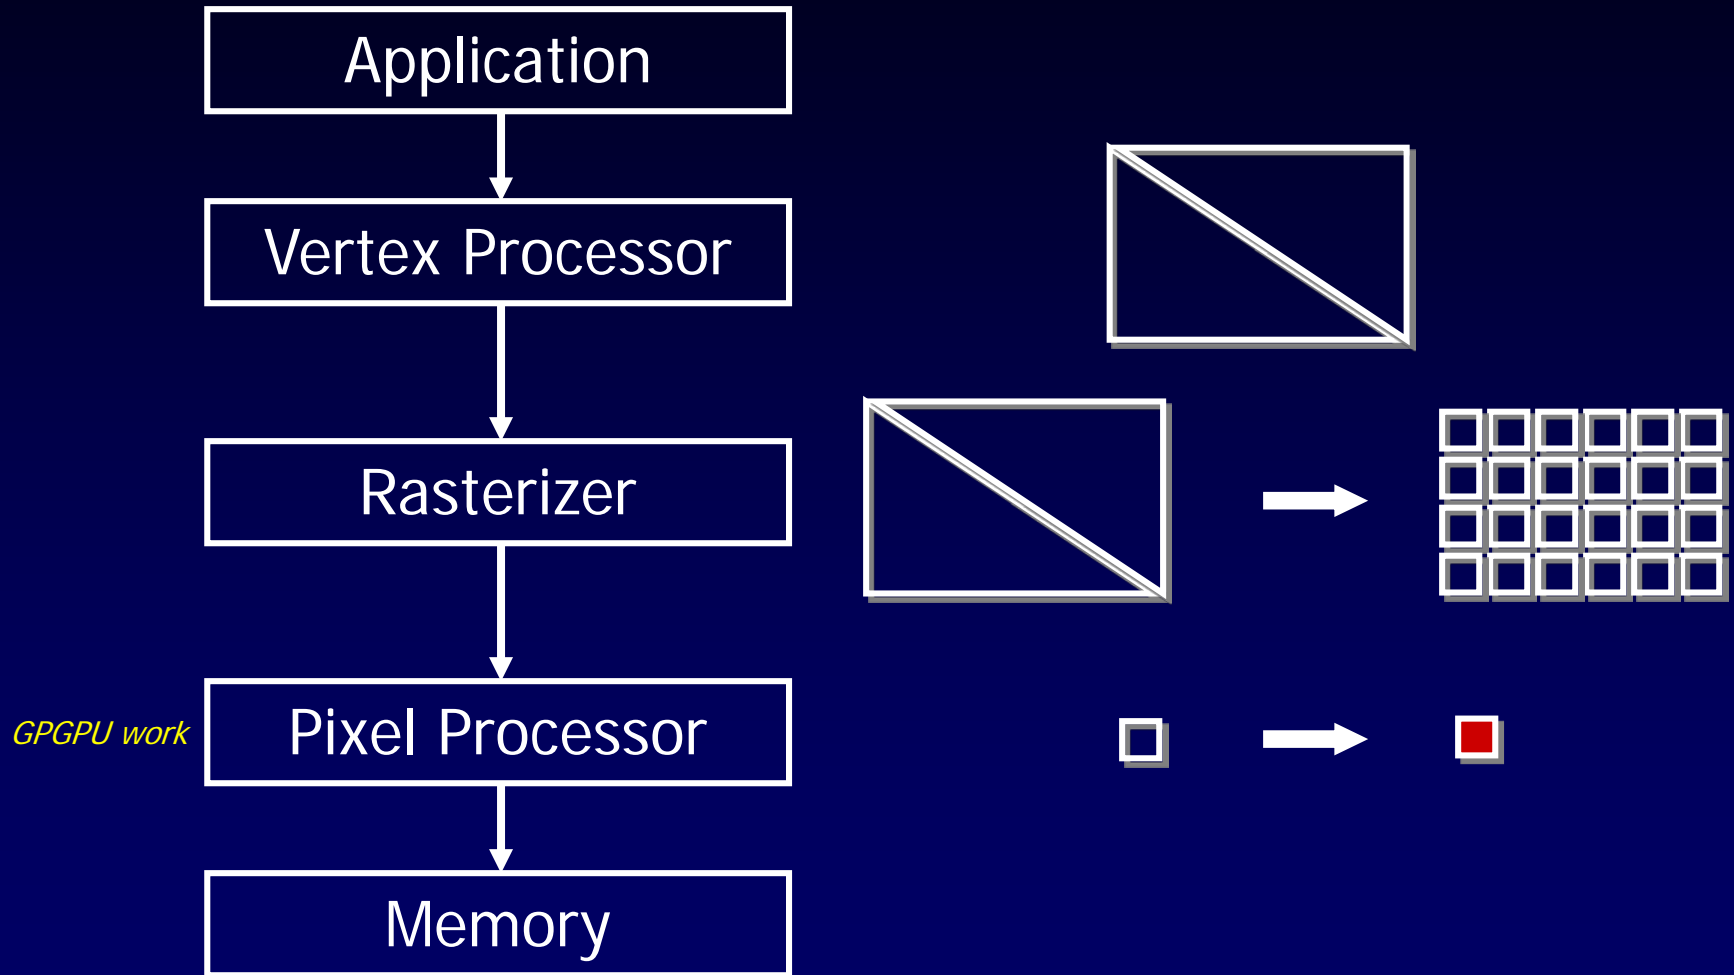
- No longer has dedicated vertex and pixel processors
- All programmable computation shares the same set of processors
- A complex scheduler determines their allocation when rendering a frame
- GPU “shaders” processes vertices and pixels

GPU Shading languages



- GLSL (OpenGL standard library)
- HLSL (Microsoft DirectX)
- CG (Nvidia proprietary)

- All C-like languages
- 2D-4D vector and matrix data types
 - Enough to represent geometry and color

GPGPU usage scenario




GPGPU Programming Model



Physical simulations on the GPU

Physics is a data parallel task



Mapping concepts to GPU


- Properties:
 - Algorithm must be data-parallel
 - Computations should ideally be local
- Now, how do we map it?

Data Streams & Kernels

- Streams
 - Collection of records requiring similar computation
 - Vertex positions, Voxels, etc.
 - Provide data parallelism
- Kernels
 - Functions applied to each element in stream
 - transforms, PDE, ...
 - Few dependencies between stream elements
 - Encourage high Arithmetic Intensity


Example: Fluid simulation

- Uses a simulation grid
- Common GPGPU rendering pipeline computation style
 - Textures represent computational grids = streams
- Many computations map to grids
 - Matrix algebra
 - Image & Volume processing
 - Global Illumination
 - Physically-based simulation
- Non-grid streams can be mapped to grids



Fluid simulation

- Navier stokes equations for fluid simulation



Algorithm

advect

accelerate

water/thermo

divergence

jacobi

jacobi

jacobi

jacobi

⋮

jacobi

u-grad(p)

Kernels

CPU

advect

GPU

```
for (int j = 1; j < height - 1; ++j)
{
    for (int i = 1; i < width - 1; ++i)
    {
        // get velocity at this cell
        Vec2f v = grid(x, y);

        // trace backwards along velocity field
        float x = (i - (v.x * timestep / dx));
        float y = (j - (v.y * timestep / dy));

        grid(x, y) = grid.bilerp(x, y);
    }
}
```

C++

```
void advect(float2 uv : WPOS,
            out float4 xNew : COLOR,

            uniform float dt, // timestep
            uniform float dx, // grid scale
            uniform samplerRECT u, // velocity
            uniform samplerRECT x) // state
{
    // trace backwards along velocity field
    float2 pos = uv - dt * f2texRECT(u, uv) / dx;

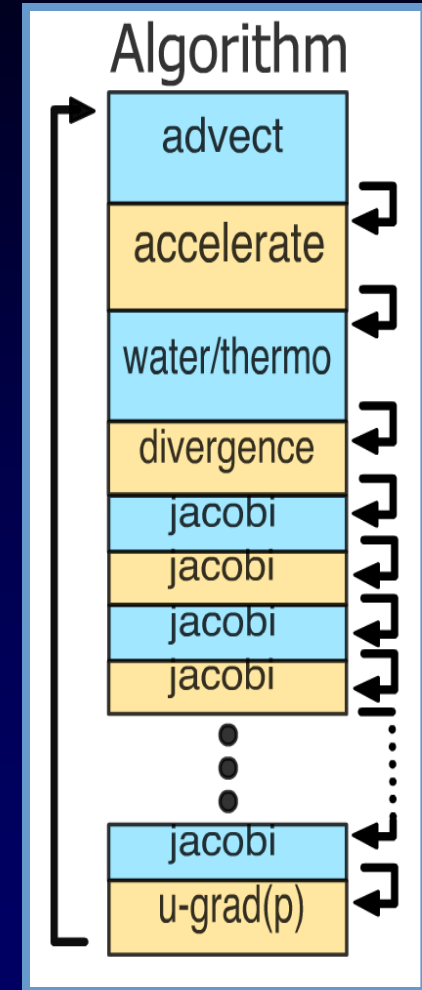
    xNew = f4texRECTbilerp(x, pos);
}
```

Cg

Kernel / loop body / algorithm step = Fragment Program

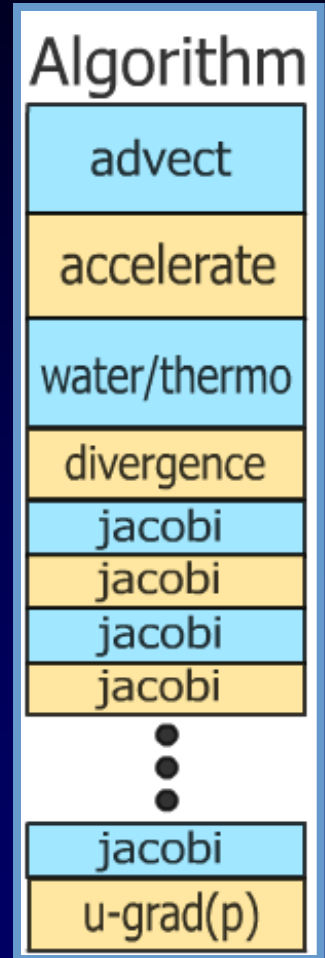
Feedback

- Each algorithm step depends on the results of previous steps
- Each time step depends on the results of the previous time step



GPU Simulation

- Algorithm steps are pixel shaders
 - Computational *kernels*
- Current state is stored in textures
- Feedback via render to texture



Invoking Computation


- Full screen quadrilateral invokes computation
- So, rasterization = kernel invocation

Examples:

<https://www.youtube.com/watch?v=7EOjAdmURY4>

<https://haxiomic.github.io/GPU-Fluid-Experiments/html5/?q=Medium>


Example: N-Body Simulation



- <https://www.youtube.com/watch?v=CPuVfiWLIHI>

GPU computation

- Brute force ☹
- $N = 8192$ bodies
- N^2 gravity computations
- 64M force comps. / frame
- ~25 flops per force




*Nyland, Harris, Prins,
GP² 2004 poster*

Computing Gravitational Forces

- Each body attracts all other bodies
 - N bodies, so N^2 forces
- Draw into an $N \times N$ buffer
 - Pixel (i,j) computes force between i and j
 - Very simple fragment program
 - More than 8192 bodies makes it trickier
 - Limited by max buffer size...

Computing Gravitational Forces



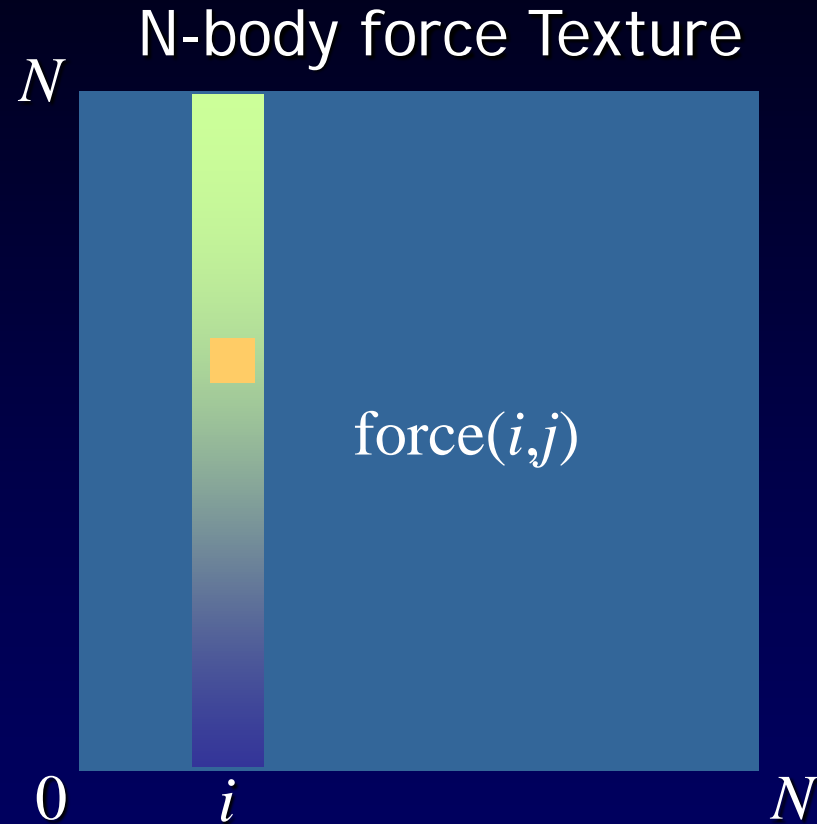
Force is proportional to the inverse square
of the distance between bodies

Computing Gravitational Forces

```
float4 force(float2 ij      : WPOS,  
             uniform sampler2D pos) : COLOR0  
{  
    // Pos texture is 2D, not 1D, so we need to  
    // convert body index into 2D coords for pos tex  
    float4 iCoords = getBodyCoords(ij);  
    float4 iPosMass = texture2D(pos, iCoords.xy);  
    float4 jPosMass = texture2D(pos, iCoords.zw);  
    float3 dir = iPos.xyz - jPos.xyz;  
    float r2 = dot(dir, dir);  
    dir = normalize(dir);  
    return dir * g * iPosMass.w * jPosMass.w / r2;  
}
```

Computing Total Force

- Have: array of (i,j) forces
- Need: total force on each particle i
 - Sum of each column of the force array
- Can do all N columns in parallel



Parallel Reduction

Parallel Reductions


- 1D parallel reduction:
 - sum N columns or rows in parallel
 - add two $N/2$ together
 - repeated
 - Until we're left with a single row of

texels

$N \times (N/2)$

$N \times (N/4)$

$N \times 1$



Requires $\log_2 N$ steps

Update Positions and Velocities

- Now we have a 1-D array of total forces
 - One per body
- Update Velocity
 - $u(i, t+dt) = u(i, t) + F_{total}(i) * dt$
 - Simple pixel shader reads previous velocity and force textures, creates new velocity texture
- Update Position
 - $x(i, t+dt) = x(i, t) + u(i, t) * dt$
 - Simple pixel shader reads previous position and velocity textures, creates new position texture

PhysX

- Nvidia middle-ware SDK
- Works on any CUDA-ready GPU
- Supports rigid body dynamics, soft body dynamics, vehicle dynamics, particles, volumetric fluid simulation, and cloth simulation
- <https://www.youtube.com/watch?v=6vipmar3wS4>
- <https://www.youtube.com/watch?v=pEX13W-luLA>

Optimize by force: Dual and Triple GPU Physics

Second (or third) GPU can be used for graphics or physics simulation, or share the primary card

