

Heterogeneous Parallel Programming



COMP4901D

Relational Query Processing on the GPU

Overview

- Relational query processing
 - Relational operators: select, project, join, aggr.
 - Access methods: table scan, B+-tree, hashing
 - Physical query operators
- Using the GPU for query co-processing
 - What can be done on the GPU and what not
 - On the GPU, how operators are implemented
 - Is co-processing worthwhile? If so, how?

Database Management Systems



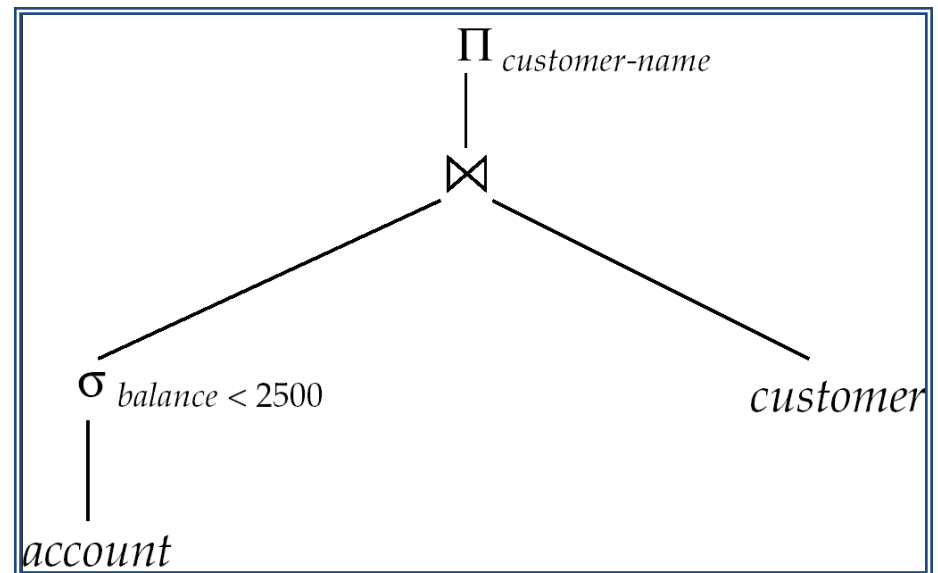
DBMSs have been a 40-year success in various applications.

SQL and Query Operators

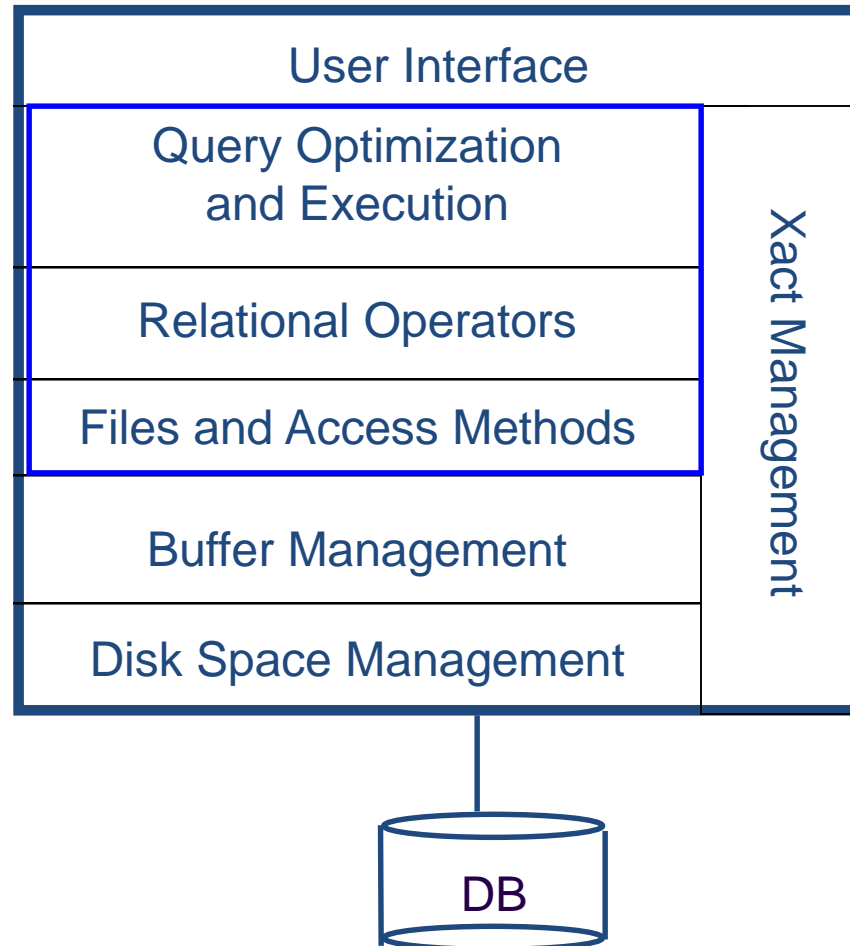
SELECT select-clause
FROM from-clause
[WHERE where-clause]
[ORDER BY order-by-expression]
[GROUP BY group-by-attributes
[HAVING condition-for-each-group]]

- Projection, aggregation
- Join
- Selection and/or Join condition
- Sort
- Partitioning
- Selection on partitions

SELECT customer-name
FROM account, customer
WHERE account.balance < 2500 AND
account.customer-ID =
customer.customer-ID



DBMS Architecture



Relational Query Processing

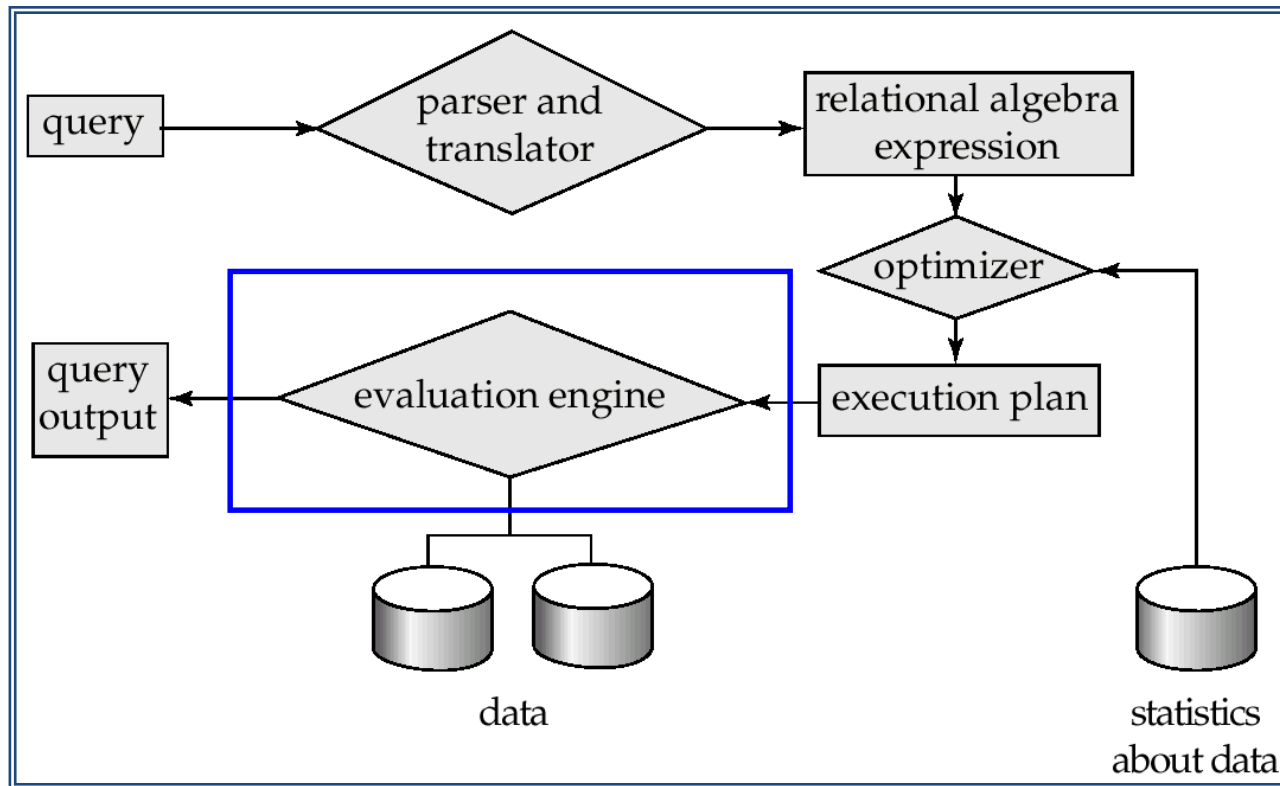



Figure source: Silberschatz et al.

Query Evaluation on the GPU

- Bring data into the GPU memory
- Construct GPU-suitable index structures
- Implement operators using primitives
- Determine result size lock-free
- Handle data skew for better parallelism


GPUQP Architecture



GPU-Based Join Algorithms

- Non-indexed nested-loop join (NINLJ)
 - Use the map primitive on both tables
- Indexed nested-loop join (INLJ)
 - Use the map primitive on the outer table
 - Adopt CSS-Tree [Rao99] to index the inner table
- Sort-merge join (SMJ)
 - Use sort on both tables and map for merging
- Hash join (HJ)
 - Adopt radix join [Boncz99]
 - Use split on both tables for partitioning

Nested-Loops Join on the GPU



Result Output Lock-free

- Problem: Join result size unknown
- Solution: Count result size before output
 - Each thread **counts** the number of join results for the partitioned join.
 - **Prefix sum** for write locations for each thread and the total number of join results.
 - Each thread **outputs** the join results in parallel.

Skew Handling in SMJ & HJ

- Identify the partitions that do not fit into the local memory.
 - Given an array storing partition sizes, we **split** it into two groups.
 - Partitions larger than the local memory
 - Partitions not larger than the local memory
- Decompose each of the large partitions into multiple small chunks.

Experimental Setup

Implementations

CPU: OpenMP

GPU: **CUDA** and DirectX


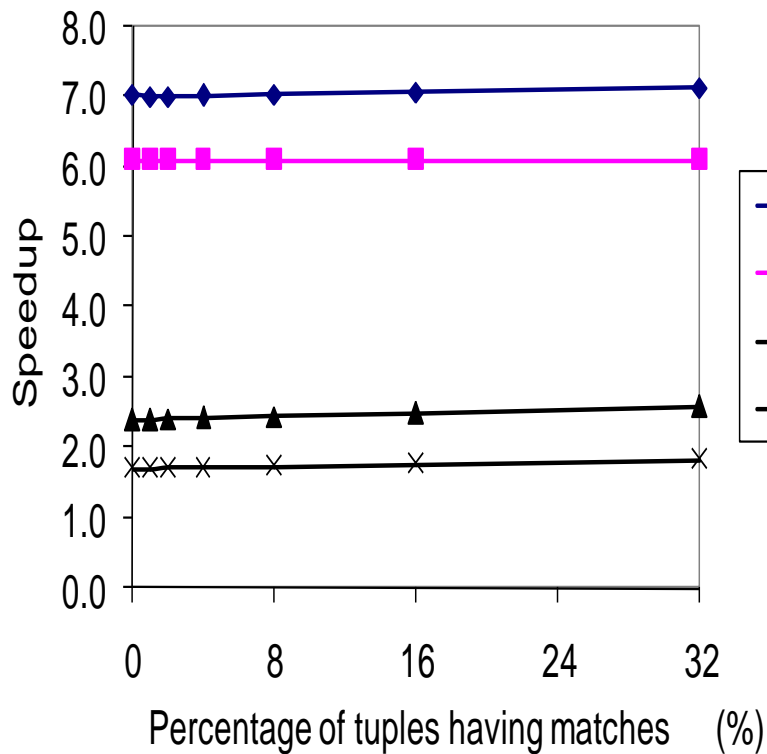
| | CMP (P4 Quad) | GPU (NV G80) |
|-----------------------|---------------|------------------|
| Processors (HZ) | 2.66G*4 | 1.35G*128 |
| Cache size | 8MB | 256KB |
| Bandwidth (GB/sec) | 10.4 | 86.4 |

Performance on Uniform Data

| Joins | CPU (sec) | GPU (sec) | Speedup |
|--------------|------------------|------------------|----------------|
| NINLJ | 528.0 | 75.0 | 7.0 |
| INLJ | 4.2 | 0.7 | 6.1 |
| SMJ | 5.0 | 2.0 | 2.4 |
| HJ | 2.5 | 1.3 | 1.9 |

The GPU measurements **include** the time for data transfer between the GPU memory and the main memory.

Performance on Skewed Data



GDB: Beyond GPUQP

- Co-processing between the CPU and the GPU
 - The CPU handles disk IO, cost estimation, workload partitioning, and runs as a worker for a query when selected.
 - The GPU runs as a worker for a query when selected.
 - The cost model estimates the execution time of a query including memory stalls and computation.
 - Each operator can be (1) on the CPU only, (2) on the GPU only, and (3) on both processors.

Results from GDB

- The query cost model for the GPU was accurate.
- For TPC-H queries on disk-resident data, GDB's performance was similar to a commercial DBMS.
- For queries on in-memory data,
 - With data transfer time excluded, the GPU was 2-27 times faster than the CPU worker.
 - With data transfer time included, the GPU was 2-7 times faster on complex queries but 2-4 times slower on simple selections.
- GDB's co-processing achieved the best performance by making the right decision on where to execute a query.

Summary

- Relational query processing is a data-parallel task, suitable for GPU processing.
- Using data-parallel primitives as building blocks in GPU-based query processing simplifies programming and improves performance.
- Utilizing both the CPU and the GPU for query co-processing gets the best of both worlds.

<http://www.cse.ust.hk/gpuqp>