

# COMP 4901Q: High Performance Computing (HPC)

## Lecture 1: Introduction

Instructor: Shaohuai SHI ([shaohuais@cse.ust.hk](mailto:shaohuais@cse.ust.hk))

Teaching assistants: Mingkai TANG ([mtangag@connect.ust.hk](mailto:mtangag@connect.ust.hk))

Yazhou XING ([yxingag@connect.ust.hk](mailto:yxingag@connect.ust.hk))

# Intended Learning Outcomes

- ▶ Be able to describe the **elements** and **principles** of HPC.
- ▶ Be able to master the **system architecture** of high performance computers.
- ▶ Be able to learn **parallel programming** languages and frameworks.
- ▶ Be able to use one of programming languages or frameworks to **design and develop HPC systems** for real-world science or engineering problems.

# Course Schedule

1. Introduction to high performance computing. (1 lecture) (4 Feb, Fri)
  - ▶ Course structure, basic concepts of HPC, examples that are using HPC.
2. Introduction to parallel computer architecture. (1 lecture) (9 Feb, Wed)
  - ▶ Parallel computer organization, processor (e.g., CPU and GPU), memory, storage, interconnect, etc.
3. Revisit C/C++ programming and common Linux commands. (1 lecture) (11 Feb, Fri)
  - ▶ For the preparation of the following classes that will have extensive code examples.
4. Introduction to parallel programming. (1 lecture) (16 Feb, Wed)
  - ▶ Overview of parallel programming and related frameworks/libraries/languages.
5. Introduction to OpenMP. (2 lectures) (18 Feb, Fri and 23 Feb, Wed) **(Lab Tutorial 1, 21 Feb, Mon)**
  - ▶ Multi-threading programming for CPUs and the framework of OpenMP and its syntax.
6. OpenMP 3.0 and tasking. (2 lectures) (25 Feb, Fri and 2 Mar, Wed)
  - ▶ Introduce main features of OpenMP 3.0
7. Introduction to GPU computing. (2 lectures) (4 Mar, Fri and 9 Mar, Wed)
  - ▶ Basic concepts of GPU computing, GPU architecture (e.g., stream processors, memory, scheduler, etc.)

# Course Schedule

8. CUDA programming on the GPU. (4 lectures) (11, 16, 18, and 23 Mar) **(Lab Tutorial 2, 14 Mar, Mon)**
  - ▶ Basic concepts in CUDA programming, how to program CUDA code on the GPU, and some advanced topics in optimizing CUDA code.
9. Introduction to MPI. (2 lectures) (25 Mar, Fri and 30 Mar, Wed) **(Lab Tutorial 3, 28 Mar, Mon)**
  - ▶ Basic concepts of MPI and distributed memory programming.
10. Collective communication algorithms in MPI. (4 lectures) (1, 6, 8, and 20 Apr)
  - ▶ Some advanced topics in optimized algorithms for collective communication.
11. Combining shared-memory and distributed-memory computing with OpenMP, CUDA, and MPI. (2 lectures) (22 Apr, Fri and 27 Apr, Wed) **(Lab Tutorial 4, 25 Apr, Mon)**
  - ▶ Practices and examples to demonstrate how to combine these three technologies to achieve higher performance.
12. A case study: large-scale machine learning in HPC. (2 lectures) (29 Apr, Fri and 4 May, Fri)
  - ▶ A detailed example to show how to exploit HPC to support large-scale machine learning
13. Introduction to Hadoop and Spark. (4 May)
  - ▶ Introduce some basic concepts of Hadoop and Spark

# Lab Sessions

1. Lab Tutorial 1: Programming with OpenMP (**21 Feb, Mon**)
  - ▶ Programming environment setup
  - ▶ Examples of writing, debugging, compiling, and running code
2. Lab Tutorial 2: Programming with CUDA on GPUs (**14 Mar, Mon**)
  - ▶ Programming environment setup
  - ▶ Examples of writing, debugging, compiling, and running code
3. Lab Tutorial 3: Programming with MPI on clusters (**28 Mar, Mon**)
  - ▶ Programming environment setup
  - ▶ Examples of writing, debugging, compiling, and running code
4. Lab Tutorial 4: Examples of combining OpenMP, CUDA, and MPI (**25 Apr, Mon**)
  - ▶ Some examples to demonstrate how to combine different parallel programming techniques

# Grading (No Final Exam)

- ▶ Homework: 2 times, 10% each.
  - ▶ Non-programming homework
- ▶ Programming assignments: 4 times, 10% each.
  - ▶ OpenMP
  - ▶ CUDA
  - ▶ MPI
  - ▶ Combine two or three of the above
- ▶ One final project (documentation, code, and presentation): 40%.
  - ▶ Proposal
  - ▶ Design
  - ▶ Coding
  - ▶ Presentation

Discussion is encouraged, plagiarism is prohibited!

Academic integrity: <http://ugadmin.ust.hk/integrity/student-1.html>

# Final Project

- ▶ Find a challenging problem
  - ▶ should be computing-intensive/time-consuming to solve with a single processor
  - ▶ can be any science or engineering area you like
- ▶ Literature review
  - ▶ search related work to understand how they address the problem
  - ▶ what are the main issues of the existing methods
- ▶ Methodology
  - ▶ propose your own method to address the problem
  - ▶ compare your method with existing ones
    - ▶ theoretically analyze the improvement
    - ▶ empirically compare performance on specific hardware (e.g., multi-core CPU, GPU, or clusters)
- ▶ Code
  - ▶ implement your methods using existing frameworks to run on specific hardware
- ▶ Documentation
- ▶ Presentation

# What is HPC

- ▶ HPC most generally refers to the practice of aggregating **computing power** in a way that delivers **higher performance** (thus shorter time) to **solve large problems** in science and engineering
  - ▶ **Computing power**: many-core processors to supercomputers
  - ▶ **Large problems**: can be divided into small problems and solved by parallel computing
  - ▶ **Higher performance**: solve the problem as fast as possible
- ▶ It is typically talking about two main areas:
  - ▶ Hardware
    - ▶ From multi-core processors to supercomputers
    - ▶ The top500 list: <https://www.top500.org/>
  - ▶ Software stack
    - ▶ Frameworks to utilize the computing resources
    - ▶ Algorithms to reduce the time/memory complexity

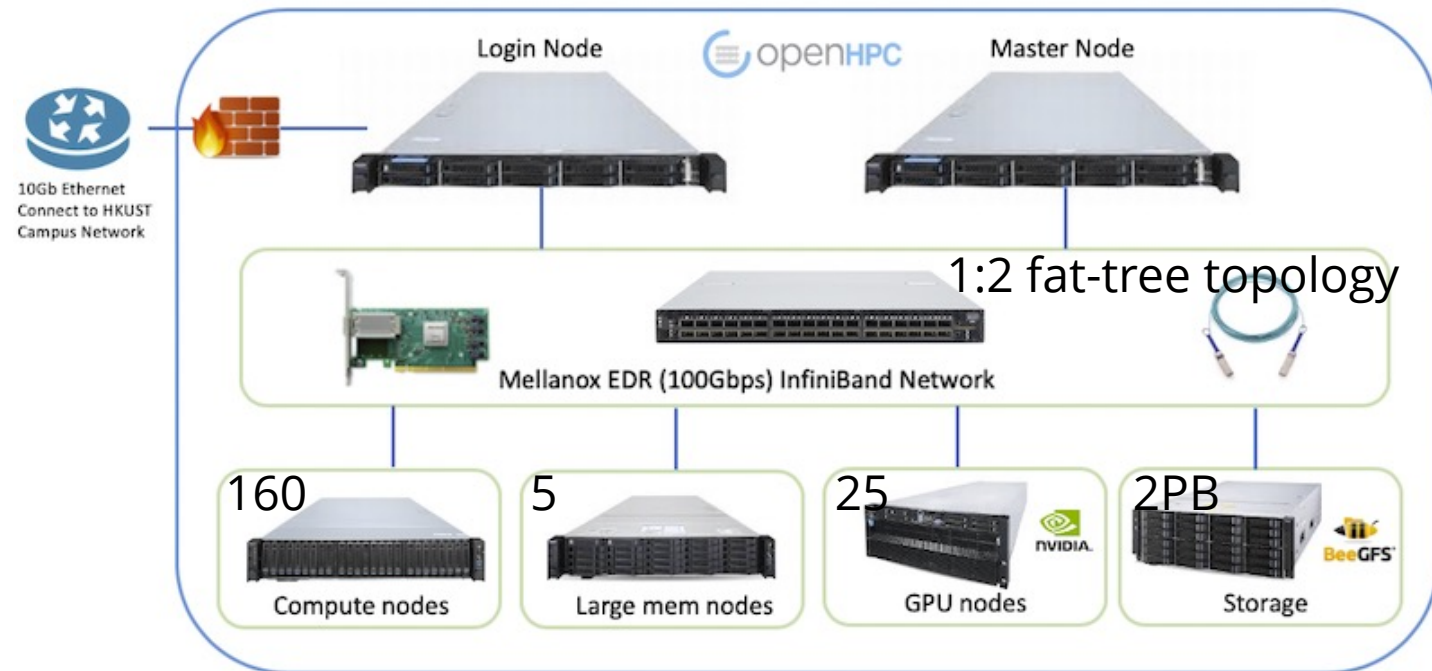


# Hardware of Supercomputers

- ▶ On a single node
  - ▶ Processors: e.g., CPU, GPU, etc.
  - ▶ Memory: e.g., DRAM, HBM
  - ▶ Storage: HDD, SSD, etc.
  - ▶ Interconnects: QPI, UPI, PCIe, NVLink, ...
- ▶ On a cluster
  - ▶ Interconnect between nodes:
    - ▶ Ethernet or InfiniBand
  - ▶ Network switch
    - ▶ Topology of interconnects
    - ▶ Fat-tree, BCube, etc.
  - ▶ Management nodes
    - ▶ Login, Monitor, ...



#1 Supercomputer Fugaku: 158,976 nodes, 7,630,848 Cores



HPC cluster at HKUST

Image credit:

[https://www.riken.jp/en/news\\_pubs/news/2020/20201117\\_2/index.html](https://www.riken.jp/en/news_pubs/news/2020/20201117_2/index.html)

<https://itsc.ust.hk/services/academic-teaching-support/high-performance-computing/hpc3-cluster/overview>

# Performance Measure

- ▶ FLOP: floating point operation
  - ▶ double (default, 8 bytes) or float precision (4 bytes)
  - ▶ measure the workload of the problem
- ▶ FLOP/s (or FLOPS): floating point operations per second
  - ▶ measure the performance of the system at addressing a target problem
- ▶ Commonly used units
  - ▶ Kilo:       KFLOP/s =  $10^3$  FLOPS
  - ▶ Mega:       MFLOP/s =  $10^6$  FLOPS
  - ▶ Giga:       GFLOP/s =  $10^9$  FLOPS
  - ▶ Tera:       TFLOP/s =  $10^{12}$  FLOPS
  - ▶ Peta:       PFLOP/s =  $10^{15}$  FLOPS
  - ▶ Exa:        EFLOP/s =  $10^{18}$  FLOPS
  - ▶ Zetta:       ZFLOP/s =  $10^{21}$  FLOPS

World-wide # 1 (Supercomputer Fugaku): ~537 PFLOPS

World-wide # 2 (Summit): ~200 PFLOPS

The top500 list: <https://www.top500.org/>

# Supercomputer Fugaku (#1)

- ▶ Calculate peak performance
  - ▶ # of cores: 48
  - ▶ SVE 512-bit × 2 vector extensions per core
  - ▶ Frequency per core: 2.2 GHz
  - ▶ Double precision: 64 bits
  - ▶ 512-bit addition or multiply (fma) per cycle
    - ▶  $512/64=8$  double-precision FLOP per cycle
  - ▶ Peak performance =  $2.2 \times 10^9$  (Frequency) ×  $8 \times 2 \times 2$  (FLOP per cycle) ×  $10^{12}$  (Tera) × 48 (Cores) = 3.3792 TFLOPS




(Source: [https://www.riken.jp/en/news\\_pubs/news/2020/20201117\\_2/index.html](https://www.riken.jp/en/news_pubs/news/2020/20201117_2/index.html))

- ▶ Supercomputer Fugaku (#1 supercomputer)
  - ▶ Peak performance: ~537 PFLOPS
  - ▶ 158,976 nodes

|              |                         |  |
|--------------|-------------------------|--|
| Architecture |                         | Armv8.2-A SVE 512bit<br>With the following Fujitsu's extensions: Hardware barrier, Sector cache, and Prefetch  |
| Core         |                         | 48 cores for compute and 2 or 4 cores for OS activities<br>4 CMGs (NUMA nodes)   |
| Performance  | Normal Mode:<br>2.0 GHz | DP: 3.072 TF, SP: 6.144 TF, HP: 12.288 TF  |
|              | Boost Mode:<br>2.2 GHz  | DP: 3.3792 TF, SP: 6.7584 TF, HP: 13.5168 TF   |
| Cache*1 *2   |                         | L1D/core: 64 KiB, 4way, 256 GB/s (load), 128 GB/s (store)<br>L2/CMG: 8 MiB, 16way<br>L2/node: 4 TB/s (load), 2 TB/s (store)<br>L2/core: 128 GB/s (load), 64 GB/s (store) |
| Memory       |                         | HBM2 32 GiB, 1024 GB/s   |
| Interconnect |                         | Tofu Interconnect D (28 Gbps x 2 lane x 10 port)   |
| I/O          |                         | PCIe Gen3 x16  |
| Technology   |                         | 7nm FinFET   |

(Source: <https://www.r-ccs.riken.jp/en/fugaku/project/outline>)

# Top500




- Graphic Processing Units (GPUs) processors play an important role in supercomputers

| Rank | System   | Cores      | Rmax (TFlop/s) | Rpeak (TFlop/s) | Power (kW) |
|------|--|------------|----------------|-----------------|------------|
| 1    | <b>Supercomputer Fugaku</b> - Supercomputer Fugaku, A64FX 48C 2.2GHz, Tofu interconnect D, Fujitsu<br>RIKEN Center for Computational Science<br>Japan                                  | 7,630,848  | 442,010.0      | 537,212.0       | 29,899     |
| 2    | <b>Summit</b> - IBM Power System AC922, IBM POWER9 22C 3.07GHz, NVIDIA Volta GV100, Dual-rail Mellanox EDR<br>Infiniband, IBM<br>DOE/SC/Oak Ridge National Laboratory<br>United States | 2,414,592  | 148,600.0      | 200,794.9       | 10,096     |
| 3    | <b>Sierra</b> - IBM Power System AC922, IBM POWER9 22C 3.1GHz, NVIDIA Volta GV100, Dual-rail Mellanox EDR<br>Infiniband, IBM / NVIDIA / Mellanox<br>DOE/NNSA/LLNL<br>United States     | 1,572,480  | 94,640.0       | 125,712.0       | 7,438      |
| 4    | <b>Sunway TaihuLight</b> - Sunway MPP, Sunway SW26010 260C 1.45GHz, Sunway, NRCPC<br>National Supercomputing Center in Wuxi<br>China   | 10,649,600 | 93,014.6       | 125,435.9       | 15,371     |
| 5    | <b>Perlmutter</b> - HPE Cray EX235n, AMD EPYC 7763 64C 2.45GHz, NVIDIA A100 SXM4 40 GB, Slingshot-10, HPE<br>DOE/SC/LBNL/NERSC<br>United States  | 761,856    | 70,870.0       | 93,750.0        | 2,589      |
| 6    | <b>Selene</b> - NVIDIA DGX A100, AMD EPYC 7742 64C 2.25GHz, NVIDIA A100, Mellanox HDR Infiniband, Nvidia<br>NVIDIA Corporation<br>United States  | 555,520    | 63,460.0       | 79,215.0        | 2,646      |

# Top500

- ▶ ~92% systems are with GPUs in the Top500 list (June 2020)




Source: <https://www.top500.org/statistics/overtime/>



# Development of Processors


- ▶ Moore's law, 1965 by Intel CEO and co-founder Gordon Moore
  - ▶ the number of transistors in a dense integrated circuit (IC) doubles about every two years.
- ▶ Huang's law, 2018 by Nvidia CEO Jensen Huang
  - ▶ the performance of GPUs will more than double every two years



Source: <https://arxiv.org/pdf/1911.11313.pdf>

Transistor count comparison

# Development of Processors



Performance comparison: CPU vs. GPU

# Why needs HPC

- ▶ A large problem: input->process->output
  - ▶ Size: too big or Time: too slow
- ▶ A large number of HPC applications from science and engineering
  - ▶ Engineering: e.g., Computational Fluid Dynamics (CFD) simulation
  - ▶ Geoscience
  - ▶ Molecular Dynamics
  - ▶ Physics
  - ▶ Quantum Mechanics
  - ▶ Data Analytics
  - ▶ ...









Image credits:

<https://www.npd-solutions.com/cfd.html>

<http://www.onegeology.org/>

<https://www.epcc.ed.ac.uk/blog/2014/10/06/improving-tinker>

<http://www.sci-news.com/physics/naked-singularities-saddle-shaped-universe-04886.html>

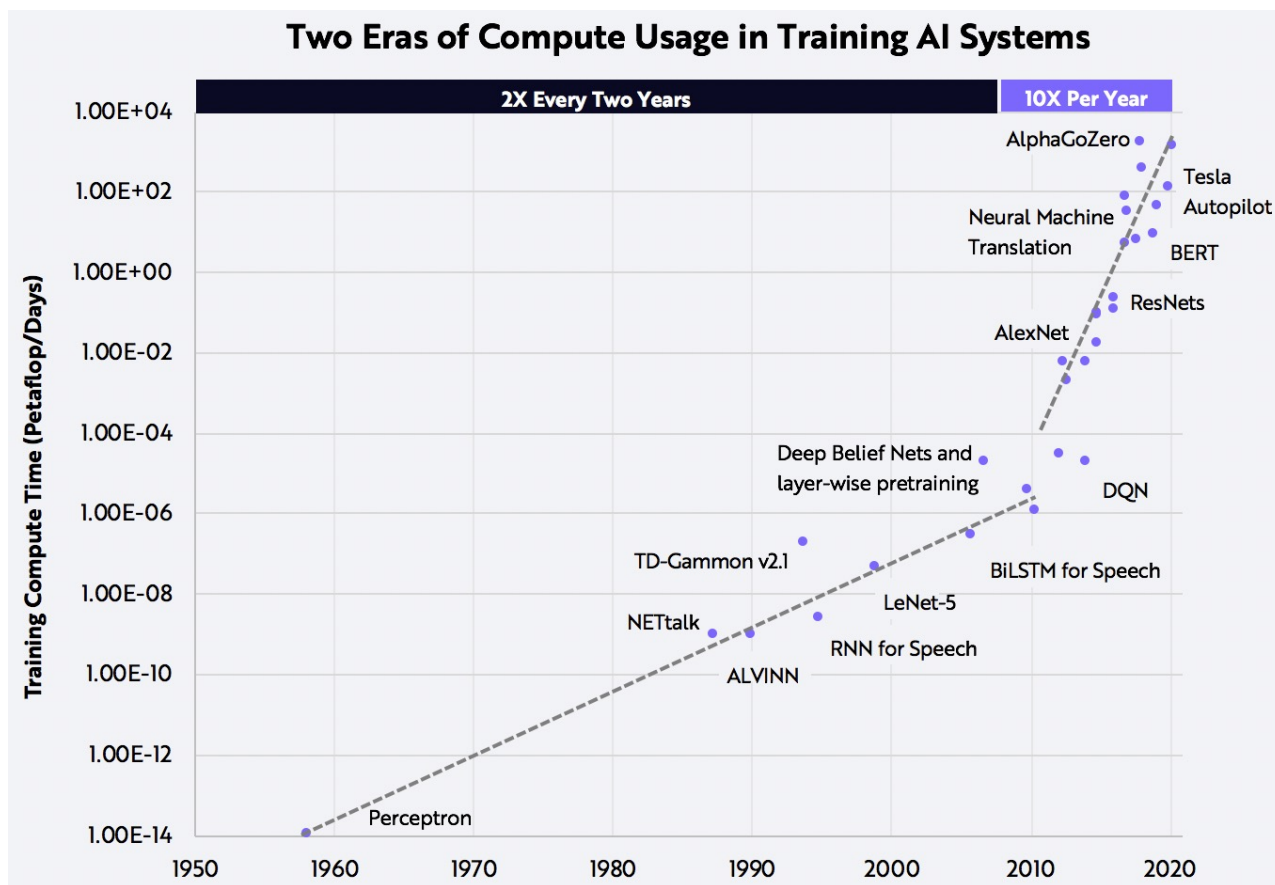
<https://online.stanford.edu/courses/ee222-applied-quantum-mechanics-i>

<https://rishi30-mehta.medium.com/object-detection-with-yolo-giving-eyes-to-ai-7a3076c6977e>




# HPC in AI

- AI models are computing hungry



Source: <https://ark-invest.com/articles/analyst-research/ai-training/>



- #1 supercomputer: 2.146176e18 FLOPS  
(Half precision for AI)!

# HPC in AI

- ▶ AI models are computing hungry
- ▶ Training large models is still very slow on single GPU/TPU
  - ▶ E.g., training a ResNet-50 model with a Tesla P100 GPU takes **11 days**;
  - ▶ BERT pre-training with Google TPUv3 takes more than **1.5 months**

Training ResNet-50 to 75+% accuracy (90 epochs)

1 hour

20 minutes

6.6 minutes

1.2 minutes

45.6 seconds

28.2 seconds

June 2017  
Facebook  
256 GPUs

September 2017  
UC Berkeley  
2048 KNLs

July 2018  
Tencent & HKBU  
2048 P40 GPUs

April 2019  
Fujitsu  
2048 V100 GPUs

July 2020  
Nvidia  
1840 A100 GPUs

July 2020  
Google  
4096 TPUv3

December 2021

ResNet-50  
4320 A100 GPUs, NVIDIA: ~20s

BERT  
4320 A100 GPUs, NVIDIA: ~14s

# How to use HPC: Software Stack

|  |  |                                      |                               |           |                |                  |                 |
|--|--|--------------------------------------|-------------------------------|-----------|----------------|------------------|-----------------|
| HPC<br>Programming<br>Tools                  | Performance Monitoring                 | HPCC                                 | IOR                           | PAPI/IPM  | NPB            | Netperf          |                 |
|  | Development Tools                      | Alliena DDT/TAU                      | Intel Cluster Studio/IBM XC   |           | PGI (PGI SDK)  | GNU Compiler     |                 |
|  | Application Libraries                  | Ferret/GRADS/PARA view/VISIT         | MVAPICH2/ OpenMPI             | ACML/ESSL | MPSS/CUDA      | BLAS, LAPACK     |                 |
| Middleware<br>Applications and<br>Management | Resource Management/<br>Job Scheduling | SLURM                                | Grid Engine                   | MOAB      | Altair PBS Pro | IBM Platform LSF | Torque/<br>Maui |
|  | File System                            | NFS                                  | Local FS<br>(ext3, ext4, XFS) |           | GPFS           | Lustre           |                 |
|  | Provisioning                           | XCAT / ROCKS / C-DAC Developed tools |                               |           |                |                  |                 |
|  | Cluster Monitoring                     | XCAT / ROCKS / C-DAC Developed tools |                               |           |                |                  |                 |
| Operating<br>Systems                         | Operating System                       | Linux (Red Hat, CentOS, SUSE)        |                               |           |                |                  |                 |

# How to use HPC: Parallel Programming

- ▶ Parallel programming
  - ▶ Multi-threading: shared memory
    - ▶ **OpenMP** for CPUs, **CUDA** for GPUs
  - ▶ Multi-processing: distributed memory
    - ▶ **Message passing interface (MPI)**, Spark, etc.
  - ▶ Storage
    - ▶ Hadoop
  - ▶ Interconnect protocol
    - ▶ TCP/IP, RDMA, etc.
- ▶ Design lower complexity algorithms
  - ▶ FFT: Fast Fourier transform with complexity of  $O(n \log n)$ 
    - ▶ A fast algorithm for computing discrete Fourier transform (DFT)
    - ▶ DFT complexity:  $O(n^2)$
    - ▶  $n$  is the data size
  - ▶ Matrix multiplication:  $A \times B$ , dimension in:  $n \times n$ 
    - ▶ Typically  $O(n^3)$
    - ▶ It can be reduced to  $O(n^{2.37548})$  [1] -  $O(n^{2.3728596})$  [2]
  - ▶ ...

[1] Don et al., "Matrix multiplication via arithmetic progressions," Proceedings of the nineteenth annual ACM symposium on Theory of computing, 1987.

[2] Josh et al., "A Refined Laser Method and Faster Matrix Multiplication," SODA 2021