

COMP 4901Q: High Performance Computing (HPC)

Lecture 3: Revisit C/C++ Programming and Common Linux Commands

Instructor: Shaohuai SHI (shaohuais@cse.ust.hk)

Teaching assistants: Mingkai TANG (mtangag@connect.ust.hk)

Yazhou XING (yxingag@connect.ust.hk)

Course website: <https://course.cse.ust.hk/comp4901q/>

Outline

- ▶ C/C++ Programming: Basic Syntax
- ▶ Common Linux Commands

High-level Programming Languages in HPC

- ▶ Fortran

- ▶ Since 1954 for HPC
- ▶ Good for handling arrays
- ▶ Less flexible in complex and highly dynamic data structure

- ▶ **C/C++**

- ▶ Comparable performance with Fortran
- ▶ More flexible

- ▶ Java

- ▶ May have performance issues

C/C++ Hello World

```
#include <iostream>
```

```
using namespace std;
```

```
int main() {
```

```
    cout << "Hello World" << endl; // prints Hello World
```

```
    return 0;
```

```
}
```

Data Types

- ▶ Integer types
- ▶ Floating-point types
- ▶ The **void** type: no value is available

Type	Storage size	Value range
char	1 byte	-128 to 127 or 0 to 255
unsigned char	1 byte	0 to 255
signed char	1 byte	-128 to 127
int	2 or 4 bytes	-32,768 to 32,767 or -2,147,483,648 to 2,147,483,647
unsigned int	2 or 4 bytes	0 to 65,535 or 0 to 4,294,967,295
short	2 bytes	-32,768 to 32,767
unsigned short	2 bytes	0 to 65,535
long	8 bytes or (4bytes for 32 bit OS)	-9223372036854775808 to 9223372036854775807
unsigned long	8 bytes	0 to 18446744073709551615

Type	Storage size	Value range	Precision
float	4 bytes	1.2E-38 to 3.4E+38	6 decimal places
double	8 bytes	2.3E-308 to 1.7E+308	15 decimal places
long double	10 bytes	3.4E-4932 to 1.1E+4932	19 decimal places

Data Types

- ▶ **typedef** Declarations

- ▶ typedef type newname;

```
typedef int feet;  
feet distance;
```

- ▶ Enumerated Types

- ▶ enum enum-name { list of names } var-list;

```
enum color { red, green, blue };  
color c;
```

Variables

- ▶ **type** identifier;
 - ▶ “identifier” is called a **variable**
 - ▶ Only the following characters may appear in an identifier
 - ▶ 0-9, a-z, A-Z, _
 - ▶ The first character cannot be a digit (0–9)
 - ▶ Keyword (reserved words) are not allowed
 - ▶ Case-sensitive
 - ▶ **type** must be a valid data type
 - ▶ E.g., char, int, float, ...

```
int i, j, k;  
char c, ch;  
float f, salary;  
double d;
```

- ▶ A variable is a named memory location
- ▶ Declaration, definition, and initialization
 - ▶ `int a = 10;`

```
// Variable declaration:  
extern int a, b;  
extern int c;  
extern float f;
```

```
// Variable definition:  
int a, b;  
int c;  
float f;
```

```
// Variable initialization:  
a = 10;  
b = 20;  
c = a + b;
```

asm	else	new	this
auto	enum	operator	throw
bool	explicit	private	true
break	export	protected	try
case	extern	public	typedef
catch	false	register	typeid
char	float	reinterpret_cast	typename
class	for	return	union
const	friend	short	unsigned
const_cast	goto	signed	using
continue	if	sizeof	virtual
default	inline	static	void
delete	int	static_cast	volatile
do	long	struct	wchar_t
double	mutable	switch	while
dynamic_cast	namespace	template	

reserved words

Constants

▶ Fixed values that the program may not alter

▶ Integer

- ▶ 212 // Legal
- ▶ 215u // Legal
- ▶ 0xFeeL // Legal
- ▶ 078 // Illegal: 8 is not an octal digit
- ▶ 032UU // Illegal: cannot repeat a suffix
- ▶ 85 // decimal
- ▶ 0213 // octal
- ▶ 0x4b // hexadecimal
- ▶ 30 // int
- ▶ 30u // unsigned int
- ▶ 30l // long
- ▶ 30ul // unsigned long

▶ Floating-point

- ▶ 3.14159 // Legal
- ▶ 314159E-5L // Legal
- ▶ 510E // Illegal: incomplete exponent
- ▶ 210f // Illegal: no decimal or exponent
- ▶ .e55 // Illegal: missing integer or fraction

▶ Boolean

- ▶ true or false

▶ Character

- ▶ in single quotes, e.g., 'x'

▶ String

- ▶ are enclosed in double quotes, e.g., "hello, hpc"

▶ Defining constants

▶ #define preprocessor:

- ▶ #define identifier value

```
#define LENGTH 10
#define WIDTH 5
#define NEWLINE '\n'
```

▶ const keyword

- ▶ const type variable = value;

```
const int LENGTH = 10;
const int WIDTH = 5;
const char NEWLINE = '\n';
```


Operators

► Assignment: =

```
int A = 10;  
int B = 20;  
int c;  
c = A + B;  
c = c * A;
```

► Arithmetic operators

► Relational operators

Operator	Description	Example
+	Adds two operands	A + B will give 30
-	Subtracts second operand from the first	A - B will give -10
*	Multiplies both operands	A * B will give 200
/	Divides numerator by de-numerator	B / A will give 2
%	Modulus Operator and remainder of after an integer division	B % A will give 0
++	Increment operator , increases integer value by one	A++ will give 11
--	Decrement operator , decreases integer value by one	A-- will give 9

Operator	Description	Example
==	Checks if the values of two operands are equal or not, if yes then condition becomes true.	(A == B) is not true.
!=	Checks if the values of two operands are equal or not, if values are not equal then condition becomes true.	(A != B) is true.
>	Checks if the value of left operand is greater than the value of right operand, if yes then condition becomes true.	(A > B) is not true.
<	Checks if the value of left operand is less than the value of right operand, if yes then condition becomes true.	(A < B) is true.
>=	Checks if the value of left operand is greater than or equal to the value of right operand, if yes then condition becomes true.	(A >= B) is not true.
<=	Checks if the value of left operand is less than or equal to the value of right operand, if yes then condition becomes true.	(A <= B) is true.

Operators cont.

▶ Logical Operators

```
bool A = true;
bool B = false;
```

▶ Bitwise Operators

- ▶ perform bit-by-bit operation

```
int A = 60; //A = 0011 1100
int B = 13; //B = 0000 1101
```

A = 0011 1100

B = 0000 1101

A&B = 0000 1100

A|B = 0011 1101

A^B = 0011 0001

~A = 1100 0011

Operator	Description	Example
&&	Called Logical AND operator. If both the operands are non-zero, then condition becomes true.	(A && B) is false.
	Called Logical OR Operator. If any of the two operands is non-zero, then condition becomes true.	(A B) is true.
!	Called Logical NOT Operator. Use to reverses the logical state of its operand. If a condition is true, then Logical NOT operator will make false.	!(A && B) is true.

Operator	Description	Example
&	Binary AND Operator copies a bit to the result if it exists in both operands.	(A & B) will give 12 which is 0000 1100
	Binary OR Operator copies a bit if it exists in either operand.	(A B) will give 61 which is 0011 1101
^	Binary XOR Operator copies the bit if it is set in one operand but not both.	(A ^ B) will give 49 which is 0011 0001
~	Binary Ones Complement Operator is unary and has the effect of 'flipping' bits.	(~A) will give -61 which is 1100 0011 in 2's complement form due to a signed binary number.
<<	Binary Left Shift Operator. The left operands value is moved left by the number of bits specified by the right operand.	A << 2 will give 240 which is 1111 0000
>>	Binary Right Shift Operator. The left operands value is moved right by the number of bits specified by the right operand.	A >> 2 will give 15 which is 0000 1111

Operators cont.

▶ Operators with assignment operator

- ▶ $B \text{ op} = A$ is equivalent to $B = B \text{ op } A$;
- ▶ For example
 - ▶ $B += A$ is equivalent to $B = B + A$;
 - ▶ $B /= A$ is equivalent to $B = B / A$;
 - ▶ $B |= 2$ is equivalent to $B = B | 2$;

▶ Operators precedence

- ▶ determines the grouping of terms in an expression

```
int A = 60;
int B = 13;
int C = 5;
int D = 15;
int E;
E = A + B * C / D;
E = (A + B) * C / D;
E = (A + B) * (C / D);
E = A++;
E = ++A;
```

Precedence	Operator	Description	Associativity
1	++ --	Suffix/postfix increment and decrement	Left-to-right
	()	Function call	
	[]	Array subscripting	
	.	Structure and union member access	
	-> (type){list}	Structure and union member access through pointer Compound literal(C99)	
2	++ --	Prefix increment and decrement	Right-to-left
	+ -	Unary plus and minus	
	! ~	Logical NOT and bitwise NOT	
	(type)	Cast	
	*	Indirection (dereference)	
	&	Address-of	
	sizeof	Size-of	
	_Alignof	Alignment requirement(C11)	
3	* / %	Multiplication, division, and remainder	Left-to-right
4	+ -	Addition and subtraction	
5	<< >>	Bitwise left shift and right shift	
6	< <=	For relational operators < and ≤ respectively	
	> >=	For relational operators > and ≥ respectively	
7	== !=	For relational = and ≠ respectively	
8	&	Bitwise AND	
9	^	Bitwise XOR (exclusive or)	
10		Bitwise OR (inclusive or)	
11	&&	Logical AND	
12		Logical OR	
13	?:	Ternary conditional	Right-to-left
14	=	Simple assignment	
	+= -=	Assignment by sum and difference	
	*= /= %=	Assignment by product, quotient, and remainder	
	<<= >>=	Assignment by bitwise left shift and right shift	
	&= ^= =	Assignment by bitwise AND, XOR, and OR	
15	,	Comma	Left-to-right

Expression and Statement

▶ Expression

- ▶ has a value which is the result of some operation(s) on its(theirs) operands
- ▶ Examples
 - ▶ 4
 - ▶ $x - y$
 - ▶ `!x`

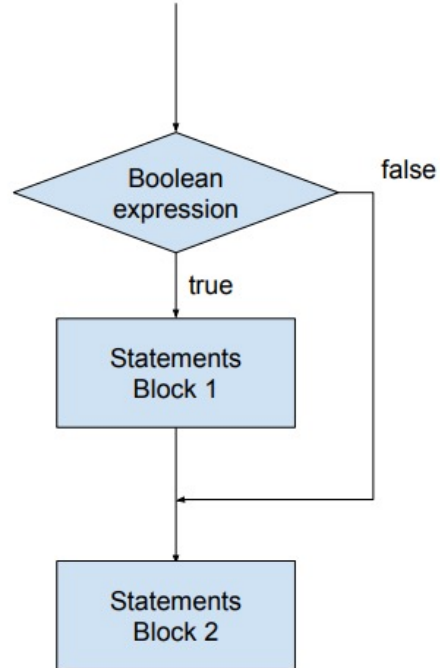
▶ Statement

- ▶ a sentence that acts as a command
- ▶ It does not have a value
- ▶ Ends in a “;”
- ▶ Examples
 - ▶ Input statement: `cin >> x;`
 - ▶ Output statement: `cout << x;`
 - ▶ Assignment statement: `x = 5;`
 - ▶ Variable definition: `int x;`

Flow Control

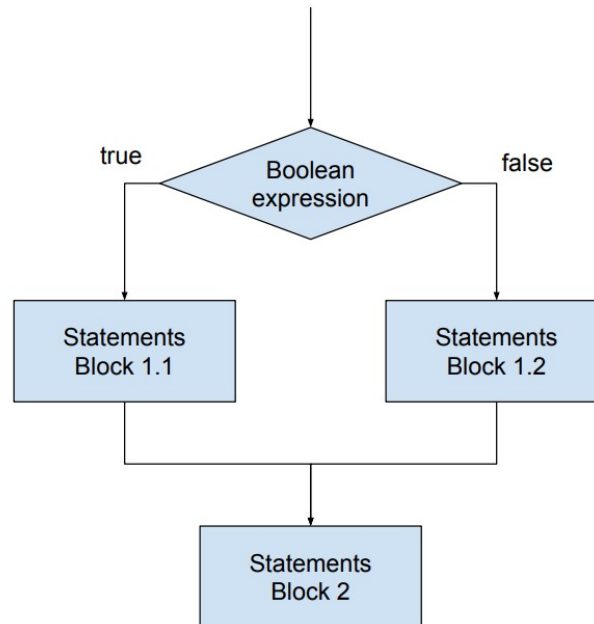
▶ if statement

- ▶ if (boolean expression) statement;
- ▶ if (boolean expression) { statements; }



▶ if-else statement

- ▶ if (boolean expression) statement;
else statement;
- ▶ if (boolean expression) {
statements;
} else { statements; }



▶ if-else-if statement

- ▶ if (boolean exp1) state;
else if (boolean exp2) state;
. . .
else state;

▶ The ? : Operator

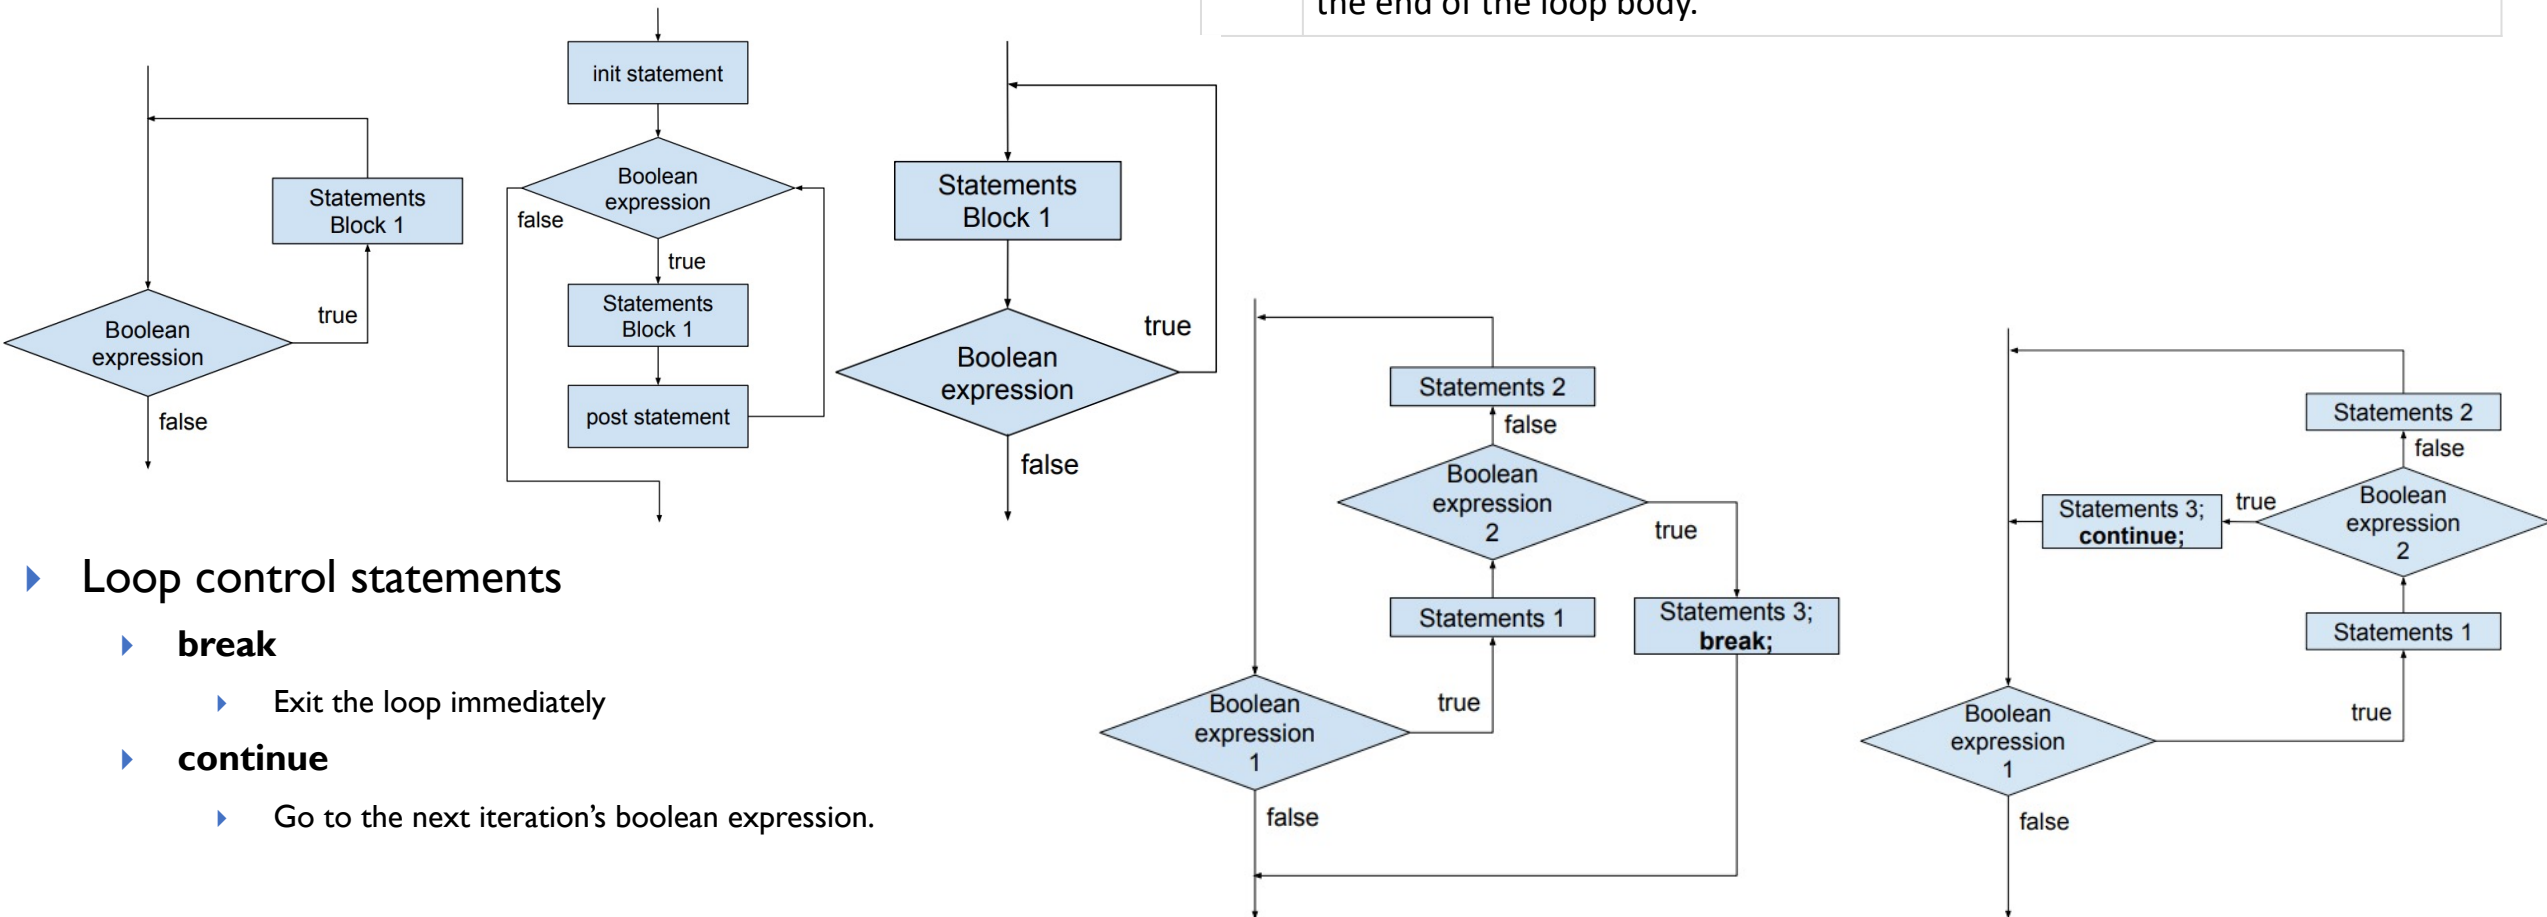
- ▶ is equivalent to if-else
- ▶ boolean exp1 ? exp2 : exp3;
- ▶ The same as
 - ▶ if (exp1) exp2; else exp3;

Loops

▶ Loop statements

- ▶ `while (boolean expression) { statements; }`
- ▶ `for(init statement; boolean expression; post statement) { statements; }`
- ▶ `do { statements; } while (boolean expression);`

No	Loop Type & Description
1	while loop Repeats a statement or group of statements while a given condition is true. It tests the condition before executing the loop body.
2	for loop Execute a sequence of statements multiple times and abbreviates the code that manages the loop variable.
3	do...while loop Like a 'while' statement, except that it tests the condition at the end of the loop body.



▶ Loop control statements

- ▶ **break**
 - ▶ Exit the loop immediately
- ▶ **continue**
 - ▶ Go to the next iteration's boolean expression.

Functions

- ▶ Function declaration

- ▶ tells the compiler about a function's name, return type, and parameters.
- ▶ `return_type function_name(formal parameter list);`

- ▶ Function definition

- ▶ provides the actual body of the function
- ▶ `return_type function_name(formal parameter list) {
 body of the function
}`

- ▶ Function call

- ▶ `function_name(actual parameter list);`

- ▶ **void** as the return type of a function

- ▶ return statement can be avoided.

- ▶ **void** as the formal parameter list of a function

- ▶ no arguments are needed when calling a function

```
int max(int num1, int num2);
```

```
int max(int num1, int num2) {  
    int result;  
    if (num1 > num2) result = num1; else result = num2;  
    return result;  
}
```

```
int main () {  
    int a = 10;  
    int b = 12;  
    int larger;  
    larger = max(a, b);  
    return 0;  
}
```

Function Arguments

- ▶ When calling a function, there are three ways that arguments can be passed to a function
 - ▶ Call by value
 - ▶ The value should be copied
 - ▶ Changes inside the function to the parameter have no effect on the argument
 - ▶ Call by reference
 - ▶ A reference variable is an alias of another variable
 - ▶ Copy the reference of argument instead of copying values
 - ▶ Changes made to the parameter affect the argument.
 - ▶ Call by pointer
 - ▶ Copy the address of an argument into the formal parameter
 - ▶ The address is used to access the actual argument
 - ▶ Changes made to the parameter affect the argument.

```
int max1(int num1, int num2) {  
    if (num1 > num2) num1 = num1; else num1 = num2;  
    return num1;  
}
```

```
int max2(int& num1, int& num2) {  
    if (num1 > num2) num1 = num1; else num1 = num2;  
    return num1;  
}
```

```
int max3(int* num1, int* num2) {  
    if (*num1 > *num2) *num1 = *num1; else *num1 = *num2;  
    return *num1;  
}
```

```
int main () {  
    int a = 10, b = 12;  
    max1(a, b);  
    max2(a, b);  
    max2(&a, &b);  
    return 0;  
}
```


Arrays

▶ Definition of a 1D array

- ▶ **type arrayName [arraySize];**
- ▶ arraySize should be a positive constant or constant expression
- ▶ Examples
 - ▶ `double balance[10];`
 - ▶ `double balance[5] = {1000.0, 2.0, 3.4, 17.0, 50.0};`
 - ▶ `double balance[] = {1000.0, 2.0, 3.4, 17.0, 50.0};` //omit the size of the array
- ▶ Successive elements are stored in contiguous memory

▶ Access

- ▶ Use subscript operator: `[]` with an array index
- ▶ The index is from 0, 1, 2, ..., N-1 for an N-element array

▶ Multi-dimensional array

- ▶ **type name[size1][size2]...[sizeN];**

```
int a[3][4] = { {0, 1, 2, 3} , /* initializers for row indexed by 0 */  
               {4, 5, 6, 7} , /* initializers for row indexed by 1 */  
               {8, 9, 10, 11} /* initializers for row indexed by 2 */ };
```

String

- ▶ The C-Style Character String

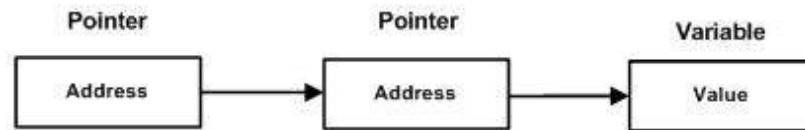
- ▶ For a string of length N , add `'\0'` as the $(N + 1)$ th element of its char array
- ▶ char arrays
 - ▶ `char greeting[6] = {'H', 'e', 'l', 'l', 'o', '\0'};`
- ▶ double quotes
 - ▶ `char greeting[] = "Hello";`

- ▶ The String Class in C++

- ▶ `#include <string>`
- ▶ `string str1 = "Hello";`

Pointers

- ▶ Every variable is a memory location and every memory location has its address
- ▶ The ampersand (&) operator denotes an address in memory
- ▶ A **pointer** is a variable whose value is the address of another variable
 - ▶ **type *var-name;**
- ▶ Pointer to Pointer
 - ▶ **type **var-name;**



```
int sum = 255;  
short int age = -1;  
double average = 4.45015e-308;  
int *ptrSum = &sum;
```

Computer		Programmers		
Address	Content	Name	Type	Value
90000000	00	sum	int (4 bytes)	000000FF (255 ₁₀)
90000001	00			
90000002	00			
90000003	FF			
90000004	FF	age	short (2 bytes)	FFFF (-1 ₁₀)
90000005	FF			
90000006	1F	average	double (8 bytes)	1FFFFFFFFFFFFFFF (4.45015E-308 ₁₀)
90000007	FF			
90000008	FF			
90000009	FF			
9000000A	FF			
9000000B	FF			
9000000C	FF			
9000000D	FF			
9000000E	90	ptrSum	int* (4 bytes)	90000000
9000000F	00			
90000010	00			
90000011	00			

Note: All numbers in hexadecimal

Source: <https://astrocoke.tistory.com/32>

C++ Classes and Objects

▶ Class

- ▶ define a blueprint for a data type

Sr.No	Concept & Description
1	Class Member Functions : A member function of a class is a function that has its definition or its prototype within the class definition like any other variable.
2	Class Access Modifiers : A class member can be defined as public, private or protected. By default members would be assumed as private.
3	Constructor & Destructor : A class constructor is a special function in a class that is called when a new object of the class is created. A destructor is also a special function which is called when created object is deleted.

```
class Box {  
    public:  
        double length;    // Length of a box  
        double breadth;   // Breadth of a box  
        double height;    // Height of a box  
        Box() { length = 1; } // Constructor  
        ~Box() {}          // Destructor  
        double getVolume(void) {  
            return length * breadth * height;  
        }  
};
```

Files and Streams

- ▶ **iostream** standard library

- ▶ **cin**: standard input
- ▶ **cout**: standard output

- ▶ **fstream**

- ▶ **ofstream**: output file stream
- ▶ **ifstream**: input file stream
- ▶ **fstream**: represents the file stream, and has the capabilities of both ofstream and ifstream

- ▶ Open/close a File

- ▶ `void open(const char *filename, ios::openmode mode);`

Sr.No	Mode Flag & Description
1	ios::app : Append mode. All output to that file to be appended to the end.
2	ios::ate : Open a file for output and move the read/write control to the end of the file.
3	ios::in : Open a file for reading.
4	ios::out : Open a file for writing.
5	ios::trunc : If the file already exists, its contents will be truncated before opening the file.

- ▶ `void close();`

C++ Multithreading

- ▶ Multithreading for multitasking that allows your computer to run two or more programs concurrently

- ▶ Usage

- ▶ **#include <pthread.h>**

- ▶ **pthread_create (thread, attr, start_routine, arg)**

- ▶ **thread:** an opaque, unique identifier for the new thread
 - ▶ **attr:** an opaque attribute object that may be used to set thread attributes
 - ▶ **start_routine:** the C++ routine that the thread will execute once it is created
 - ▶ **arg:** A single argument that may be passed to start_routine

- ▶ **pthread_exit (status)**

```
#include <iostream>
#include <cstdlib>
#include <pthread.h>
using namespace std;
#define NUM_THREADS 5
struct thread_data {
    int thread_id;
    char *message;
};

void *PrintHello(void *threadarg) {
    struct thread_data *my_data;
    my_data = (struct thread_data *) threadarg;
    cout << "Thread ID : " << my_data->thread_id ;
    cout << " Message : " << my_data->message << endl;
    pthread_exit(NULL);
}

int main () {
    pthread_t threads[NUM_THREADS];
    struct thread_data td[NUM_THREADS];
    int rc, i;
    for( i = 0; i < NUM_THREADS; i++ ) {
        cout << "main() : creating thread, " << i << endl;
        td[i].thread_id = i;
        td[i].message = "This is message";
        rc = pthread_create(&threads[i], NULL, PrintHello, (void *)&td[i]);
        if (rc) {
            cout << "Error:unable to create thread," << rc << endl;
            exit(-1);
        }
    }
    pthread_exit(NULL);
}
```

Common Linux Commands

cmd	Description
ping	Check the connectivity status to a server
ssh	Login a remote server
scp	Copy files from/to a server

cmd	Description
pwd	Path of the current directory
cd	Go to a directory \$cd folder; \$cd - \$cd \$cd ..
ls	View the contents of a directory

cmd	Description
cat	List the contents of a file
cp	Copy files \$cp file.txt file-copy.txt; \$cd -r folder folder-copy;
mv	Move files
mkdir	Create a new directory \$mkdir new-directory \$mkdir -p d1/d2/d3
rm	Remove files \$rm file.txt \$rm -r d1
touch	Create a blank new file
find	Searches for files and directories \$find . -name notes.txt \$find / -type d -name notes
grep	search through all the text in a given file \$grep blue notepad.txt

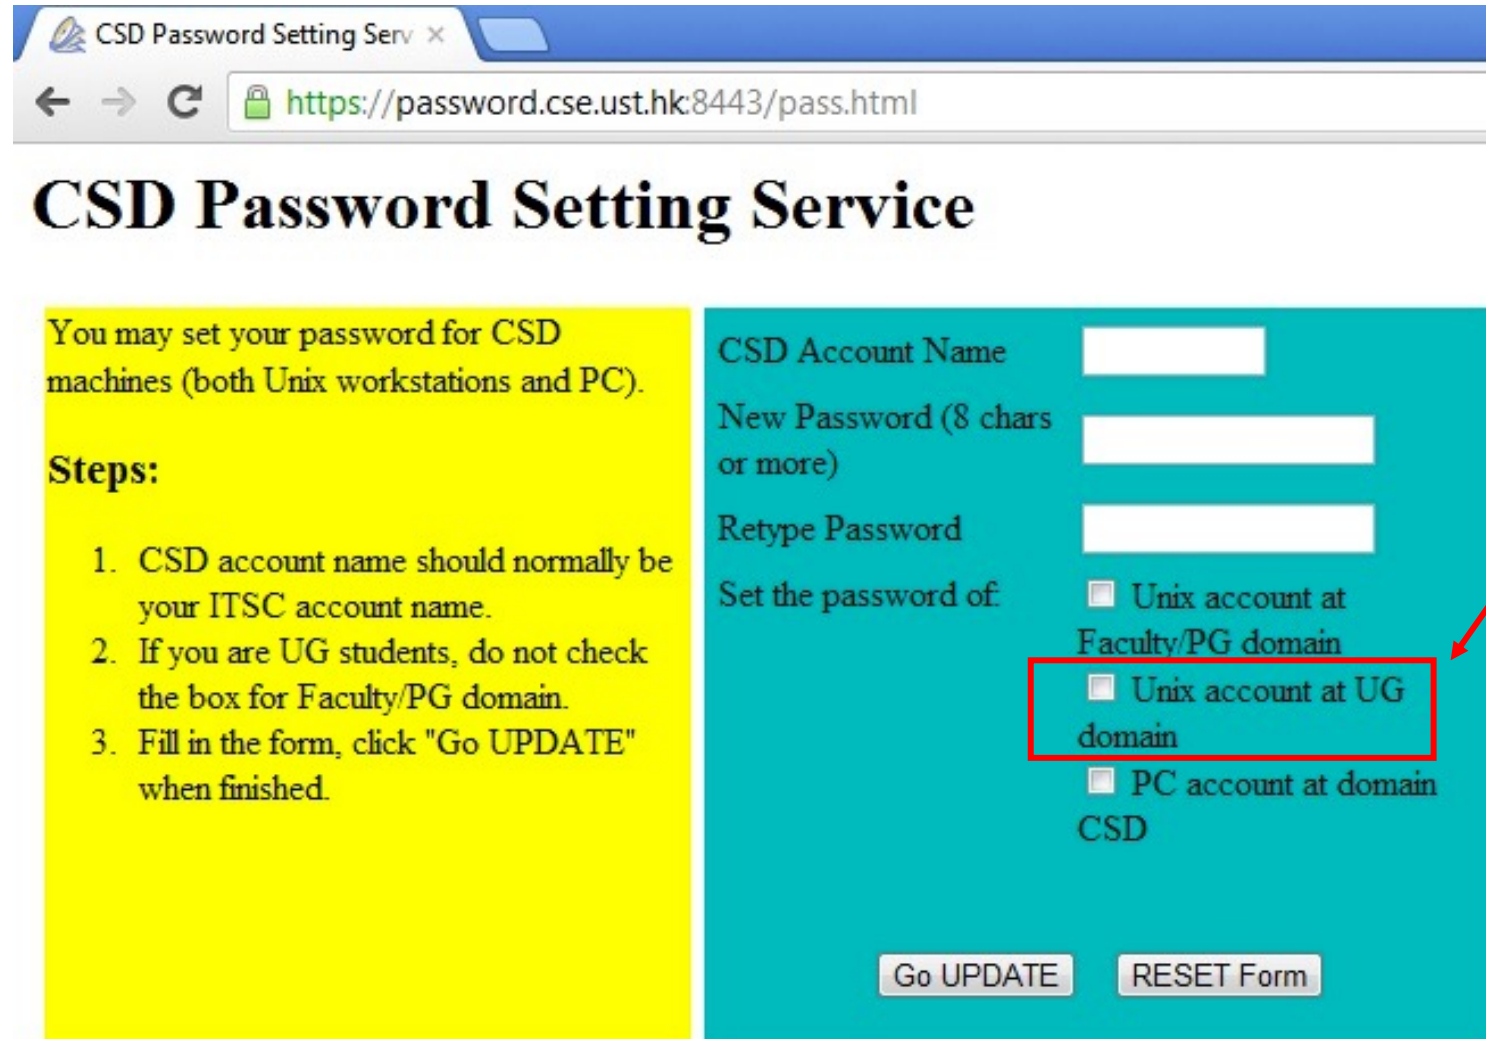
cmd	Description
kill	Terminate a process
wget	Download files
top	Display a list of running processes
lscpu	Check CPU information
df	Check disk information \$df -h
ifconfig	Display network configuration
free	Display system memory \$free -h
vi	Default editor of Unix/Linux systems

Ubuntu is easy to use!

<http://mally.stanford.edu/~sr/computing/basic-unix.html>

Access Our GPU Cluster

Follow the guide: <https://cssystem.cse.ust.hk/UGuides/activation.html>



CSD Password Setting Serv x

← → ↻ <https://password.cse.ust.hk:8443/pass.html>

CSD Password Setting Service

You may set your password for CSD machines (both Unix workstations and PC).

Steps:

1. CSD account name should normally be your ITSC account name.
2. If you are UG students, do not check the box for Faculty/PG domain.
3. Fill in the form, click "Go UPDATE" when finished.

CSD Account Name

New Password (8 chars or more)

Retype Password

Set the password of:

- ☐ Unix account at Faculty/PG domain
- ☒ Unix account at UG domain
- ☐ PC account at domain CSD

Go UPDATE RESET Form

Remember to check this!

Summary

- ▶ **C/C++ programming basics**
- ▶ **Common Linux commands**

Practice makes perfect!