

# COMP 4901Q: High Performance Computing (HPC)

## Lecture 13: Introduction to Hadoop and Spark

Instructor: Shaohuai SHI ([shaohuais@cse.ust.hk](mailto:shaohuais@cse.ust.hk))

Teaching assistants: Mingkai TANG ([mtangag@connect.ust.hk](mailto:mtangag@connect.ust.hk))

Yazhou XING ([yxingag@connect.ust.hk](mailto:yxingag@connect.ust.hk))

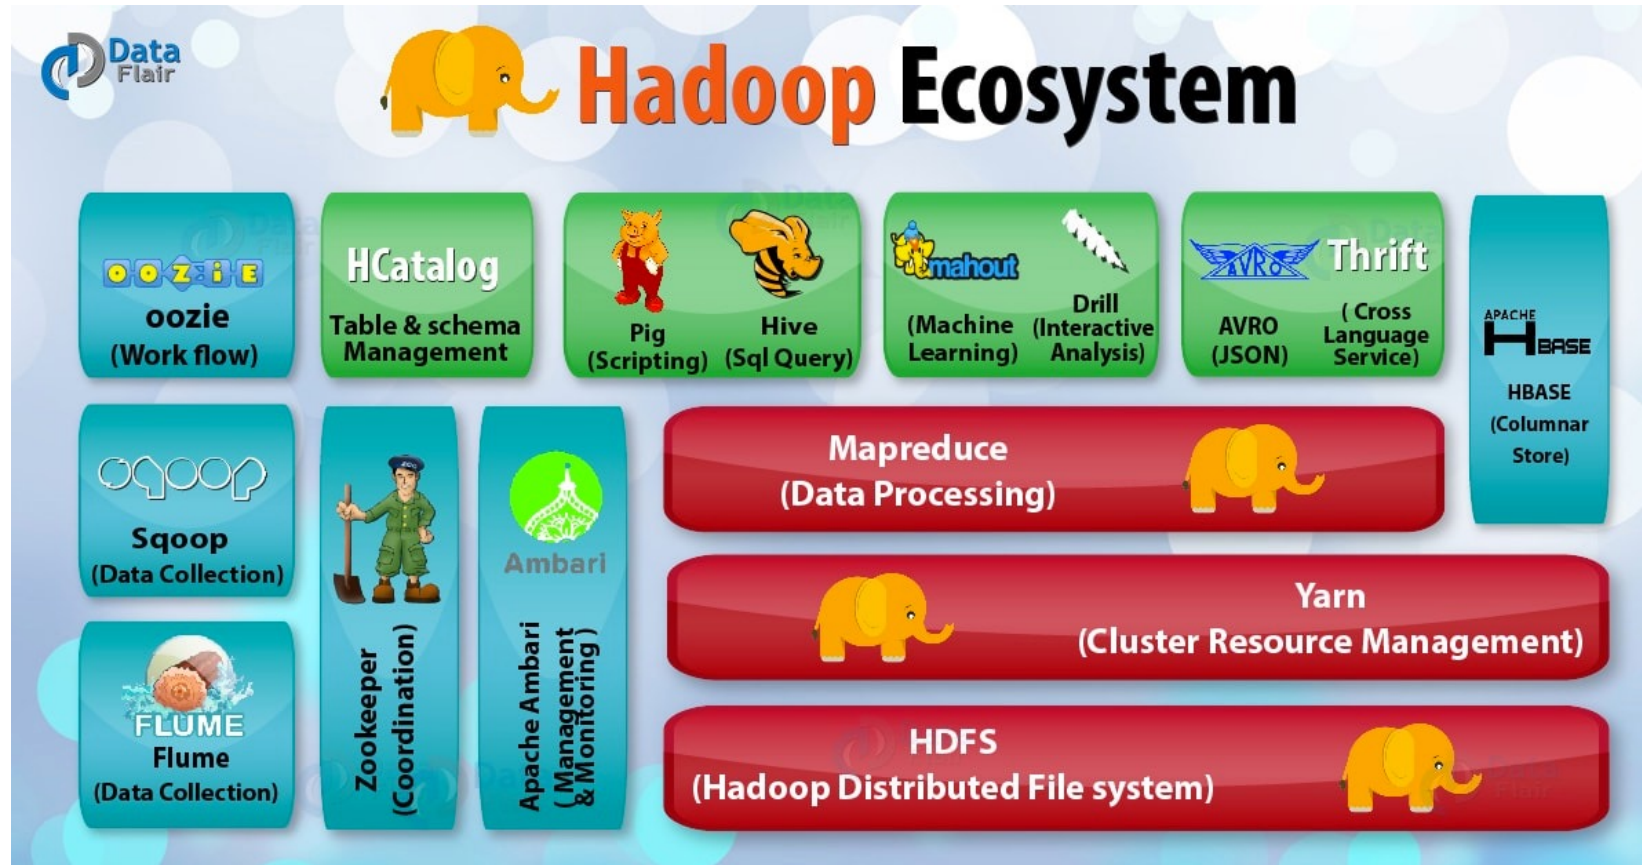
Course website: <https://course.cse.ust.hk/comp4901q/>

# Outline

- ▶ Hadoop Ecosystem
- ▶ MapReduce and Spark

# Hadoop Ecosystem

- ▶ The Apache Hadoop software library is a framework that allows for the distributed processing of large data sets across clusters of computers using simple programming models.
  - ▶ <https://hadoop.apache.org/>



Source: <https://data-flair.training/blogs/hadoop-ecosystem-components/>

# Apache Hadoop Basic Modules


- ▶ Hadoop Common
  - ▶ The common utilities that support the other Hadoop modules
- ▶ Hadoop Distributed File System (HDFS)
  - ▶ A distributed file system that provides high-throughput access to application data
- ▶ Hadoop YARN
  - ▶ A framework for job scheduling and cluster resource management
- ▶ Hadoop MapReduce
  - ▶ A YARN-based system for parallel processing of large data sets
- ▶ Hadoop Ozone
  - ▶ An object store for Hadoop.

# Hadoop MapReduce

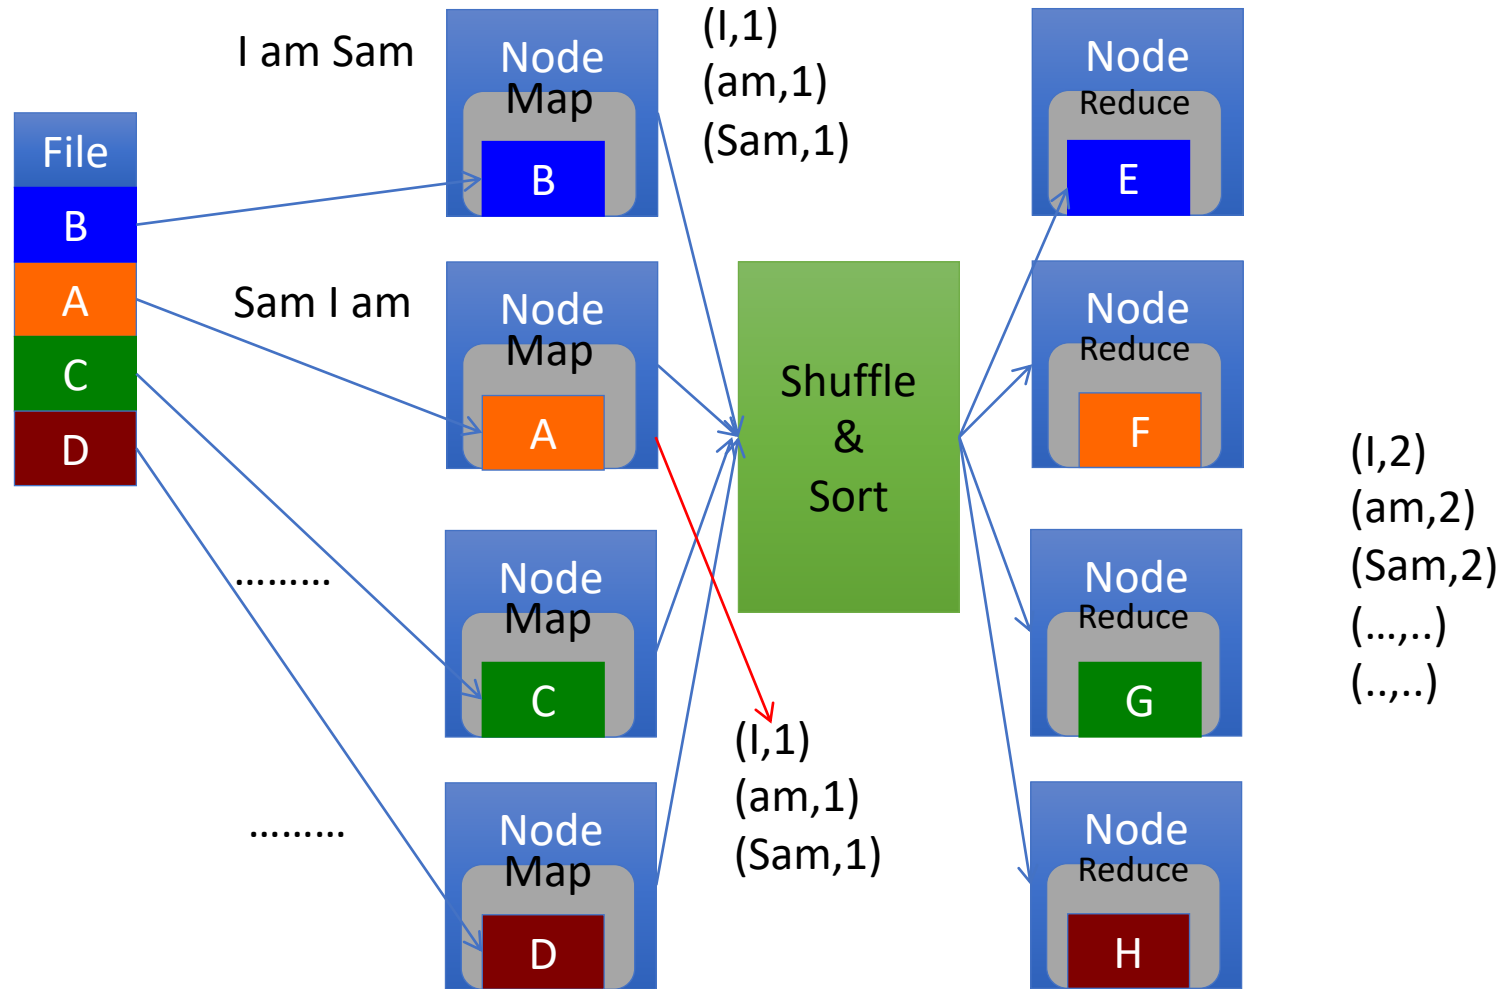
- ▶ MapReduce is simple programming paradigm for the Hadoop
- ▶ Traditional parallel programming requires expertise of different computing/systems concepts
  - ▶ Examples: multithreads, synchronization mechanisms (locks, semaphores, and monitors )
  - ▶ Incorrect use: can crash your program, get incorrect results, or severely impact performance
  - ▶ Usually not fault tolerant to hardware failure
- ▶ The MapReduce programming model greatly simplifies running code in parallel

# MapReduce Paradigm

- ▶ Map and Reduce are based on functional programming
- ▶ Breaking the processing into two phases:
  - ▶ Map phase: takes a set of data and converts it into another set of data, where individual elements are broken down into tuples (key/value pairs)
  - ▶ Reduce phase: takes the output from the Map as an input and combines those data tuples based on the key and accordingly modifies the value of the key



# MapReduce: Word Count Example




# Shortcoming of MapReduce

- ▶ Forces your data processing into Map and Reduce
- ▶ Based on “Acyclic Data Flow” from Disk to Disk (HDFS)
  - ▶ Read and write to Disk before and after Map and Reduce (stateless machine)
  - ▶ Not efficient for iterative tasks, i.e. Machine Learning
- ▶ Only Java natively supported
- ▶ Only for Batch processing
  - ▶ Interactivity, streaming data



# Apache Spark

- ▶ Apache Spark is a unified analytics engine for large-scale data processing.
  - ▶ Runs on many frameworks
  - ▶ High speed: ~100x faster
  - ▶ Ease-of-use: supports Java, Scala, Python, R, and SQL.
  - ▶ A stack of libraries: SQL and DataFrames, MLlib for machine learning, GraphX, and Spark Streaming



# Word Count

```
text_file = sc.textFile("hdfs://input/book.txt")
counts = text_file.flatMap(lambda line: line.split(" ")) \
    .map(lambda word: (word, 1)) \
    .reduceByKey(lambda a, b: a + b)
counts.saveAsTextFile("hdfs://output/wordCount.txt")
```

```
public class WordCount {

    public static class TokenizerMapper
        extends Mapper<Object, Text, Text, IntWritable>{

        private final static IntWritable one = new IntWritable(1);
        private Text word = new Text();


        public void map(Object key, Text value, Context context
            ) throws IOException, InterruptedException {
            StringTokenizer itr = new StringTokenizer(value.toString());
            while (itr.hasMoreTokens()) {
                word.set(itr.nextToken());
                context.write(word, one);
            }
        }
    }

    public static class IntSumReducer
        extends Reducer<Text, IntWritable, Text, IntWritable> {
        private IntWritable result = new IntWritable();


        public void reduce(Text key, Iterable<IntWritable> values,
            Context context
            ) throws IOException, InterruptedException {

            int sum = 0;
            for (IntWritable val : values) {
                sum += val.get();
            }
            result.set(sum);
            context.write(key, result);
        }
    }
}
```

# Spark vs. MPI



Performance comparison [1]



Deep Learning research [2]

[1] Abuín, J. M., Lopes, N., Ferreira, L., Pena, T. F., & Schmidt, B. (2020). Big Data in metagenomics: Apache Spark vs MPI. *Plos one*, 15(10), e0239741.

[2] Ben-Nun, T., & Hoefler, T. (2019). Demystifying parallel and distributed deep learning: An in-depth concurrency analysis. *ACM Computing Surveys (CSUR)*, 52(4), 1-43.