

# Machine Learning

## Lecture 01-1: Basics of Probability Theory

Nevin L. Zhang

`lzhang@cse.ust.hk`

Department of Computer Science and Engineering  
The Hong Kong University of Science and Technology

# Outline

1 Basic Concepts in Probability Theory

2 Interpretation of Probability

3 Bayes' Theorem

4 Parameter Estimation

# Random Experiments

- Probability associated with a **random experiment** — a process with uncertain outcomes
- Often kept implicit



Tail




Head

In machine learning, we often assume that data are generated by a hypothetical process (or a model), and task is to determine the structure and parameters of the model from data.

# Sample Space


- **Sample space (aka population)  $\Omega$ :** Set of possible outcomes and a random experiment.
- Example: Rolling two dice.



- Elements in a sample space are outcomes.

# Events

- **Event:** A subset of the sample space.

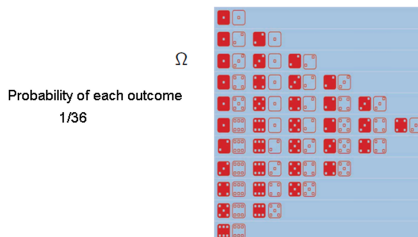


- **Example:** The two results add to 4.



# Probability Weight Function

- A **probability weight**  $P(\omega)$  is assigned to each outcome.



In Machine Learning, we often need to determine the probability weights, or related parameters, from data. This task is called **parameter learning**.

# Probability measure

- Probability  $P(E)$  of an event  $E$ :  $P(E) = \sum_{\omega \in E} P(\omega)$
- A **probability measure** is a mapping from the set of **events** to  $[0, 1]$

$$P : 2^{\Omega} \rightarrow [0, 1]$$

that satisfies Kolmogorov's axioms:

- 1  $P(\Omega) = 1$ .
- 2  $P(A) \geq 0 \ \forall A \subseteq \Omega$
- 3 **Additivity**:  $P(A \cup B) = P(A) + P(B)$  if  $A \cap B = \emptyset$ .

In a more advanced treatment of Probability Theory, we would start with the concept of probability measure, instead of probability weights.

# Random Variables

- A **random variable** is a function over the sample space.
  - Example:  $X = \text{sum of the two results}$ .  $X((2, 5)) = 7$ ;  $X((3, 1)) = 4$

DICE CHART		$P(X=x)$
$\Omega_X$		PROBABILITY
2		1/36
3		2/36
4		3/36
5		4/36
6		5/36
7		6/36
8		5/36
9		4/36
10		3/36
11		2/36
12		1/36

- Why is it random? The experiment.
- **Domain** of a random variable: Set of all its possible values.

$$\Omega_X = \{2, 3, \dots, 12\}$$



# Random Variables and Event

- A random variable  $X$  taking a specific value  $x$  is an event:

$$\Omega_{X=x} = \{\omega \in \Omega | X(\omega) = x\}$$


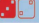
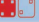
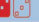
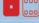
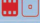
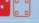
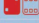
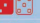
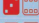
DICE CHART	
$\Omega_X$	$P(X=x)$
2	1/36
3	2/36
4	3/36
5	4/36
6	5/36
7	6/36
8	5/36
9	4/36
10	3/36
11	2/36
12	1/36

- $\Omega_{X=4} = \{(1, 3), (2, 2), (3, 1)\}$ .

# Probability Mass Function (Distribution)

- **Probability mass function**  $P(X): \Omega_X \rightarrow [0, 1]$

$$P(X = x) = P(\Omega_{X=x})$$

$\Omega_X$	DICE CHART		$P(X=x)$
	PROBABILITY		
2			1/36
3	 		2/36
4	  		3/36
5	   		4/36
6	    		5/36
7	     		6/36
8	    		5/36
9	   		4/36
10	  		3/36
11	 		2/36
12			1/36

- $P(X = 4) = P(\{(1, 3), (2, 2), (3, 1)\}) = \frac{3}{36}$ .
- If  $X$  is continuous, we have a **density function**  $p(X)$ .

# Outline

- 1 Basic Concepts in Probability Theory
- 2 Interpretation of Probability**
- 3 Bayes' Theorem
- 4 Parameter Estimation

# Frequentist interpretation

- Probabilities are **long term relative frequencies**.

- Example:

- $X$  is result of coin tossing.  $\Omega_X = \{H, T\}$
- $P(X=H) = 1/2$  means that
  - *the relative frequency of getting heads* will almost surely approach  $1/2$  as the number of tosses goes to infinite.
- Justified by the Law of Large Numbers:
  - $X_i$ : result of the  $i$ -th tossing; 1 – H, 0 – T
  - Law of Large Numbers:

$$\lim_{n \rightarrow \infty} \frac{\sum_{i=1}^n X_i}{n} = \frac{1}{2} \quad \text{with probability 1}$$

- The frequentist interpretation is meaningful only when experiment can be repeated under the same condition.

# Bayesian interpretation

- Probabilities are **logically consistent degrees of beliefs**.
- Applicable when experiment not repeatable.
- Depends on a person's state of knowledge.
- Example: “probability that Suez canal is longer than the Panama canal”.
  - Doesn't make sense under frequentist interpretation.
  - Subjectivist: degree of belief based on state of knowledge
    - Primary school student: 0.5
    - Me: 0.8
    - Geographer: 1 or 0
- Arguments such as **Dutch book** are used to explain why one's probability beliefs must satisfy Kolmogorov's axioms.

# Interpretations of Probability

- Now both interpretations are accepted. In practice, subjective beliefs and statistical data complement each other.
  - We rely on subjective beliefs (**prior probabilities**) when data are scarce.
  - As more and more data become available, we rely less and less on subjective beliefs.
  - Often, we also use **prior probabilities** to impose some **bias** on the kind of results we want from a machine learning algorithm.
- The subjectivist interpretation makes concepts such as conditional independence easy to understand.

# Outline

- 1 Basic Concepts in Probability Theory
- 2 Interpretation of Probability
- 3 Bayes' Theorem**
- 4 Parameter Estimation

# Prior, posterior, and likelihood

- Three important concepts in Bayesian inference.
- With respect to a piece of evidence:  $E$
- **Prior probability**  $P(H)$ : belief about a hypothesis before observing evidence.
  - Example: Suppose 10% of people suffer from Hepatitis B. A doctor's prior probability about a new patient suffering from Hepatitis B is 0.1.
- **Posterior probability**  $P(H|E)$ : belief about a hypothesis after obtaining the evidence.
  - If the doctor finds that the eyes of the patient are yellow, his belief about patient suffering from Hepatitis B would be  $> 0.1$ .



# Prior, posterior, and likelihood

- Suppose a patient is observed to have yellow eyes ( $E$ ).
- Consider two possible explanations:
  - 1 The patient has Hepatitis B ( $H_1$ ),
  - 2 The patient does not have Hepatitis B ( $H_2$ )
- Obviously,  $H_1$  is a better explanation because  $P(E|H_1) > P(E|H_2)$ . To state it another way, we say that  $H_1$  is more **likely** than  $H_2$  given  $E$ .
- In general, the **likelihood** of a hypothesis  $H$  given evidence  $E$  is a measure of how well  $H$  explains  $E$ . Mathematically, it is

$$L(H|E) = P(E|H)$$

- In Machine Learning, we often talk about the likelihood of a model  $M$  given data  $D$ . It is a measure of how well the model  $M$  explains the data  $D$ . Mathematically, it is

$$L(M|D) = P(D|M)$$

# Bayes' Theorem/Bayes Rule

- **Bayes' Theorem:** relates prior probability, likelihood, and posterior probability:

$$P(H|E) = \frac{P(H)P(E|H)}{P(E)} \propto P(H)L(H|E)$$

where  $P(E)$  is normalization constant to ensure  $\sum_{h \in \Omega_H} P(H = h|E) = 1$ .

That is:                  posterior  $\propto$  prior  $\times$  likelihood


# Outline

- 1 Basic Concepts in Probability Theory
- 2 Interpretation of Probability
- 3 Bayes' Theorem
- 4 Parameter Estimation**

# A Simple Problem

- Let  $X$  be the result of tossing a thumbtack and  $\Omega_X = \{H, T\}$ .
- Data instances:  
 $D_1 = H, D_2 = T, D_3 = H, \dots, D_m = H$
- Data set:  $\mathcal{D} = \{D_1, D_2, D_3, \dots, D_m\}$
- Task: To estimate parameter  $\theta = P(X=H)$ .

X: result of tossing a thumbtack



# Likelihood

- Data:  $\mathcal{D} = \{H, T, H, T, T, H, T\}$
- As possible values of  $\theta$ , which of the following is the most likely?  
Why?
  - $\theta = 0$
  - $\theta = 0.01$
  - $\theta = 0.5$
- $\theta = 0$  contradicts data because  $P(\mathcal{D}|\theta = 0) = 0$ . It cannot explain the data at all.
- $\theta = 0.01$  almost contradicts with the data. It does not explain the data well.  
However, it is more consistent with the data than  $\theta = 0$  because  $P(\mathcal{D}|\theta = 0.01) > P(\mathcal{D}|\theta = 0)$ .
- So  $\theta = 0.5$  is more consistent with the data than  $\theta = 0.01$  because  $P(\mathcal{D}|\theta = 0.5) > P(\mathcal{D}|\theta = 0.01)$   
It explains the data the best, and is hence the most likely.

# Maximum Likelihood Estimation


- In general, the larger  $P(\mathcal{D}|\theta)$  is, the more likely the value  $\theta$  is.
- Likelihood of parameter  $\theta$  given data set:

$$L(\theta|\mathcal{D}) = P(\mathcal{D}|\theta)$$

- The **maximum likelihood estimation (MLE)**  $\theta^*$  is

$$L(\theta^*|\mathcal{D}) = \arg \max_{\theta} L(\theta|\mathcal{D}).$$

MLE best explains data or best fits data.



# i.i.d and Likelihood

- Assume the data instances  $D_1, \dots, D_m$  are independent given  $\theta$ :

$$P(D_1, \dots, D_m | \theta) = \prod_{i=1}^m P(D_i | \theta)$$

- Assume the data instances are identically distributed:

$$P(D_i = H) = \theta, P(D_i = T) = 1 - \theta \quad \text{for all } i$$

(Note: i.i.d means independent and identically distributed)

- Then

$$\begin{aligned} L(\theta | \mathcal{D}) &= P(\mathcal{D} | \theta) = P(D_1, \dots, D_m | \theta) \\ &= \prod_{i=1}^m P(D_i | \theta) = \theta^{m_h} (1 - \theta)^{m_t} \end{aligned} \quad (1)$$

where  $m_h$  is the number of heads and  $m_t$  is the number of tail.

**Binomial likelihood.**

# Example of Likelihood Function

■ Example:  $\mathcal{D} = \{D_1 = H, D_2 = T, D_3 = H, D_4 = H, D_5 = T\}$

$$\begin{aligned} L(\theta|\mathcal{D}) &= P(\mathcal{D}|\theta) \\ &= P(D_1 = H|\theta)P(D_2 = T|\theta)P(D_3 = H|\theta)P(D_4 = H|\theta)P(D_5 = T|\theta) \\ &= \theta(1 - \theta)\theta\theta(1 - \theta) \\ &= \theta^3(1 - \theta)^2. \end{aligned}$$



# Sufficient Statistic

- A **sufficient statistic** is a function  $s(\mathcal{D})$  of data that summarizing the relevant information for computing the likelihood. That is

$$s(\mathcal{D}) = s(\mathcal{D}') \Rightarrow L(\theta|\mathcal{D}) = L(\theta|\mathcal{D}')$$

- Sufficient statistics tell us all there is to know about data.
- Since  $L(\theta|\mathcal{D}) = \theta^{m_h}(1 - \theta)^{m_t}$ ,  
the pair  $(m_h, m_t)$  is a **sufficient statistic**.

# Loglikelihood

## ■ Loglikelihood:

$$l(\theta|\mathcal{D}) = \log L(\theta|\mathcal{D}) = \log \theta^{m_h} (1 - \theta)^{m_t} = m_h \log \theta + m_t \log(1 - \theta)$$

Maximizing likelihood is the same as maximizing loglikelihood. The latter is easier.

- Taking the derivative of  $\frac{dl(\theta|\mathcal{D})}{d\theta}$  and setting it to zero, we get

$$\theta^* = \frac{m_h}{m_h + m_t} = \frac{m_h}{m}$$

- MLE is intuitive.
- It also has nice properties:
  - E.g. **Consistence**:  $\theta^*$  approaches the true value of  $\theta$  with probability 1 as  $m$  goes to infinity.

# Drawback of MLE

- Thumbtack tossing:

- $(m_h, m_t) = (3, 7)$ . MLE:  $\theta = 0.3$ .
- Reasonable. Data suggest that the thumbtack is biased toward tail.

- Coin tossing:

- Case 1:  $(m_h, m_t) = (3, 7)$ . MLE:  $\theta = 0.3$ .
  - Not reasonable.
  - Our experience (prior) suggests strongly that coins are fair, hence  $\theta=1/2$ .
  - The size of the data set is too small to convince us this particular coin is biased.
  - The fact that we get  $(3, 7)$  instead of  $(5, 5)$  is probably due to randomness.
- Case 2:  $(m_h, m_t) = (30,000, 70,000)$ . MLE:  $\theta = 0.3$ .
  - Reasonable.
  - Data suggest that the coin is after all biased, overshadowing our prior.
- MLE does not differentiate between those two instances. It does not take prior information into account.

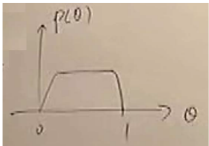
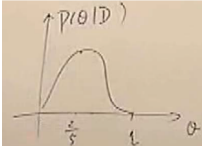
# Two Views on Parameter Estimation

## MLE:

- Assumes that  $\theta$  is unknown but fixed parameter.
- Estimates it using  $\theta^*$ , the value that maximizes the likelihood function
- Makes prediction based on the estimation:  $P(D_{m+1} = H|\mathcal{D}) = \theta^*$

## Bayesian Estimation:

- Treats  $\theta$  as a random variable.
- Assumes a prior probability of  $\theta$ :  $p(\theta)$
- Uses data to get posterior probability of  $\theta$ :  $p(\theta|\mathcal{D})$

	Before Seeing Data	After Seeing Data: {2H, 3T}
MLE	?	2/5
Bayesian Estimation		

# Two Views on Parameter Estimation

## Bayesian Estimation:

### ■ Predicting $D_{m+1}$

$$\begin{aligned}P(D_{m+1} = H|\mathcal{D}) &= \int P(D_{m+1} = H, \theta|\mathcal{D})d\theta \\&= \int P(D_{m+1} = H|\theta, \mathcal{D})p(\theta|\mathcal{D})d\theta \\&= \int P(D_{m+1} = H|\theta)p(\theta|\mathcal{D})d\theta \\&= \int \theta p(\theta|\mathcal{D})d\theta.\end{aligned}$$

**Full Bayesian:** Take expectation over  $\theta$ .

### ■ Bayesian MAP:

$$P(D_{m+1} = H|\mathcal{D}) = \theta^* = \arg \max p(\theta|\mathcal{D})$$

# Calculating Bayesian Estimation

- Posterior distribution:

$$\begin{aligned} p(\theta|\mathcal{D}) &\propto p(\theta)L(\theta|\mathcal{D}) \\ &= \theta^{m_h}(1-\theta)^{m_t}p(\theta) \end{aligned}$$

where the equation follows from (1)


- To facilitate analysis, assume prior has **Beta distribution**  $B(\alpha_h, \alpha_t)$

$$p(\theta) \propto \theta^{\alpha_h-1}(1-\theta)^{\alpha_t-1}$$

- Then

$$p(\theta|\mathcal{D}) \propto \theta^{m_h+\alpha_h-1}(1-\theta)^{m_t+\alpha_t-1} \quad (2)$$

# Beta Distribution



- The normalization constant for the Beta distribution  $B(\alpha_h, \alpha_t)$

$$\frac{\Gamma(\alpha_t + \alpha_h)}{\Gamma(\alpha_t)\Gamma(\alpha_h)}$$

where  $\Gamma(\cdot)$  is the **Gamma function**. For any integer  $\alpha$ ,  $\Gamma(\alpha) = (\alpha - 1)!$ . It is also defined for non-integers.

- Density function of prior Beta distribution  $B(\alpha_h, \alpha_t)$ ,

$$p(\theta) = \frac{\Gamma(\alpha_t + \alpha_h)}{\Gamma(\alpha_t)\Gamma(\alpha_h)} \theta^{\alpha_h-1} (1 - \theta)^{\alpha_t}$$

- The **hyperparameters**  $\alpha_h$  and  $\alpha_t$  can be thought of as "imaginary" counts from our prior experiences.
- Their sum  $\alpha = \alpha_h + \alpha_t$  is called **equivalent sample size**.
- The larger the equivalent sample size, the more confident we are in our prior.

# Conjugate Families

- Binomial Likelihood:  $\theta^{m_h}(1 - \theta)^{m_t}$
- Beta Prior:  $\theta^{\alpha_h-1}(1 - \theta)^{\alpha_t-1}$
- Beta Posterior:  $\theta^{m_h+\alpha_h-1}(1 - \theta)^{m_t+\alpha_t-1}$ .
- Beta distributions are hence called a **conjugate family** for Binomial likelihood.
- Conjugate families allow closed-form for posterior distribution of parameters and closed-form solution for prediction.



# Calculating Prediction

- We have

$$\begin{aligned}
 P(D_{m+1} = H|\mathcal{D}) &= \int \theta p(\theta|\mathcal{D}) d\theta \\
 &= c \int \theta \theta^{m_h + \alpha_h - 1} (1 - \theta)^{m_t + \alpha_t - 1} d\theta \\
 &= \frac{m_h + \alpha_h}{m + \alpha}
 \end{aligned}$$

where  $c$  is the normalization constant,  $m = m_h + m_t$ ,  $\alpha = \alpha_h + \alpha_t$ .

- Consequently,

$$P(D_{m+1} = T|\mathcal{D}) = \frac{m_t + \alpha_t}{m + \alpha}$$

- After taking data  $\mathcal{D}$  into consideration, now our **updated belief** on  $X = T$  is  $\frac{m_t + \alpha_t}{m + \alpha}$ .

# MLE and Bayesian estimation

- As  $m$  goes to infinity,  $P(D_{m+1} = H|\mathcal{D})$  approaches the MLE  $\frac{m_h}{m_h + m_t}$ , which approaches the true value of  $\theta$  with probability 1.

- Coin tossing example revisited:

- Suppose  $\alpha_h = \alpha_t = 100$ . Equivalent sample size: 200

- In case 1,

$$P(D_{m+1} = H|\mathcal{D}) = \frac{3 + 100}{10 + 100 + 100} \approx 0.5$$

Our prior prevails.

- In case 2,

$$P(D_{m+1} = H|\mathcal{D}) = \frac{30,000 + 100}{100,000 + 100 + 100} \approx 0.3$$

Data prevail.

# MLE vs Bayesian Estimation

- Much of Machine Learning is about parameter estimation.
- In all case, both MLE and Bayesian estimations can used, although the latter is harder mathematically.
- In this course, we will focus on MLE.