

# Advanced Deep Learning Architectures


## *COMP 5214 & ELEC 5680*

Instructor: Dr. Qifeng Chen  
<https://cqf.io>


# Logistics

- Assignment 1 is released today
- Pay attention to project proposal


Given an image of a ball,  
can you predict where it will go next?



Given an image of a ball,  
can you predict where it will go next?




Given an image of a ball,  
can you predict where it will go next?



Given an image of a ball,  
can you predict where it will go next?



# Sequences in the wild



Audio

# Sequences in the wild

character:

6.S191 Introduction to Deep Learning

word:

Text



# A Sequence Modeling Problem: Predict the Next Word

# A sequence modeling problem: predict the next word

“This morning I took my cat for a walk.”

# A sequence modeling problem: predict the next word

“This morning I took my cat for a walk.”

given these words

# A sequence modeling problem: predict the next word

“This morning I took my cat for a walk.”

given these words

predict the  
next word

# Idea #1: use a fixed window

“This morning I took my cat for a walk.”

given these	predict the
two words	next word

# Idea #1: use a fixed window

“This morning I took my cat for a walk.”

given these      predict the  
two words      next word

One-hot feature encoding: tells us what each word is

[ 1 0 0 0 0 0 1 0 0 0 ]

for


a



prediction

# Problem #1: can't model long-term dependencies

“France is where I grew up, but I now live in Boston. I speak fluent \_\_\_\_.”



We need information from the distant past to accurately predict the correct word.

## Idea #2: use entire sequence as set of counts

“This morning I took my cat for a”



“bag of words”

[ 0 1 0 0 1 0 0 ... 0 0 1 1 0 0 0 1 ]



prediction



## Problem #2: counts don't preserve order



The food was good, not bad at all.

vs.

The food was bad, not good at all.



# Idea #3: use a really big fixed window

“This morning I took my cat for a walk.”

given these  
words

predict the  
next word

[ 1 0 0 0 0 0 0 0 0 1 0 0 1 0 0 0 1 0 0 0 0 0 0 1 0 ... ]

morning I took this cat

prediction

# Problem #3: no parameter sharing

[ 1 0 0 0 0 0 0 0 0 1 0 0 1 0 0 0 1 0 0 0 0 0 0 1 0 ... ]  
this morning took the cat

Each of these inputs has a separate parameter:

# Problem #3: no parameter sharing

[ 1 0 0 0 0 0 0 0 0 1 0 0 1 0 0 0 1 0 0 0 0 0 1 0 ... ]  
this morning took the cat

Each of these inputs has a separate parameter:

[ 0 0 0 1 0 0 0 1 0 0 0 1 0 0 0 1 0 0 0 0 0 0 0 0 0 1 ... ]  
this morning

# Problem #3: no parameter sharing

[ 1 0 0 0 0 0 0 0 0 1 0 0 1 0 0 0 1 0 0 0 0 0 1 0 ... ]

this morning took the cat

Each of these inputs has a separate parameter:

[ 0 0 0 1 0 0 0 1 0 0 0 1 0 0 0 1 0 0 0 0 0 0 0 0 0 1 ... ]


this morning

Things we learn about the sequence won't transfer if they appear elsewhere in the sequence.

# Sequence modeling: design criteria

To model sequences, we need to:


1. Handle variable-length sequences
2. Track long-term dependencies
3. Maintain information about order
4. Share parameters across the sequence



Today: Recurrent Neural Networks (RNNs) as  
an approach to sequence modeling problems

# Recurrent Neural Networks (RNNs)


# Standard feed-forward neural network




One to One  
“Vanilla” neural network



# Recurrent neural networks: sequence modeling




One to One  
"Vanilla" neural network




Many to One  
*Sentiment Classification*


# Recurrent neural networks: sequence modeling



One to One  
“Vanilla” neural network



Many to One  
*Sentiment Classification*




Many to Many  
*Music Generation*




**6.SI91 Lab!**


# Recurrent neural networks: sequence modeling



One to One  
“Vanilla” neural network



Many to One  
*Sentiment Classification*




Many to Many  
*Music Generation*

... and many other  
architectures and  
applications




**6.SI91 Lab!**


# A standard “vanilla” neural network




# A recurrent neural network (RNN)



# A recurrent neural network (RNN)




# A recurrent neural network (RNN)



Apply a recurrence relation at every time step to process a sequence:

# A recurrent neural network (RNN)




Apply a recurrence relation at every time step to process a sequence:

$$\boxed{h_t} = \boxed{f_W}(\boxed{h_{t-1}}, \boxed{x_t})$$

cell state      function parameterized by  $W$       old state      input vector at time step  $t$



# A recurrent neural network (RNN)




Apply a recurrence relation at every time step to process a sequence:

$$\boxed{h_t} = \boxed{f_W}(\boxed{h_{t-1}}, \boxed{x_t})$$


cell state      function parameterized by  $W$       old state      input vector at time step  $t$

Note: the same function and set of parameters are used at every time step

# RNN state update and output




# RNN state update and output



InputVector

# RNN state update and output




Update Hidden State

$$h_t = \tanh(\mathbf{W}_{hh}h_{t-1} + \mathbf{W}_{xh}x_t)$$

InputVector

# RNN state update and output



Output Vector


$$\hat{y}_t = \mathbf{W}_{hy}h_t$$

Update Hidden State

$$h_t = \tanh(\mathbf{W}_{hh}h_{t-1} + \mathbf{W}_{xh}x_t)$$

Input Vector


# RNNs: computational graph across time




=

Represent as computational graph unrolled across time

# RNNs: computational graph across time




# RNNs: computational graph across time

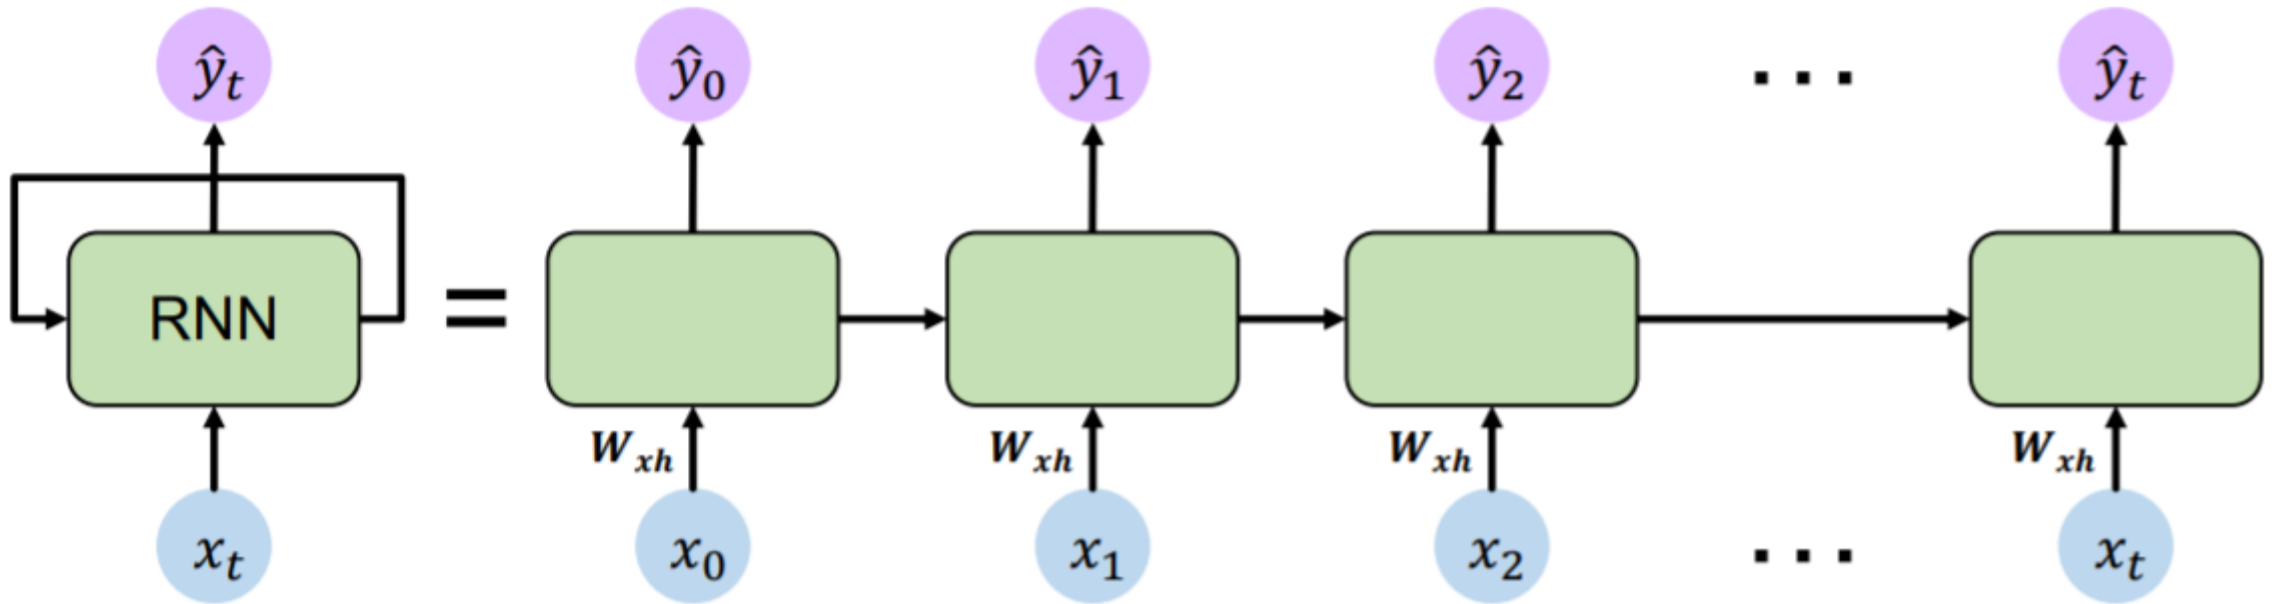




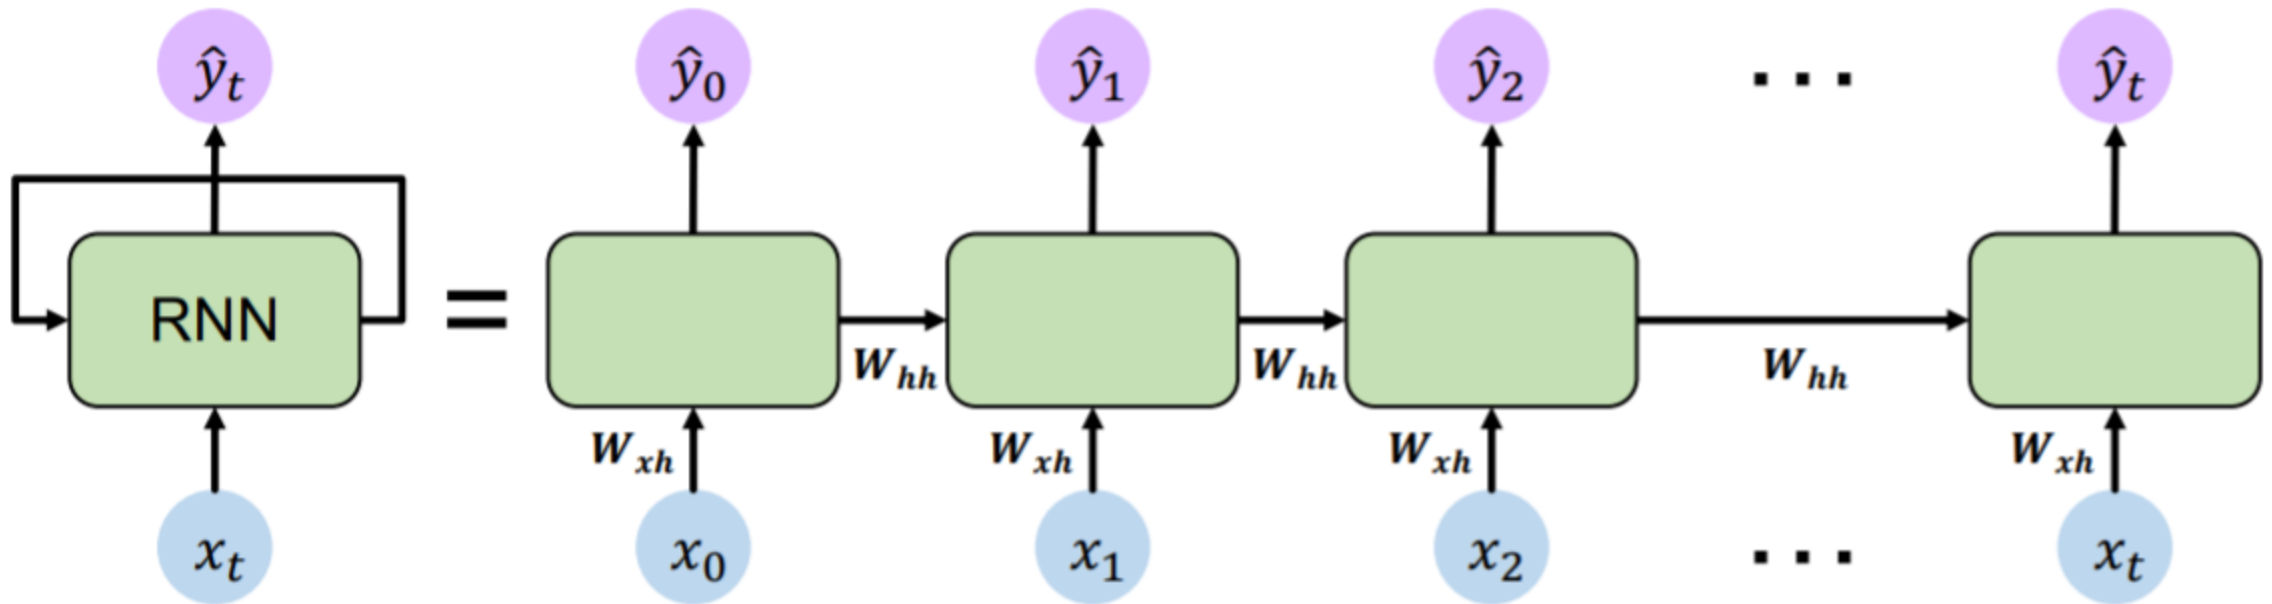
# RNNs: computational graph across time



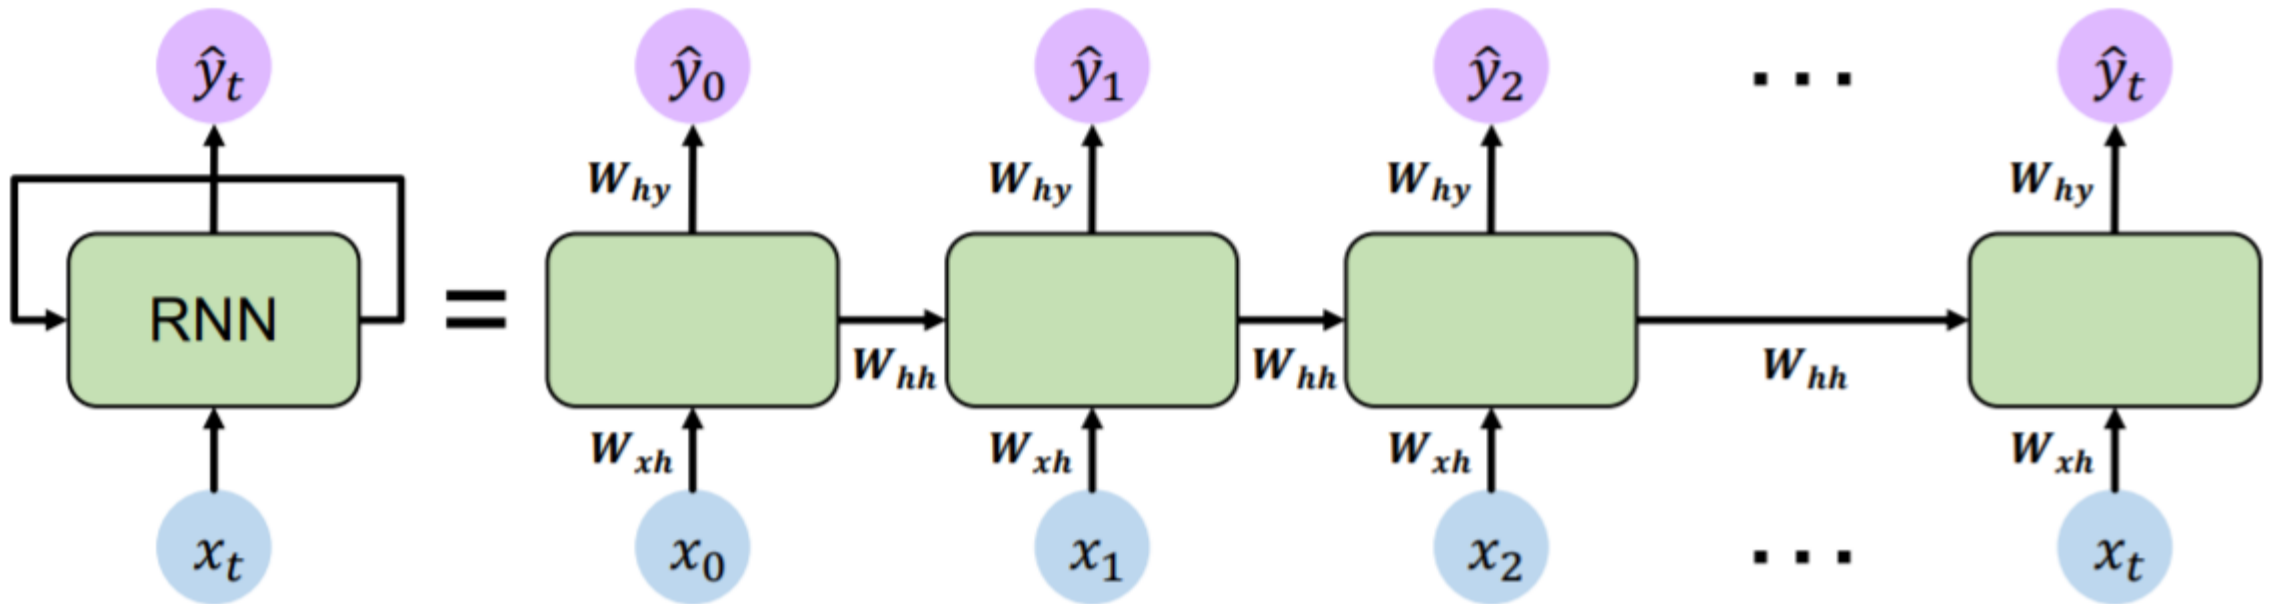
# RNNs: computational graph across time



# RNNs: computational graph across time

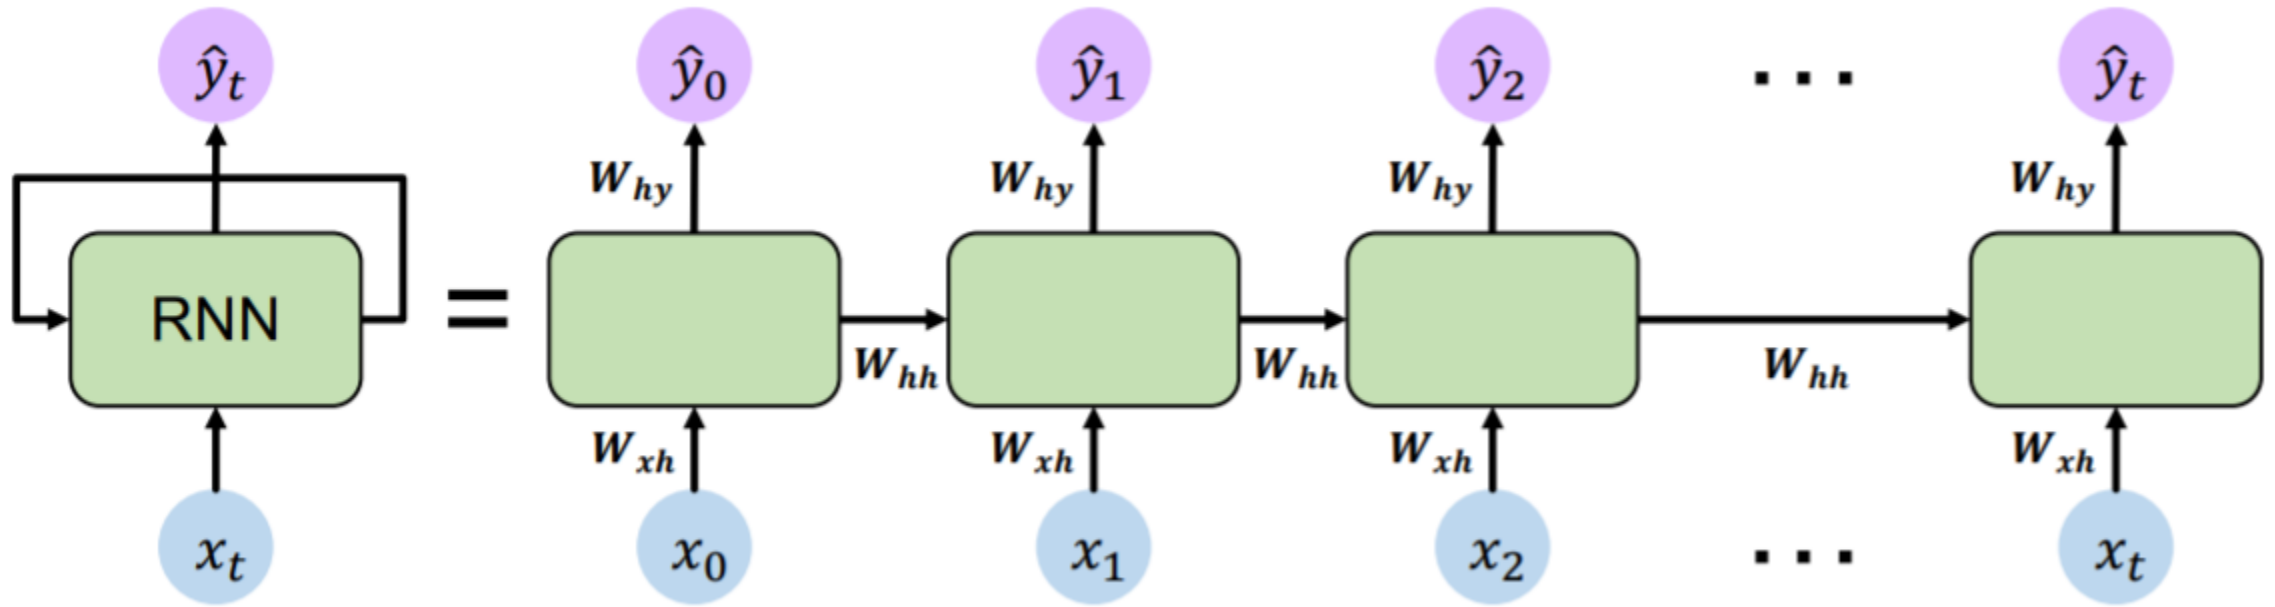


# RNNs: computational graph across time



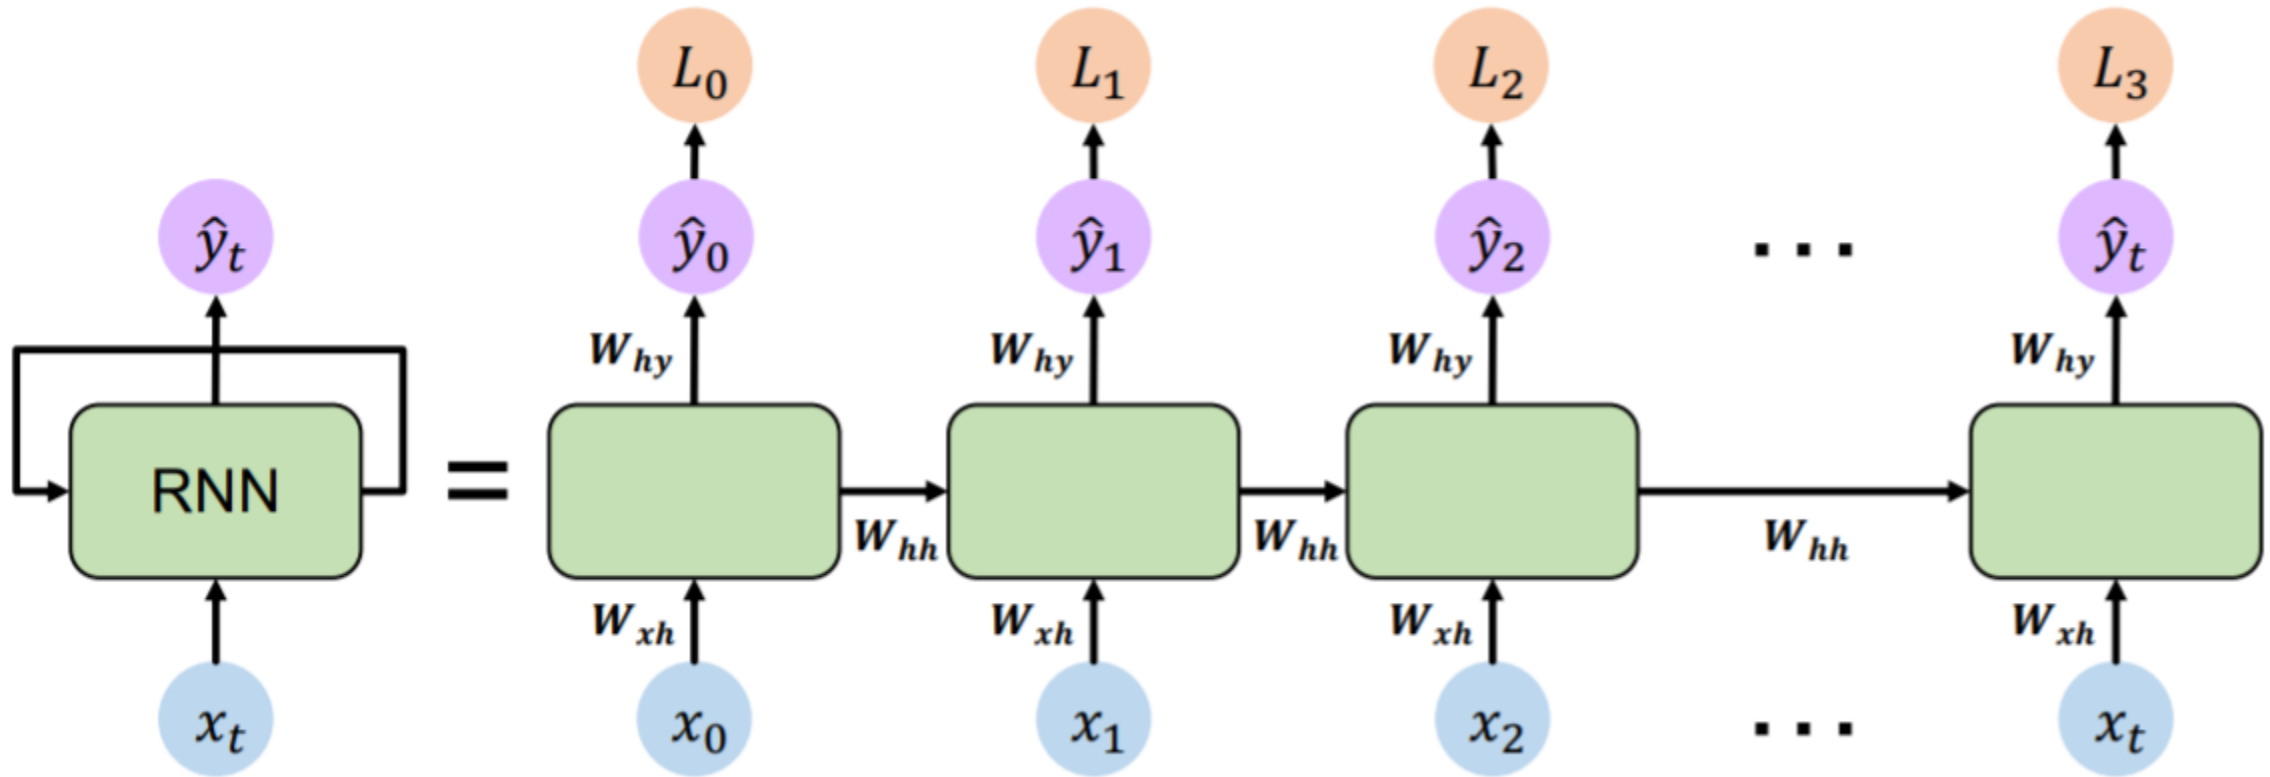
# RNNs: computational graph across time

Re-use the same weight matrices at every time step




# RNNs: computational graph across time

→ Forward pass




# RNNs: computational graph across time



# Backpropagation Through Time (BPTT)



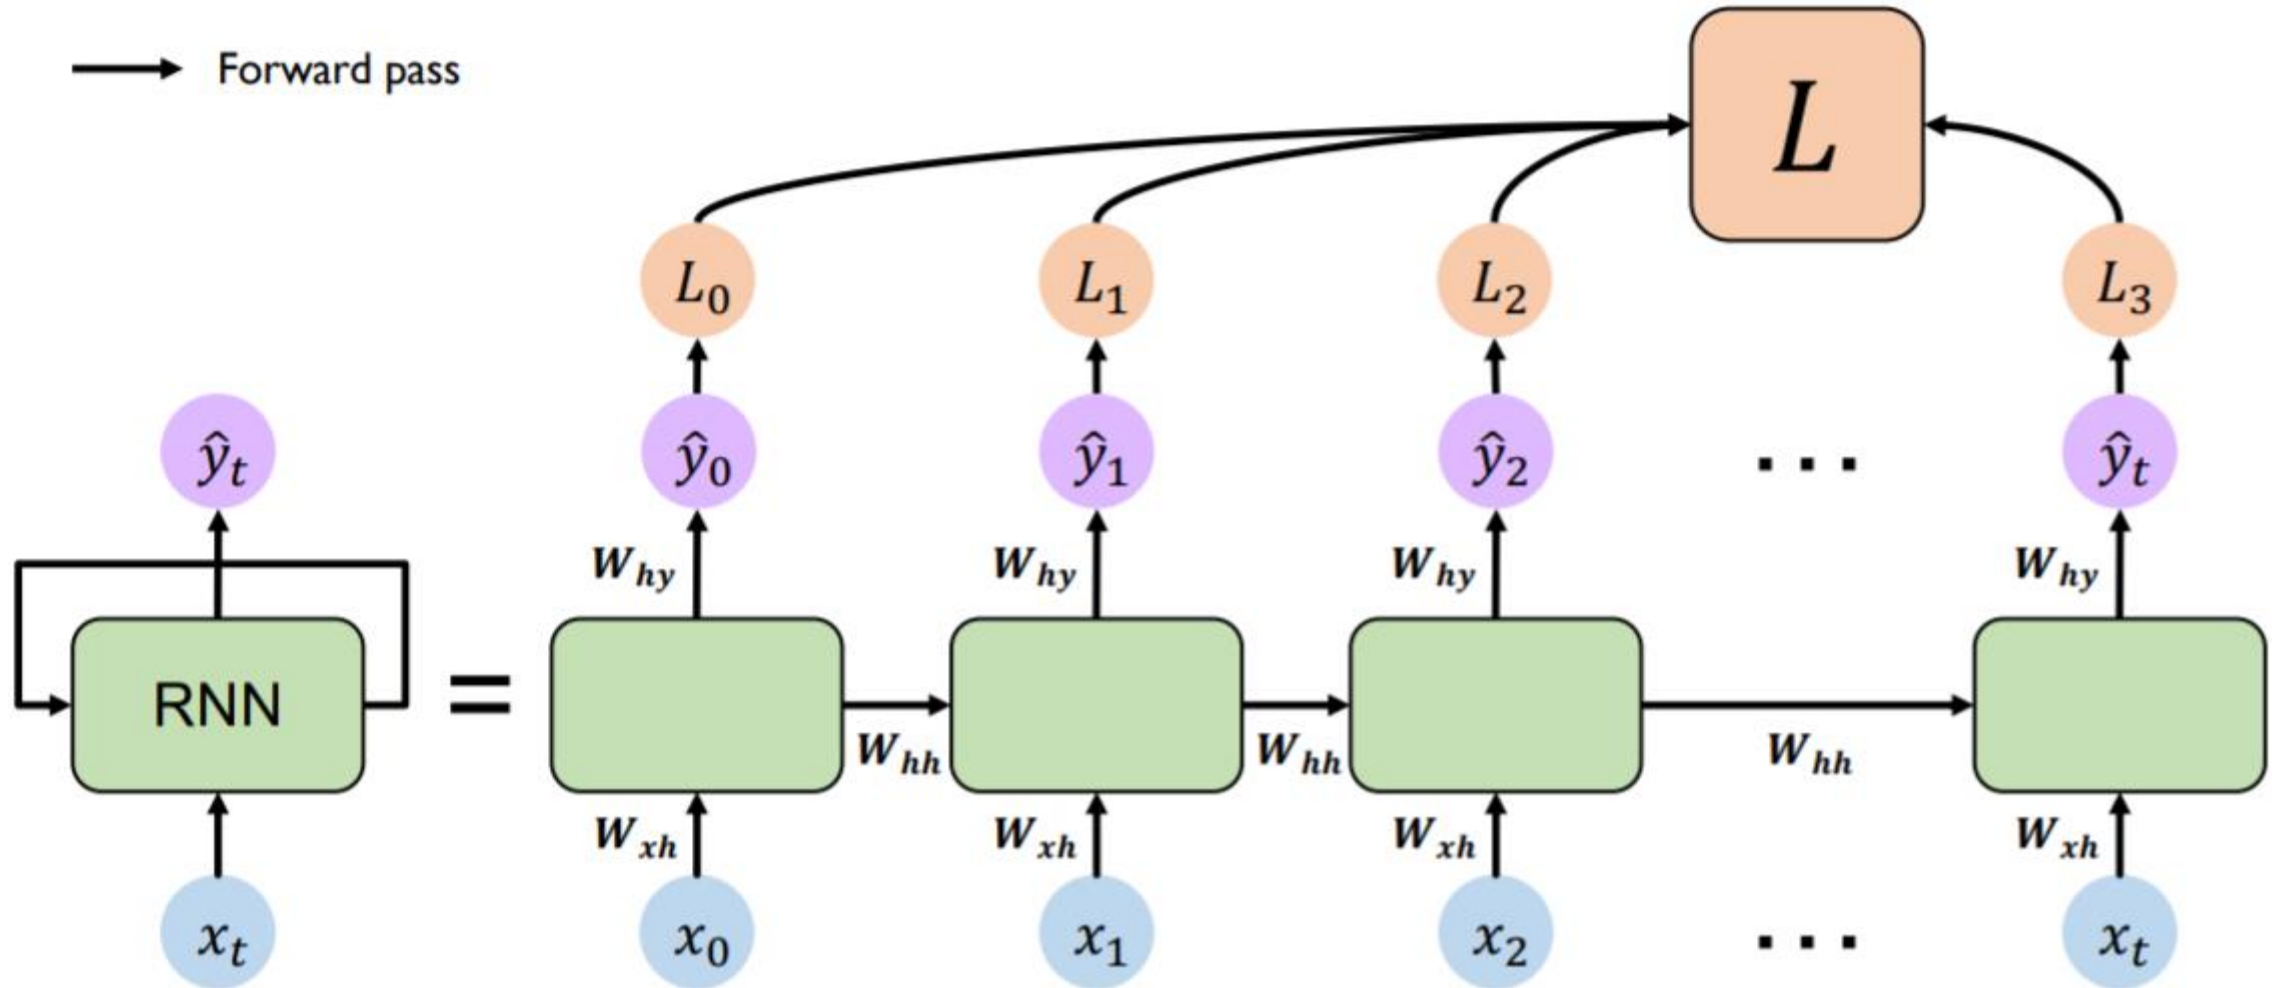
# Recall: backpropagation in feed forward models




Backpropagation algorithm:

1. Take the derivative (gradient) of the loss with respect to each parameter
2. Shift parameters in order to minimize loss

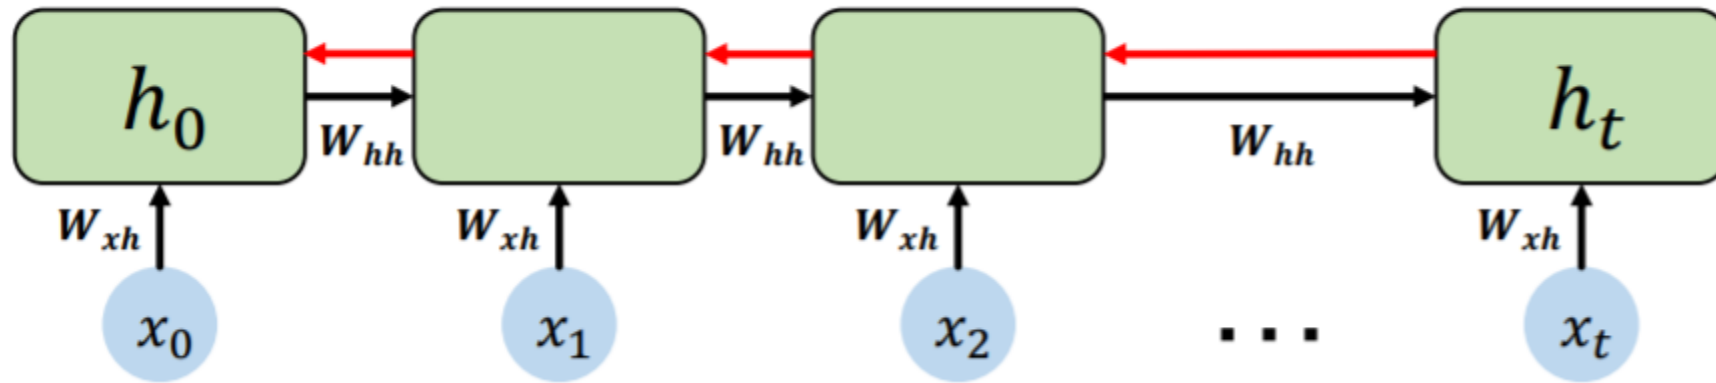
# RNNs: computational graph across time




# RNNs: backpropagation through time



# Standard RNN gradient flow




# Standard RNN gradient flow



Computing the gradient wrt  $h_0$  involves **many factors of  $W_{hh}$**  (and repeated  $f'$ !)


# Standard RNN gradient flow: exploding gradients



Computing the gradient wrt  $h_0$  involves **many factors of  $W_{hh}$**  (and repeated  $f'$ !)

Many values  $> 1$ :  
**exploding gradients**

# Standard RNN gradient flow: exploding gradients

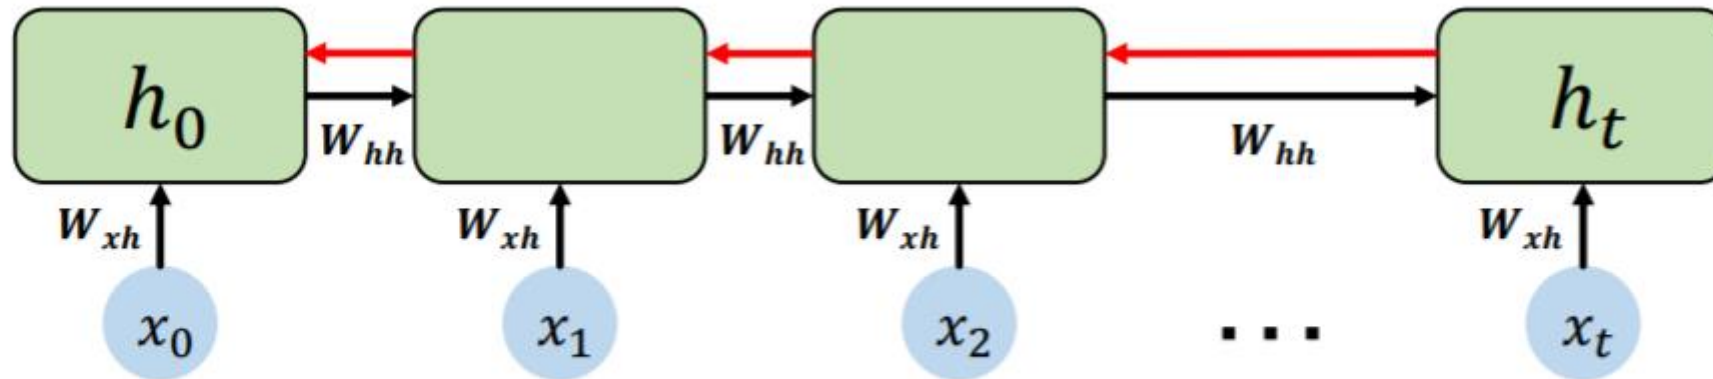


Computing the gradient wrt  $h_0$  involves **many factors of  $W_{hh}$**  (and repeated  $f'!$ )

Many values  $> 1$ :  
**exploding gradients**

**Gradient clipping** to  
scale big gradients

# Standard RNN gradient flow: vanishing gradients



Computing the gradient wrt  $h_0$  involves **many factors of  $W_{hh}$**  (and repeated  $f'$ !)


Many values  $> 1$ :  
exploding gradients

Gradient clipping to  
scale big gradients

Many values  $< 1$ :  
**vanishing gradients**



# Standard RNN gradient flow: vanishing gradients



Computing the gradient wrt  $h_0$  involves **many factors of  $W_{hh}$**  (and repeated  $f'$ !)

Largest singular value  $> 1$ :  
exploding gradients

Gradient clipping to  
scale big gradients

Largest singular value  $< 1$ :  
**vanishing gradients**

1. Activation function
2. Weight initialization
3. Network architecture

# The problem of long-term dependencies

Why are vanishing gradients a problem?

# The problem of long-term dependencies

Why are vanishing gradients a problem?

Multiply many small numbers together

# The problem of long-term dependencies

Why are vanishing gradients a problem?

Multiply many small numbers together



Errors due to further back time steps  
have smaller and smaller gradients

# The problem of long-term dependencies

Why are vanishing gradients a problem?

Multiply many small numbers together



Errors due to further back time steps  
have smaller and smaller gradients



Bias network to capture short-term  
dependencies

# The problem of long-term dependencies

“The clouds are in the \_\_\_\_”

Why are vanishing gradients a problem?

Multiply many small numbers together



Errors due to further back time steps  
have smaller and smaller gradients



Bias network to capture short-term  
dependencies

# The problem of long-term dependencies

Why are vanishing gradients a problem?

Multiply many small numbers together




Errors due to further back time steps  
have smaller and smaller gradients



Bias parameters to capture short-term  
dependencies

“The clouds are in the \_\_\_\_”



# The problem of long-term dependencies

Why are vanishing gradients a problem?

Multiply many small numbers together




Errors due to further back time steps  
have smaller and smaller gradients



Bias parameters to capture short-term  
dependencies

“The clouds are in the \_\_\_\_”



“I grew up in France,... and I I speak fluent\_\_\_\_”



# The problem of long-term dependencies

Why are vanishing gradients a problem?

Multiply many small numbers together




Errors due to further back time steps  
have smaller and smaller gradients




Bias parameters to capture short-term  
dependencies


“The clouds are in the \_\_\_\_”



“I grew up in France,... and I I speak fluent\_\_\_\_”



# Trick #1: activation functions



Using ReLU prevents  $f'$  from shrinking the gradients when  $x > 0$

# Trick #2: parameter initialization

Initialize weights to identity matrix


Initialize biases to zero

$$I_n = \begin{pmatrix} 1 & 0 & 0 & \dots & 0 \\ 0 & 1 & 0 & \dots & 0 \\ 0 & 0 & 1 & \dots & 0 \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & 0 & \dots & 1 \end{pmatrix}$$

This helps prevent the weights from shrinking to zero.


# Solution #3: gated cells

Idea: use a more complex recurrent unit with gates to control what information is passed through



# Solution #3: gated cells

Idea: use a more complex recurrent unit with gates to control what information is passed through




Long Short Term Memory (LSTMs) networks rely on a gated cell to track information throughout many time steps.

# Long Short Term Memory (LSTM) Networks


# Standard RNN

In a standard RNN, repeating modules contain a simple computation node



# Long Short Term Memory (LSTMs)

LSTM repeating modules contain interacting layers that control information flow




LSTM cells are able to track information throughout many timesteps




# Long Short Term Memory (LSTMs)

LSTMs maintain a cell state  $c_t$  where it's easy for information to flow



# Long Short Term Memory (LSTMs)

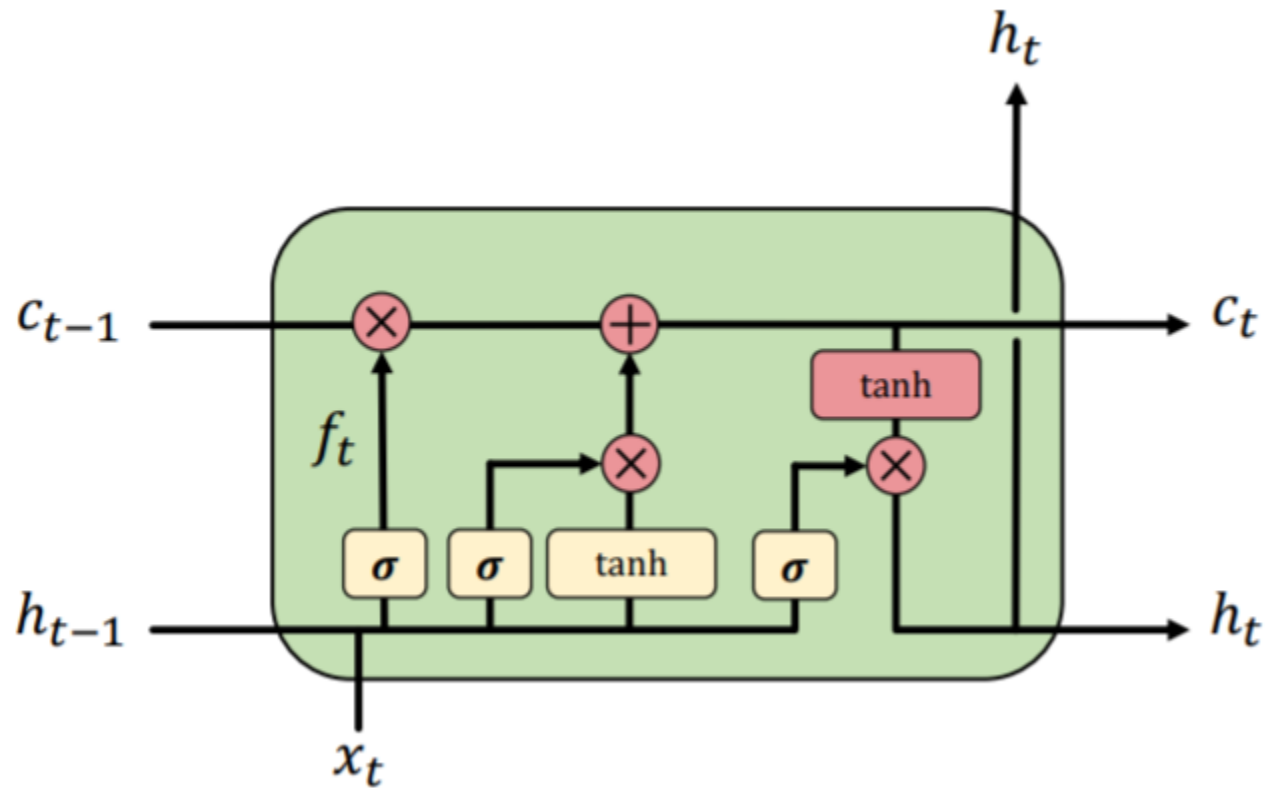
Information is added or removed to cell state through structures called gates



Gates optionally let information through, via a sigmoid neural net layer and pointwise multiplication


# Long Short Term Memory (LSTMs)

How do LSTMs work?




# Long Short Term Memory (LSTMs)

LSTMs forget irrelevant parts of the previous state




# Long Short Term Memory (LSTMs)

LSTMs selectively update cell state values



# Long Short Term Memory (LSTMs)


LSTMs use an output gate to output certain parts of the cell state




# Long Short Term Memory (LSTMs)

How do LSTMs work?

1) Forget 2) Update 3) Output



# LSTMs: forget irrelevant information




$$f_t = \sigma(\mathbf{W}_i[ h_{t-1}, x_t ] + b_f)$$

- Use previous cell output and input
- Sigmoid: value 0 and 1 – “completely forget” vs. “completely keep”

*ex: Forget the gender pronoun of previous subject in sentence.*



# LSTMs: identify new information to be stored




$$i_t = \sigma(\mathbf{W}_i[ h_{t-1}, x_t ] + b_i)$$
$$\tilde{C}_t = \tanh(\mathbf{W}_c[ h_{t-1}, x_t ] + b_c)$$

- Sigmoid layer: decide what values to update
- Tanh layer: generate new vector of "candidate values" that could be added to the state

*ex: Add gender of new subject to replace that of old subject.*

# LSTMs: update cell state




$$C_t = f_t * C_{t-1} + i_t * \tilde{C}_t$$

- Apply forget operation to previous internal cell state:  $f_t * C_{t-1}$
- Add new candidate values, scaled by how much we decided to update:  $i_t * \tilde{C}_t$

*ex: Actually drop old information and add new information about subject's gender.*

# LSTMs: output filtered version of cell state



$$o_t = \sigma(W_o [h_{t-1}, x_t] + b_o)$$


$$h_t = o_t * \tanh(C_t)$$

- Sigmoid layer: decide what parts of state to output
- Tanh layer: squash values between -1 and 1
- $o_t * \tanh(C_t)$ : output filtered version of cell state

*ex: Having seen a subject, may output information relating to a verb.*

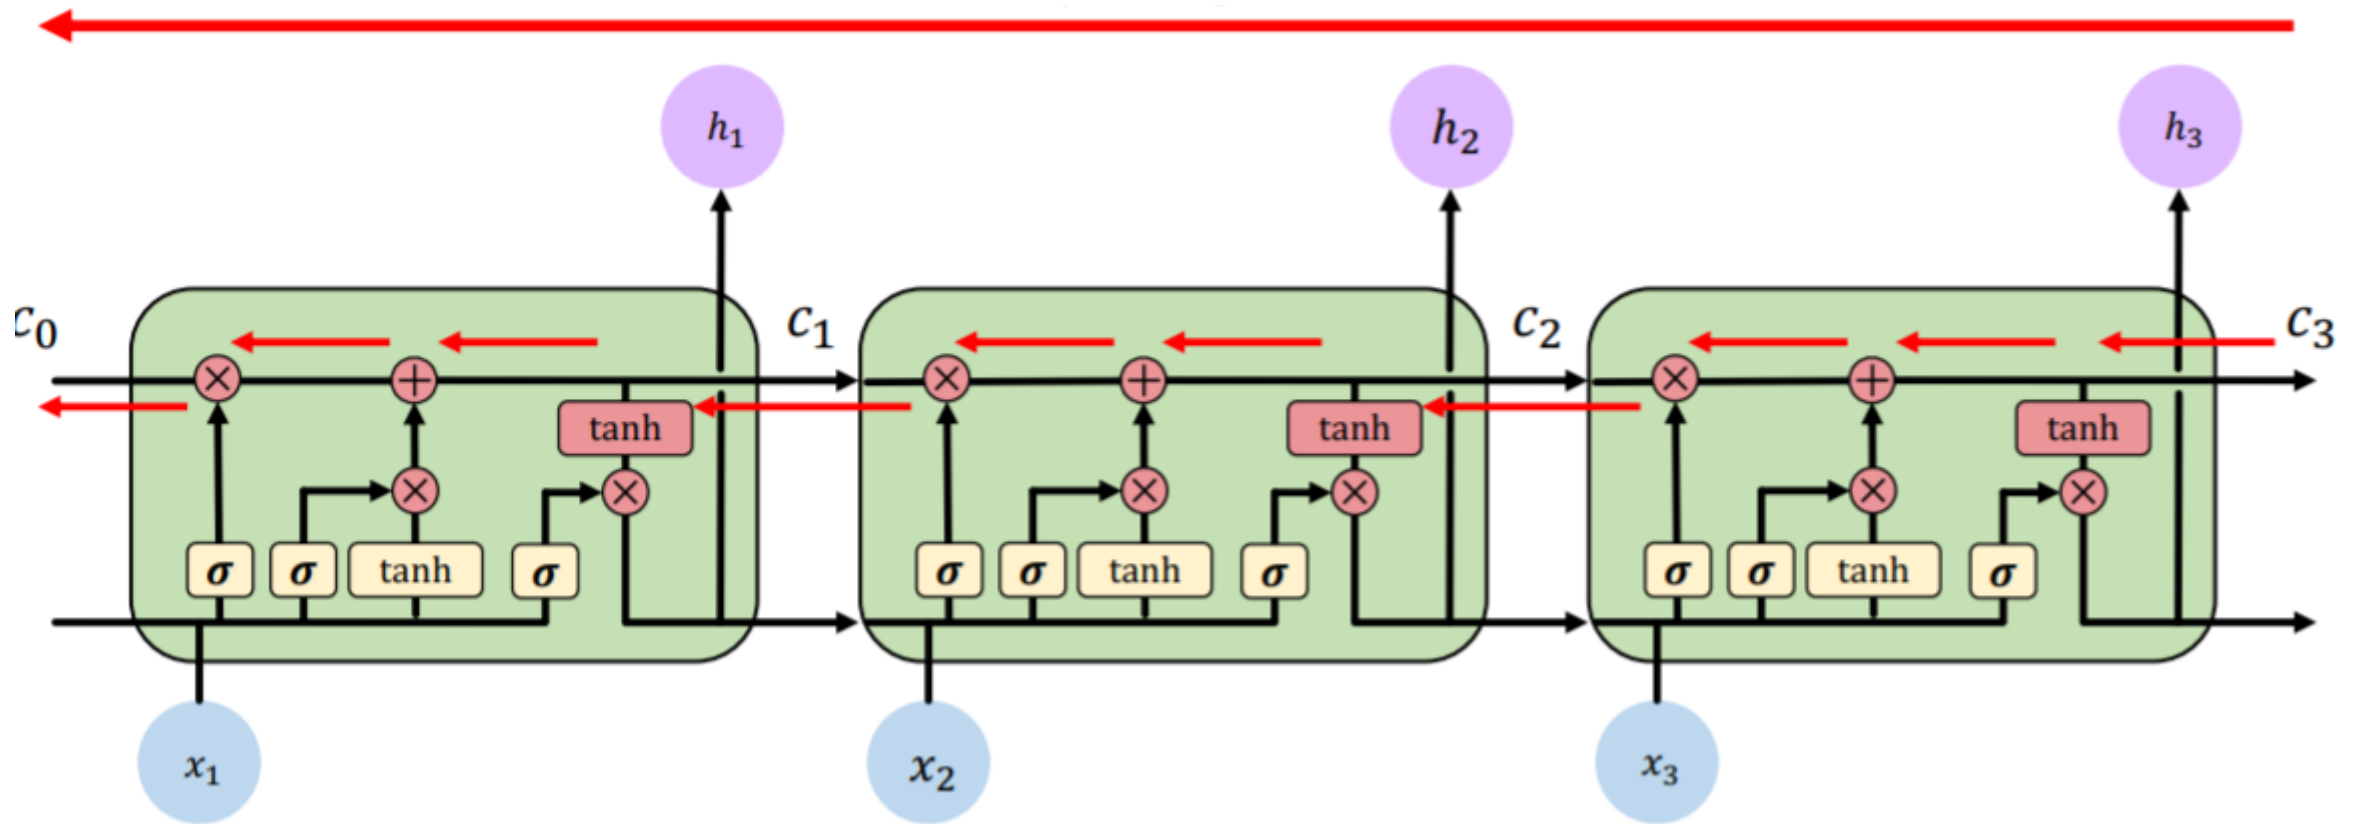
# LSTM gradient flow

Backpropagation from  $C_t$  to  $C_{t-1}$  requires only elementwise multiplication!  
No matrix multiplication  $\rightarrow$  avoid vanishing gradient problem.



# LSTM gradient flow

Uninterrupted gradient flow!




# LSTMs: key concepts

1. Maintain a separate cell state from what is outputted
2. Use gates to control the flow of information
  - Forget gate gets rid of irrelevant information
  - Selectively update cell state
  - Output gate returns a filtered version of the cell state
3. Backpropagation from  $C_t$  to  $C_{t-1}$  doesn't require matrix multiplication: uninterrupted gradient flow


# RNN Applications

# Example task: music generation




**Input:** sheet music

**Output:** next character in sheet music






# Example task: sentiment classification




**Input:** sequence of words

**Output:** probability of having positive sentiment

```
 loss = tf.nn.softmax_cross_entropy_with_logits(  
    labels=model.y, logits=model.pred  
)
```

# Example task: sentiment classification



## Tweet sentiment classification



Ivar Hagendoorn  
@IvarHagendoorn

Follow



The @MIT Introduction to #DeepLearning is definitely one of the best courses of its kind currently available online [introtodeeplearning.com](http://introtodeeplearning.com)

12:45 PM - 12 Feb 2018



Angels-Cave  
@AngelsCave

Follow

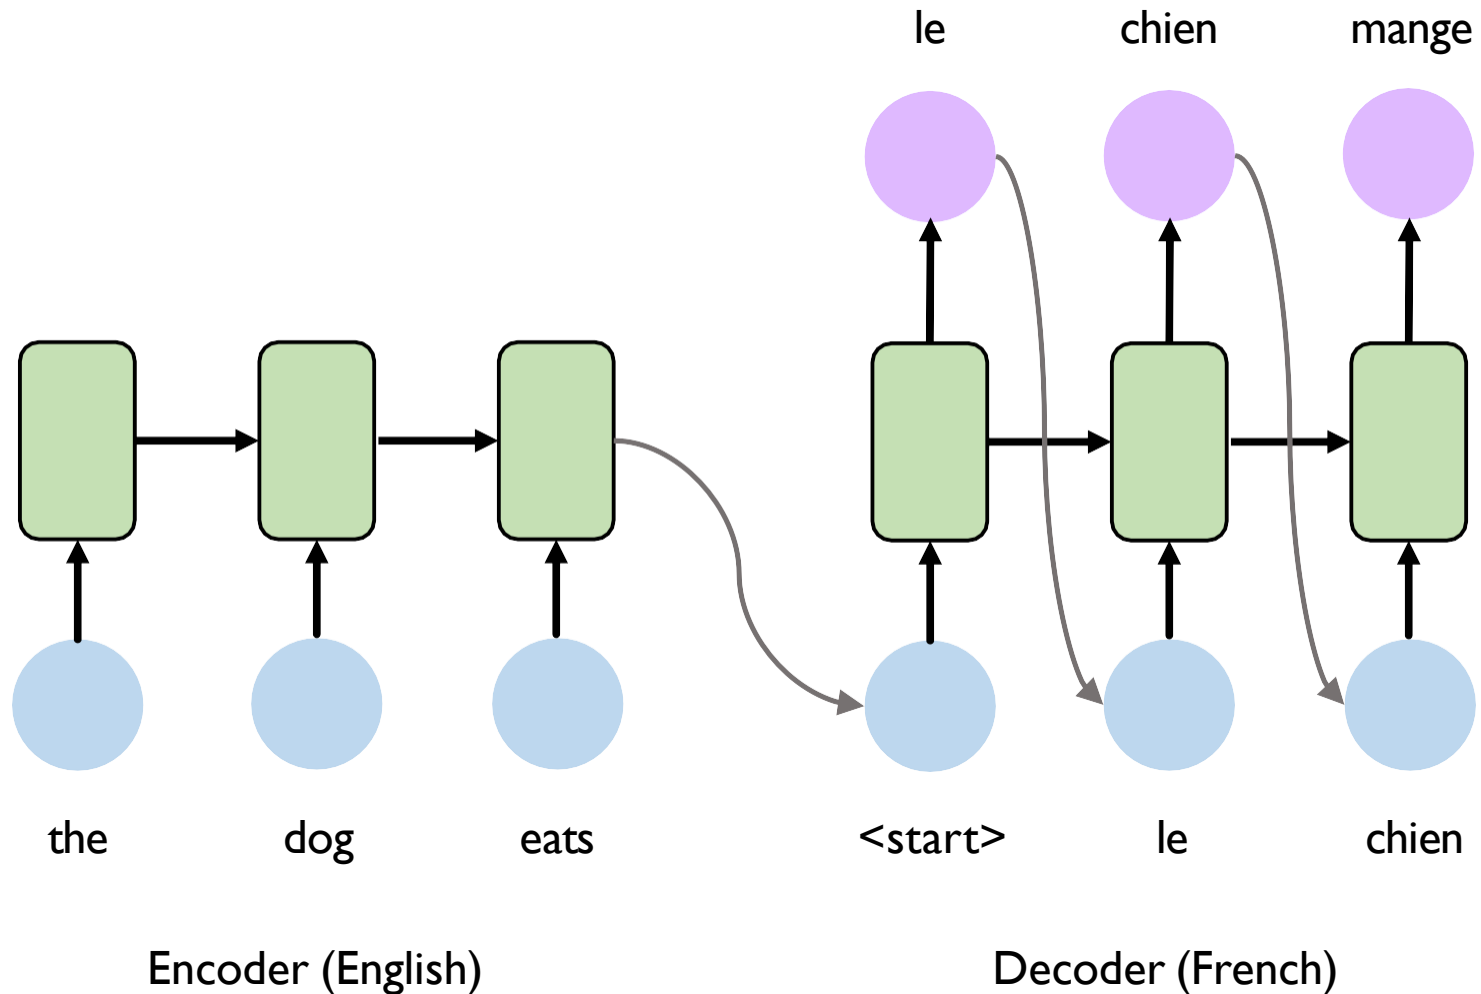


Replying to @Kazuki2048

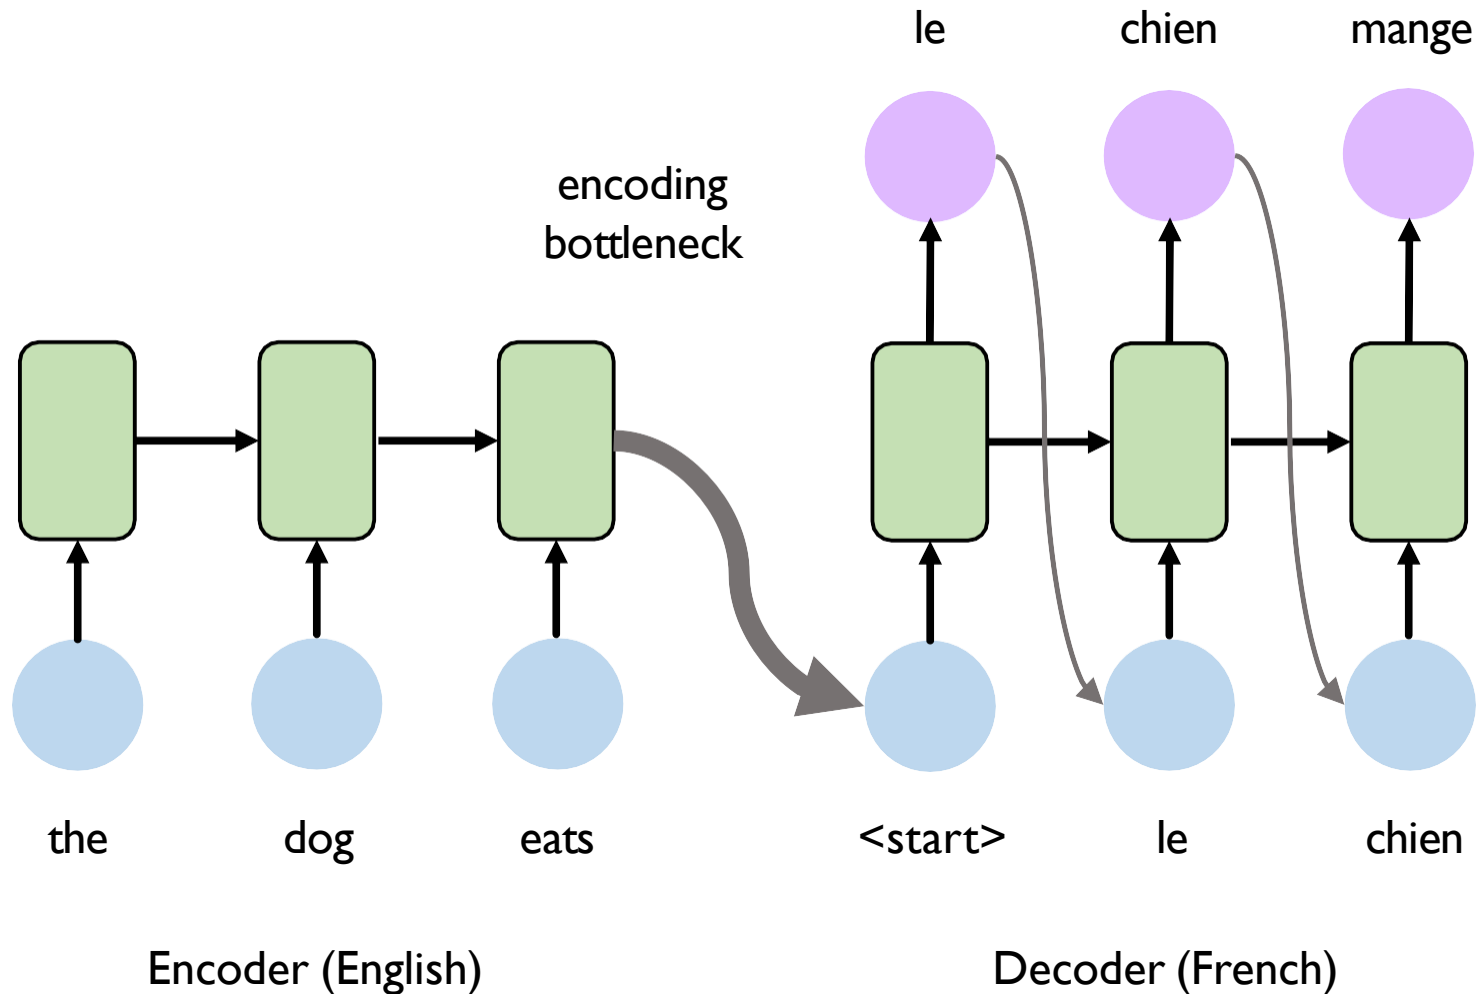
I wouldn't mind a bit of snow right now. We haven't had any in my bit of the Midlands this winter! :(

2:19 AM - 25 Jan 2019

# Example task: machine translation

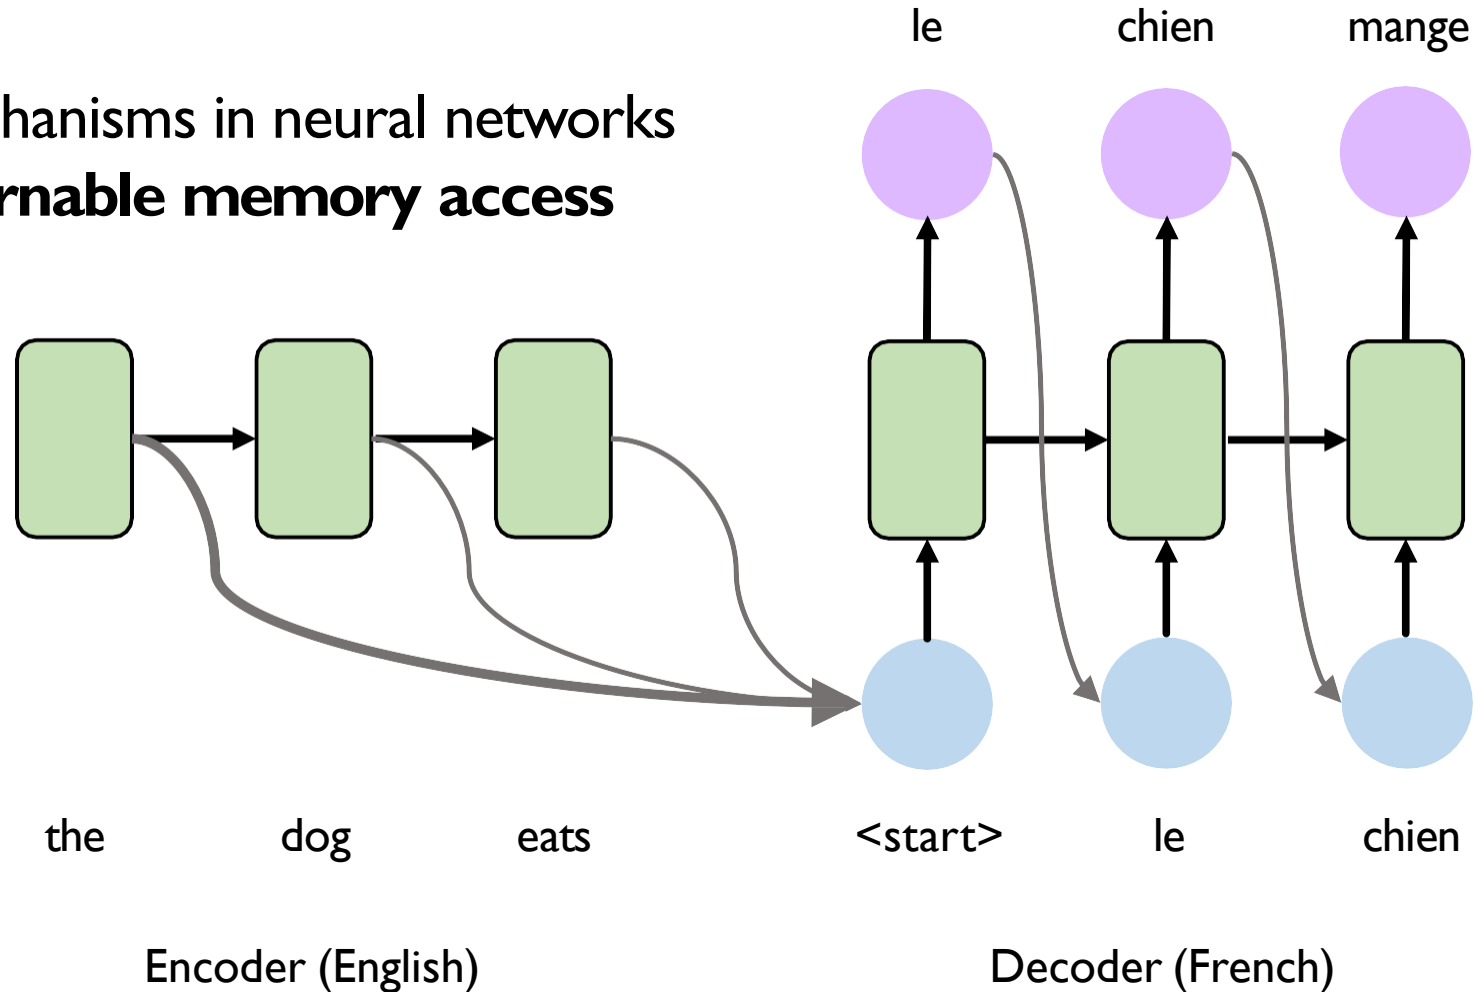


# Example task: machine translation



# Attention mechanisms

Attention mechanisms in neural networks provide **learnable memory access**



# Gated Recurrent Unit (GRU)



# Paying attention to a sequence

- Not all observations are equally relevant

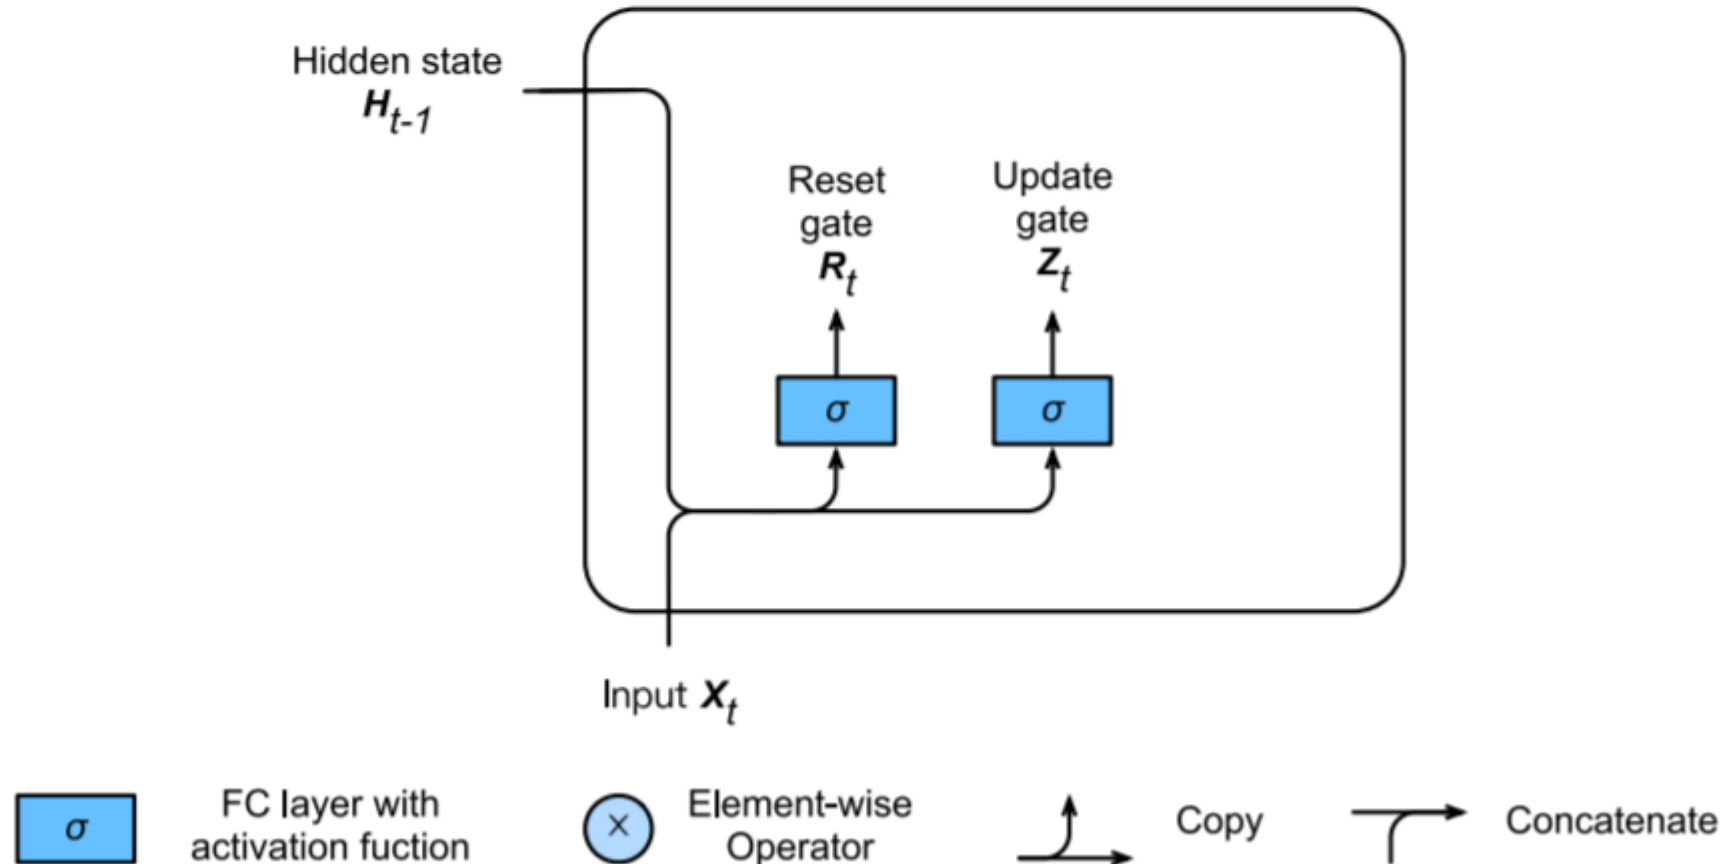


- Only remember the relevant ones
  - Need mechanism to **pay attention (update gate)**
  - Need mechanism to **forget (reset gate)**

# Gating

$$R_t = \sigma(X_t W_{xr} + H_{t-1} W_{hr} + b_r),$$

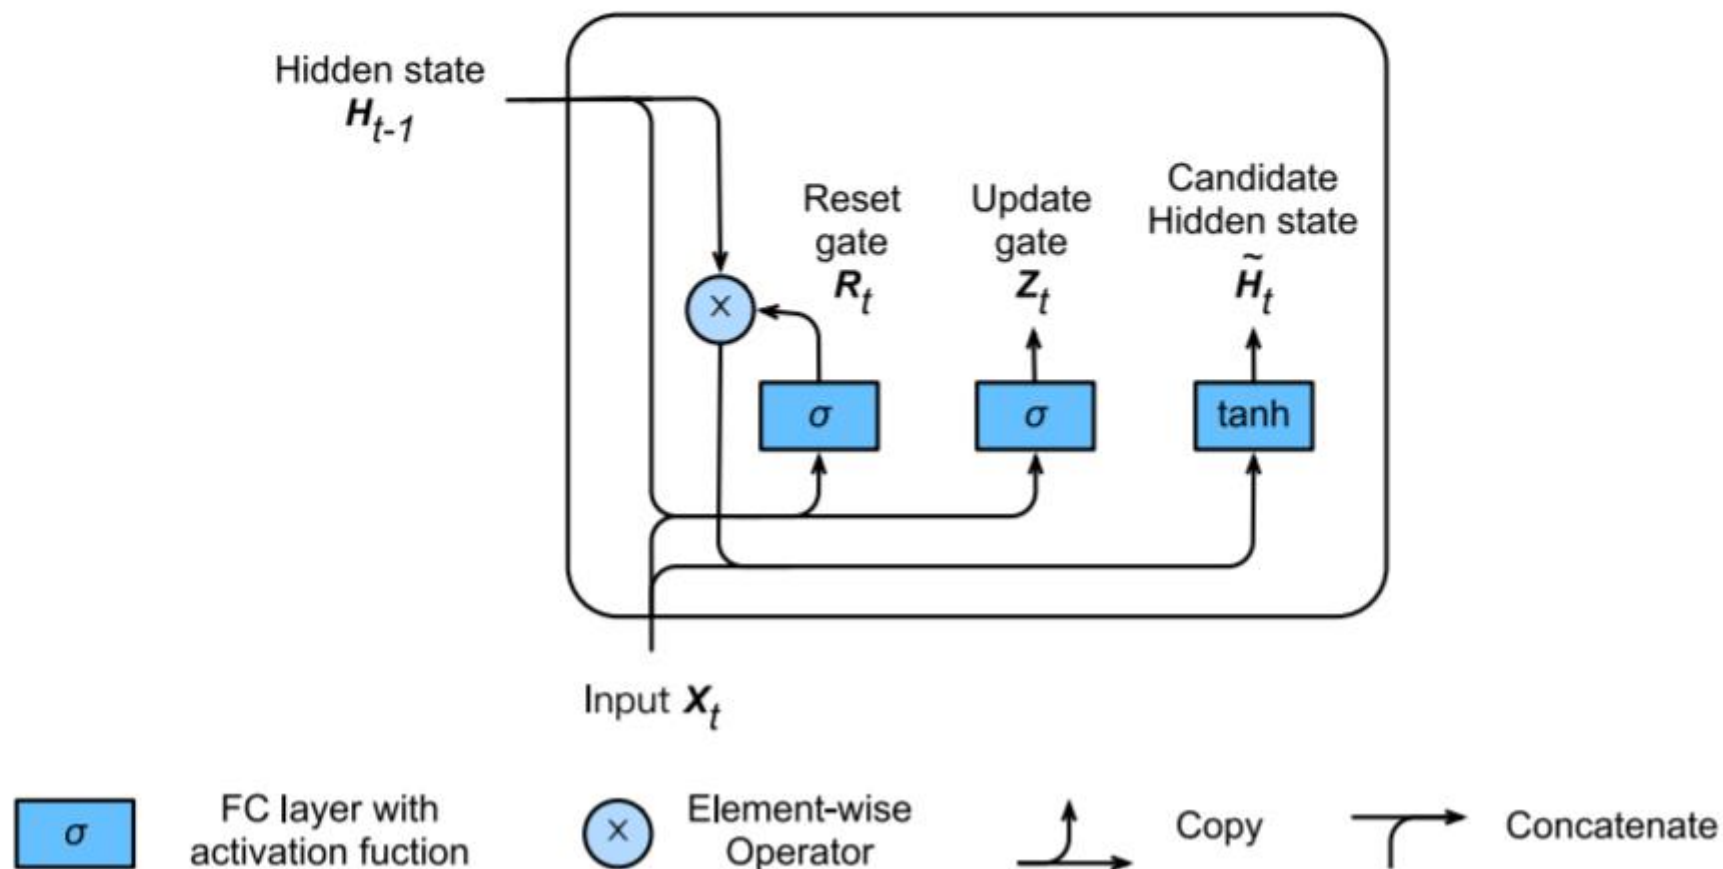
$$Z_t = \sigma(X_t W_{xz} + H_{t-1} W_{hz} + b_z)$$





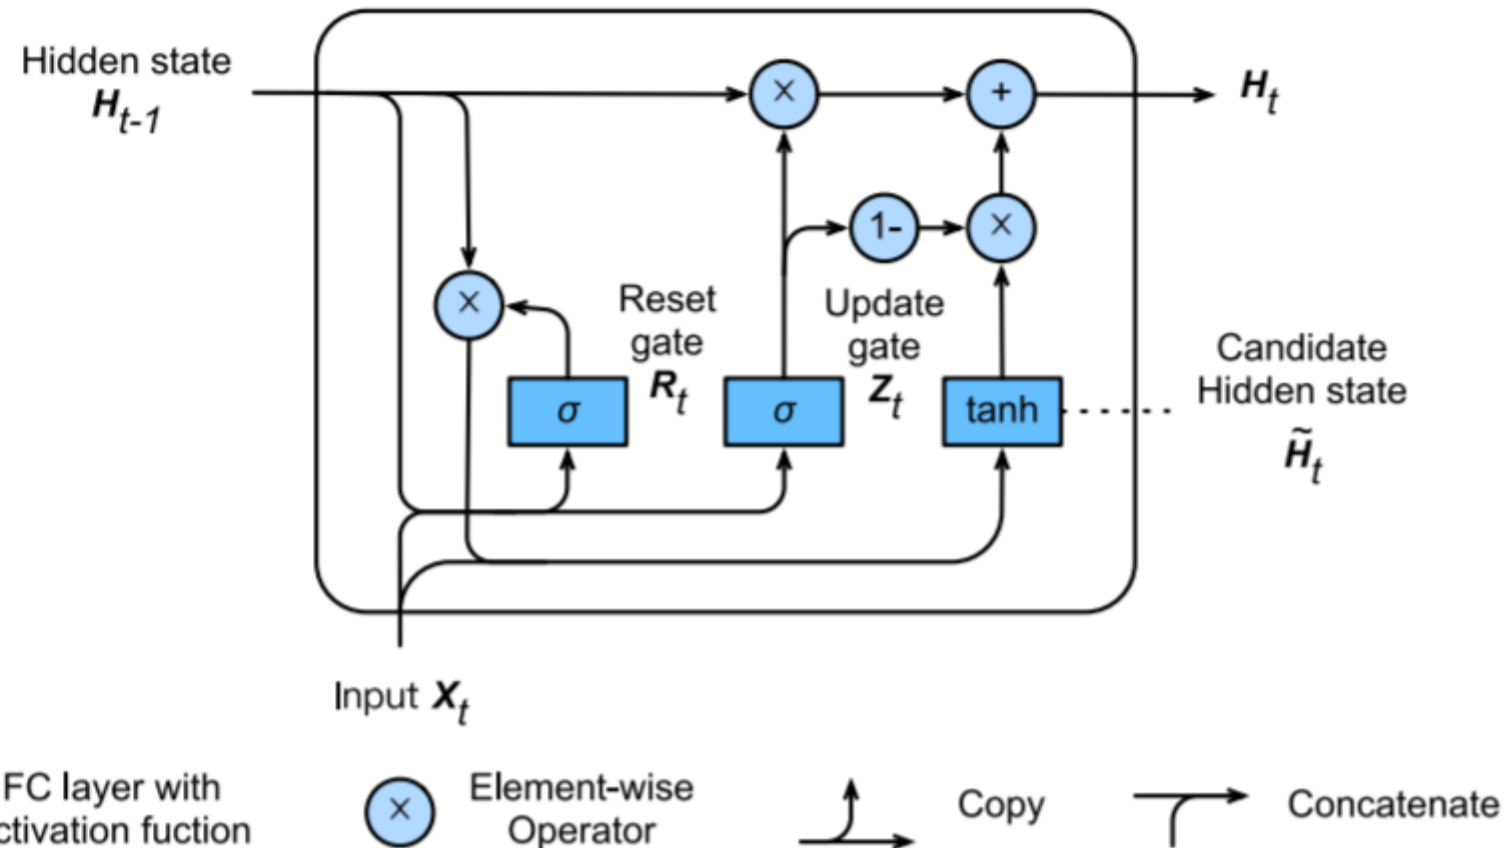
# Candidate Hidden State

$$\tilde{H}_t = \tanh(X_t W_{xh} + (R_t \odot H_{t-1}) W_{hh} + b_h)$$



# Hidden State

$$H_t = Z_t \odot H_{t-1} + (1 - Z_t) \odot \tilde{H}_t$$



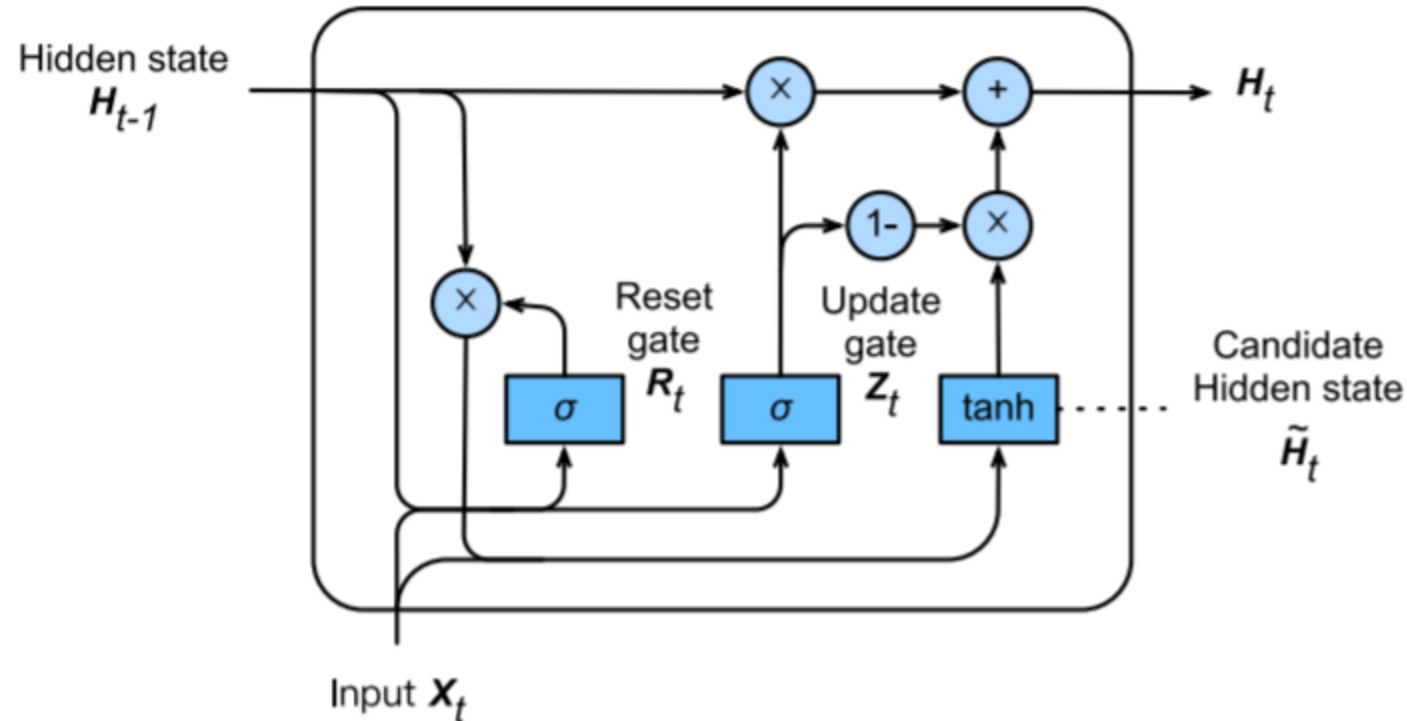
# Summary

$$R_t = \sigma(X_t W_{xr} + H_{t-1} W_{hr} + b_r),$$


$$Z_t = \sigma(X_t W_{xz} + H_{t-1} W_{hz} + b_z)$$

$$\tilde{H}_t = \tanh(X_t W_{xh} + (R_t \odot H_{t-1}) W_{hh} + b_h)$$


$$H_t = Z_t \odot H_{t-1} + (1 - Z_t) \odot \tilde{H}_t$$



# LSTM vs. GRU



LSTM  
(Long-Short Term Memory)



GRU  
(Gated Recurrent Unit)

The GRU is like a long short-term memory (LSTM) with a forget gate, but has fewer parameters than LSTM, as it lacks an output gate.

# Recurrent neural networks (RNNs)

1. RNNs are well suited for sequence modeling tasks
2. Model sequences via a recurrence relation
3. Training RNNs with backpropagation through time
4. Gated cells like LSTMs & GRUs let us model long-term dependencies
5. Models for music generation, classification, machine translation

