


# Spring 2022 COMP 3511

## Review #5

---





# Coverages

---

- Semaphore (mutex part will be covered in the POSIX mutex example)
- POSIX Synchronization
- Questions related to synchronization





# Semaphore

- Semaphore **S** – non-negative integer variable, can be considered as a generalized lock
  - First defined by Dijkstra in late 1960s. It can behave similarly as mutex lock, but more sophisticated in its usage - the main synchronization primitive used in original UNIX
- Two standard operations modify **S**: **wait()** and **signal()**
  - Originally called **P()** and **V()**, where **P()** stands for “**proberen**” (to test) and **V()** stands for “**verhogen**” (to increment) in Dutch
- It is critical that semaphore operations are executed atomically, which guarantees that no more than one process can execute **wait()** and **signal()** operations on the same semaphore at the same time.
- The semaphore can only be accessed via these two atomic operations except initialization

```
wait (S) {  
    while (S <= 0)  
        ; // busy wait  
    S--;  
}  
  
signal (S) {  
    S++;  
}
```





# Semaphore Usage

- **Counting semaphore** – An integer value can range over an unrestricted domain
  - Counting semaphore can be used to control access to a given resource consisting of a finite number of instances; semaphore value is initialized to the number of resource available
- **Binary semaphore** – integer value can range only between 0 and 1
  - This can behave similar to [mutex locks](#), can also be used in different ways
- This can also be used to solve various synchronization problems
- Consider  $P_1$  and  $P_2$  that shares a common semaphore **synch**, initialized to 0; it ensures that  $P_1$  process executes  $S_1$  before  $P_2$  process executes  $S_2$

P1:

$S_1$ ;

signal(synch);

P2:

wait(synch);

$S_2$ ;





# Semaphore Implementation with no Busy waiting

- With each semaphore there is an associated **waiting queue**
- Each entry in a waiting queue has two data items:
  - value (of type integer)
  - pointer to next record on the queue
- Two operations:
  - **block** – place the process invoking the operation on the appropriate waiting queue
  - **wakeup** – remove one of processes in the waiting queue and place it in the ready queue
- Semaphore values may become negative, whereas this value can never be negative under the classical definition of semaphores with busy waiting.
- If a semaphore value is negative, its magnitude is the number of processes currently waiting on the semaphore.





# Semaphore Implementation with no Busy waiting (Cont.)

```
typedef struct{
    int value;
    struct process *list;
} semaphore;
wait(semaphore *S) {
    S->value--;
    if (S->value < 0) {
        add this process to S->list;
        block();
    }
}
signal(semaphore *S) {
    S->value++;
    if (S->value <= 0) {
        remove a process P from S->list;
        wakeup(P);
    }
}
```

Noticing that

- Increment and decrement are done before checking the semaphore value, unlike the busy waiting implementation
- The **block()** operation suspends the process that invokes it.
- The **wakeup(P)** operation resumes the execution of a suspended process P.





# A Sample Synchronization Question

---

- Is binary semaphore equivalent as a mutex lock (Yes/No)? Briefly explain your answer:





# Answer

---

- Is binary semaphore equivalent as a mutex lock (Yes/No)? Briefly explain your answer:
- **Answer:** No (1 mark). (Explanation: 1 mark) Binary semaphore initialized to 1 can be used as a mutex lock, but it can be used for other purposes (e.g., when initialized to 0).







# A Sample Synchronization Question

---

- Suppose that there are three processes  $P_0$ ,  $P_1$  and  $P_2$  that need to access a critical section in turn strictly following the order:

$P_0, P_1, P_2, P_0, P_1, P_2, P_0, P_1, P_2, \dots$

- In other words, we want to execute  $P_0$ 's critical section, and then  $P_1$ 's critical section, and then  $P_2$ 's critical section, and the pattern repeats
- In the following slides, you will see a program with missing BLANKS





# A Sample Synchronization Question

- Goal: access a critical section in turn strictly following the order: P0, P1, P2, P0, P1, P2...

```
// Initialize          // Process P0          // Process P1
// S0, S1, S2 are      while(true) {          while(true){
// shared               wait(S0);              BLANK5;
                       // critical section    // critical section
S0 = BLANK1;          BLANK4;              BLANK6;
S1 = BLANK2;          }                      }
S2 = BLANK3;

                       // Process P2
                       while(true) {
                           BLANK7;
                           // critical section
                           signal(S0);
                       }
```





# A Sample Synchronization Answers

- Goal: access a critical section in turn strictly following the order: P0, P1, P2, P0, P1, P2...

```
// Initialize
// S0, S1, S2 are
// shared
```

```
S0 = 1;
S1 = 0;
S2 = 0;
```

```
// Process P0
while(true) {
    wait(S0);
    // critical section
    signal(S1);
}
```

```
// Process P2
while(true) {
    wait(S2);
    // critical section
    signal(S0);
}
```

```
// Process P1
while(true){
    wait(S1);
    // critical section
    signal(S2);
}
```





# Deadlock and Starvation

- **Deadlock** – two or more processes are waiting indefinitely for an event that can be caused by only one of the waiting processes (details in Chapter 7)
- Let  $S$  and  $Q$  be two semaphores initialized to 1

$P_0$	$P_1$
<code>wait(S) ;</code>	<code>wait(Q) ;</code>
<code>wait(Q) ;</code>	<code>wait(S) ;</code>
...	...
<code>signal(S) ;</code>	<code>signal(Q) ;</code>
<code>signal(Q) ;</code>	<code>signal(S) ;</code>

- Consider if  $P_0$  executes `wait(S)` and  $P_1$  `wait(Q)`. When  $P_0$  executes `wait(Q)`, it must wait until  $P_1$  executes `signal(Q)`. However,  $P_1$  is waiting until  $P_0$  execute `signal(S)`. Since these `signal()` operations will never be executed,  $P_0$  and  $P_1$  are **deadlocked**, extremely difficult to debug
- **Starvation** – **indefinite blocking**
  - A process may never be removed from the semaphore queue, in which it is suspended. For instance, if we remove processes from the queue associated with a semaphore using LIFO (last-in, first-out) order or based on certain priorities.





# POSIX Synchronization

- POSIX API provides
  - Mutex locks (Note: pthread library supports mutex locks)
  - Semaphores (Note: pthread library don't have a direct support on semaphores)
  - condition variables (Note: `pthread_cond_wait` and `pthread_cond_signal` are used to provide a waiting based on a condition)
- Widely used on UNIX, Linux, and MacOS





# POSIX Mutex Locks

## ■ Creating and initializing the lock

```
#include <pthread.h>

pthread_mutex_t mutex;

/* create and initialize the mutex lock */
pthread_mutex_init(&mutex, NULL);
```

## ■ Acquiring and releasing the lock

```
/* acquire the mutex lock */
pthread_mutex_lock(&mutex);

/* critical section */

/* release the mutex lock */
pthread_mutex_unlock(&mutex);
```





# Sample Mutex in Pthread (Question)

```
#include <stdio.h>
#include <stdlib.h>
#include <pthread.h>
#include <unistd.h>
unsigned int counter = 0;

// This constant controls
// the number of iteration
#define TOTAL_ITERATIONS 5000

void *add(void *arg) {
    int i;
    pthread_mutex_t* m = BLANK1;
    for (i=0; i<TOTAL_ITERATIONS; i++) {
        BLANK2;
        ++counter;
        BLANK3;
    }
}
```

```
int main() {
    pthread_t tid1, tid2;
    pthread_mutex_t mutex;
    pthread_mutex_init(&mutex, NULL);
    printf("The original counter value is %d\n", counter);
    pthread_create(&tid1, NULL, add, &mutex);
    pthread_create(&tid2, NULL, add, &mutex);
    pthread_join(tid1, NULL);
    pthread_join(tid2, NULL);
    printf("The final counter value is %d\n", counter);
    return 0;
}
```

What are the missing BLANKs?  
What is the purpose of this program?





# Sample Mutex in Pthread (Solution)

```
#include <stdio.h>
#include <stdlib.h>
#include <pthread.h>
#include <unistd.h>
unsigned int counter = 0;

// This constant controls
// the number of iteration
#define TOTAL_ITERATIONS 5000

void *add(void *arg) {
    int i;
    pthread_mutex_t* m = (pthread_mutex_t*) arg;
    for (i=0; i<TOTAL_ITERATIONS; i++) {
        pthread_mutex_lock(m);
        ++counter;
        pthread_mutex_unlock(m);
    }
}
```

```
int main() {
    pthread_t tid1, tid2;
    pthread_mutex_t mutex;
    pthread_mutex_init(&mutex, NULL);
    printf("The original counter value is %d\n", counter);
    pthread_create(&tid1, NULL, add, &mutex);
    pthread_create(&tid2, NULL, add, &mutex);
    pthread_join(tid1, NULL);
    pthread_join(tid2, NULL);
    printf("The final counter value is %d\n", counter);
    return 0;
}
```

The purpose of this program is to protect the shared variable counter. Both threads will increase the counter by 5,000 times, so the final counter value is 10,000.







# POSIX Condition Variables

---

- POSIX condition variables are associated with a POSIX mutex lock to provide mutual exclusion: Creating and initializing the condition variable:

```
pthread_mutex_t mutex;  
pthread_cond_t cond_var;
```

```
pthread_mutex_init(&mutex, NULL);  
pthread_cond_init(&cond_var, NULL);
```





# POSIX Condition Variables

- Thread waiting for the condition  $a == b$  to become true:

```
pthread_mutex_lock(&mutex);  
while (a != b)  
    pthread_cond_wait(&cond_var, &mutex);  
  
pthread_mutex_unlock(&mutex);
```

- Thread signaling another thread waiting on the condition variable:

```
pthread_mutex_lock(&mutex);  
a = b;  
pthread_cond_signal(&cond_var);  
pthread_mutex_unlock(&mutex);
```

