

COMP4211-Tutorial 8: Policy Iteration and Value Iteration

Weiyu CHEN, wchenbx@connect.ust.hk

OpenAI Gym is a cool python package that provides ready-to-use environments to test RL algorithms on. As a python package, it is easy to install:


In [1]:

```
pip install gym==0.17.3
```

```
Requirement already satisfied: gym==0.17.3 in /Users/weiyu/opt/anaconda3/lib/python3.9/site-packages (0.17.3)
Requirement already satisfied: numpy>=1.10.4 in /Users/weiyu/opt/anaconda3/lib/python3.9/site-packages (from gym==0.17.3) (1.20.3)
Requirement already satisfied: pygame<=1.5.0,>=1.4.0 in /Users/weiyu/opt/anaconda3/lib/python3.9/site-packages (from gym==0.17.3) (1.5.0)
Requirement already satisfied: scipy in /Users/weiyu/opt/anaconda3/lib/python3.9/site-packages (from gym==0.17.3) (1.7.1)
Requirement already satisfied: cloudpickle<1.7.0,>=1.2.0 in /Users/weiyu/opt/anaconda3/lib/python3.9/site-packages (from gym==0.17.3) (1.6.0)
Requirement already satisfied: future in /Users/weiyu/opt/anaconda3/lib/python3.9/site-packages (from pygame<=1.5.0,>=1.4.0->gym==0.17.3) (0.18.2)
Note: you may need to restart the kernel to use updated packages.
```

Environment

Let's start with a simple example introduced in the lecture. You are an agent on an $M \times N$ grid and your goal is to reach the terminal state at the top left or the bottom right corner. You can take actions in each direction (UP=0, RIGHT=1, DOWN=2, LEFT=3). Actions going off the edge leave you in your current state. You receive a reward of -1 at each step until you reach a terminal state.



In [2]:

```
import io
import numpy as np
import sys
from gym.envs.toy_text import discrete
```

```

UP = 0
RIGHT = 1
DOWN = 2
LEFT = 3

class GridworldEnv(discrete.DiscreteEnv):

    metadata = {'render.modes': ['human', 'ansi']}

    def __init__(self, shape=[4,4]):
        if not isinstance(shape, (list, tuple)) or not len(shape) == 2:
            raise ValueError('shape argument must be a list/tuple of length 2')

        self.shape = shape

        nS = np.prod(shape)
        nA = 4

        MAX_Y = shape[0]
        MAX_X = shape[1]

        P = {}
        grid = np.arange(nS).reshape(shape)
        it = np.nditer(grid, flags=['multi_index'])

        while not it.finished:
            s = it.iterindex
            y, x = it.multi_index

            # P[s][a] = (prob, next_state, reward, is_done)
            P[s] = {a : [] for a in range(nA)}

            is_done = lambda s: s == 0 or s == (nS - 1)
            reward = 0.0 if is_done(s) else -1.0

            # We're stuck in a terminal state
            if is_done(s):
                P[s][UP] = [(1.0, s, reward, True)]
                P[s][RIGHT] = [(1.0, s, reward, True)]
                P[s][DOWN] = [(1.0, s, reward, True)]
                P[s][LEFT] = [(1.0, s, reward, True)]
            # Not a terminal state
            else:
                ns_up = s if y == 0 else s - MAX_X
                ns_right = s if x == (MAX_X - 1) else s + 1
                ns_down = s if y == (MAX_Y - 1) else s + MAX_X
                ns_left = s if x == 0 else s - 1
                P[s][UP] = [(1.0, ns_up, reward, is_done(ns_up))]
                P[s][RIGHT] = [(1.0, ns_right, reward, is_done(ns_right))]
                P[s][DOWN] = [(1.0, ns_down, reward, is_done(ns_down))]
                P[s][LEFT] = [(1.0, ns_left, reward, is_done(ns_left))]

            it.iternext()

        # Initial state distribution is uniform
        isd = np.ones(nS) / nS

        self.P = P

        super(GridworldEnv, self).__init__(nS, nA, P, isd)

```

Policy Evaluation

Input π , the policy to be evaluated
Initialize $V(s) = 0$, for all $s \in \mathcal{S}^+$
Repeat
 $\Delta \leftarrow 0$
 For each $s \in \mathcal{S}$:
 $v \leftarrow V(s)$
 $V(s) \leftarrow \sum_a \pi(s, a) \sum_{s'} \mathcal{P}_{ss'}^a [\mathcal{R}_{ss'}^a + \gamma V(s')]$
 $\Delta \leftarrow \max(\Delta, |v - V(s)|)$
until $\Delta < \theta$ (a small positive number)
Output $V \approx V^\pi$

In [3]:

```
env = GridworldEnv()
def policy_eval(policy, env, discount_factor=1.0, theta=0.00001):
    """
    Evaluate a policy given an environment and a full description of the environment.

    Args:
        policy: [S, A] shaped matrix representing the policy.
        env: OpenAI env. env.P represents the transition probabilities of the environment.
            env.P[s][a] is a list of transition tuples (prob, next_state, reward, done).
            env.nS is a number of states in the environment.
            env.nA is a number of actions in the environment.
        theta: We stop evaluation once our value function change is less than theta.
        discount_factor: Gamma discount factor.

    Returns:
        Vector of length env.nS representing the value function.
    """
    # Start with a random (all 0) value function
    V = np.zeros(env.nS)
    while True:
        delta = 0
        # For each state, perform a "full backup"
        for s in range(env.nS):
            v = 0
            # Look at the possible next actions
            for a, action_prob in enumerate(policy[s]):
                # For each action, look at the possible next states...
                for prob, next_state, reward, done in env.P[s][a]:
                    # Calculate the expected value.
                    v += action_prob * prob * (reward + discount_factor * V[next_state])
            # How much our value function changed (across any states)
            delta = max(delta, np.abs(v - V[s]))
            V[s] = v
        # Stop evaluating once our value function change is below a threshold
        if delta < theta:
            break
    return np.array(V)
```

In [4]:

```
random_policy = np.ones([env.nS, env.nA]) / env.nA
v = policy_eval(random_policy, env)
print(v)

[  0.          -13.99993529 -19.99990698 -21.99989761 -13.99993529
 -17.9999206  -19.99991379 -19.99991477 -19.99990698 -19.99991379
 -17.99992725 -13.99994569 -21.99989761 -19.99991477 -13.99994569
   0.          ]
```

In [5]:

```
# Test: Make sure the evaluated policy is what we expected
expected_v = np.array([0, -14, -20, -22, -14, -18, -20, -20, -20, -18])
np.testing.assert_array_almost_equal(v, expected_v, decimal=2)
```

Policy Iteration

Policy iteration iterates through policies until it converges on the optimal policy.

1. Initialization

$V(s) \in \mathbb{R}$ and $\pi(s) \in \mathcal{A}(s)$ arbitrarily for all $s \in \mathcal{S}$

2. Policy Evaluation

Repeat

$\Delta \leftarrow 0$

For each $s \in \mathcal{S}$:

$v \leftarrow V(s)$

$V(s) \leftarrow \sum_{s'} \mathcal{P}_{ss'}^{\pi(s)} [\mathcal{R}_{ss'}^{\pi(s)} + \gamma V(s')]$

$\Delta \leftarrow \max(\Delta, |v - V(s)|)$

until $\Delta < \theta$ (a small positive number)

3. Policy Improvement

policy-stable \leftarrow true

For each $s \in \mathcal{S}$:

$b \leftarrow \pi(s)$

$\pi(s) \leftarrow \arg \max_a \sum_{s'} \mathcal{P}_{ss'}^a [\mathcal{R}_{ss'}^a + \gamma V(s')]$

If $b \neq \pi(s)$, then *policy-stable* \leftarrow false

If *policy-stable*, then stop; else go to 2

In [6]:

```
def policy_improvement(env, policy_eval_fn=policy_eval, discount_factor=1.0):
    """
```

Policy Improvement Algorithm. Iteratively evaluates and improves a policy until an optimal policy is found.

Args:

env: The OpenAI environment.
policy_eval_fn: Policy Evaluation function that takes 3 arguments:
policy, env, discount_factor.
discount_factor: gamma discount factor.

Returns:

A tuple (policy, V).
policy is the optimal policy, a matrix of shape [S, A] where each row contains a valid probability distribution over actions.
V is the value function for the optimal policy.

"""

```
def one_step_lookahead(state, V):
```

```
    """
```

```
    Helper function to calculate the value for all action in a given state
```

```
    Args:
```

```
        state: The state to consider (int)  
        V: The value to use as an estimator, Vector of length env.nS
```

```
    Returns:
```

```
        A vector of length env.nA containing the expected value of each action
```

```
    """
```

```
    A = np.zeros(env.nA)
```

```
    for a in range(env.nA):
```

```
        for prob, next_state, reward, done in env.P[state][a]:
```

```
            A[a] += prob * (reward + discount_factor * V[next_state])
```

```
    return A
```

```
# Start with a random policy
```

```
policy = np.ones([env.nS, env.nA]) / env.nA
```

```
iteration = 0
```

```
while True:
```

```
    # Evaluate the current policy
```

```
    V = policy_eval_fn(policy, env, discount_factor)
```

```
    # Will be set to false if we make any changes to the policy
```

```
    policy_stable = True
```

```
    # For each state...
```

```
    for s in range(env.nS):
```

```
        # The best action we would take under the current policy
```

```
        chosen_a = np.argmax(policy[s])
```

```
        # Find the best action by one-step lookahead
```

```
        # Ties are resolved arbitrarily
```

```
        action_values = one_step_lookahead(s, V)
```

```
        best_a = np.argmax(action_values)
```

```
        # Greedily update the policy
```

```
        if chosen_a != best_a:
```

```
            policy_stable = False
```

```
        policy[s] = np.eye(env.nA)[best_a]
```

```

print(f"[Iteration {iteration}]")
print("Reshaped Grid Policy (0=up, 1=right, 2=down, 3=left):")
print(np.reshape(np.argmax(policy, axis=1), env.shape))
print("Reshaped Grid Value Function:")
print(V.reshape(env.shape))
print("")
iteration += 1

# If the policy is stable we've found an optimal policy. Return it
if policy_stable:
    return policy, V

```

In [7]:

```

env = GridworldEnv()
policy, V = policy_improvement(env)

```

```

[Iteration 0]
Reshaped Grid Policy (0=up, 1=right, 2=down, 3=left):
[[0 3 3 3]
 [0 0 3 2]
 [0 0 1 2]
 [0 1 1 0]]
Reshaped Grid Value Function:
[[ 0.          -13.99993529 -19.99990698 -21.99989761]
 [-13.99993529 -17.9999206  -19.99991379 -19.99991477]
 [-19.99990698 -19.99991379 -17.99992725 -13.99994569]
 [-21.99989761 -19.99991477 -13.99994569  0.          ]]

[Iteration 1]
Reshaped Grid Policy (0=up, 1=right, 2=down, 3=left):
[[0 3 3 2]
 [0 0 0 2]
 [0 0 1 2]
 [0 1 1 0]]
Reshaped Grid Value Function:
[[ 0. -1. -2. -3.]
 [-1. -2. -3. -2.]
 [-2. -3. -2. -1.]
 [-3. -2. -1.  0.]]

[Iteration 2]
Reshaped Grid Policy (0=up, 1=right, 2=down, 3=left):
[[0 3 3 2]
 [0 0 0 2]
 [0 0 1 2]
 [0 1 1 0]]
Reshaped Grid Value Function:
[[ 0. -1. -2. -3.]
 [-1. -2. -3. -2.]
 [-2. -3. -2. -1.]
 [-3. -2. -1.  0.]]

```

In [8]:

```

# Test the value function
expected_v = np.array([ 0, -1, -2, -3, -1, -2, -3, -2, -2, -3, -2, -1, -3,
np.testing.assert_array_almost_equal(V, expected_v, decimal=2)

```


Value Iteration

Initialize V arbitrarily, e.g., $V(s) = 0$, for all $s \in \mathcal{S}^+$

Repeat

$\Delta \leftarrow 0$

For each $s \in \mathcal{S}$:

$v \leftarrow V(s)$

$V(s) \leftarrow \max_a \sum_{s'} \mathcal{P}_{ss'}^a [\mathcal{R}_{ss'}^a + \gamma V(s')]$

$\Delta \leftarrow \max(\Delta, |v - V(s)|)$

until $\Delta < \theta$ (a small positive number)

Output a deterministic policy, π , such that

$\pi(s) = \arg \max_a \sum_{s'} \mathcal{P}_{ss'}^a [\mathcal{R}_{ss'}^a + \gamma V(s')]$

In [9]:

```
def value_iteration(env, theta=0.0001, discount_factor=1.0):
    """
    Value Iteration Algorithm.

    Args:
        env: OpenAI env. env.P represents the transition probabilities of the environment.
            env.P[s][a] is a list of transition tuples (prob, next_state, reward, done).
            env.nS is a number of states in the environment.
            env.nA is a number of actions in the environment.
        theta: We stop evaluation once our value function change is less than theta.
        discount_factor: Gamma discount factor.

    Returns:
        A tuple (policy, V) of the optimal policy and the optimal value function.
    """

    def one_step_lookahead(state, V):
        """
        Helper function to calculate the value for all action in a given state.

        Args:
            state: The state to consider (int)
            V: The value to use as an estimator, Vector of length env.nS

        Returns:
            A vector of length env.nA containing the expected value of each action.
        """
        A = np.zeros(env.nA)
        for a in range(env.nA):
            for prob, next_state, reward, done in env.P[state][a]:
                A[a] += prob * (reward + discount_factor * V[next_state])
        return A
```

```

V = np.zeros(env.nS)
iteration = 0
while True:
    # Stopping condition
    delta = 0
    # Update each state...
    for s in range(env.nS):
        # Do a one-step lookahead to find the best action
        A = one_step_lookahead(s, V)
        best_action_value = np.max(A)
        # Calculate delta across all states seen so far
        delta = max(delta, np.abs(best_action_value - V[s]))
        # Update the value function.
        V[s] = best_action_value

    # Create a deterministic policy using the optimal value function
    policy = np.zeros([env.nS, env.nA])
    for s in range(env.nS):
        # One step lookahead to find the best action for this state
        A = one_step_lookahead(s, V)
        best_action = np.argmax(A)
        # Always take the best action
        policy[s, best_action] = 1.0

    print(f"[Iteration {iteration}]")
    print("Reshaped Grid Policy (0=up, 1=right, 2=down, 3=left):")
    print(np.reshape(np.argmax(policy, axis=1), env.shape))
    print("Reshaped Grid Value Function:")
    print(V.reshape(env.shape))
    print("")
    iteration += 1
    # Check if we can stop
    if delta < theta:
        break

return policy, V

```

```

In [10]: env = GridworldEnv()
         policy, V = value_iteration(env)

```



```

[Iteration 0]
Reshaped Grid Policy (0=up, 1=right, 2=down, 3=left):
[[0 3 0 0]
 [0 0 0 0]
 [0 0 0 2]
 [0 0 1 0]]
Reshaped Grid Value Function:
[[ 0. -1. -1. -1.]
 [-1. -1. -1. -1.]
 [-1. -1. -1. -1.]
 [-1. -1. -1.  0.]]

[Iteration 1]
Reshaped Grid Policy (0=up, 1=right, 2=down, 3=left):
[[0 3 3 0]
 [0 0 0 2]
 [0 0 1 2]
 [0 1 1 0]]
Reshaped Grid Value Function:
[[ 0. -1. -2. -2.]
 [-1. -2. -2. -2.]
 [-2. -2. -2. -1.]
 [-2. -2. -1.  0.]]

[Iteration 2]
Reshaped Grid Policy (0=up, 1=right, 2=down, 3=left):
[[0 3 3 2]
 [0 0 0 2]
 [0 0 1 2]
 [0 1 1 0]]
Reshaped Grid Value Function:
[[ 0. -1. -2. -3.]
 [-1. -2. -3. -2.]
 [-2. -3. -2. -1.]
 [-3. -2. -1.  0.]]

[Iteration 3]
Reshaped Grid Policy (0=up, 1=right, 2=down, 3=left):
[[0 3 3 2]
 [0 0 0 2]
 [0 0 1 2]
 [0 1 1 0]]
Reshaped Grid Value Function:
[[ 0. -1. -2. -3.]
 [-1. -2. -3. -2.]
 [-2. -3. -2. -1.]
 [-3. -2. -1.  0.]]

```

In [11]:

```

# Test the value function
expected_v = np.array([ 0, -1, -2, -3, -1, -2, -3, -2, -2, -3, -2, -1, -3,
np.testing.assert_array_almost_equal(V, expected_v, decimal=2)

```

Reference: <https://github.com/dennybritz/reinforcement-learning>