

Data Mining

Classification: Basic Concepts, Decision Trees, and Model Evaluation

Lecture Notes for Chapter 4 Part III

Introduction to Data Mining by


Tan, Steinbach, Kumar

Adapted by Qiang Yang (2010)

Practical Issues of Classification

- | Underfitting and Overfitting
- | Missing Values
- | Costs of Classification

Underfitting and Overfitting (Example)



500 circular and 500 triangular data points.

Circular points:


$$0.5 \leq \text{sqrt}(x_1^2 + x_2^2) \leq 1$$

Triangular points:

$$\text{sqrt}(x_1^2 + x_2^2) > 0.5 \text{ or}$$


$$\text{sqrt}(x_1^2 + x_2^2) < 1$$

Underfitting and Overfitting



Underfitting: when model is too simple, both training and test errors are large

Overfitting due to Noise



Decision boundary is distorted by noise point

Notes on Overfitting

- | Overfitting results in decision trees that are more complex than necessary
- | Training error no longer provides a good estimate of how well the tree will perform on previously unseen records
- | Need new ways for estimating errors

Estimating Generalization Errors

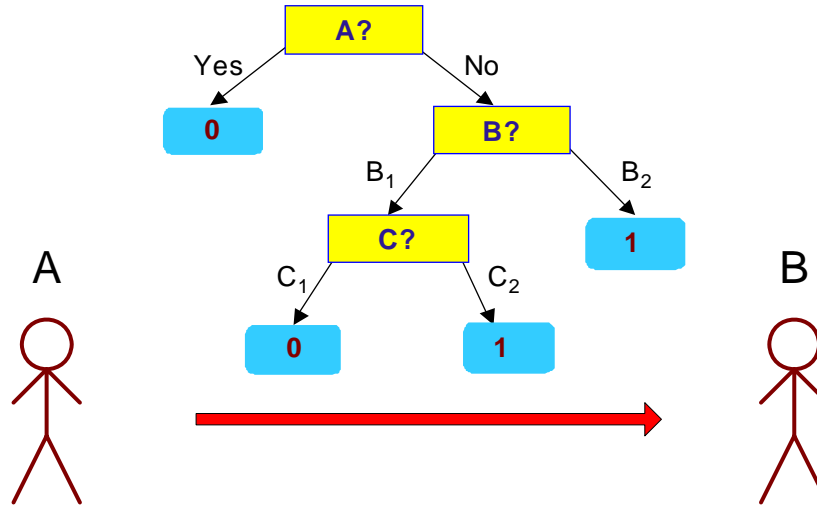
- | **Re-substitution errors:** error on training ($\sum e(t)$)
- | **Generalization errors:** error on testing ($\sum e'(t)$)
- | Methods for estimating generalization errors:
 - **Optimistic approach:** $e'(t) = e(t)$
 - **Pessimistic approach:**
 - ◆ For each leaf node: $e'(t) = (e(t)+0.5)$
 - ◆ Total error counts: $e'(T) = e(T) + N \times 0.5$ (N: number of leaf nodes)
 - ◆ For a tree with 30 leaf nodes and 10 errors on training (out of 1000 instances):
Training error = $10/1000 = 1\%$
Generalization error = $(10 + 30 \times 0.5)/1000 = 2.5\%$
 - **Reduced error pruning (REP):**
 - ◆ uses validation data set to estimate generalization error

Occam's Razor

- | Given two models of similar generalization errors, one should prefer the **simpler** model over the more complex model
 - For complex models, there is a greater chance that it was fitted accidentally by errors in data
 - Therefore, one should include **model complexity** when evaluating a model

Minimum Description Length (MDL)

X	y
X ₁	1
X ₂	0
X ₃	0
X ₄	1
...	...
X _n	1



X	y
X ₁	?
X ₂	?
X ₃	?
X ₄	?
...	...
X _n	?

- | $\text{Cost}(\text{Model}, \text{Data}) = \text{Cost}(\text{Data} | \text{Model}) + \text{Cost}(\text{Model})$
 - Cost is the **number of bits** needed for encoding.
 - We should search for the least costly model.
- | $\text{Cost}(\text{Data} | \text{Model})$ encodes the errors on training data.
- | $\text{Cost}(\text{Model})$ estimates model complexity, or future error...

How to Address Overfitting in Decision Trees

| Pre-Pruning (Early Stopping Rule)

- Stop the algorithm before it becomes a fully-grown tree
- Typical stopping conditions for a node:
 - ◆ Stop if all instances belong to the same class
 - ◆ Stop if all the attribute values are the same
- More restrictive conditions:
 - ◆ Stop if number of instances is less than some user-specified threshold
 - ◆ Stop if class distribution of instances are independent of the available features (e.g., using χ^2 test)
 - ◆ Stop if expanding the current node does not improve impurity measures (e.g., Gini or information gain).

How to Address Overfitting...

| Post-pruning

- Grow decision tree to its entirety
- Trim the nodes of the decision tree in a bottom-up fashion
- **If generalization error improves after trimming, replace sub-tree by a leaf node.**
 - ◆ Heuristic: Class label of leaf node is determined from majority class of instances in the sub-tree
 - ◆ **generalization error count = error count + $0.5 * N$, where N is the number of leaf nodes,**
 - ◆ This is a heuristic used in some algorithms, but there are other ways using statistics

Post-Pruning based on |leaves|

Class = Yes	20
Class = No	10
Error = 10/30	

Training Error (Before splitting) = 10/30

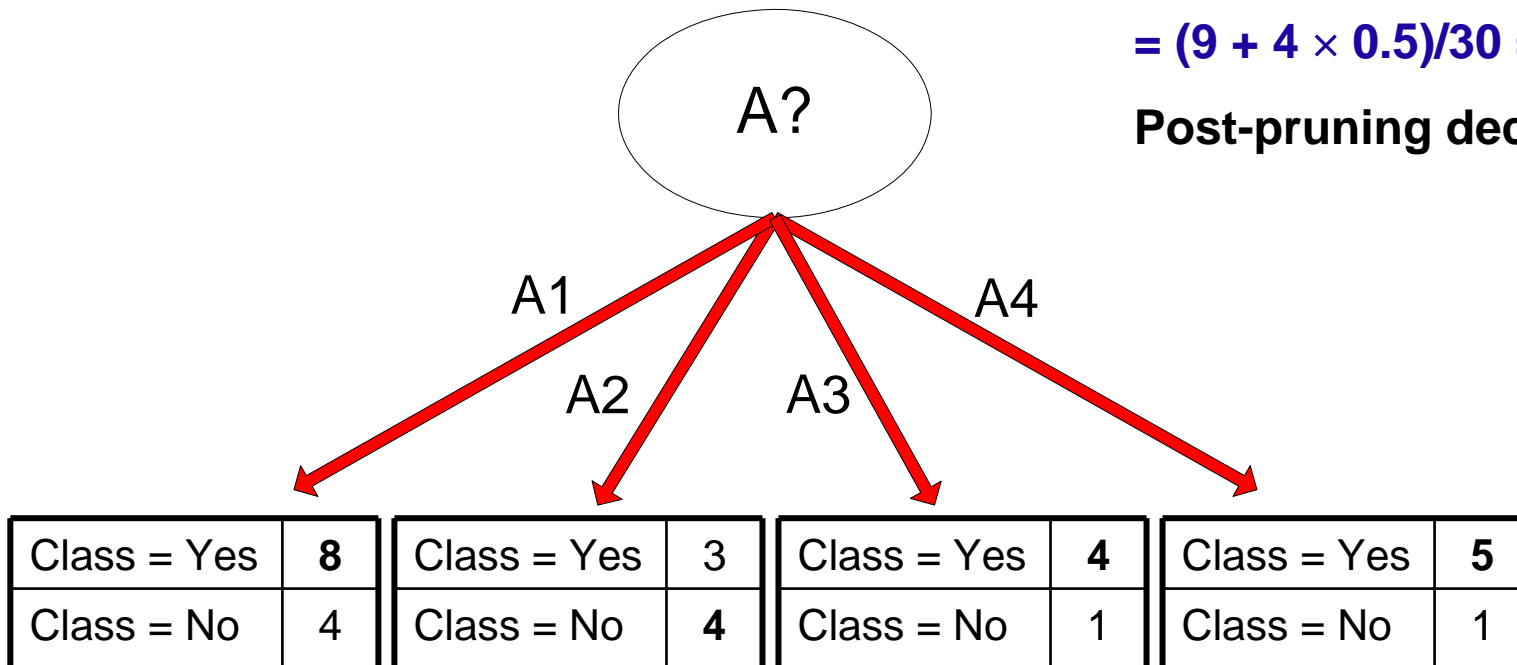
Pessimistic error (Before splitting) = $(10 + 1 \times 0.5)/30 = 10.5/30$

Training Error (After splitting) = 9/30

Pessimistic error (After splitting)

$$= (9 + 4 \times 0.5)/30 = 11/30$$

Post-pruning decision: **PRUNE!**



Examples of Post-pruning


- Optimistic error?

Don't prune for both cases


- Pessimistic error?

Don't prune case 1, prune case 2

Case 1:





Case 2:



Data Fragmentation

- | Number of instances gets smaller as you traverse down the tree
- | Number of instances at the leaf nodes could be too small to make any statistically significant decision
- | Solution: limit number of instances per leaf node \geq a user given value n .

Decision Trees: Feature Construction



- **Test condition may involve multiple attributes, but hard to automate!**
- **Finding better node test features is a difficult research issue**

Model Evaluation

- | **Metrics** for Performance Evaluation
 - How to evaluate the performance of a model?
- | **Methods** for Performance Evaluation
 - How to obtain reliable estimates?
- | Methods for Model **Comparison**
 - How to compare the relative performance among competing models?

Model Evaluation

- | Metrics for Performance Evaluation
 - How to evaluate the performance of a model?
- | Methods for Performance Evaluation
 - How to obtain reliable estimates?
- | Methods for Model Comparison
 - How to compare the relative performance among competing models?

Metrics for Performance Evaluation

- | Focus on the **predictive capability** of a model
 - Rather than how fast it takes to classify or build models, scalability, etc.
- | Confusion Matrix: count or percentage

	PREDICTED CLASS		
ACTUAL CLASS		Class=Yes	Class=No
	Class=Yes	a	b
	Class=No	c	d

- a: TP (true positive)
- b: FN (false negative)
- c: FP (false positive)
- d: TN (true negative)

Metrics for Performance Evaluation...

ACTUAL CLASS	PREDICTED CLASS	
	Class=Yes	Class=No
Class=Yes	a (TP)	b (FN)
	c (FP)	d (TN)

| Most widely-used metric:

$$\text{Accuracy} = \frac{a + d}{a + b + c + d} = \frac{TP + TN}{TP + TN + FP + FN}$$

Limitation of Accuracy

- | Consider a 2-class problem
 - Number of Class 0 examples = 9990
 - Number of Class 1 examples = 10
- | If model predicts everything to be class 0, accuracy is $9990/10000 = 99.9\%$
 - **Accuracy** is **misleading** because model does not detect any class 1 example

Cost Matrix

	PREDICTED CLASS		
	$C(i j)$	Class=Yes	Class=No
	Class=Yes	$C(\text{Yes} \text{Yes})$	$C(\text{No} \text{Yes})$
	Class=No	$C(\text{Yes} \text{No})$	$C(\text{No} \text{No})$

$C(i|j)$: Cost of misclassifying class j example as class i
- medical diagnosis, customer segmentation

Computing Cost of Classification

Cost Matrix	PREDICTED CLASS		
ACTUAL CLASS	C(i j)	+	-
	+	-1	100
	-	1	0

Confusion matrix

Model M_1	PREDICTED CLASS		
ACTUAL CLASS		+	-
	+	150	40
	-	60	250

Accuracy = 80%

Cost = 3910

Model M_2	PREDICTED CLASS		
ACTUAL CLASS		+	-
	+	250	45
	-	5	200

Accuracy = 90%

Cost = 4255

Information Retrieval Measures

$$\text{Precision : } p = \frac{a}{a + c}$$

$$\text{Recall : } r = \frac{a}{a + b}$$

$$\text{F - measure (F)} = \frac{2rp}{r + p} = \frac{2a}{2a + b + c}$$

	PREDICTED CLASS	
	Class=Yes	Class=No
	Class=Yes	Class=No
ACTUAL CLASS	a	b
	c	d

- | Let C be cost (can be count in our example)
- | Precision is biased towards C(Yes|Yes) & C(Yes|No)
- | Recall is biased towards C(Yes|Yes) & C(No|Yes)
- | F-measure is biased towards all except C(No|No)

Model Evaluation

- | Metrics for Performance Evaluation
 - How to evaluate the performance of a model?
- | **Methods for Performance Evaluation**
 - How to obtain reliable estimates?
- | Methods for Model Comparison
 - How to compare the relative performance among competing models?

Methods of Estimation

- | Holdout
 - Reserve $2/3$ for training and $1/3$ for testing
- | Cross validation
 - Partition data into k disjoint subsets
 - k -fold: train on $k-1$ partitions, test on the remaining one
 - Leave-one-out: $k=n$

Test of Significance (Sections 4.5,4.6 of TSK Book)

- | Given two models:
 - Model M1: accuracy = 85%, tested on 30 instances
 - Model M2: accuracy = 75%, tested on 5000 instances
- | Can we say M1 is better than M2?
 - How much **confidence** can we place on accuracy of M1 and M2?
 - Can the difference in performance measure be explained as a result of random fluctuations in the test set?

Confidence Interval for Accuracy

- | Prediction can be regarded as a Bernoulli trial
 - A Bernoulli trial has 2 possible outcomes
 - Possible outcomes for prediction: correct or wrong
 - Collection of Bernoulli trials has a Binomial distribution:
 - ◆ $x \sim \text{Bin}(N, p)$ x : number of correct predictions
 - ◆ e.g: Toss a fair coin 50 times, how many heads would turn up?
Expected number of heads = $N \times p = 50 \times 0.5 = 25$
- | Given x (# of correct predictions) or equivalently, $\text{acc} = x/N$, and N = # of test instances,
 - Can we predict p (true accuracy of model)?

Confidence Interval for Accuracy


- For large N , let $1-\alpha$ be confidence

- acc has a normal distribution with mean p and variance $p(1-p)/N$

$$P(Z_{\alpha/2} < \frac{acc - p}{\sqrt{p(1-p)/N}} < Z_{1-\alpha/2}) = 1 - \alpha$$

- Confidence Interval for p :

$$p = \frac{2 \times N \times acc + Z_{\alpha/2}^2 \pm \sqrt{Z_{\alpha/2}^2 + 4 \times N \times acc - 4 \times N \times acc^2}}{2(N + Z_{\alpha/2}^2)}$$



Confidence Interval for Accuracy

- Consider a model that produces an accuracy of 80% when evaluated on 100 test instances:
 - $N=100$, $\text{acc} = 0.8$
 - Let $1-\alpha = 0.95$ (95% confidence)
 - From probability table, $Z_{\alpha/2}=1.96$

N	50	100	500	1000	5000
p(lower)	0.670	0.711	0.763	0.774	0.789
p(upper)	0.888	0.866	0.833	0.824	0.811

$1-\alpha$	Z
0.99	2.58
0.98	2.33
0.95	1.96
0.90	1.65

ROC (Receiver Operating Characteristic)

- | Page 298 of TSK book.
- | Many applications care about ranking (give a queue from the most likely to the least likely)
- | Examples...
- | Which ranking order is better?
- | ROC: Developed in 1950s for signal detection theory to analyze noisy signals
 - Characterize the trade-off between positive hits and false alarms
- | ROC curve plots TP (on the y-axis) against FP (on the x-axis)
- | Performance of each classifier represented as a point on the ROC curve
 - changing the threshold of algorithm, sample distribution or cost matrix changes the location of the point

How to Construct an ROC curve

Instance	$P(+ A)$	True Class
1	0.95	+
2	0.93	+
3	0.87	-
4	0.85	-
5	0.85	-
6	0.85	+
7	0.76	-
8	0.53	+
9	0.43	-
10	0.25	+

Predicted by classifier


This is the ground truth

- Use classifier that produces posterior probability for each test instance $P(+|A)$ for instance A
- Sort the instances according to $P(+|A)$ in decreasing order
- Apply threshold at each unique value of $P(+|A)$
- Count the number of TP, FP, TN, FN at each threshold
- TP rate, $TPR = TP/(TP+FN)$
- FP rate, $FPR = FP/(FP + TN)$


How to construct an ROC curve

Class	+	-	+	-	-	-	+	-	+	+	
Threshold \geq	0.25	0.43	0.53	0.76	0.85	0.85	0.85	0.87	0.93	0.95	1.00
TP	5	4	4	3	3	3	3	2	2	1	0
FP	5	5	4	4	3	2	1	1	0	0	0
TN	0	0	1	1	2	3	4	4	5	5	5
FN	0	1	1	2	2	2	2	3	3	4	5
TPR	1	0.8	0.8	0.6	0.6	0.6	0.6	0.4	0.4	0.2	0
FPR	1	1	0.8	0.8	0.6	0.4	0.2	0.2	0	0	0

ROC Curve:



Using ROC for Model Comparison



- | No model consistently outperform the other
 - | M_1 is better for small FPR
 - | M_2 is better for large FPR
- | Area Under the ROC curve: AUC
 - | Ideal:
 - Area = 1
 - | Random guess:
 - Area = 0.5

ROC Curve

(TP,FP):

- | (0,0): declare everything to be negative class
- | (1,1): declare everything to be positive class
- | (1,0): ideal
- | Diagonal line:
 - Random guessing
 - Below diagonal line:
 - ◆ prediction is opposite of the true class

