

# Dijkstra's Shortest Path Algorithm

DPV 4.4, CLRS 24.3

Revised, October 23, 2014

## Outline of this Lecture

- Recalling the BFS solution of the shortest path problem for unweighted (di)graphs.
- The shortest path problem for **weighted** digraphs.
- Dijkstra's algorithm.  
Given for digraphs but easily modified to work on undirected graphs.

## Recall: Shortest Path Problem for Graphs

Let  $G = (V, E)$  be a (di)graph.

- The shortest path between two vertices is a path with the shortest **length** (least number of edges). Call this the **link-distance**.
- Breadth-first-search is an algorithm for finding shortest (link-distance) paths from a **single source vertex** to all other vertices.
- BFS processes vertices in increasing order of their distance from the root vertex.
- BFS has running time  $O(|V| + |E|)$ .


## Shortest Path Problem for **Weighted** Graphs

Let  $G = (V, E)$  be a **weighted digraph**, with weight function  $w : E \mapsto \mathbb{R}$  mapping edges to real-valued weights. If  $e = (u, v)$ , we write  $w(u, v)$  for  $w(e)$ .

- The **length** of a path  $p = \langle v_0, v_1, \dots, v_k \rangle$  is the **sum** of the weights of its constituent edges:

$$\text{length}(p) = \sum_{i=1}^k w(v_{i-1}, v_i).$$

- The **distance** from  $u$  to  $v$ , denoted  $\delta(u, v)$ , is the length of the **minimum length path** if there is a path from  $u$  to  $v$ ; and is  $\infty$  otherwise.



$\text{length}(\langle a, b, c, e \rangle) = 6$   
distance from  $a$  to  $e$  is 6

## Single-Source Shortest-Paths Problem


**The Problem:** Given a digraph with **positive** edge weights  $G = (V, E)$  and a distinguished **source vertex**,  $s \in V$ , determine the **distance** and a **shortest path** from the source vertex to every vertex in the digraph.

**Question:** How do you design an efficient algorithm for this problem?

## Single-Source Shortest-Paths Problem

**Important Observation:** Any subpath of a shortest path must also be a shortest path. Why?

**Example:** In the following digraph,  $\langle a, b, c, e \rangle$  is a shortest path. The subpath  $\langle a, b, c \rangle$  is also a shortest path.



**Observation** Extending this idea we observe the existence of a *shortest path tree* in which distance from source to vertex  $v$  is length of shortest path from source to vertex in original tree.

## Intuition behind Dijkstra's Algorithm

- Report the vertices in increasing order of their distance from the source vertex.
- Construct the shortest path tree edge by edge; at each step adding one new edge, corresponding to construction of shortest path to the current new vertex.

## The Rough Idea of Dijkstra's Algorithm

- Maintain an *estimate*  $d[v]$  of the length  $\delta(s, v)$  of the shortest path for each vertex  $v$ .
- Always  $d[v] \geq \delta(s, v)$  and  $d[v]$  equals the length of a known path ( $d[v] = \infty$  if we have no paths so far).
- Initially  $d[s] = 0$  and all the other  $d[v]$  values are set to  $\infty$ . The algorithm will then **process** the vertices one by one in **some order**.  
The processed vertex's estimate will be validated as being real shortest distance, i.e.  $d[v] = \delta(s, v)$ .

Here “processing a vertex  $u$ ” means **finding** new paths and **updating**  $d[v]$  for all  $v \in Adj[u]$  if necessary. The process by which an estimate is updated is called **relaxation**.

When all vertices have been processed,  
 $d[v] = \delta(s, v)$  for all  $v$ .

## The Rough Idea of Dijkstra's Algorithm

**Question 1:** How does the algorithm find new paths and do the **relaxation**?

**Question 2:** In which order does the algorithm **process** the vertices one by one?



## Answer to Question 1

- **Finding new paths.** When processing a vertex  $u$ , the algorithm will examine all vertices  $v \in Adj[u]$ . For each vertex  $v \in Adj[u]$ , a new path from  $s$  to  $v$  is found (path from  $s$  to  $u$  + new edge).
- **Relaxation.** If the new path from  $s$  to  $v$  is shorter than  $d[v]$ , then update  $d[v]$  to the length of this new path.

**Remark:** Whenever we set  $d[v]$  to a finite value, there exists a path of that length. Therefore  $d[v] \geq \delta(s, v)$ .

(Note: If  $d[v] = \delta(s, v)$ , then further relaxations cannot change its value.)

## Implementing the Idea of Relaxation


Consider an edge from a vertex  $u$  to  $v$  whose weight is  $w(u, v)$ . Suppose that we have already processed  $u$  so that we know  $d[u] = \delta(s, u)$  and also computed a current estimate for  $d[v]$ . Then

- There is a (shortest) path from  $s$  to  $u$  with length  $d[u]$ .
- There is a path from  $s$  to  $v$  with length  $d[v]$ .

Combining this path from  $s$  to  $u$  with the edge  $(u, v)$ , we obtain another path from  $s$  to  $v$  with length  $d[u] + w(u, v)$ .

If  $d[u] + w(u, v) < d[v]$ , then we replace the old path  $\langle s, \dots, w, v \rangle$  with the new shorter path  $\langle s, \dots, u, v \rangle$ . Hence we update

- $d[v] = d[u] + w(u, v)$
- $pred[v] = u$  (originally,  $pred[v] == w$ ).



## The Algorithm for Relaxing an Edge

Relax( $u, v$ )

```
{  
    if ( $d[u] + w(u, v) < d[v]$ )  
    {  
         $d[v] = d[u] + w(u, v);$   
         $pred[v] = u;$   
    }  
}
```

**Remark:** The predecessor pointer  $pred[]$  is for determining the shortest paths.

## Idea of Dijkstra's Algorithm: Repeated Relaxation

- Dijkstra's algorithm operates by maintaining a subset of vertices,  $S \subseteq V$ , for which we **know** the true distance, that is  $d[v] = \delta(s, v)$ .
- Initially  $S = \emptyset$ , the empty set, and we set  $d[s] = 0$  and  $d[v] = \infty$  for all others vertices  $v$ . One by one we **select** vertices from  $V \setminus S$  to add to  $S$ .
- The set  $S$  can be implemented using an array of vertex colors. Initially all vertices are white, and we set  $color[v] = \text{black}$  to indicate that  $v \in S$ .

## The Selection in Dijkstra's Algorithm

**Recall Question 2:** What is the best order in which to process vertices, so that the estimates are guaranteed to converge to the true distances.

That is, how does the algorithm **select** which vertex among the vertices of  $V \setminus S$  to process next?

**Answer:** We use a **greedy** algorithm. For each vertex in  $u \in V \setminus S$ , we have computed a distance estimate  $d[u]$ . The next vertex processed is always a vertex  $u \in V \setminus S$  for which  $d[u]$  is minimum, that is, we take the unprocessed vertex that is closest (by our estimate) to  $s$ .

**Question:** How do we implement this selection of vertices **efficiently**?

## The Selection in Dijkstra's Algorithm

**Question:** How do we perform this selection efficiently?

**Answer:** We store the vertices of  $V \setminus S$  in a *priority queue*, where the key value of each vertex  $v$  is  $d[v]$ .

[Note: if we implement the priority queue using a heap, we can perform the operations **Insert()**, **Extract\_Min()**, and **Decrease\_Key()**, each in  $O(\log n)$  time.]

## Review of Priority Queues

A **Priority Queue** is a data structure (can be implemented as a heap) which supports the following operations:

**insert( $u, key$ ):** Insert  $u$  with the key value  $key$  in  $Q$ .

**$u = \text{extractMin}()$ :** Extract the item with the minimum key value in  $Q$ .

**decreaseKey( $u, new\text{-}key$ ):** Decrease  $u$ 's key value to  $new\text{-}key$ .

**Remark:** *Priority Queues can be implemented such that each operation takes time  $O(\log |Q|)$ .*

*In Class we only saw insert, extractMin and Delete. DecreaseKey can either be implemented directly in  $O(\log |Q|)$  time or could be implemented by a delete followed by an insert.*

## Description of Dijkstra's Algorithm

Dijkstra( $G, w, s$ )

```
{
    for (each  $u \in V$ )
    {
         $d[u] = \infty$ ;
         $color[u] = \text{white}$ ;
    }
     $d[s] = 0$ ;
     $pred[s] = \text{NIL}$ ;
     $Q = (\text{queue with all vertices})$ ;

    while (Non-Empty( $Q$ ))
    {
         $u = \text{Extract-Min}(Q)$ ;
        for (each  $v \in Adj[u]$ )
        {
            if ( $d[u] + w(u, v) < d[v]$ )
            {
                 $d[v] = d[u] + w(u, v)$ ;
                Decrease-Key( $Q, v, d[v]$ );
                 $pred[v] = u$ ;
            }
        }
         $color[u] = \text{black}$ ;
    }
}
```

**% Initialize**

**% Process all vertices**

**% Find new vertex**


**% If estimate improves**

**relax**



## Dijkstra's Algorithm

**Example:**



**Step 0:** Initialization.


$v$	s	a	b	c	d
$d[v]$	0	$\infty$	$\infty$	$\infty$	$\infty$
$pred[v]$	nil	nil	nil	nil	nil
$color[v]$	W	W	W	W	W

**Priority Queue:**

$v$	s	a	b	c	d
$d[v]$	0	$\infty$	$\infty$	$\infty$	$\infty$

## Dijkstra's Algorithm

**Example:**



**Step 1:** As  $Adj[s] = \{a, b\}$ , work on  $a$  and  $b$  and update information.


$v$	s	a	b	c	d
$d[v]$	0	2	7	$\infty$	$\infty$
$pred[v]$	nil	s	s	nil	nil
$color[v]$	B	W	W	W	W

**Priority Queue:**

$v$	a	b	c	d
$d[v]$	2	7	$\infty$	$\infty$

## Dijkstra's Algorithm

**Example:**



**Step 2:** After Step 1,  $a$  has the minimum key in the priority queue. As  $Adj[a] = \{b, c, d\}$ , work on  $b, c, d$  and update information.


$v$	s	a	b	c	d
$d[v]$	0	2	5	10	7
$pred[v]$	nil	s	a	a	a
$color[v]$	B	B	W	W	W

**Priority Queue:**

$v$	b	c	d
$d[v]$	5	10	7

## Dijkstra's Algorithm

**Example:**



**Step 3:** After Step 2,  $b$  has the minimum key in the priority queue. As  $Adj[b] = \{a, c\}$ , work on  $a, c$  and update information.


$v$	s	a	b	c	d
$d[v]$	0	2	5	6	7
$pred[v]$	nil	s	a	b	a
$color[v]$	B	B	B	W	W

**Priority Queue:**

$v$	c	d
$d[v]$	6	7

## Dijkstra's Algorithm

**Example:**



**Step 4:** After Step 3,  $c$  has the minimum key in the priority queue. As  $Adj[c] = \{d\}$ , work on  $d$  and update information.


$v$	s	a	b	c	d
$d[v]$	0	2	5	6	7
$pred[v]$	nil	s	a	b	a
$color[v]$	B	B	B	B	W

**Priority Queue:**

$v$	d
$d[v]$	7

## Dijkstra's Algorithm

**Example:**



**Step 5:** After Step 4,  $d$  has the minimum key in the priority queue. As  $Adj[d] = \{c\}$ , work on  $c$  and update information.

$v$	s	a	b	c	d
$d[v]$	0	2	5	6	7
$pred[v]$	nil	s	a	b	a
$color[v]$	B	B	B	B	B

**Priority Queue:**  $Q = \emptyset$ .


We are done.

## Dijkstra's Algorithm

**Shortest Path Tree:**  $T = (V, A)$ , where

$$A = \{(pred[v], v) | v \in V \setminus \{s\}\}.$$

The array  $pred[v]$  is used to build the tree.



**Example:**


$v$	s	a	b	c	d
$d[v]$	0	2	5	6	7
$pred[v]$	nil	s	a	b	a

## Correctness of Dijkstra's Algorithm

**Lemma:** When a vertex  $u$  is added to  $S$  (i.e., dequeued from the queue),  $d[u] = \delta(s, u)$ .

**Proof:** Suppose to the contrary that at some point Dijkstra's algorithm *first* attempts to add a vertex  $u$  to  $S$  for which  $d[u] \neq \delta(s, u)$ . By our observations about relaxation,  $d[u] > \delta(s, u)$ .

Consider the situation just prior to the insertion of  $u$ . Consider the *true shortest path* from  $s$  to  $u$ . Because  $s \in S$  and  $u \in V \setminus S$ , at some point this path must first take a jump out of  $S$ . Let  $(x, y)$  be the edge taken by the path, where  $x \in S$  and  $y \in V \setminus S$  (it may be that  $x = s$  and/or  $y = u$ ).





## Correctness of Dijkstra's Algorithm – Continued

We now prove that  $d[y] = \delta(s, y)$ . We have done relaxation when processing  $x$ , so

$$d[y] \leq d[x] + w(x, y). \quad (1)$$

Since  $x$  is added to  $S$  earlier, by hypothesis,

$$d[x] = \delta(s, x). \quad (2)$$

Since  $\langle s, \dots, x, y \rangle$  is subpath of a shortest path, by (2)

$$\delta(s, y) = \delta(s, x) + w(x, y) = d[x] + w(x, y). \quad (3)$$

By (1) and (3),

$$d[y] \leq \delta(s, y).$$

Hence

$$d[y] = \delta(s, y).$$

So  $y \neq u$  (because we suppose  $d[u] > \delta(s, u)$ ).

Now observe that since  $y$  appears midway on the path from  $s$  to  $u$ , and all subsequent edges are positive, we have  $\delta(s, y) < \delta(s, u)$ , and thus

$$d[y] = \delta(s, y) < \delta(s, u) \leq d[u].$$

Thus  $y$  would have been added to  $S$  *before*  $u$ , in contradiction to our assumption that  $u$  is the next vertex to be added to  $S$ .

## Proof of the Correctness of Dijkstra's Algorithm

- By the lemma,  $d[v] = \delta(s, v)$  when  $v$  is added into  $S$ , that is when we set  $color[v] = \text{black}$ .
- At the end of the algorithm, all vertices are in  $S$ , then all distance estimates are correct.

## Analysis of Dijkstra's Algorithm:

**Insert:** Performed once for each vertex.  $n$  times. .

**Non-Empty():** Called one for each vertex.  $n$  times.  
Each time does one **Extract-Min()** and one color setting

inner loop **for (each  $v \in Adj[u]$ ) :**  
is called once for each of the  $m$  edges in the graph.  
 $O(m)$  times.  
Each call of the inner loop does  $O(1)$  work plus, possibly, one **Decrease-Key** operation.

Recall that all of the priority queue operations require  $O(\log |Q|) = O(\log n)$  time  
we have that the algorithm uses

$nO(\log n) + nO(\log n) + O(e \log n) = O((n+e) \log n)$   
time.

### Further Analysis of Dijkstra's Algorithm:

Dijkstra's original implementation did not use a priority queue but, instead, ran through the entire list of white vertices at each step, finding the one with smallest weight. Each run through requires  $O(n)$  time and there are only  $n$  steps so his implementation runs in  $O(n^2)$  time.

This is slightly better than the priority queue based algorithm for *dense* graphs but much worse for *sparse* ones.

## Further Improvements to Dijkstra's Algorithm:

A more advanced priority queue data structure called a **Fibonacci Heap** implements

- **Insert**: in  $O(1)$  time
- **Extract-Min()**:  $O(\log n)$  time
- **Decrease-Key**:  $O(1)$  (amortized) time.

This reduces the running time down to

$nO(\log n) + nO(\log n) + O(e) = O((n \log n + e))$   
time.

This is a huge improvement for dense graphs

Prove: Dijkstra's algorithm processes vertices in non-decreasing order of their actual distance from the source vertex.