


Convolutional Neural Net


COMP4211




Sparse Connectivity



Sparse
connections
due to small
convolution
kernel




Dense
connections




- receptive field

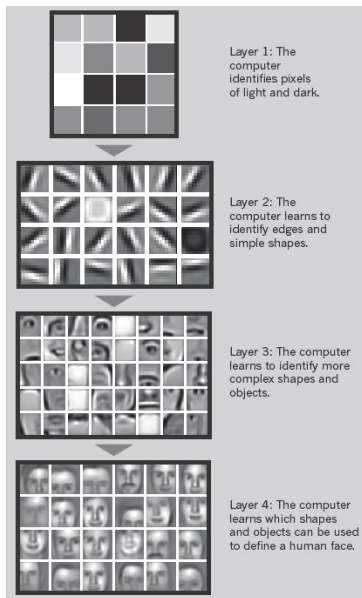
Feature Hierarchy




- hidden units are connected to a **local** subset of units in the previous layer



Example: Face Recognition




Growing Receptive Field




Shared Weights

- each local receptive field is replicated across the entire image




- weights of the same color are **shared** (constrained to be identical)




Parameter Sharing

Convolution
shares the same
parameters
across all spatial
locations




Traditional
matrix
multiplication
does not share
any parameters




- allows for features to be detected regardless of their position in the image
 - robustness to shifts of the input

Convolutional Layer

- multiple feature maps look at the same region of the input



- stack the activation maps for all filters along the depth dimension



- 1 × 1 convolution**
 - perform convolution **without** looking at neighboring pixels
 - dimension reduction**

Efficiency of Convolution

- parameter sharing greatly reduces the number of free parameters to learn

Input size: 320 by 280


Kernel size: 2 by 1

Output size: 319 by 280

	Convolution	Dense matrix	Sparse matrix
Stored floats	2	$319 \times 280 \times 320 \times 280$ $> 8e9$	$2 \times 319 \times 280 =$ 178,640


- Convolution is a linear operation
- need nonlinearity
 - otherwise 2 convolution layers would be no more powerful than 1
- common to apply a rectified linear unit (ReLU): $y = \max(z, 0)$

Zero-Padding



- representation shrink at each layer
- limits the number of layers

Zero-padding



- adding zeros to each layer
- allows the use of an arbitrarily deep convolutional network

Pooling Layer

motivation

once a feature has been detected, only its **approximate** position relative to other features is relevant

Example

the input image contains


- 1 the endpoint of a roughly horizontal segment in the upper left area
- 2 a corner in the upper right area
- 3 the endpoint of a roughly vertical segment in the lower portion

the input image is a seven

- positions are likely to vary for different instances of the character
- **spatial invariance**


Pooling

- **max-pooling**
 - for each such sub-region (e.g., over a 2×2 area in the previous layer), outputs the **maximum** value



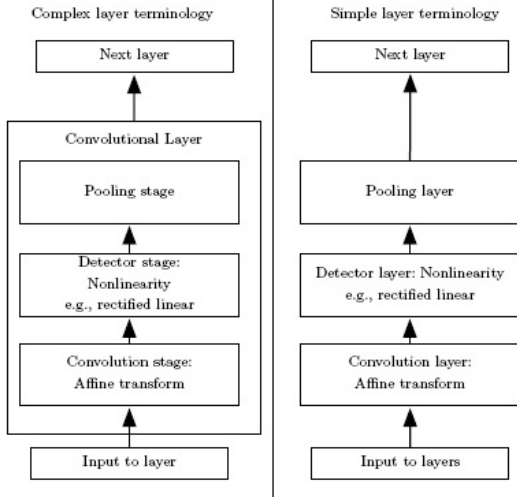
- can also have **average pooling**

Pooling with Downsampling




- **stride** of two
- reduces the representation size by a factor of two
- reduces the computational and statistical burden on the next layer


Convolutional Network Components



Example Classification Architecture



Example



- lower-layers: alternating convolution and max-pooling layers
- fully-connected (traditional MLP)
- classification error