

# Heterogeneous Parallel Programming

## COMP4901D

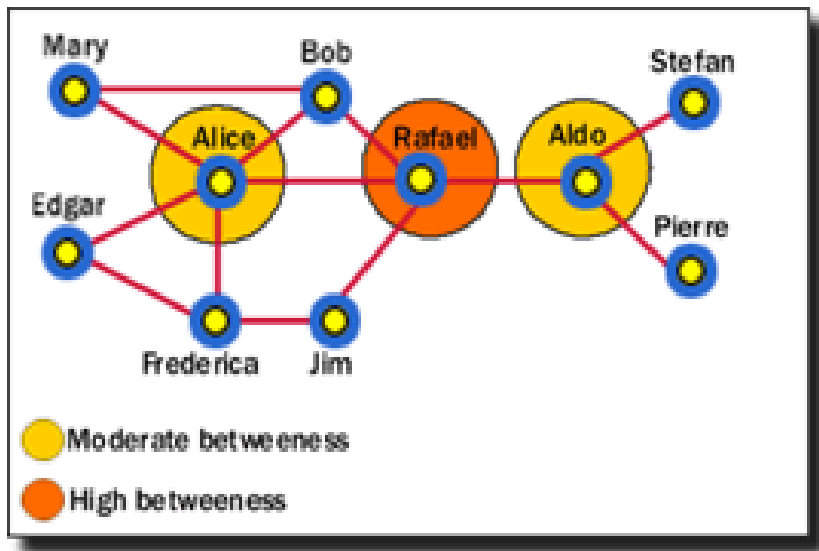
Betweenness Centrality Computation  
On Heterogeneous Processors

# Overview

- Betweenness Centrality measure
- Sequential algorithms
- Existing parallel algorithms for weighted and unweighted graphs
- Our GPU-based algorithms

# BC Measure

## Centrality Metric



- importance of individual nodes
- wide range of applications

# Betweenness Centrality Measure

## Freeman [1]

- Denote  $\sigma_{st}$   
# shortest paths between  $s$  and  $t$
- Denote  $\sigma_{st}(v)$   
# shortest paths between  $s$  and  $t$  through  $v$

$$BC(v) = \sum_{s \neq t \neq v} \delta_{st}(v) = \sum_{s \neq t \neq v} \frac{\sigma_{st}(v)}{\sigma_{st}}$$

## Brandes [2]


$$\begin{aligned} \delta_s(v) &= \sum_t \delta_{st}(v) \\ &= \sum_{w: v \in \text{prev}(s, w)} \frac{\sigma_{sv}}{\sigma_{sw}} (1 + \delta_s(w)) \end{aligned}$$

$$BC(v) = \sum_{s \neq v} \delta_s(v)$$

[1] L. Freeman. A Set of Measures of Centrality Based on Betweenness. *Sociometry*, 40:35–41, 1977.


[2] U. Brandes. A faster algorithm for betweenness centrality. *Journal of Mathematical Sociology*, 25:163–177, 2001.

# Sequential Algorithm




An example of BC computation rooted at A on unweighted graph

# Sequential Algorithm




An example of BC computation rooted at A on unweighted graph

# Sequential Algorithm



An example of BC computation rooted at A on unweighted graph

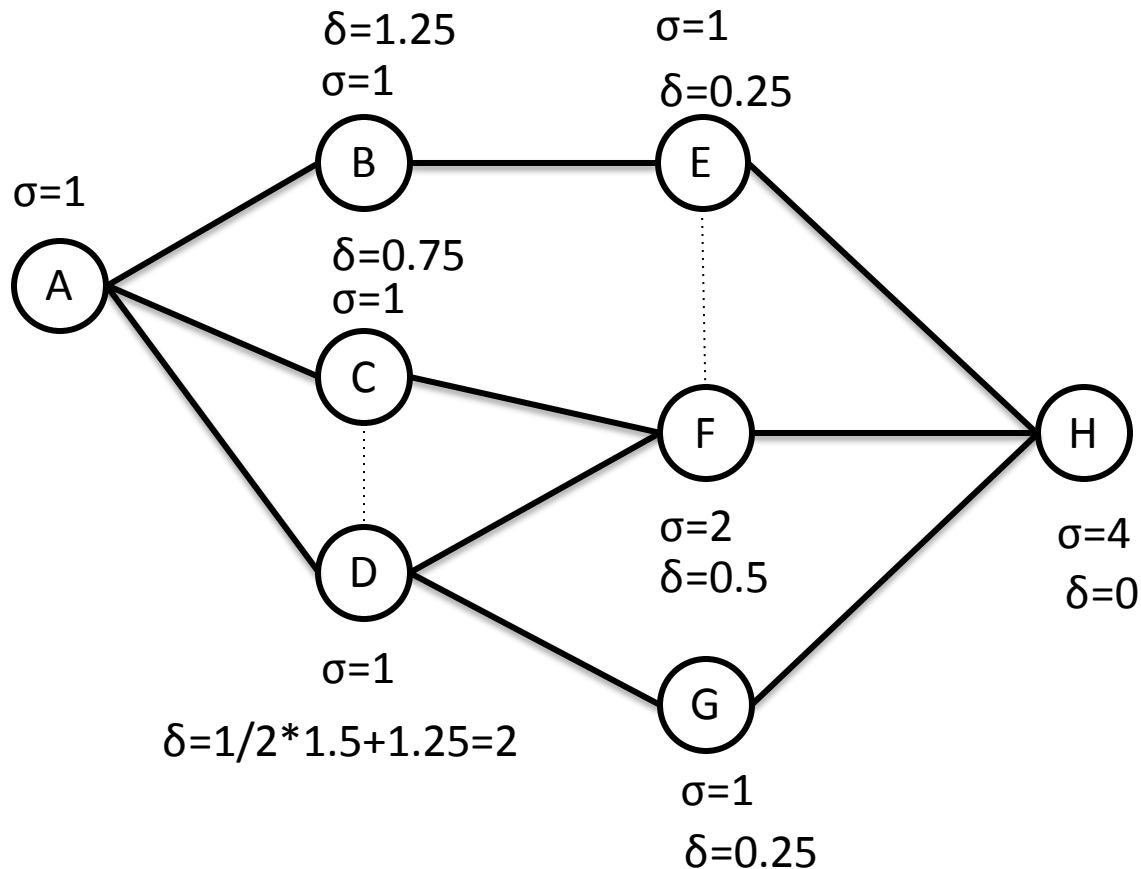
# Sequential Algorithm



An example of BC computation rooted at A on unweighted graph




# Sequential Algorithm



An example of BC computation rooted at A on unweighted graph


# Algorithmic Redundancy in BC

one-degree vertex



+

=



$\sigma_{v_1 v_2}$




is root



is the rest

two-degree vertex



$\times 2$

for each  $v \in V$  &  $v \neq s$  do

$d1[v] \leftarrow d1[v] + \text{weight1}$

$d2[v] \leftarrow d2[v] + \text{weight2}$

if  $d1[v] < d2[v]$  then

$ds[v] \leftarrow d1[v],$

$\sigma s[v] \leftarrow \sigma1[v]$

if  $d1[v] = d2[v]$  then

$ds[v] \leftarrow d1[v],$

$\sigma s[v] \leftarrow \sigma1[v] + \sigma2[v]$

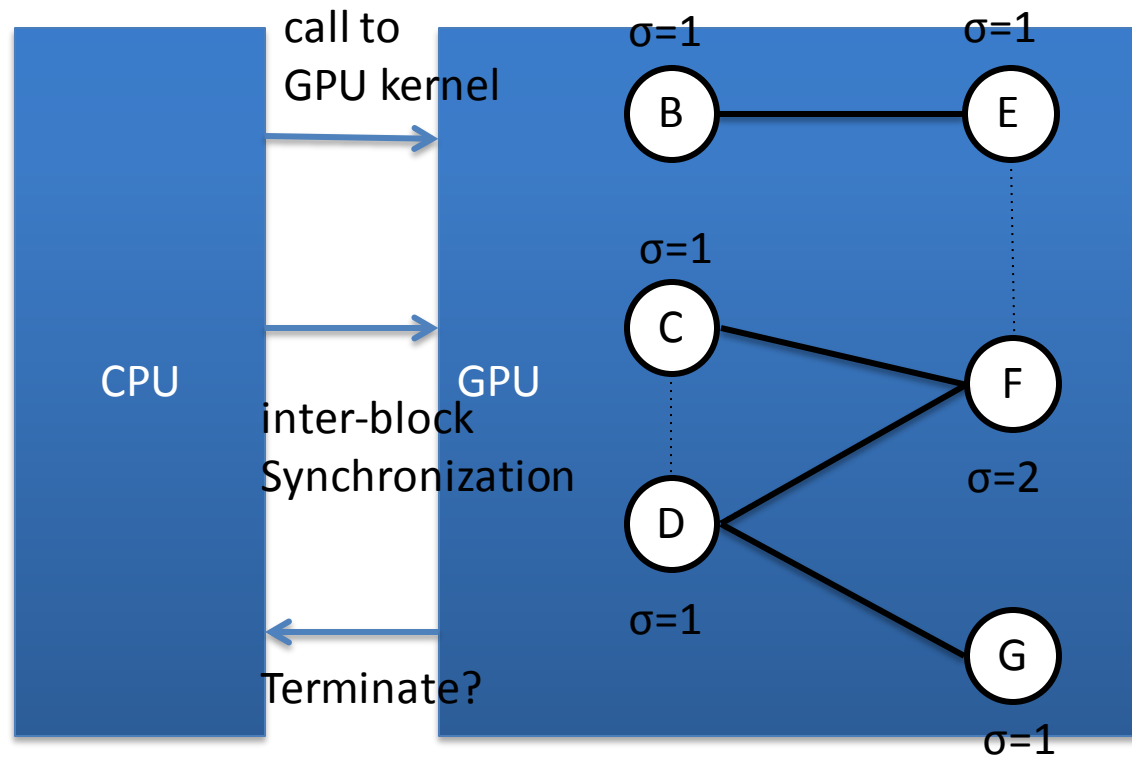
if  $d1[v] > d2[v]$  then

$ds[v] \leftarrow d2[v]$

$\sigma s[v] \leftarrow \sigma2[v]$

# Parallel BC on Unweighted Graphs

- Level-synchronization mechanism




# Parallel BC on Unweighted Graphs

- Vertex-oriented BC[1]
- Edge-oriented BC [2]
- Virtualization-based BC [3]

[1] Pawan Harish and P. J. Narayanan. Accelerating Large Graph Algorithms on the GPU using CUDA. In Srinivas Aluru, Manish Parashar, Ramamurthy Badrinath, and Viktor K. Prasanna, editors, *HiPC*, volume 4873 of *Lecture Notes in Computer Science*, pages 197–208. Springer, 2007


[2] Zhiao Shi and Bing Zhang 0003. Fast network centrality analysis using gpus. *BMC Bioinformatics*, 12:149, 2011.

[3] Ahmet Erdem Sariyüce, Kamer Kaya, Erik Saule, and Umit V. Catalyürek. Betweenness centrality on gpus and heterogeneous architectures. In *Proceedings of the 6th Workshop on General Purpose Processor Using Graphics Processing Units*, pages 76–85. ACM, 2013.



# Parallel Shortest-Path Search on Weighted Graphs

- Distance-sensitive GPU-based BC




$U_i$ : the unsettled nodes in  $i$ th iteration

Limited value  $\Delta_i$ : the shortest distance of vertices in  $U_i$  plus the minimum of edge weights

$F_i$ : the set of frontier nodes in  $i$ th iteration

# Parallel BC on Unweighted Graphs

- Edge-oriented BC with filter
  - CPU-based filter
  - GPU-based filter
- Dynamic parallel BC

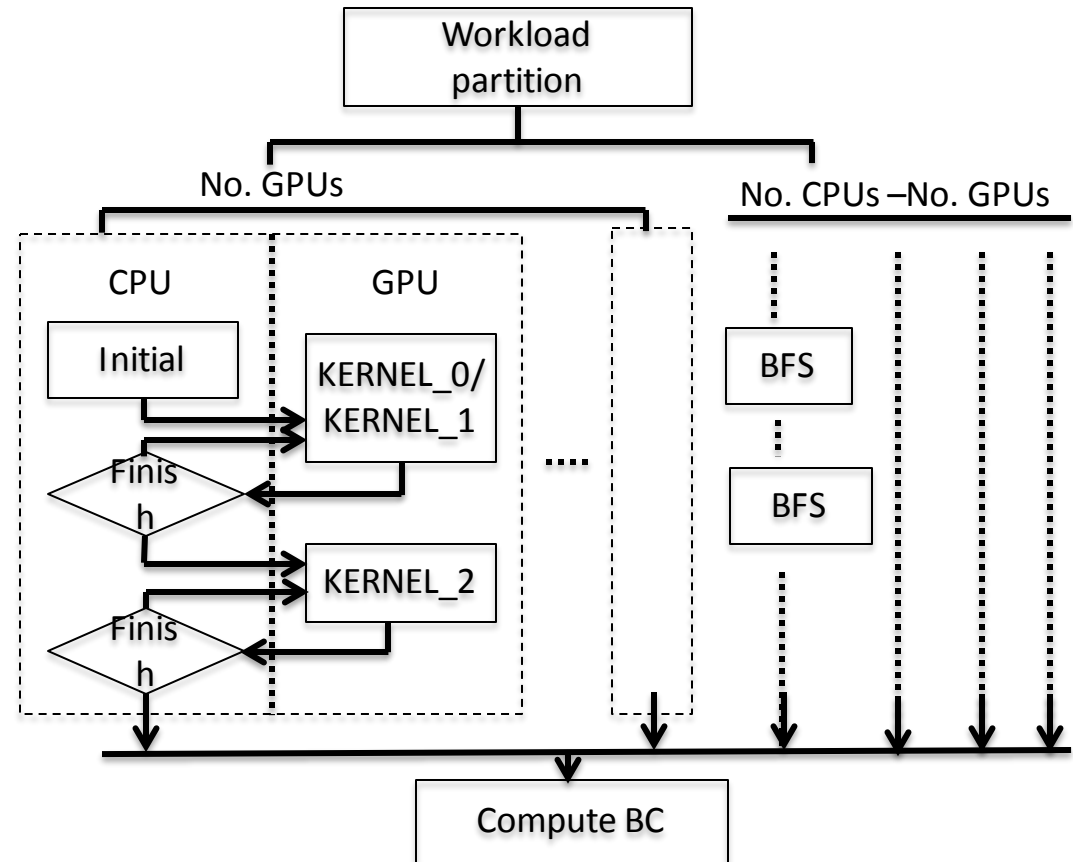


# Parallel BC on Unweighted Graphs

- GPU-based implementation issues
  - Maximize memory throughput
    - Segmented sorting end array
    - Mark array
  - Reduce thread workload
    - Special first F-STEP kernel
    - Vertex-backward
    - DAG-traversal merging

# Hetero-BC for Unweighted Graphs

1. discard the one-degree vertices
2. assign the computation of the two-degree vertices and their neighbors to the GPUs
3. partition the rest of the work on heterogeneous processors



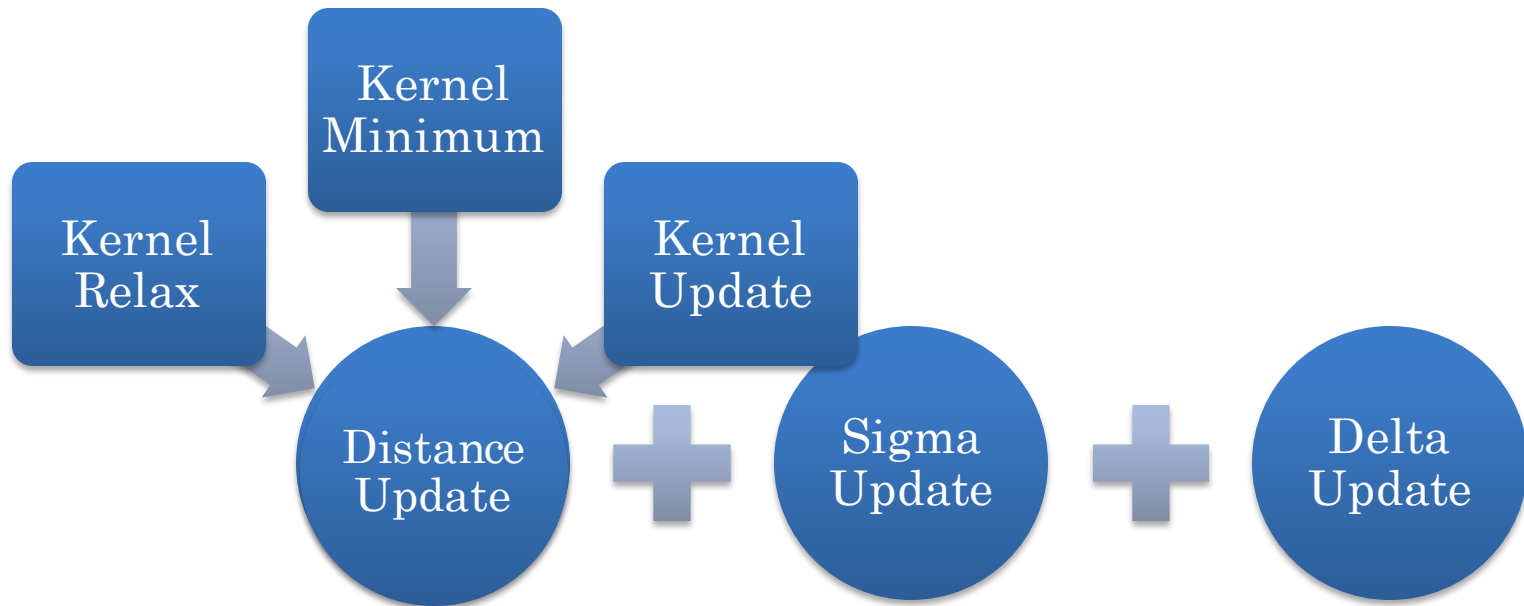


# Parallel BC on Weighted Graphs

- Distance-sensitive GPU-based BC
- Distance-insensitive GPU-based BC
- Hetero-BC for Weighted Graphs

# Parallel BC on Weighted Graphs

- Distance-sensitive GPU-based BC



# Parallel BC on Weighted Graphs


- Distance-insensitive GPU-based BC
  - cons
    - redundant relaxation in distance update kernel
  - pros
    - remove Kernel Minimum & Update
    - fewer iterations in level-synchronization framework

# Parallel BC on Weighted Graphs

## Hetero-BC for Weighted Graphs

- reduce write conflict

```
w ← head[getThreadID]
v ← end[getThreadID]
if level[w] = current level then
  if distance[v] ≥ distance[w] + weight[w][v] then
    atomicMin d[v] ← distance[w] + weight[w][v]
    level[v] ← current level
  d cont ← TRUE
```



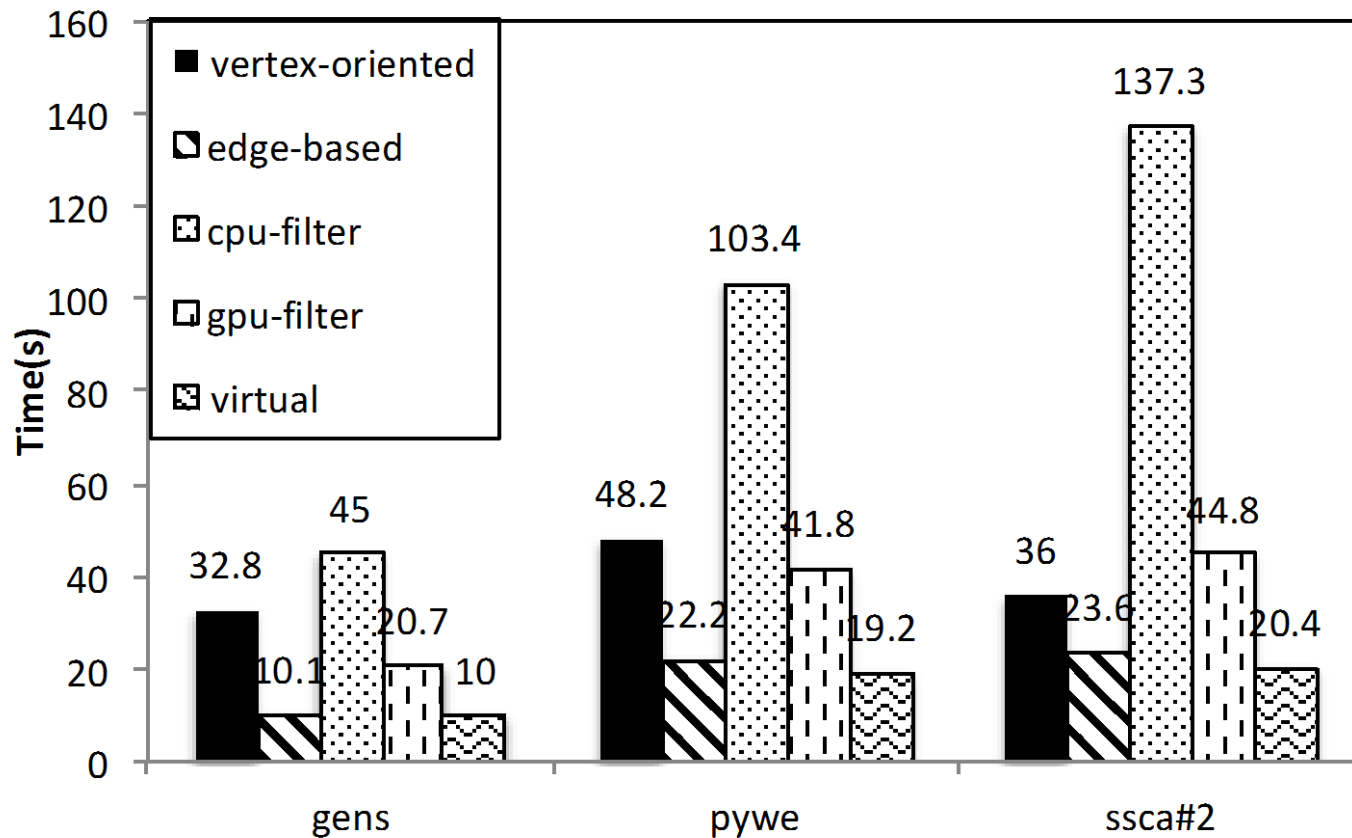
# Experimental Evaluation

- Setup

Processor	Intel E5-2650	Tesla M2090	Tesla K20m
No. cores	8	512	2496
Processor clock	2 GHz	1.3 GHz	3.52 GHz
Cache size	20 M	L2: 768 KB	L2: 1.3 MB
Memory size	768 GB	6 GB	5 GB
Memory bandwidth	51.2 GB/s	138 GB/s	208 GBy/s
Bandwidth between CPU and GPU	2973 MB/s		

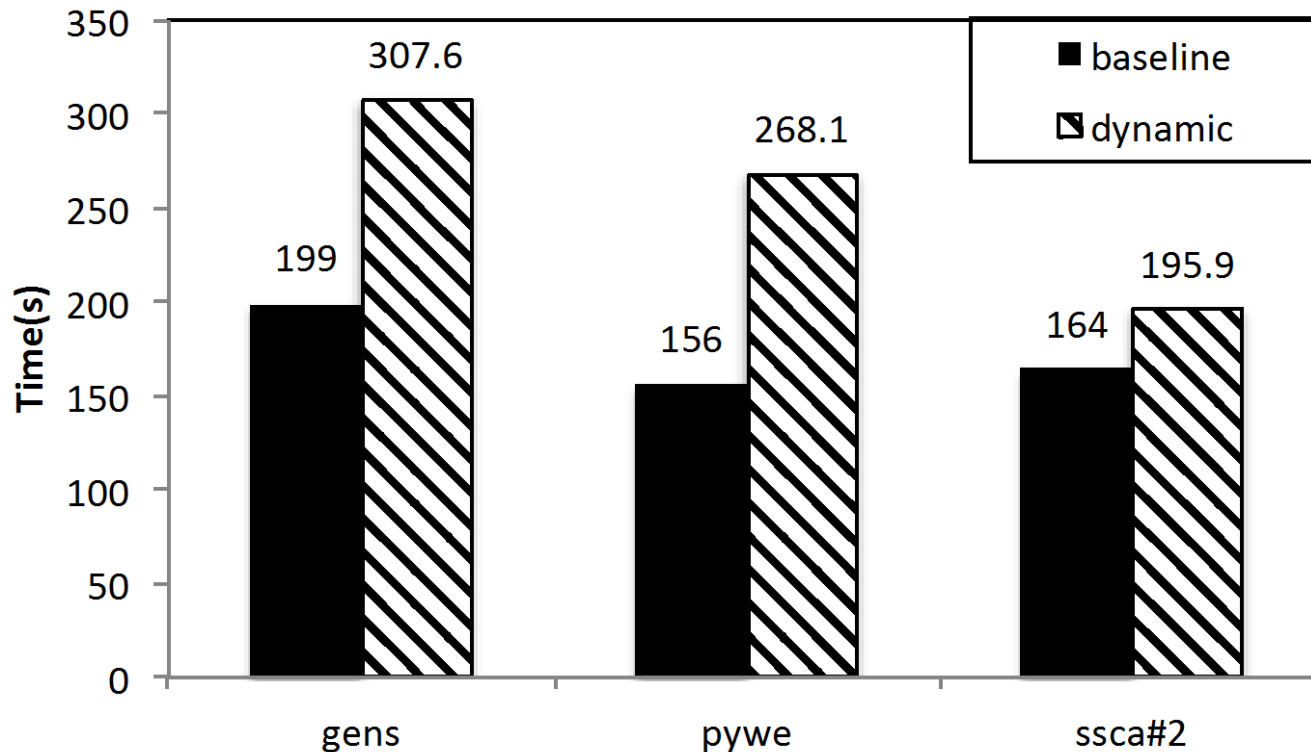
Software :CentOS 6.3 64-bit with gcc 4.4.6 and CUDA Toolkit 4.2.9

# Experimental Evaluation




Performance of GPU-based BC on unweighted graphs  
with five basic F-STEP kernel programs

# Experimental Evaluation



Performance of dynamic Parallel BC on unweighted graphs


# Experimental Evaluation



Performance of GPU-based BC on unweighted graphs with various optimizations

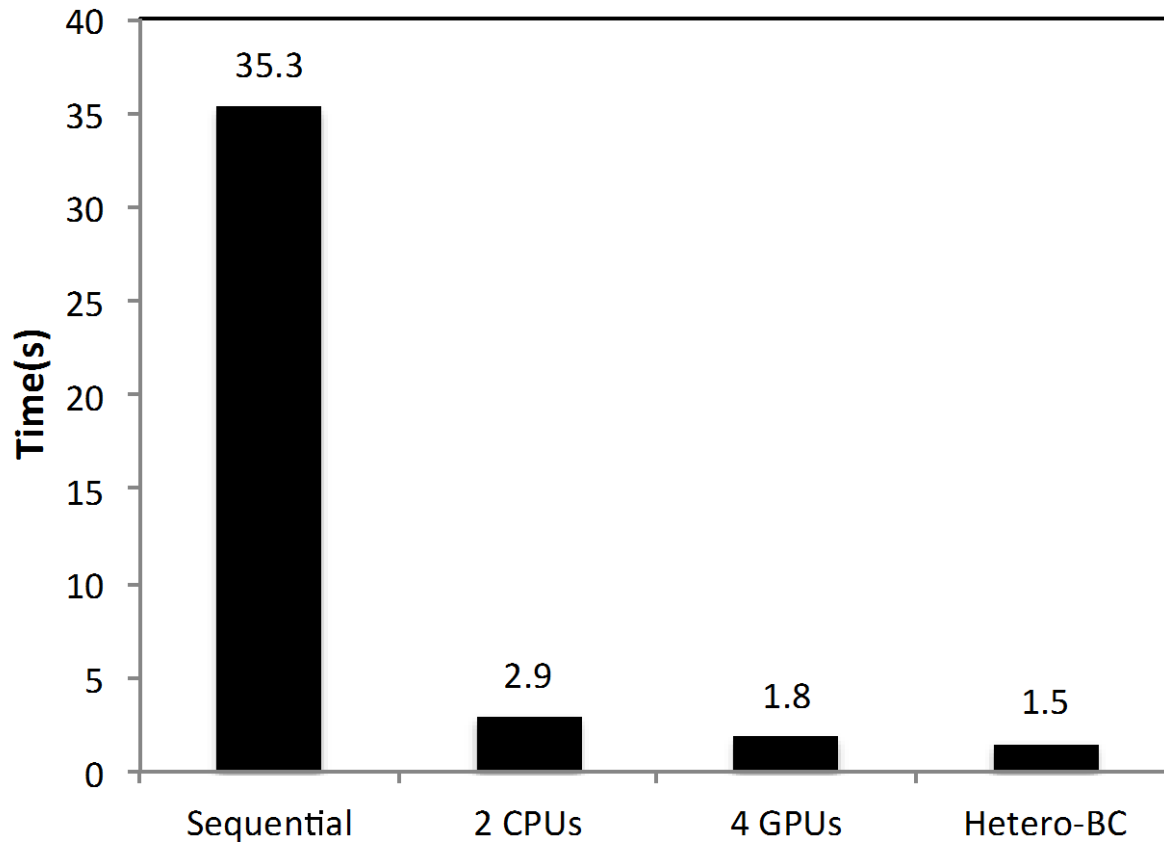


# Experimental Evaluation




Comparison of distance-sensitive/insensitive CPU-based and GPU-based BC on weighted graphs

# Experimental Evaluation



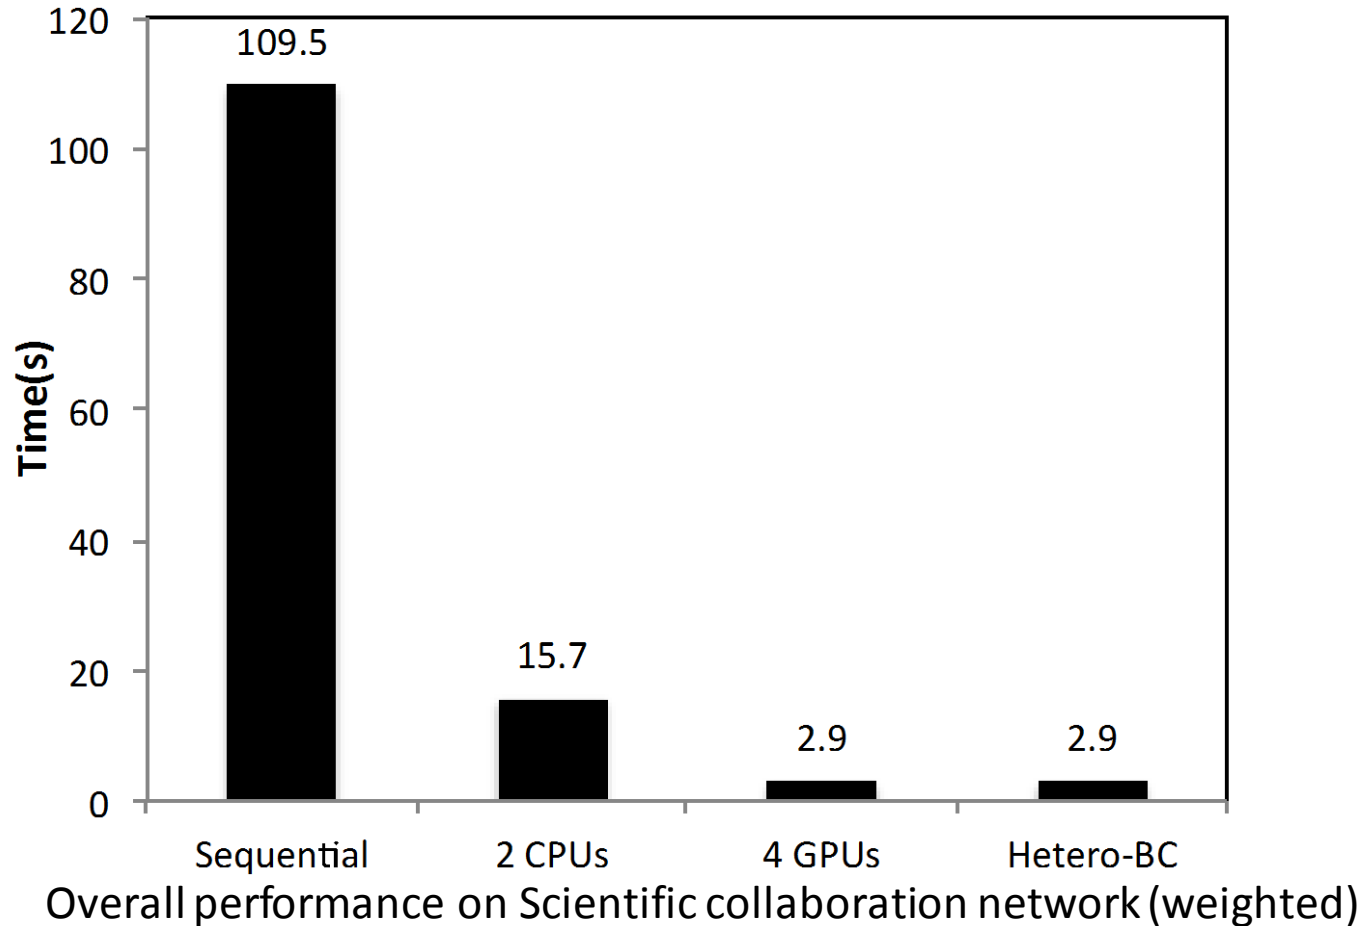
Overall performance on Protein interaction network (unweighted)

# Experimental Evaluation



Overall performance on Co-author and Citation Networks (unweighted)

# Experimental Evaluation



# Summary

- We design a GPU-based BC computation algorithm for unweighted graphs, where the BFS is parallelized by edges for small scale graphs or by virtual vertices for large scale graphs.
- We design a novel distance-insensitive edge-based BC algorithm for weighted graphs on the GPU.
- We perform general optimizations such as for low-degree nodes as well as GPU-specific optimizations such as splitting different computation stages into separate kernel programs.