

Heterogeneous Parallel Programming

COMP4901D

Data-Parallel Primitives:
Scatter and Gather

Overview

- Data-Parallel Primitives
 - Map, Prefix Scan, Scatter, Gather, Split, Sort
 - Others: Reduce, Filter, Search...
- GPU-Based Implementations
 - Gather and Scatter

Processing a Large Number of Data Items

```
//sequential
```

```
for (i = 0; i < N; i++)
```

```
    h_C[i] = h_A[i] + h_B[i];
```

```
//data-parallel
```

```
__global__ void VecAdd(int* A, int* B, int* C)
```

```
{  
    int i = blockDim.x * blockIdx.x + threadIdx.x;  
    C[i] = A[i] + B[i];  
}
```

Map and Prefix Scan

Primitive: Map

Input: $R_{in}[1, \dots, n]$, a map function fcn .

Output: $R_{out}[1, \dots, n]$.

Function: $R_{out}[i] = fcn(R_{in}[i])$.

Primitive: Prefix Scan

Input: $R_{in}[1, \dots, n]$, binary operator \oplus .

Output: $R_{out}[1, \dots, n]$.

Function: $R_{out}[i] = \bigoplus_{j < i} R_{in}[j]$.

Scatter and Gather

Primitive: Scatter

Input: $R_{in}[1, \dots, n], L[1, \dots, n]$.

Output: $R_{out}[1, \dots, n]$.

Function: $R_{out}[L[i]] = R_{in}[i], i=1, \dots, n$.

Primitive: Gather

Input: $R_{in}[1, \dots, n], L[1, \dots, n]$.

Output: $R_{out}[1, \dots, n]$.

Function: $R_{out}[i] = R_{in}[L[i]], i=1, \dots, n$.

Split and Sort

Primitive: Split

Input: $R_{in}[1, \dots, n]$, $func(R_{in}[i]) \in [1, \dots, F]$, $i=1, \dots, n$.

Output: $R_{out}[1, \dots, n]$.

Function: $\{R_{out}[i], i=1, \dots, n\} = \{R_{in}[i], i=1, \dots, n\}$
and $func(R_{out}[i]) \leq func(R_{out}[j]), \forall i, j \in [1, \dots, n], i \leq j$.

Primitive: Sort

Input: $R_{in}[1, \dots, n]$.

Output: $R_{out}[1, \dots, n]$.


Function: $\{R_{out}[i], i=1, \dots, n\} = \{R_{in}[i], i=1, \dots, n\}$ and
 $R_{out}[i] \leq R_{out}[j], \forall i, j \in [1, \dots, n]$ and $i \leq j$.

Scatter and Gather: Overview


- Widely supported
 - Parallel programming languages, e.g., MPI, NESL, ZPL.
 - Supercomputers, e.g., Cray MTA, Stanford Merrimac
 - Commodity co-processors (IBM Cell, GPUs)
- Irregular access patterns
 - Sparse matrix computations, hashing, searching, etc.
- Performance is memory bandwidth limited
 - Require high bandwidth architectures
 - HPC benchmarks (HPC Challenge, NAS PB, etc.)

Access Patterns

- Scatter: sequential reads and random writes.
- Gather: random reads and sequential writes.




(a) Scatter



(b) Gather

Scatter and Gather on the GPU

- Access pattern makes a 30X difference in performance.



Single-pass Scatter

R

0	3	1	2	3	1	0	2	0	3	1	2	3	0	1	2
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

L

0	12	4	8	13	5	1	9	2	14	6	10	15	3	7	11
---	----	---	---	----	---	---	---	---	----	---	----	----	---	---	----

4 partitions
4 concurrent threads.



Cache

Single-pass Scatter

R

0	3	1	2	3	1	0	2	0	3	1	2	3	0	1	2
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

L

0	12	4	8	13	5	1	9	2	14	6	10	15	3	7	11
---	----	---	---	----	---	---	---	---	----	---	----	----	---	---	----

R

0	3	1	2	3	1	0	2	0	3	1	2	3	0	1	2
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

R_{out}


--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--

4 partitions
4 concurrent threads.



Cache

Single-pass Scatter



4 partitions
4 concurrent threads.



Cache

Cache Misses = 4
Cache Hits = 0

Single-pass Scatter

R

0	3	1	2	3	1	0	2	0	3	1	2	3	0	1	2
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

L

0	12	4	8	13	5	1	9	2	14	6	10	15	3	7	11
---	----	---	---	----	---	---	---	---	----	---	----	----	---	---	----

R

0	3	1	2	3	1	0	2	0	3	1	2	3	0	1	2
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

R_{out}

0				1				2				3			
---	--	--	--	---	--	--	--	---	--	--	--	---	--	--	--

4 partitions
4 concurrent threads.

0				1			
---	--	--	--	---	--	--	--

Cache

Cache Misses = 4

Cache Hits = 0

Single-pass Scatter

R

0	3	1	2	3	1	0	2	0	3	1	2	3	0	1	2
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

L

0	12	4	8	13	5	1	9	2	14	6	10	15	3	7	11
---	----	---	---	----	---	---	---	---	----	---	----	----	---	---	----

R

0	3	1	2	3	1	0	2	0	3	1	2	3	0	1	2
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

R_{out}

0	0			1	1			2	2			3	3		
---	---	--	--	---	---	--	--	---	---	--	--	---	---	--	--

4 partitions
4 concurrent threads.

2				3			
---	--	--	--	---	--	--	--

Cache

Cache Misses = 6

Cache Hits = 2

Single-pass Scatter

R	0	3	1	2	3	1	0	2	0	3	1	2	3	0	1	2
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

L	0	12	4	8	13	5	1	9	2	14	6	10	15	3	7	11
---	---	----	---	---	----	---	---	---	---	----	---	----	----	---	---	----

R	0	3	1	2	3	1	0	2	0	3	1	2	3	0	1	2
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

R _{out}	0	0			1	1			2	2			3	3		
------------------	---	---	--	--	---	---	--	--	---	---	--	--	---	---	--	--

4 partitions
4 concurrent threads.

2				3			
---	--	--	--	---	--	--	--

Cache

Cache Misses = 6
Cache Hits = 2

Single-pass Scatter

R

0	3	1	2	3	1	0	2	0	3	1	2	3	0	1	2
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

L

0	12	4	8	13	5	1	9	2	14	6	10	15	3	7	11
---	----	---	---	----	---	---	---	---	----	---	----	----	---	---	----

R

0	3	1	2	3	1	0	2	0	3	1	2	3	0	1	2
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

R_{out}

0	0	0		1	1	1		2	2	2		3	3	3	
---	---	---	--	---	---	---	--	---	---	---	--	---	---	---	--

4 partitions
4 concurrent threads.

0				3			
---	--	--	--	---	--	--	--

Cache

Cache Misses = 8

Cache Hits = 4

Single-pass Scatter

R

0	3	1	2	3	1	0	2	0	3	1	2	3	0	1	2
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

L

0	12	4	8	13	5	1	9	2	14	6	10	15	3	7	11
---	----	---	---	----	---	---	---	---	----	---	----	----	---	---	----

R

0	3	1	2	3	1	0	2	0	3	1	2	3	0	1	2
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

R_{out}

0	0	0		1	1	1		2	2	2		3	3	3	
---	---	---	--	---	---	---	--	---	---	---	--	---	---	---	--

4 partitions
4 concurrent threads.

0				3			
---	--	--	--	---	--	--	--

Cache

Cache Misses = 8

Cache Hits = 4

Single-pass Scatter

R

0	3	1	2	3	1	0	2	0	3	1	2	3	0	1	2
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

L

0	12	4	8	13	5	1	9	2	14	6	10	15	3	7	11
---	----	---	---	----	---	---	---	---	----	---	----	----	---	---	----

R

0	3	1	2	3	1	0	2	0	3	1	2	3	0	1	2
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

R_{out}

0	0	0	0	1	1	1	1	2	2	2	2	3	3	3	3
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

4 partitions
4 concurrent threads.

0				3			
---	--	--	--	---	--	--	--

Cache

Cache Misses = 10

Cache Hits = 6

Cache miss rate = 62.5%
Effective bandwidth = 0.4 B_{seq}

Multi-pass Scheme

- The entire scatter is performed in multiple passes.
- Each pass writes to a small chunk

Two-pass Scatter


R

0	3	1	2	3	1	0	2	0	3	1	2	3	0	1	2
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

L

0	12	4	8	13	5	1	9	2	14	6	10	15	3	7	11
---	----	---	---	----	---	---	---	---	----	---	----	----	---	---	----

4 partitions
4 concurrent threads.



Cache

Two-pass Scatter

R

0	3	1	2	3	1	0	2	0	3	1	2	3	0	1	2
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

L

0	12	4	8	13	5	1	9	2	14	6	10	15	3	7	11
---	----	---	---	----	---	---	---	---	----	---	----	----	---	---	----

R

0	3	1	2	3	1	0	2	0	3	1	2	3	0	1	2
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

R_{out}


--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--

4 partitions
4 concurrent threads.



Cache

Two-pass Scatter



4 partitions
4 concurrent threads.



Cache

Cache Misses = 2
Cache Hits = 0

Two-pass Scatter

R

0	3	1	2	3	1	0	2	0	3	1	2	3	0	1	2
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

L

0	12	4	8	13	5	1	9	2	14	6	10	15	3	7	11
---	----	---	---	----	---	---	---	---	----	---	----	----	---	---	----

R

0	3	1	2	3	1	0	2	0	3	1	2	3	0	1	2
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

R_{out}

0				1											
---	--	--	--	---	--	--	--	--	--	--	--	--	--	--	--


4 partitions
4 concurrent threads.

0				1			
---	--	--	--	---	--	--	--


Cache

Cache Misses = 2
Cache Hits = 0

Two-pass Scatter



4 partitions
4 concurrent threads.



Cache

Cache Misses = 2
Cache Hits = 2

Two-pass Scatter

R

0	3	1	2	3	1	0	2	0	3	1	2	3	0	1	2
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

L

0	12	4	8	13	5	1	9	2	14	6	10	15	3	7	11
---	----	---	---	----	---	---	---	---	----	---	----	----	---	---	----

R

0	3	1	2	3	1	0	2	0	3	1	2	3	0	1	2
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

R_{out}

0	0			1	1										
---	---	--	--	---	---	--	--	--	--	--	--	--	--	--	--


4 partitions
4 concurrent threads.

0				1			
---	--	--	--	---	--	--	--


Cache

Cache Misses = 2
Cache Hits = 2

Two-pass Scatter



4 partitions
4 concurrent threads.



Cache

Cache Misses = 2
Cache Hits = 4

Two-pass Scatter

R

0	3	1	2	3	1	0	2	0	3	1	2	3	0	1	2
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

L

0	12	4	8	13	5	1	9	2	14	6	10	15	3	7	11
---	----	---	---	----	---	---	---	---	----	---	----	----	---	---	----

R

0	3	1	2	3	1	0	2	0	3	1	2	3	0	1	2
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

R_{out}

0	0	0		1	1	1									
---	---	---	--	---	---	---	--	--	--	--	--	--	--	--	--

4 partitions
4 concurrent threads.

0				1			
---	--	--	--	---	--	--	--

Cache

Cache Misses = 2
Cache Hits = 4

Two-pass Scatter

R

0	3	1	2	3	1	0	2	0	3	1	2	3	0	1	2
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

L

0	12	4	8	13	5	1	9	2	14	6	10	15	3	7	11
---	----	---	---	----	---	---	---	---	----	---	----	----	---	---	----

R

0	3	1	2	3	1	0	2	0	3	1	2	3	0	1	2
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

R_{out}

0	0	0	0	1	1	1	1								
---	---	---	---	---	---	---	---	--	--	--	--	--	--	--	--

4 partitions
4 concurrent threads.



Cache

Cache Misses = 2
Cache Hits = 6

Two-pass Scatter

R

0	3	1	2	3	1	0	2	0	3	1	2	3	0	1	2
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

L

0	12	4	8	13	5	1	9	2	14	6	10	15	3	7	11
---	----	---	---	----	---	---	---	---	----	---	----	----	---	---	----

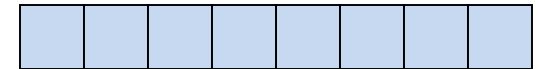
R

0	3	1	2	3	1	0	2	0	3	1	2	3	0	1	2
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

R_{out}

0	0	0	0	1	1	1	1	2	2	2	2	3	3	3	3
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

4 partitions
4 concurrent threads.



Cache

Cache Misses = 4
Cache Hits = 12

Cache miss rate = 25%
Effective bandwidth = B_{seq}

Cost Model


- Estimate the performance of different access patterns
 - Sequential bandwidth
 - Random bandwidth
- Estimate the total cost of sequential access and random access in the multi-pass scheme.

$$T_{\text{scatter}} = (|R| + |L|) * \text{npasses} / B_{\text{seq}} + |R| / B_{\text{rand}}$$

- Determine the optimal number of passes.

Performance Results

-- Multi-pass Scatter




The optimal number of passes is 16.

Applications and Analysis


- Applications
 - Radix sort, hash search, and sparse-matrix vector multiplication
- Platforms
 - CPUs: Intel Quad, or two AMD dual-core processors.
 - GPU: Nvidia 8800 GTX.
- Overall results
 - The cost model has an accuracy of over 85%.
 - The multipass scheme improves the application 10%~50%.
 - The GPU-based algorithm outperforms the CPU-based algorithm by 2-7X.

Sequential Scatter/Gather

Scatter



Gather



Accuracy of the cost model on the sequential gather and the scatter: 87%


$$\text{Accuracy} = 1 - \frac{| \text{Measurement} - \text{Estimation} |}{\text{Measurement}}$$

Random Scatter/Gather



Accuracy of the cost model on the random gather and the scatter: 90%.

Performance Impact of Multi-Pass Scatter



- (1) The speedup is 7-13X and 2-4X on Intel and AMD, respectively.
- (2) The multi-pass scheme improves the GPU-based scatter by 2-4X.

Summary

- Data-parallel primitives are an efficient way of utilizing GPU's parallelism.
- Scatter and gather are memory-bound and can be optimized through multi-pass schemes.

References:

Bingsheng He, Naga K. Govindaraju, Qiong Luo, and Burton Smith. Efficient Gather and Scatter Operations on Graphics Processors. ACM/IEEE SuperComputing (SC), Nov 2007.

<http://www.cse.ust.hk/gpuqp>