

Heterogeneous Parallel Programming

COMP4901D

Host-Device Data Transfer

Overview

- CUDA memory copy between host and device
- Pinned memory
- CUDA streams
- Overlapping memory copy and computation

Why

- Bandwidth between GPU and device memory: 178GB/second on M2090
- Bandwidth between host memory and device memory: 8 GB/second on the PCIe x16 Gen2
- Host-device data transfer often the bottleneck in the overall performance

Simple Ways of Optimizing Host-Device Data Transfer

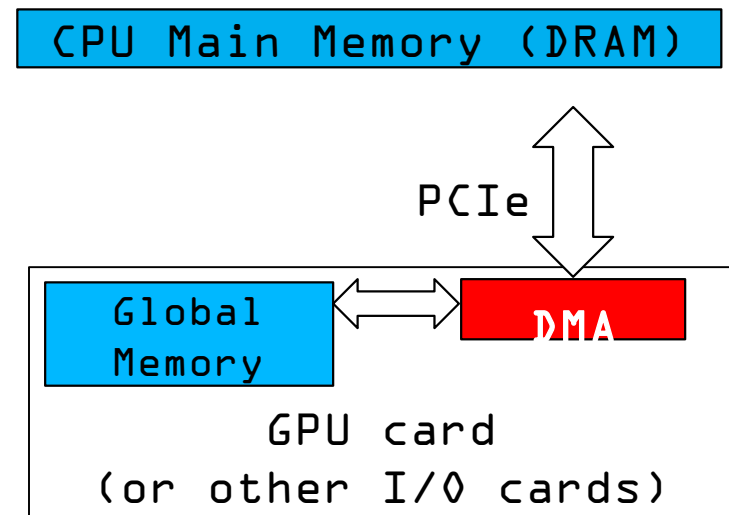
- Minimize transfers
 - Intermediate data can be allocated, operated on, and deallocated without ever copying to host memory
 - Run most of the code on the GPU
- Group transfers
 - One large transfer much better than many small ones; needs a large buffer for batching

More Advanced Ways of Optimizing Host-Device Data Transfer

- Pinned memory (page-locked memory)
- Asynchronous and overlapping transfers

Behind cudaMemcpy ()

- DMA (Direct Memory Access) hardware is used for cudaMemcpy() for efficiency
 - Frees CPU for other tasks
 - Transfers a number of bytes requested by OS
 - Uses PCI-e bus



Virtual Memory Management

- Modern computers use virtual memory management
 - Many virtual memory spaces mapped into a single physical memory
 - Virtual addresses (pointer values) are translated into physical addresses
 - Not all variables and data structures are always in the physical memory
 - Each virtual address space is divided into pages when mapped into physical memory
 - Memory pages can be paged out to make room
 - Whether a variable is in the physical memory is checked at address translation time

DMA and Virtual Memory

- DMA uses physical addresses
- When `cudaMemcpy()` copies an array, it is implemented as one or more DMA transfers
- Address is translated and page presence checked at the beginning of each DMA transfer
- No address translation for the rest of the same DMA transfer so that high efficiency can be achieved
- The OS could accidentally page-out the data that is being read or written by a DMA and page-in another virtual page into the same physical location

Pinned Memory and DMA

- Pinned memory, also called Page Locked Memory, Locked Pages, etc., are virtual memory pages that are specially marked so that they cannot be paged out
- Allocated with a special system API function call
- CPU memory that serve as the source or destination of a DMA transfer must be allocated as pinned memory

Data Transfer

Using Pinned Memory

- If a source or destination of a `cudaMemcpy()` in the host memory is not allocated in pinned memory, it needs to be first copied to a pinned memory – extra overhead
- `cudaMemcpy()` is faster if the host memory source or destination is allocated in pinned memory since no extra copy is needed

Allocate/Free Pinned Memory

- `cudaHostAlloc()`, three parameters
 - Address of pointer to the allocated memory
 - Size of the allocated memory in bytes
 - Option – use `cudaHostAllocDefault` for now
- `cudaFreeHost()`, one parameter
 - Pointer to the memory to be freed

Using Pinned Memory in CUDA

- Use the allocated pinned memory and its pointer the same way as those returned by `malloc()`;
- The only difference is that the allocated memory cannot be paged by the OS
- The `cudaMemcpy()` function is about 2X faster with pinned memory
- Pinned memory is a limited resource
 - over-subscription can have serious consequences


Vector Add Example with Pinned Memory

```
int main()
{
    float *h_A, *h_B, *h_C;
    ...
    cudaHostAlloc((void **) &h_A, N* sizeof(float), cudaHostAllocDefault);
    cudaHostAlloc((void **) &h_B, N* sizeof(float), cudaHostAllocDefault);
    cudaHostAlloc((void **) &h_C, N* sizeof(float), cudaHostAllocDefault);
    ...
    vecAdd(h_A, h_B, h_C, N);
}
```

Data Transfer and Computation

Serialized

- So far, the way we use cudaMemcpy serializes data transfer and GPU computation



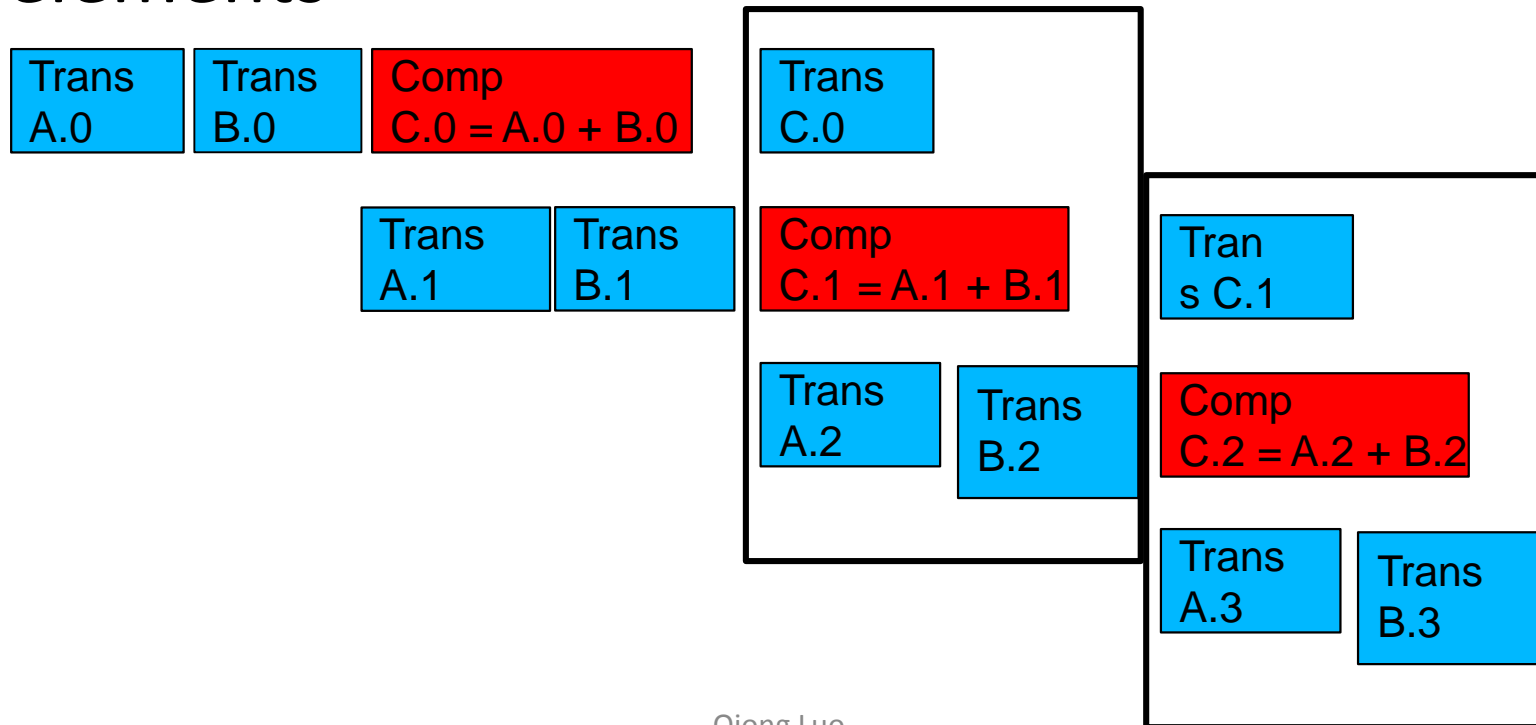
Device Capability of Overlapping Data Transfer and Execution

- Some CUDA devices support device overlap
 - Simultaneously execute a kernel while copying data between device and host memory

```
int dev_count; cudaDeviceProp prop;  
  
cudaGetDeviceCount( &dev_count);  
for (int i = 0; i < dev_count; i++) {  
    cudaGetDeviceProperties(&prop, i);  
    if (prop.deviceOverlap) ...
```

Ideal Timing Pipeline

- Divide large vectors into segments
- Overlap transfer and compute of adjacent elements




CUDA Streams

- CUDA supports parallel execution of kernels and Memcpy with “Streams”
 - Each stream is a queue of operations (kernel launches and Memcpy’s)
 - Operations (tasks) in different streams can go in parallel: “Task parallelism”


Queues Behind CUDA Streams

- Requests made from the host code are put into First-In-First-Out queues
- Queues are read and processed asynchronously by the driver and device
- Driver ensures that commands in a queue are processed in sequence, e.g., Memory copies end before kernel executes

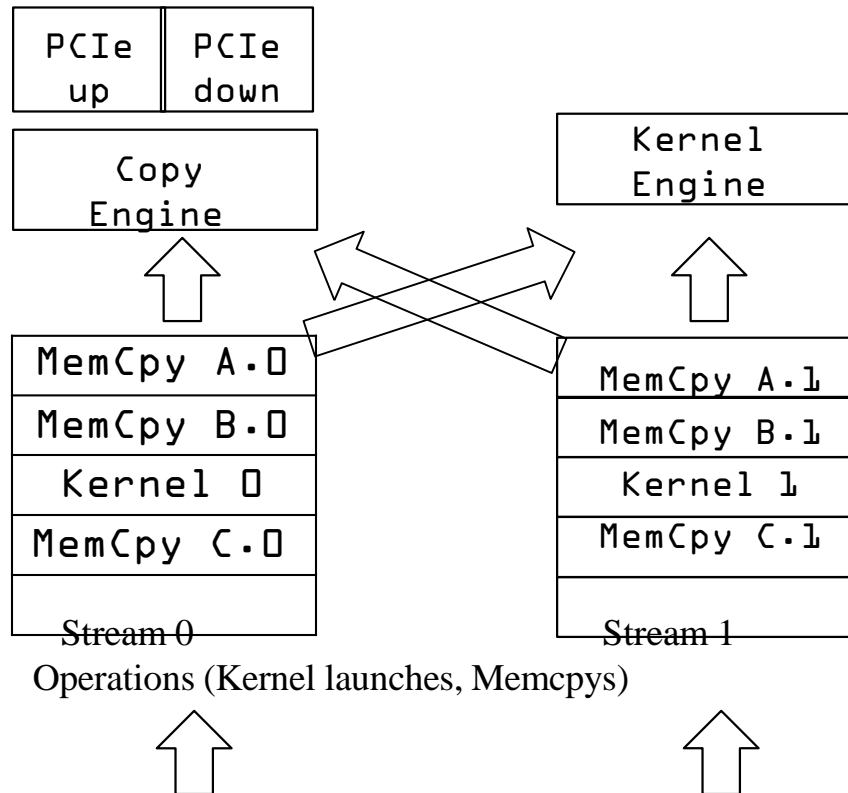


Multiple Streams for Concurrent Copying and Kernel Execution

- To allow concurrent copying and kernel execution, use multiple CUDA streams
 - CUDA “events” allow the host thread to query and synchronize with individual streams.



Conceptual View of Multiple Streams



Example of Multiple CUDA Streams

1.cudaStream_t stream0, stream1;

2.cudaStreamCreate(&stream0);

3.cudaStreamCreate(&stream1);

4.float *d_A0, *d_B0, *d_C0;// device memory for stream 0

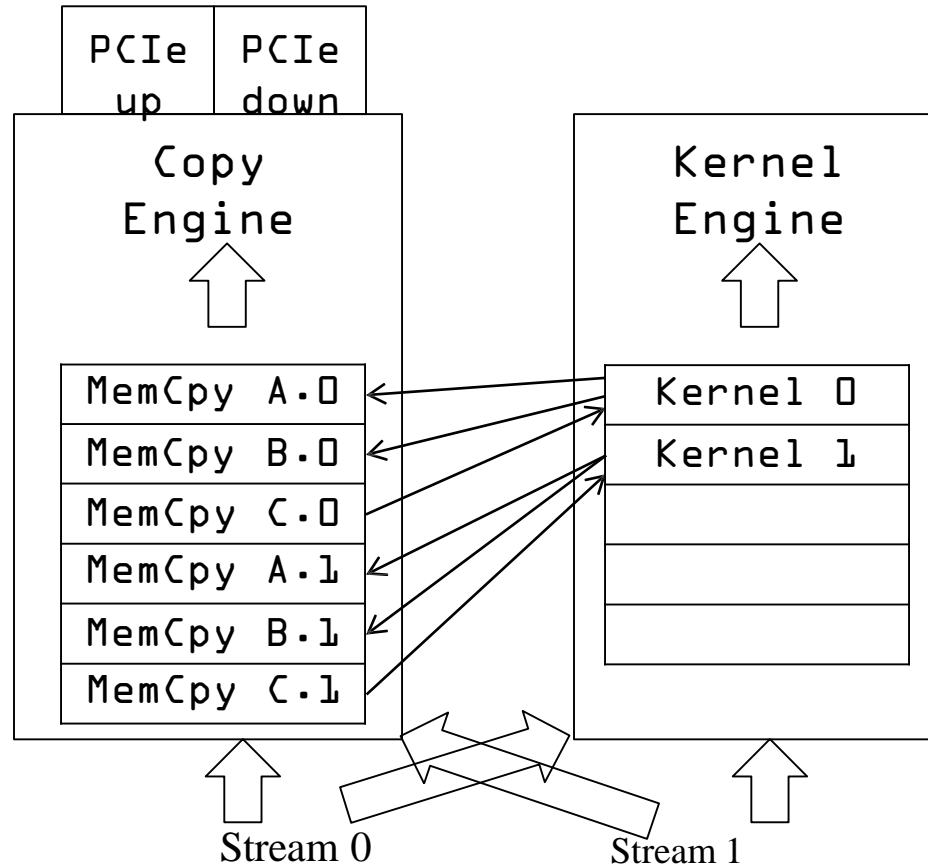
5.float *d_A1, *d_B1, *d_C1;// device memory for stream 1

// cudaMalloc for d_A0, d_B0, d_C0, d_A1, d_B1, d_C1 go here

Example of Multiple CUDA Streams (cont.)

```
6. for (int i=0; i<n; i+=SegSize*2) {  
7.   cudaMemcpyAsync(d_A0, h_A+i, SegSize*sizeof(float), ..., stream0);  
8.   cudaMemcpyAsync(d_B0, h_B+i, SegSize*sizeof(float), ..., stream0);  
9.   vecAdd<<<SegSize/256, 256, 0, stream0>>>(d_A0, d_B0, ...);  
10.  cudaMemcpyAsync(h_C+i, d_C0, SegSize*sizeof(float), ..., stream0);  
11.  cudaMemcpyAsync(d_A1, h_A+i+SegSize, SegSize*sizeof(float), ...,  
    stream1);  
12.  cudaMemcpyAsync(d_B1, h_B+i+SegSize, SegSize*sizeof(float) , ...,  
    stream1);  
13.  vecAdd<<<SegSize/256, 256, 0, stream1>>>(d_A1, d_B1, ...);  
14.  cudaMemcpyAsync(d_C1, h_C+i+SegSize, SegSize*sizeof(float), ...,  
    stream1);  
}
```


Conceptual View



Operations (Kernel launches, Memcpys)

We Want Finer Grain Concurrency

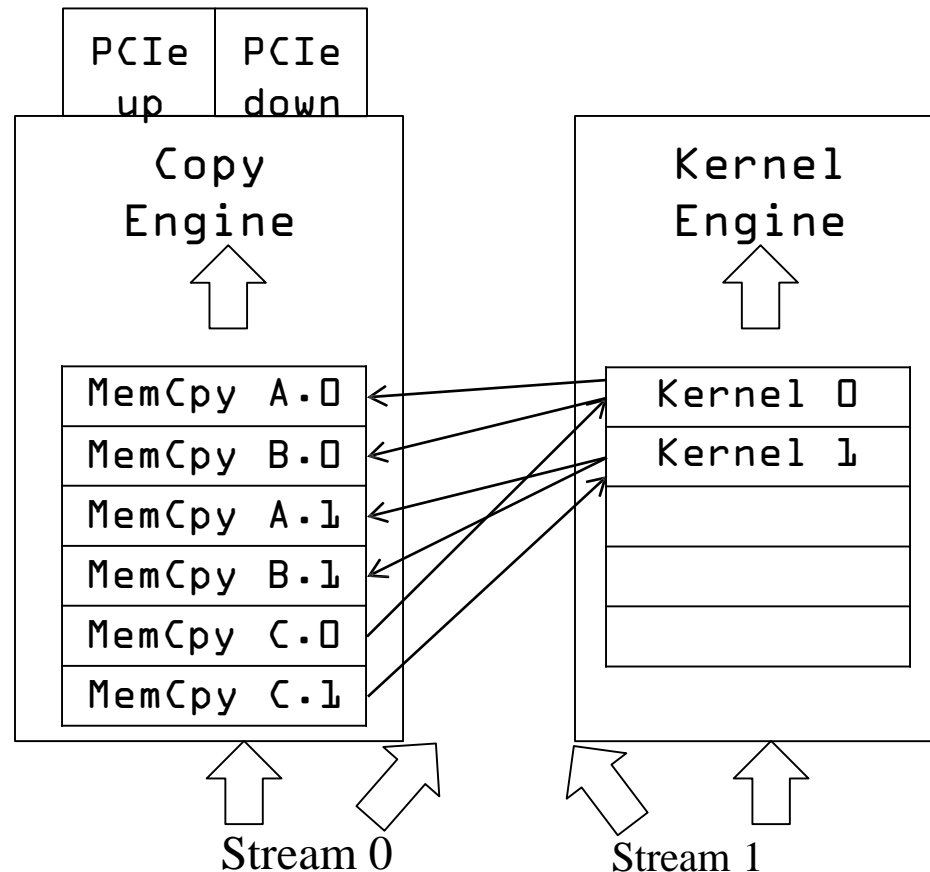
- C.0 blocks A.1 and B.1 in the copy engine queue



Better Concurrency with Streams

```
    for (int i=0; i<n; i+=SegSize*2) {  
        cudaMemcpyAsync(d_A0,  h_A+i,  SegSize*sizeof(float),...,  stream0);  
        cudaMemcpyAsync(d_B0,  h_B+i,  SegSize*sizeof(float),...,  stream0);  
        cudaMemcpyAsync(d_A1, h_A+i+SegSize, SegSize*sizeof(float),...,  
                        stream1);  
        cudaMemcpyAsync(d_B1, h_B+i+SegSize, SegSize*sizeof(float),...,  
                        stream1);  
  
        vecAdd<<<SegSize/256, 256, 0, stream0>>>(d_A0, d_B0, ...);  
        vecAdd<<<SegSize/256, 256, 0, stream1>>>(d_A1, d_B1, ...);  
  
        cudaMemcpyAsync(h_C+i,  d_C0,  SegSize*sizeof(float),...,  stream0);  
        cudaMemcpyAsync(h_C+i+SegSize, d_C1, SegSize*sizeof(float),...,  
                        stream1);  
    }
```

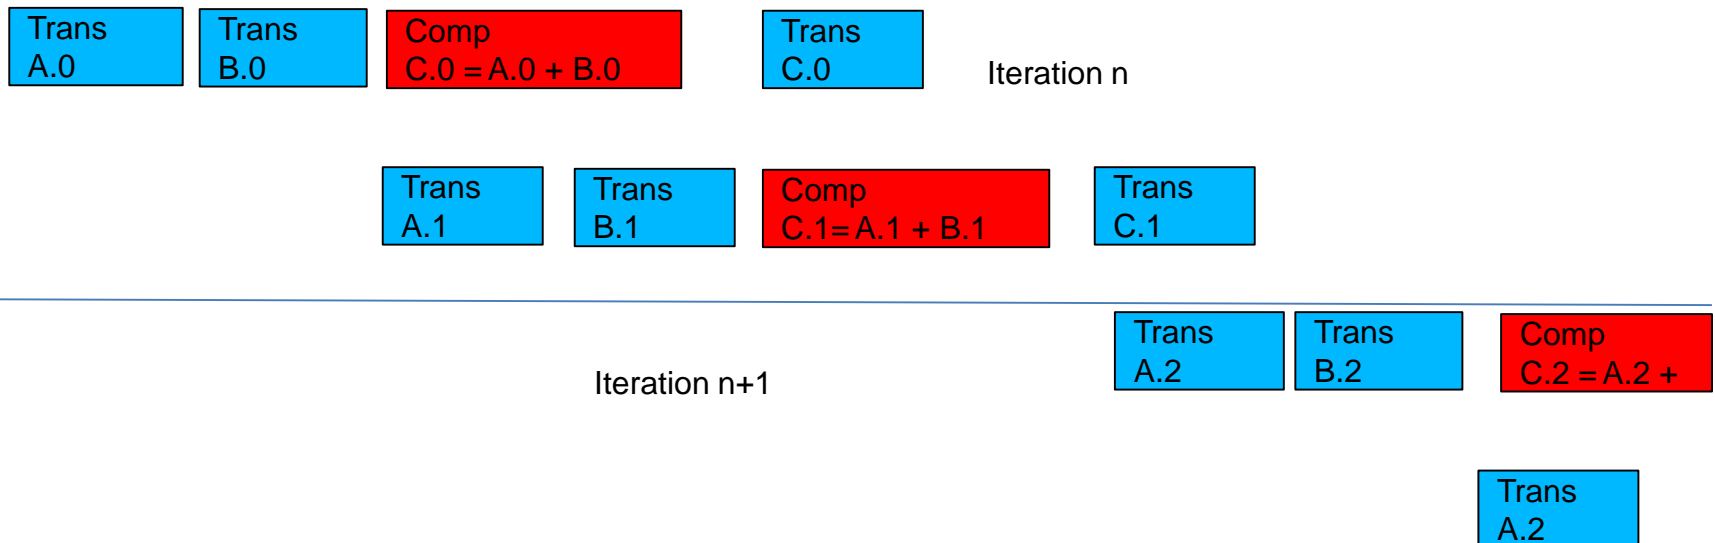
C.0 no longer blocks A.1 or B.1



Operations (Kernel launches, Memcpys)

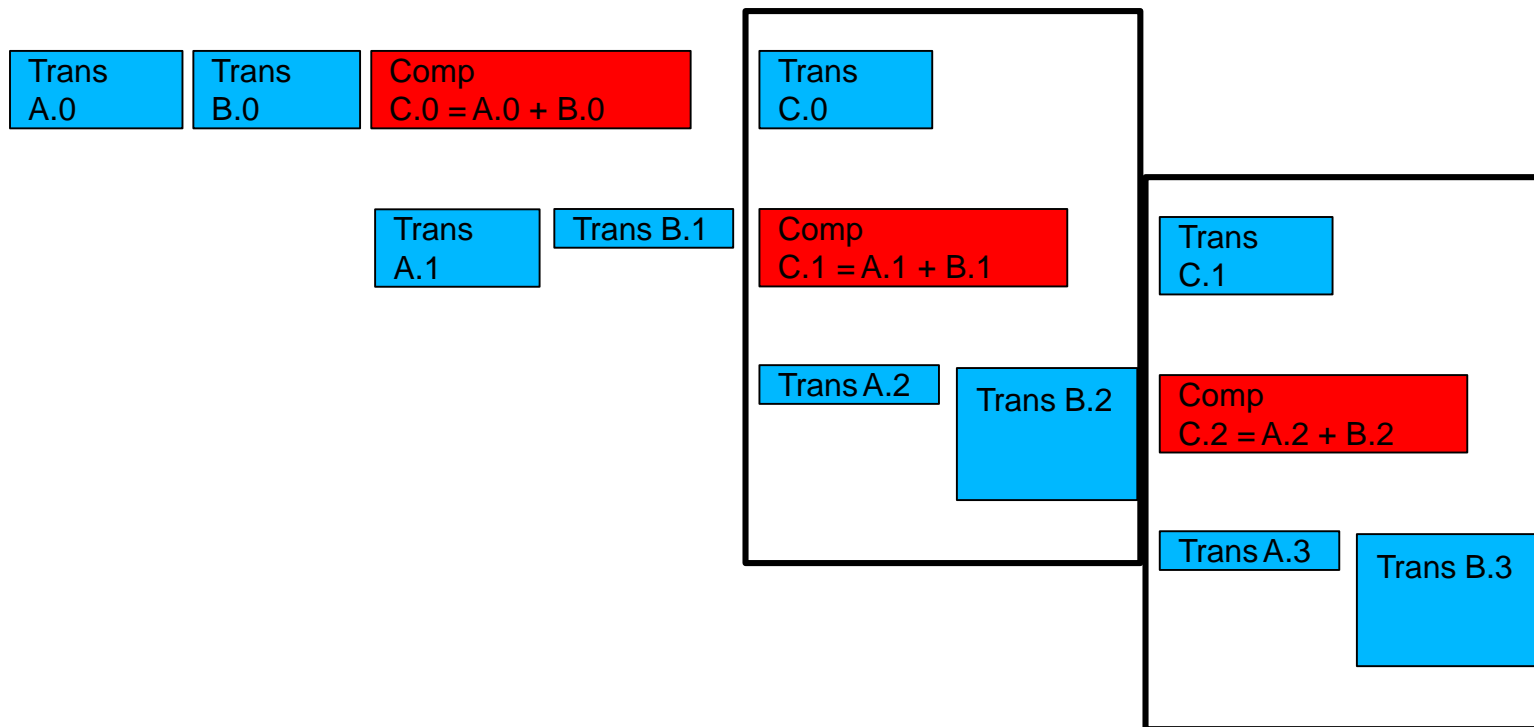
Not the Best Overlap Yet

- C.1 blocks the next iteration's A.0 and B.0 in the copy engine queue



Ideal, Pipelined Timing

- Will need at least three buffers for each original A, B, and C, code is more complicated



Wait until all tasks have completed

- `cudaStreamSynchronize(stream_id)`
 - Used in host code
 - Takes one parameter – stream identifier
 - Wait until all tasks in a stream have completed
- `cudaDeviceSynchronize()`
 - Also used in host code
 - No parameter
 - Wait until all tasks in all streams have completed for the current device

Summary

- Memory transfer between host and device is usually the performance bottleneck.
- General ways of optimizing memory transfer
 - Minimize transfer
 - Batch transfer
- Two more advanced optimizations in CUDA
 - Pinned memory
 - Overlapping transfer and computation