

Chapter 17: Protection



Chapter 17: Protection

- Goals of Protection
- Principles of Protection
- Protection Rings
- Domain of Protection
- Access Matrix
- Implementation of Access Matrix



Objectives

- Discuss the goals and principles of protection in modern computer systems
- Explain how protection domains combined with an access matrix are used to specify the resources that a process may access
- Examine capability-based protection system



Goals of Protection

- In a protection model, computer system consists of a collection of **objects**, hardware or software
 - Hardware objects: CPU, memory segments, printers, disks, and tape
 - Software objects: files, programs, and semaphores
- Each object has a unique name and can be accessed through a **well-defined set of operations**
- **Protection problem** is to ensure that each object is accessed correctly and only by those processes allowed to do so
- **Mechanisms** are distinct from **policies**, in which mechanisms determine how something will be done and policies decide what will be done.
 - The separation is important for flexibility, as policies are likely to change from place to place or from time to time.
 - The separation ensures that not every change in policy would require a change in the underlying mechanism.



Principles of Protection

- The guiding principle – **principle of least privilege**
 - Programs, users and systems should be given just enough privileges to perform their tasks - mitigate the attack
 - In file permissions, this principle dictates that a user have read access but not write or execute access to a file. **The principle of least privilege** would require that the OS provides a mechanism to only allow read access but not write or execute access
- Properly set **permissions** (i.e., the access rights to an object) can limit damage if entity has a bug or gets abused



Principles of Protection (Cont.)


- **Audit trail** – the collection of activities in a log for monitoring review
 - An audit trail is a record in the system logs
 - It can reveal early warnings of an attack or provide clues as to which attack were used and assess the damage caused
- No single principle is a panacea for security vulnerabilities – need **defense in depth**
 - The theory states more layers of defense provide stronger defense than fewer layers





Protection Rings

- **User mode** and **kernel mode** – **privilege separation**
- Hardware support required to support the notion of separate execution
- Let D_i and D_j be any two domain rings
- If $j < i \Rightarrow D_i \subseteq D_j$
- The innermost ring, ring 0, provides the full set of privileges

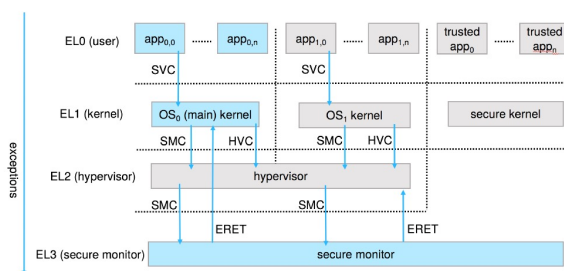


Protection Rings (Cont.)

- Components ordered by amount of privilege and protected from each other
 - For example, the kernel is in one ring and user applications in another
 - This privilege separation requires hardware support
 - “Gates” used to transfer between rings, for example the **syscall** Intel instruction, also **traps** and **interrupts**
- **Hypervisors** (Intel) is introduced (another ring) - virtual machine managers, which create and run virtual machines, and have more capabilities than the kernels of the guest operating systems
- ARM processors added **TrustZone** or **TZ** ring to protect crypto functions with access (more privileged than kernel)
 - This most privileged execution environment has exclusive access to hardware-backed cryptographic features, such as the NFC Secure Element and an on-chip cryptographic key, that make handling passwords and sensitive information more secure.



ARM CPU Architecture



Domain of Protection

- Protection rings separate functions into different domains and order them hierarchically
- **Domain** can be considered as a generalization of rings without a hierarchy
- A computer system can be treated as processes and objects
 - **Hardware objects** (such as CPU, memory, disk) and **software objects** (such as files, programs, semaphores)
- Process for example should only have access to objects it currently requires to complete its task – the **need-to-know** principle (policy)
- Implementation can be via process operating in a **protection domain**
 - Protection domain specifies the set of resources a process may access
 - Each domain specifies set of objects and types of operations may be invoked on each object



Domain of Protection (Cont.)

- Ability to execute an operation on an object is an **access right**
- A domain is a collection of access rights, each of which is an ordered pair <object-name, rights-set>
 - An example: if domain D has the access right <file F, {read,write}>, then a process executing in domain D can both read and write file F. It cannot, however, perform any other operation on that object.
- Domains may share access rights
- Associations between processes and domains can be **static** if the set of resources available to the process is fixed throughout the process's lifetime, or can be **dynamic**
- If dynamic, a mechanism is available to allow **domain switching**, enabling the process to switch from one domain to another during different stage of execution



Domain of Protection (Cont.)

Domain can be realized in a variety of ways:


- Each **user** may be a domain - the set of objects that can be accessed depends on the identity of the user. Domain switching occurs when the user is changed
- Each **process** may be a domain - the set of objects that can be accessed depends on the identity of the process. Domain switching occurs when one process sends a message to another process and then waits for a response.
- Each **procedure** may be a domain - the set of objects that can be accessed corresponds to the local variables defined within the procedure. Domain switching occurs when a procedure call is made





Domain Structure

- Access-right = $\langle \text{object-name, rights-set} \rangle$
where rights-set is a subset of all valid operations that can be performed on the object
- Domain = set of access-rights
- The access right $\langle O_4, \{\text{print}\} \rangle$ shared by domains D_2 and D_3 , thus, a process executing in either of these two domains can print object O_4 .



Access Matrix

- View protection as a matrix (**access matrix**)
- Rows** represent domains, and **columns** represent objects
- Access(i,j)** consists of a set of access rights - the set of operations that a process executing in Domain_i can invoke on Object_j

object domain	F ₁	F ₂	F ₃	printer
D ₁	read		read	
D ₂				print
D ₃		read	execute	
D ₄	read write		read write	



Use of Access Matrix

- The **access matrix** scheme provides the mechanism for specifying a variety of policies - mechanism and policy separation
- The **mechanism** consists of implementing the access matrix and ensuring that the semantic properties hold.
 - To ensure that a process executing in domain D_i can access only those objects specified in row i .
- The **policy** decisions specify which rights should be included in the (i,j) th entry, and determine the domain in which each process executes
- If a process in Domain D_i tries to do "op" on object O_j , then "op" must be in the access matrix
- User who creates an object can define access column for that object
 - When a user creates a new object O_j , the column O_j is added to the access matrix with the appropriate initialization entries, as dictated by the creator. The user may decide to enter some rights in some entries in column j and other rights in other entries, as needed.



Use of Access Matrix (Cont.)

This can be expanded to dynamic protection

- Operations to add, delete access rights
- Special access rights:
 - owner** of O_i - can add and remove any right in any entry in column
 - copy** op from O_i to O_j (denoted by "**") - only within the column (that is, for the object)
 - control** - D_i can modify D_j access rights - modify domain objects (a row)
 - transfer** - switch from domain D_i to D_j
- Copy** and **Owner** applicable to an object - change the entries in a column
- Control** applicable to domain object - change the entries in a row
- New objects and new domains can be created dynamically and included in the access-matrix model
- In a **dynamic** protection system, we may sometimes need to revoke access rights to objects shared by different users - **revocation** of access right



Access Matrix of Figure A with Domains as Objects

object domain	F ₁	F ₂	F ₃	laser printer	D ₁	D ₂	D ₃	D ₄
D ₁	read		read			switch		
D ₂				print			switch	switch
D ₃		read	execute					
D ₄	read write		read write		switch			



Access Matrix with Copy Rights

object domain	F ₁	F ₂	F ₃
D ₁	execute		write*
D ₂	execute	read*	execute
D ₃	execute	read	

(a)

object domain	F ₁	F ₂	F ₃
D ₁	execute		write*
D ₂	execute	read*	execute
D ₃	execute	read	

(b)





Access Matrix With Owner Rights

object domain	F_1	F_2	F_3
D_1	owner execute		write
D_2		read* owner	read* owner write
D_3	execute		

(a)

object domain	F_1	F_2	F_3
D_1	owner execute		write
D_2		owner read* write*	read* owner write
D_3		write	write

(b)



Modified Access Matrix of Figure B

object domain	F_1	F_2	F_3	laser printer	D_1	D_2	D_3	D_4
D_1	read		read			switch		
D_2				print			switch	switch control
D_3		read	execute					
D_4	write		write		switch			



Implementation of Access Matrix

- In general, the access matrix is sparse; that is, most of the entries will be empty
- Option 1 – Global Table**
 - Store ordered triples <domain, object, rights-set> in table
 - A requested operation M on object O_j within domain D_i -> search table for < D_i , O_j , R_k > with $M \in R_k$
 - But the table could be large -> might not fit in main memory, requires additional I/O – virtual memory techniques are often used
 - Difficult to group objects - For example, if everyone can read a particular object, this object must have a separate entry in every domain.



Implementation of Access Matrix (Cont.)

- Each column = **Access-control list** for one object
Defines who can perform what operation
 - Domain 1 = Read, Write
 - Domain 2 = Read
 - Domain 3 = Read
- Each row = **Capability List** (like a key)
For each domain, what operations allowed on what objects
 - Object F1 – Read
 - Object F4 – Read, Write, Execute
 - Object F5 – Read, Write, Delete, Copy



Implementation of Access Matrix (Cont.)

- Option 2 – Access lists for objects**
 - Each column implemented as an access list for one object
 - Resulting per-object list consists of ordered pairs <domain, rights-set> defining all domains with non-empty set of access rights for the object
 - Obviously, the empty entries can be discarded.
 - This can be easily extended to define **default** set of access rights -> If $M \in$ default set, also allow access (for all domains)



Implementation of Access Matrix (Cont.)

- Option 3 – Capability list for domains**
 - Instead of object-based, list is domain-based
 - A **capability list** for domain is a list of objects together with operations allowed on them
 - An object represented by its name or address, called a **capability**
 - To execute operation M on object O_j , a process requests operation M, specifying the capability (or pointer) for object O_j as a parameter
 - Possession of capability means access is allowed
- Capability list associated with a domain, but never directly accessible by a process executing in that domain
 - Rather, the capability list itself is a **protected object**, maintained by OS and accessed by users only indirectly
 - This avoids the possibility of capability list modification by users
 - If all capabilities are secure, the object they protect is also secure against unauthorized access





Implementation of Access Matrix (Cont.)

■ Option 4 – Lock-key

- Compromise between access lists and capability lists
- Each object has list of unique bit patterns, called **locks**
- Each domain as list of unique bit patterns called **keys**
- Process in a domain can only access object if domain has key that matches one of the locks of the object
- As with capability lists, the list of keys for a domain must be managed by the operating system on behalf of the domain.
- Users are not allowed to examine or modify the list of keys (or locks) directly.



Comparison of Implementations

Choosing a technique for implementing an access matrix involves various trade-offs.

- **Global table** is simple, but large, lack of grouping of objects or domains
- **Access lists** correspond directly to the needs of users
 - An access list on an object is specified when a user creates the object
 - Determining set of access rights for each domain is difficult - every access to the object must be checked, requiring a search of the access list.
- **Capability lists** useful for localizing information for a given process
 - But revocation capabilities can be inefficient
- **Lock-key** can be effective and flexible depending on the length of the keys
 - Keys can be passed freely from domain to domain, easy revocation
- Most systems use combination of access lists and capabilities



End of Chapter 17

