



# Model Selection

COMP4211



# Example



which one is the best?

# Training Error

given:

- target function  $f$
- hypothesis/model  $h$  (e.g., a particular neural network)
- distribution  $\mathcal{D}$  of the instances
- training set  $S$  (of size  $n$ ) drawn from  $\mathcal{D}$

training error

- proportion of examples in  $S$  that  $h$  misclassifies


$$error_S(h) \equiv \frac{1}{n} \sum_{x \in S}^n \delta(f(x) \neq h(x))$$

- $\delta(f(x) \neq h(x))$  is 1 if  $f(x) \neq h(x)$ , and 0 otherwise
- what we can measure

# How to Select the “Best” Hypothesis?

use how many hidden units / hidden layers?

use the network that minimizes the training error?



what we want to obtain is a model with low **testing error**

# Testing Error

## testing error

- probability that  $h$  will misclassify an instance drawn at random according to  $\mathcal{D}$


$$error_{\mathcal{D}}(h) \equiv Pr_{x \in \mathcal{D}}[f(x) \neq h(x)]$$

typically, use a test set to estimate the testing error


$$error_{test}(h) \equiv \frac{1}{|testSet|} \sum_{x \in testSet}^n \delta(f(x) \neq h(x))$$

- the test patterns should be drawn independently from the training patterns
- the testing data should not be used in any way to learn the hypothesis


# How to Select the “Best” Hypothesis?...



- intuitively, the model should be not too simple (too few hidden units) nor too complex (too many hidden units)



# Overfitting



hypothesis  $h \in H$  **overfits** the training data if there is an alternative hypothesis  $h' \in H$  such that

- $h$  has **smaller error** than  $h'$  over the **training examples**, but
- $h$  has a **larger error** than  $h'$  over the **entire distribution** of instances

# How to Select the “Best” Hypothesis?...

- the test data should **not** be used!
- partition the **available** (training) **data** into two sets
  - training set: used to form the learned hypothesis
  - **validation set**: used to estimate the accuracy of this hypothesis over subsequent data



- once the evaluation is complete, all the data can be used to train the final hypothesis (optional)
- generally,
  - the larger the training set, the better the hypothesis
  - the larger the validation set, the more accurate is error estimation

## Example


withhold one-third of the available examples for the validation set, using the other two-thirds for training



# How to make Error Estimation more reliable?

repeat the process with different subsamples

# $k$ -fold Cross-validation



- 1 the  $m$  available examples are partitioned into  $k$  **disjoint** subsets, each of size  $m/k$
- 2 the learning procedure is then run  $k$  times, each time
  - using one of these subsets as the **validation set**, and
  - combining the other subsets for the **training set**
- 3 **average** the performance on the validation sets over the  $k$  runs

# $k$ -fold Cross-validation...

each example is used

- in the validation set once, and
- in the training set for the other  $k - 1$  folds


a high proportion of the available data ( $1 - \frac{1}{k}$ ) is used for training, while making use of all the data in computing the error

how many times do we need to perform training?

- e.g.,  $k \sim 10$

leave-one-out cross-validation ( $k = m$ )

- train on  $m - 1$  examples and validate on 1 example
- useful for small data sets



# $k$ -fold Stratified Cross-validation

another problem:

- examples in the training set of each fold may not be representative
- e.g., all the examples of a certain class are missing
- $\rightarrow$  the classifier cannot learn to predict this class

how to ensure that each class is represented with **approximately equal proportions** in both the training and validation sets?

- partition the  $m$  examples into  $k$  folds such that each class is **uniformly distributed** among the  $k$  folds
- the class distribution in each fold is similar to that in the original data set
- e.g., 10-fold (stratified) cross-validation