


# Decision Tree

# Decision Tree




- each **internal node** denotes a test on an attribute
- each **branch** represents an outcome of the test
- each **leaf node** holds a class label

# Prediction

- given a new (unseen) tuple: the associated class is unknown

## Example

$\langle \text{Name, age, student, credit\_rating, buys\_computer} \rangle = \langle \text{Name, age, student, credit\_rating, buys\_computer} \rangle =$   
 $\langle \text{Stavros, youth, no, excellent, ?} \rangle \qquad \qquad \langle \text{Stavros, youth, no, excellent, no} \rangle$



- test the attribute values of the tuple against the decision tree
  - start from the root and trace a **path** to a leaf node (top-down), based on the attribute values of the tuple
- the class value included in this leaf is assigned to the tuple

# Building a Decision Tree

- there are several popular decision tree **algorithms**, including **ID3**, **C4.5** and **CART**
- we will describe a **general greedy framework** followed by the majority of the algorithms
- next, we will discuss some components of this framework (which essentially differentiate the various algorithms) in more detail, as well as some other additional issues on decision tree induction

# General Algorithm *Decision\_Tree\_Induction*

- top-down and recursive

## Input

- node  $N$ 
  - the first time the algorithm is called,  $N$  is the root of the tree
- dataset  $D$  of training tuples
  - initially the entire training set
- *attribute\_list*: holds the set of attributes
  - initially the attributes that remained after data preprocessing
- *attribute\_selection\_method*: a heuristic process for selecting the attribute that “best” discriminates the given tuples according to class

## output

- decision tree

# Example

<i>RID</i>	<i>age</i>	<i>income</i>	<i>student</i>	<i>credit_rating</i>	<i>class: buys_computer</i>
1	youth	high	no	fair	no
2	youth	high	no	excellent	no
3	middle_aged	high	no	fair	yes
4	senior	medium	no	fair	yes
5	senior	low	yes	fair	yes
6	senior	low	yes	excellent	no
7	middle_aged	low	yes	excellent	yes
8	youth	medium	no	fair	no
9	youth	low	yes	fair	yes
10	senior	medium	yes	fair	yes
11	youth	medium	yes	excellent	yes
12	middle_aged	medium	no	excellent	yes
13	middle_aged	high	yes	fair	yes
14	senior	medium	no	excellent	no

# Steps

## Step 1


- associate node  $N$  with dataset  $D$
- trivial; meant to emphasize that each decision tree node represents a subset of the original (entire) training set

## Step 2

- end this process if one of the terminating conditions is satisfied (will be discussed soon)

### Step 1

The algorithm is called with the initial single root node  $N$ , the entire training set as  $D$ , {age, income, student, credit\_rating} as the *attribute\_list*, and a certain *attribute\_selection\_method*



### Step 2

No terminating condition  
is satisfied yet

# Step 3

Call *attribute\_selection\_method*

- use a **splitting criterion** to select an attribute to test at node  $N$ 
  - try to “best” partition  $D$  into subsets, such that each subset is as “pure” as possible
  - a subset of  $D$  is **pure**, if it contains tuples belonging to the **same** class
- this test will result in creating new nodes (as children of  $N$ ), each of which representing a subset of  $D$

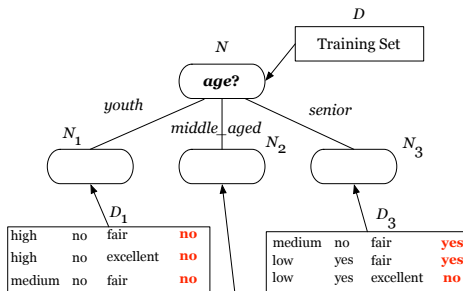
## Step 3

Suppose that the *age* is selected as the splitting attribute by the *attribute\_selection\_method*

## Step 4

Three branches and children are created for  $N$

Also  $D$  is partitioned into three datasets  $D_1$ ,  $D_2$  and  $D_3$  according to the *age* attribute values of







# Step 4

- a branch for each of the outcomes of the splitting criterion
- a new node  $N_i$  is created for each branch  $i$
- if the number of new nodes is  $m$ ,  $D$  is partitioned accordingly into  $m$  subsets  $D_1, D_2, \dots, D_m$ 
  - each  $D_i$  contains the tuples that satisfy the splitting criterion outcome of branch  $i$


**A is discrete-valued**




**Example**



**A is continuous-valued**




**Example**



**A is discrete-valued and binary**



**Example**



# Example...

## Step 3

Suppose that the *age* is selected as the splitting attribute by the *attribute\_selection\_method*

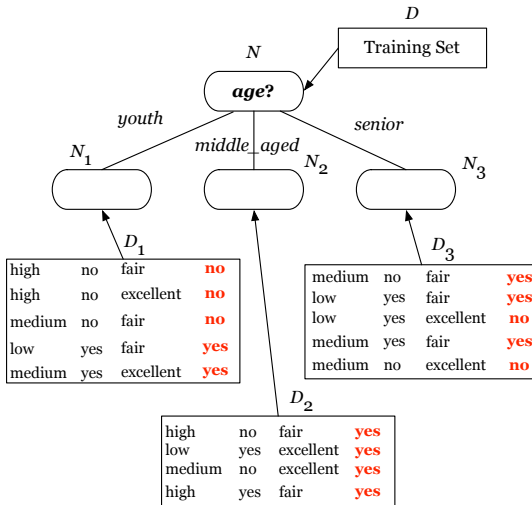
## Step 4

Three branches and children are created for *N*

Also *D* is partitioned into three datasets *D*<sub>1</sub>, *D*<sub>2</sub> and *D*<sub>3</sub> according to the *age* attribute values of the tuples

## Step 5

The algorithm is recursively called for (*N*<sub>1</sub>, *D*<sub>1</sub>), (*N*<sub>2</sub>, *D*<sub>2</sub>) and (*N*<sub>3</sub>, *D*<sub>3</sub>), after removing *age* from *attribute\_list*



The tuples in *D*<sub>1</sub>, *D*<sub>2</sub> and *D*<sub>3</sub> are of the form  
**<income, student, credit\_rating, class>**

## Step 5

- remove the splitting attribute  $A$  from *attribute\_list*
- call *Decision\_Tree\_Induction*( $N_i$ ,  $D_i$ , *attribute\_list*, *attribute\_selection\_method*) recursively for every newly created ( $N_i$ ,  $D_i$ ) pair

# Terminating Conditions


- all of the tuples in partition  $D$  belong to the same class
  - in this case,  $N$  becomes a **leaf** and is labeled with that class

## Example

In the procedure call for  $(N_2, D_2)$

### Step 2

Since all the tuples in  $D_2$  have the same class, the procedure terminates after making  $N_2$  a leaf node storing the class label ("yes")





- there are no remaining attributes in *attribute\_list* to help partitioning the tuples of  $D$  further
  - $N$  becomes a **leaf** and is labeled with the **majority class** in  $D$
- $D$  is empty
  - $N$  becomes a **leaf** and is labeled with the **majority class** of its **parent's dataset**

# Attribute Selection Measures

- ideally, the best splitting criterion is the one that decomposes  $D$  into subsets having only tuples of a single class (these subsets are called **pure**)
- since it may not be always possible to select a splitting criterion that derives only pure subsets, the attribute selection measure provides a **ranking** for each attribute
- the attribute with the **best score** is selected as the splitting attribute
- we will study three attribute selection measures
  - **Information Gain** (used in ID3)
  - **Gain Ratio** (used in C4.5)
  - **Gini Index** (used in CART)

# Which Attribute is Best?





# Example

I am thinking of an integer between 1 and 1,000 – what is it?  
What is the first question you would ask?

- “Is it 752?”, or
- “Is it a prime number between 123 and 239?”, or
- “Is it between 1 and 500?”

Which answer provides the most information?

# Idea




Want attributes that split examples into sets that are relatively **pure** in one label

How close is a set of instances to having just one label?



# Entropy: Intuitive Notion

- measures the impurity, uncertainty, irregularity, surprise
- suppose we have two discrete classes (in general, can have  $> 2$ )
  - $D$ : a sample of training examples
  - $p_{\oplus}$ : proportion of positive examples in  $D$
  - $p_{\ominus}$ : proportion of negative examples in  $D$
- **optimal purity** (impurity/uncertainty = 0): either
  - $p_{\oplus} = 1, p_{\ominus} = 0$
  - $p_{\oplus} = 0, p_{\ominus} = 1$
- **least pure** (maximum impurity/uncertainty):
  - $p_{\oplus} = 0.5, p_{\ominus} = 0.5$



# Entropy: Definition

$$\text{Entropy}(D) \equiv -p_{\oplus} \log_b p_{\oplus} - p_{\ominus} \log_b p_{\ominus}$$

- if  $p_{\oplus} = 0$ , take  $p_{\oplus} \log_b p_{\oplus} = 0$

What units is entropy measured in?

Depends on the base  $b$  of the log

- $b = e$ : **nats**
- $b = 2$ : **bits** (adopted here)

## Example

Entropy of a fair coin = 1 bit

## Example

$D$  is a collection of 14 examples, 9 positive and 5 negative

$$\text{Entropy}([9+, 5-]) = -\frac{9}{14} \log_2 \frac{9}{14} - \frac{5}{14} \log_2 \frac{5}{14} = 0.94$$

# Properties


All members of  $D$  belong to the same class

- $\rightarrow$  entropy = 0

$D$  has an equal number of +ve and -ve examples

- $\rightarrow$  entropy = 1

Otherwise, entropy is between 0 and 1





In general, if the target attribute can take on  $m$  different values

$$Entropy(D) = \sum_{i=1}^m -p_i \log_2 p_i$$

- entropy is between 0 and  $\log_2 m$

# Information Gain



What is the **uncertainty removed** by splitting on the value of  $A$ ?

## Definition

The **information gain** of  $D$  relative to attribute  $A$  is the expected **reduction** in entropy caused by knowing the value of  $A$

- $D_v$ : the set of examples in  $D$  where attribute  $A$  has value  $v$

$$Gain(D, A) \equiv Entropy(D) - \sum_{v \in Values(A)} \frac{|D_v|}{|D|} Entropy(D_v)$$

# Example

<i>RID</i>	<i>age</i>	<i>income</i>	<i>student</i>	<i>credit_rating</i>	<i>class: buys_computer</i>
1	youth	high	no	fair	no
2	youth	high	no	excellent	no
3	middle_aged	high	no	fair	yes
4	senior	medium	no	fair	yes
5	senior	low	yes	fair	yes
6	senior	low	yes	excellent	no
7	middle_aged	low	yes	excellent	yes
8	youth	medium	no	fair	no
9	youth	low	yes	fair	yes
10	senior	medium	yes	fair	yes
11	youth	medium	yes	excellent	yes
12	middle_aged	medium	no	excellent	yes
13	middle_aged	high	yes	fair	yes
14	senior	medium	no	excellent	no

●  $\text{entropy}(D) = 0.940$

● we have to compute  $\text{Gain}(A)$  for every attribute  $A$

# Example

- let us first focus on  $A = age$

$$entropy(D_{\text{youth}}) = -\frac{2}{5} \log_2 \frac{2}{5} - \frac{3}{5} \log_2 \frac{3}{5}$$

$$entropy(D_{\text{middle\_aged}}) = -\frac{4}{4} \log_2 \frac{4}{4}$$

$$entropy(D_{\text{senior}}) = -\frac{3}{5} \log_2 \frac{3}{5} - \frac{2}{5} \log_2 \frac{2}{5}$$

$$\begin{aligned} \text{weighted average} &= \frac{5}{14} \times entropy(D_{\text{youth}}) + \frac{4}{14} \times entropy(D_{\text{middle\_aged}}) \\ &\quad + \frac{5}{14} \times entropy(D_{\text{senior}}) \\ &= 0.694 \text{ bits} \end{aligned}$$

- information gain for  $age$ :

$$Gain(D, age) = 0.940 - 0.694 = 0.246 \text{ bits}$$

- Similarly,  $Gain(D, income) = 0.029$ ,

$$Gain(D, student) = 0.151, \text{ and}$$

$$Gain(D, credit\_rating) = 0.048$$

# Continuous-Valued Attributes

## Example

<i>Temperature:</i>	40	48	60	72	80	90
<i>PlayTennis:</i>	No	No	Yes	Yes	Yes	No

**Discretize** the continuous-valued attributes

- Create a **boolean** attribute  $A_c$  that is true if  $A < c$  and false otherwise

How to pick the threshold  $c$ ?

# Continuous-Valued Attributes

## Example

<i>Temperature:</i>	40	48	60	72	80	90
<i>PlayTennis:</i>	No	No	Yes	Yes	Yes	No

- 1 **sort** the attribute values in the training set ( $\{v_1, v_2, \dots, v_m\}$ )
- 2 generate a set of candidate thresholds **midway** between these values
  - cut at 44, 54, 66, 76, 85
  - there are thus  $m - 1$  candidate thresholds
- 3 evaluate these candidate thresholds by the information gain



# Example

## Example

<i>Temperature:</i>	40	48	60	72	80	90
<i>PlayTennis:</i>	No	No	Yes	Yes	Yes	No

6 items, total weight 6.0

Att A1


- Cut at 44.000 (gain 0.191):
- Cut at 54.000 (gain 0.459):
- Cut at 66.000 (gain 0.082):
- Cut at 76.000 (gain 0.000):
- Cut at 85.000 (gain 0.191):

Best cut: 54.000

# Attribute Selection Measures: Gain Ratio

## Example

- if we try to split according to a unique identifier (e.g., *product\_id*)



- there will be as many partitions as the number of tuples, and each partition will have a single tuple
- entropy after splitting = 0, which means that the gain is maximized, and *product\_id* is chosen as the splitting criterion
- obviously, such a partitioning is useless for classification

Information gain is **biased** towards tests with many outcomes

# Gain Ratio

- **gain ratio** is an extension of information gain, where a large number of partitions is **penalized**
- penalty should be
  - large when data is evenly spread
  - small when all data belong to one branch
- first compute  $SplitInfo_A(D)$ , which measures the **entropy** of the **partitioning** according to the  $v$  distinct values of  $A$ :

$$SplitInfo_A(D) = - \sum_{j=1}^v \frac{|D_j|}{|D|} \times \log_2 \frac{|D_j|}{|D|}$$

- gain ratio is:

$$GainRatio(A) = \frac{Gain(A)}{SplitInfo_A(D)}$$

- we select as the splitting criterion the  $A$  that leads to the **highest**  $GainRatio(A)$  value

# Example

RID	age	income	student	credit-rating	class:buy_computer
1	youth	high	no	fair	no
2	youth	high	no	excellent	no
3	middle-aged	high	no	fair	yes
4	senior	medium	no	fair	yes
5	senior	low	yes	fair	yes
6	senior	low	yes	excellent	no
7	middle-aged	low	yes	excellent	yes
8	youth	medium	no	fair	no
9	youth	low	yes	fair	yes
10	senior	medium	yes	fair	yes
11	youth	medium	yes	excellent	yes
12	middle-aged	medium	no	excellent	yes
13	middle-aged	high	yes	fair	yes
14	senior	medium	no	excellent	no

## Using attribute income

1<sup>st</sup> partition (low) **D1** has **4 tuples**

2<sup>nd</sup> partition (medium) **D2** has **6 tuples**

3<sup>rd</sup> partition (high) **D3** has **4 tuples**

$$Gain(income) = 0.029$$

$$GainRatio(income) = \frac{0.029}{0.926} = 0.031$$

$$SplitInfo_{income}(D) = -\frac{4}{14} \log_2 \left( \frac{4}{14} \right) - \frac{6}{14} \log_2 \left( \frac{6}{14} \right) - \frac{4}{14} \log_2 \left( \frac{4}{14} \right)$$

$$= 0.926$$

# Attribute Selection Measures: Gini Index

- suppose that the class attribute of  $D$  has  $m$  distinct class values  $C_1, C_2, \dots, C_m$
- $|D|$ : cardinality of  $D$ ,  $|C_i|$ : number of tuples in  $D$  having class label  $C_i$
- probability  $p_i$  that a tuple of  $D$  belongs to class  $C_i: |C_i|/|D|$
- Gini index

$$Gini(D) = 1 - \sum_{i=1}^m p_i^2$$

- $Gini(D)$  is small if most of the tuples belong to a few classes

# Gini Index...

- if data set  $D$  is split on  $A$  into two subsets  $D_1$  and  $D_2$ , the gini index is defined as

$$gini_A(D) = \frac{|D_1|}{|D|} gini(D_1) + \frac{|D_2|}{|D|} gini(D_2)$$

- reduction in impurity

$$\Delta gini(A) = gini(D) - gini_A(D)$$

- select as the splitting criterion the  $A$  that leads to the **highest**  $\Delta gini(A)$  value

# Example: PlayTennis

Day	Outlook	Temperature	Humidity	Wind	PlayTennis
D1	Sunny	Hot	High	Weak	No
D2	Sunny	Hot	High	Strong	No
D3	Overcast	Hot	High	Weak	Yes
D4	Rain	Mild	High	Weak	Yes
D5	Rain	Cool	Normal	Weak	Yes
D6	Rain	Cool	Normal	Strong	No
D7	Overcast	Cool	Normal	Strong	Yes
D8	Sunny	Mild	High	Weak	No
D9	Sunny	Cool	Normal	Weak	Yes
D10	Rain	Mild	Normal	Weak	Yes
D11	Sunny	Mild	Normal	Strong	Yes
D12	Overcast	Mild	High	Strong	Yes
D13	Overcast	Hot	Normal	Weak	Yes
D14	Rain	Mild	High	Strong	No

# Over-fitting in Decision Trees

Tree output

A1 = overcast: + (4.0)

A1 = sunny:

| A3 = high: - (3.0)

| A3 = normal: + (2.0)

A1 = rain:

| A4 = weak: + (3.0)

| A4 = strong: - (2.0)



# Over-fitting in Decision Trees

## Example

Consider adding noisy training example #15:

*Sunny, Hot, Normal, Strong, PlayTennis = No*




```
A1 = overcast: + (4.0)
A1 = sunny:
|   A3 = high: - (3.0)
|   A3 = normal: + (2.0)
A1 = rain:
|   A4 = weak: + (3.0)
|   A4 = strong: - (2.0)
```

```
A1 = overcast: + (4.0)
A1 = sunny:
|   A2 = hot: - (3.0)
|   A2 = cold: + (1.0)
|   A2 = mild:
|   |   A3 = high: - (1.0)
|   |   A3 = normal: + (1.0)
A1 = rain:
|   A4 = weak: + (3.0)
|   A4 = strong: - (2.0)
```


- more complex than the original tree
- fits the noisy data better than the old tree

# Overfitting

- special **anomalies** of the training set may have been incorporated in the tree



- may further affect the accuracy of the tree on the test set (i.e., during prediction)
- sometimes smaller trees are preferable because of their interpretability



# Tackling Overfitting



- early-stopping
- pruning

# Early-Stopping

- add **extra terminating conditions** at Step 2 of the decision tree induction algorithm
  - stop if the number of tuples is fewer than some user-specified threshold
- may be difficult to choose an appropriate threshold

# Pruning

- grow the whole tree, then prune



How to prune a decision node?

- remove the **subtree** rooted at that node
- make it a leaf node

The leaves that are left may not necessarily be pure

- assign it the most common classification of the training examples affiliated with that node

# Pruning...

The pruned tree usually misclassifies **more** training examples than the unpruned tree

## When to stop pruning?

- uses a validation set
- stop if the pruned tree performs worse than the original on the validation set

## Which node to prune?

- remove the one that most improves validation set accuracy
- **greedy** approach

# Pruning Algorithm

## Reduced-error pruning

Do until further pruning is harmful:

- 1 Evaluate impact on validation set of pruning each possible node
- 2 Greedily remove the one that most improves validation set accuracy