



Perceptron and Multilayer Perceptron

Perceptron

a **feed-forward** network with only **one** layer of adjustable (learnable) weights connected to one or more **threshold** units (as output units)



Perceptron Network



Single Perceptron

Model

input: l_1, l_2, \dots, l_n

- signals from the other neurons

weights: w_1, w_2, \dots, w_n


- can be negative

activation function:

- relating the input and output


$$O = \text{step}(\sum_{j=1}^n w_j l_j - \theta)$$

$$\text{step}(x) = \begin{cases} 0 & \text{if } x < 0 \\ 1 & \text{if } x \geq 0 \end{cases} \quad (\text{step function})$$



What Boolean Functions can Perceptrons Represent?



AND?



- Perceptron output: $O = \text{step}(\sum_{j=0}^n w_j I_j)$
 - **decision boundary:** $\sum_{j=0}^n w_j I_j = 0$


What Boolean Functions can Perceptrons Represent?...

OR?




What Boolean Functions can Perceptrons Represent?...

NOT?




What Boolean Functions can Perceptrons Represent?...

XOR?




Linearly Separable Functions

- a function can be represented by a single perceptron if and only if the function is **linearly separable**



Three points in a plane shattered by a half-space.



Multi-layer Feedforward Networks

- Generalization of simple perceptrons
- **Multi-layer** perceptrons (MLP)


Output units O_i

$W_{j,i}$

Hidden units a_j

$W_{k,j}$


Input units I_k




Hidden Unit Transfer (Activation) Function

Sigmoid unit


- a unit very much like a perceptron, but based on a **smoothed, differentiable** threshold function: $\sigma(x) = \frac{1}{1+e^{-x}}$



(a) Step function




(b) Sign function



(c) Sigmoid function

Rectified Linear Unit (ReLU)


- $f(x) = \max(0, x)$



- the most popular activation function for deep networks
- more efficient computation
- simple gradient
 - if > 0 , gradient = 1
 - if ≤ 0 , gradient = 0

Training: Finding the Weight

- use **gradient descent** to search the space of possible weight vectors to find the weights that **minimizes** the error
- start at any point and keep going **downhill**





Idea: Gradient descent

- start with initial value for \mathbf{w}
- repeat until convergence
 - compute the gradient vector of the error function for current \mathbf{w}
 - move in the opposite direction

Gradient Descent

gradient $\nabla E[\vec{w}]$ at \vec{w} :

$$\left[\frac{\partial E}{\partial w_0}, \frac{\partial E}{\partial w_1}, \dots, \frac{\partial E}{\partial w_n} \right]$$



move \vec{w} :

- **direction**: opposite to $\nabla E[\vec{w}]$
- **magnitude**: a small fraction of $\nabla E[\vec{w}]$


$$\vec{w} \leftarrow \vec{w} + \Delta \vec{w}$$

$$\Delta \vec{w} = -\eta \nabla E[\vec{w}]$$

$$\Delta w_i = -\eta \frac{\partial E}{\partial w_i}$$

in general, the error surface can be very complicated

Back-Propagation



- we need to “propagate error back” when computing the gradient vector